



**MPLAB[®] XC16 Assembler,
Linker and Utilities
User's Guide**

Note the following details of the code protection feature on Microchip devices:

- Microchip products meet the specifications contained in their particular Microchip Data Sheet.
- Microchip believes that its family of products is secure when used in the intended manner and under normal conditions.
- There are dishonest and possibly illegal methods being used in attempts to breach the code protection features of the Microchip devices. We believe that these methods require using the Microchip products in a manner outside the operating specifications contained in Microchip's Data Sheets. Attempts to breach these code protection features, most likely, cannot be accomplished without violating Microchip's intellectual property rights.
- Microchip is willing to work with any customer who is concerned about the integrity of its code.
- Neither Microchip nor any other semiconductor manufacturer can guarantee the security of its code. Code protection does not mean that we are guaranteeing the product is "unbreakable." Code protection is constantly evolving. We at Microchip are committed to continuously improving the code protection features of our products. Attempts to break Microchip's code protection feature may be a violation of the Digital Millennium Copyright Act. If such acts allow unauthorized access to your software or other copyrighted work, you may have a right to sue for relief under that Act.

Information contained in this publication is provided for the sole purpose of designing with and using Microchip products. Information regarding device applications and the like is provided only for your convenience and may be superseded by updates. It is your responsibility to ensure that your application meets with your specifications.

THIS INFORMATION IS PROVIDED BY MICROCHIP "AS IS". MICROCHIP MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND WHETHER EXPRESS OR IMPLIED, WRITTEN OR ORAL, STATUTORY OR OTHERWISE, RELATED TO THE INFORMATION INCLUDING BUT NOT LIMITED TO ANY IMPLIED WARRANTIES OF NON-INFRINGEMENT, MERCHANTABILITY, AND FITNESS FOR A PARTICULAR PURPOSE OR WARRANTIES RELATED TO ITS CONDITION, QUALITY, OR PERFORMANCE.

IN NO EVENT WILL MICROCHIP BE LIABLE FOR ANY INDIRECT, SPECIAL, PUNITIVE, INCIDENTAL OR CONSEQUENTIAL LOSS, DAMAGE, COST OR EXPENSE OF ANY KIND WHATSOEVER RELATED TO THE INFORMATION OR ITS USE, HOWEVER CAUSED, EVEN IF MICROCHIP HAS BEEN ADVISED OF THE POSSIBILITY OR THE DAMAGES ARE FORESEEABLE. TO THE FULLEST EXTENT ALLOWED BY LAW, MICROCHIP'S TOTAL LIABILITY ON ALL CLAIMS IN ANY WAY RELATED TO THE INFORMATION OR ITS USE WILL NOT EXCEED THE AMOUNT OF FEES, IF ANY, THAT YOU HAVE PAID DIRECTLY TO MICROCHIP FOR THE INFORMATION. Use of Microchip devices in life support and/or safety applications is entirely at the buyer's risk, and the buyer agrees to defend, indemnify and hold harmless Microchip from any and all damages, claims, suits, or expenses resulting from such use. No licenses are conveyed, implicitly or otherwise, under any Microchip intellectual property rights unless otherwise stated.

Trademarks

The Microchip name and logo, the Microchip logo, Adaptec, AnyRate, AVR, AVR logo, AVR Freaks, BesTime, BitCloud, chipKIT, chipKIT logo, CryptoMemory, CryptoRF, dsPIC, FlashFlex, flexPWR, HELDO, IGLOO, JukeBlox, KeeLoq, Kleer, LANCheck, LinkMD, maxStylus, maxTouch, MediaLB, megaAVR, Microsemi, Microsemi logo, MOST, MOST logo, MPLAB, OptoLyzer, PackTime, PIC, picoPower, PICSTART, PIC32 logo, PolarFire, Prochip Designer, QTouch, SAM-BA, SenGenuity, SpyNIC, SST, SST Logo, SuperFlash, Symmetricom, SyncServer, Tachyon, TimeSource, tinyAVR, UNI/O, Vectron, and XMEGA are registered trademarks of Microchip Technology Incorporated in the U.S.A. and other countries.

AgileSwitch, APT, ClockWorks, The Embedded Control Solutions Company, EtherSynch, FlashTec, Hyper Speed Control, HyperLight Load, IntelliMOS, Libero, motorBench, mTouch, Powermite 3, Precision Edge, ProASIC, ProASIC Plus, ProASIC Plus logo, Quiet-Wire, SmartFusion, SyncWorld, Temux, TimeCesium, TimeHub, TimePictra, TimeProvider, WinPath, and ZL are registered trademarks of Microchip Technology Incorporated in the U.S.A.

Adjacent Key Suppression, AKS, Analog-for-the-Digital Age, Any Capacitor, AnyIn, AnyOut, Augmented Switching, BlueSky, BodyCom, CodeGuard, CryptoAuthentication, CryptoAutomotive, CryptoCompanion, CryptoController, dsPICDEM, dsPICDEM.net, Dynamic Average Matching, DAM, ECAN, Espresso T1S, EtherGREEN, IdealBridge, In-Circuit Serial Programming, ICSP, INICnet, Intelligent Paralleling, Inter-Chip Connectivity, JitterBlocker, maxCrypto, maxView, memBrain, Mindi, MiWi, MPASM, MPF, MPLAB Certified logo, MPLIB, MPLINK, MultiTRAK, NetDetach, Omniscient Code Generation, PICDEM, PICDEM.net, PICKit, PICtail, PowerSmart, PureSilicon, QMatrix, REAL ICE, Ripple Blocker, RTAX, RTG4, SAM-ICE, Serial Quad I/O, simpleMAP, SimpliPHY, SmartBuffer, SMART-I.S., storClad, SQL, SuperSwitcher, SuperSwitcher II, Switchtec, SynchroPHY, Total Endurance, TSHARC, USBCheck, VariSense, VectorBlox, VeriPHY, ViewSpan, WiperLock, XpressConnect, and ZENA are trademarks of Microchip Technology Incorporated in the U.S.A. and other countries.

SQTP is a service mark of Microchip Technology Incorporated in the U.S.A.

The Adaptec logo, Frequency on Demand, Silicon Storage Technology, and Symmcom are registered trademarks of Microchip Technology Inc. in other countries.

GestIC is a registered trademark of Microchip Technology Germany II GmbH & Co. KG, a subsidiary of Microchip Technology Inc., in other countries.

All other trademarks mentioned herein are property of their respective companies.

© 2013-2021, Microchip Technology Incorporated, All Rights Reserved.

ISBN: 978-1-5224-7664-1

For information regarding Microchip's Quality Management Systems, please visit www.microchip.com/quality.

Table of Contents

Preface	7
Part 1 – MPLAB XC16 Assembler	
<hr/>	
Chapter 1. Assembler Overview	
1.1 Introduction	15
1.2 Feature Set	15
1.3 Assembler Usage	15
1.4 Input/Output Files	16
Chapter 2. Assembler Command Line Options	
2.1 Introduction	19
2.2 Command-Line Syntax	19
2.3 Options that Modify the Listing Output	19
2.4 Options that Control Informational Output	28
2.5 Options that Control Output File Creation	29
2.6 Other Options	31
Chapter 3. MPLAB XC16 Assembly Language	
3.1 Introduction	33
3.2 Internal Preprocessor	33
3.3 Source Code Format	34
3.4 Characters	37
3.5 Constants	38
3.6 Symbols	40
3.7 Expressions	44
3.8 Operators	44
3.9 Special Operators	46
Chapter 4. Assembler Directives	
4.1 Introduction	53
4.2 Directives that Define Sections	54
4.3 Directives that Fill Program Memory	62
4.4 Directives that Initialize Constants	64
4.5 Directives that Declare Symbols	68
4.6 Directives that Define Symbols	69
4.7 Directives that Modify Section Alignment	70
4.8 Directives that Format the Output Listing	76
4.9 Directives that Control Conditional Assembly	77

16-Bit Assembler, Linker and Utilities User's Guide

4.10 Directives for Substitution/Expansion	79
4.11 Miscellaneous Directives	82
4.12 Directives for Debug Information	84

Chapter 5. Assembler Listing File

5.1 Introduction	87
5.2 Generation	87
5.3 Contents	88

Chapter 6. Assembler Errors/Warnings/Messages

6.1 Introduction	91
6.2 Fatal Errors	91
6.3 Errors	92
6.4 Warnings	100
6.5 Messages	106

Part 2 – MPLAB XC16 Object Linker

Chapter 7. Linker Overview

7.1 Introduction	109
7.2 Feature Set	109
7.3 Linker Usage	109
7.4 Input/Output Files	110

Chapter 8. Linker Command Line Options

8.1 Introduction	113
8.2 Highlights	113
8.3 Syntax	113
8.4 Options that Control Output File Creation	115
8.5 Options that Control Run-time Initialization	125
8.6 Options that Control Informational Output	127
8.7 Options that Modify the Link Map Output	130
8.8 Options that Specify CodeGuard™ Security Features	131
8.9 Options that Control the Preprocessor	133

Chapter 9. Linker Scripts

9.1 Introduction	135
9.2 Highlights	135
9.3 Overview of Linker Scripts	135
9.4 Command Line Information	136
9.5 Contents of a Linker Script	136
9.6 Creating a Custom Linker Script	149
9.7 Linker Script Command Language	149
9.8 Expressions in Linker Scripts	166

Chapter 10. Linker Processing

10.1 Introduction	173
10.2 Highlights	173
10.3 Overview of Linker Processing	173

10.4 Memory Addressing	176
10.5 Linker Allocation	178
10.6 Global and Weak Symbols	184
10.7 Handles	185
10.8 Initialized Data	186
10.9 Read-only Data	189
10.10 Stack Allocation	191
10.11 Heap Allocation	192
10.12 Interrupt Vector Tables	192
10.13 Optimizing Memory Usage	194
10.14 Boot and Secure Segments	198
10.15 Co-resident Application Linking	202
10.16 Notable Symbols	204
Chapter 11. Linker Examples	
11.1 Introduction	205
11.2 Highlights	205
11.3 Memory Addresses and Relocatable Code	206
11.4 Locating a Variable at a Specific Address	207
11.5 Locating a Function at a Specific Address	207
11.6 Using More than 32K of Constants	208
11.7 Locating a Constant at a Specific Address in Program Memory	210
11.8 Locating and Accessing Data in EEPROM Memory	211
11.9 Creating an Incrementing Modulo Buffer in X Memory	212
11.10 Creating a Decrementing Modulo Buffer in Y Memory	213
11.11 Locating the Stack at a Specific Address	213
11.12 Locating and Reserving Program Memory	214
Chapter 12. Linker Map File	
12.1 Introduction	215
12.2 Generation	215
12.3 Contents	215
Chapter 13. Linker Errors/Warnings	
13.1 Introduction	223
13.2 Highlights	223
13.3 Errors	223
13.4 Warnings	229
Part 3 – 16-Bit Utilities (including the Archiver/Librarian)	
Chapter 14. MPLAB XC16 Object Archiver/Librarian	
14.1 Introduction	233
14.2 Highlights	233
14.3 Archiver/Librarian and Other Development Tools	234
14.4 Feature Set	234
14.5 Input/Output Files	234

16-Bit Assembler, Linker and Utilities User's Guide

14.6 Syntax	234
14.7 Options	235
14.8 Scripts	237
Chapter 15. Other Utilities	
15.1 Introduction	239
15.2 Highlights	239
15.3 xc16-bin2hex Utility	240
15.4 xc16-nm Utility	242
15.5 xc16-objdump Utility	245
15.6 xc16-ranlib Utility	249
15.7 xc16-strings Utility	250
15.8 xc16-strip Utility	252
Part 4 – Appendices	
Appendix A. Deprecated Features	
A.1 Introduction	257
A.2 Highlights	257
A.3 Assembler Directives that Define Sections	257
A.4 Reserved Section Names with Implied Attributes	258
A.5 Environmental Variables	258
Appendix B. Useful Tables	
B.1 Introduction	259
B.2 Highlights	259
B.3 ASCII Character Set	259
B.4 Hexadecimal to Decimal Conversion	260
Appendix C. GNU Free Documentation License	
C.1 Preamble	261
C.2 Applicability and Definitions	261
C.3 Verbatim Copying	262
C.4 Copying in Quantity	263
C.5 Modifications	264
C.6 Combining Documents	265
C.7 Collections of Documents	265
C.8 Aggregation with Independent Works	266
C.9 Translation	266
C.10 Termination	266
C.11 Future Revisions of this License	267
C.12 Relicensing	267
Appendix D. Document Revision History	
Index	271
Worldwide Sales and Service	281

Preface

NOTICE TO CUSTOMERS

All documentation becomes dated, and this manual is no exception. Microchip tools and documentation are constantly evolving to meet customer needs, so some actual dialogs and/or tool descriptions may differ from those in this document. Please refer to our web site (www.microchip.com) to obtain the latest documentation available.

Documents are identified with a “DS” number. This number is located on the bottom of each page, in front of the page number. The numbering convention for the DS number is “DSXXXXA”, where “XXXX” is the document number and “A” is the revision level of the document.

For the most up-to-date information on development tools, see the MPLAB® X IDE online help, available from the Help menu.

INTRODUCTION

This chapter contains general information that will be useful to know before using 16-bit language tools. Items discussed include:

- [Document Layout](#)
- [Conventions Used in this Guide](#)
- [Recommended Reading](#)
- [The Microchip Web Site](#)
- [myMicrochip Personalized Notification Service](#)
- [Customer Support](#)

DOCUMENT LAYOUT

This document describes how to use GNU language tools to write code for 16-bit applications. The document layout is as follows:

Part 1 – MPLAB XC16 Assembler

- [Chapter 1. “Assembler Overview”](#) – gives an overview of assembler operation.
- [Chapter 2. “Assembler Command Line Options”](#) – details command line options for the assembler.
- [Chapter 3. “MPLAB XC16 Assembly Language”](#) – describes syntax used with the assembler.
- [Chapter 4. “Assembler Directives”](#) – details the available assembler directives.
- [Chapter 5. “Assembler Listing File”](#) – describes what symbols are and how to use them.
- [Chapter 6. “Assembler Errors/Warnings/Messages”](#) – contains a descriptive list of the errors, warnings and messages generated by the 16-bit assembler.

Part 2 – MPLAB XC16 Object Linker

- [Chapter 7. “Linker Overview”](#) – gives an overview of linker operation.
- [Chapter 8. “Linker Command Line Options”](#) – details command line options for the linker.
- [Chapter 9. “Linker Scripts”](#) – describes how to generate and use linker scripts to control linker operation.
- [Chapter 10. “Linker Processing”](#) – discusses how the linker builds an application from input files.
- [Chapter 11. “Linker Examples”](#) – discusses a number of 16-bit specific linker examples and shows the equivalent syntax in C and assembly language.
- [Chapter 13. “Linker Errors/Warnings”](#) – contains a descriptive list of the errors and warnings generated by the 16-bit linker.

Part 3 – 16-Bit Utilities (including the Archiver/Librarian)

- [Chapter 14. “MPLAB XC16 Object Archiver/Librarian”](#) – details command line options for the librarian.
- [Chapter 15. “Other Utilities”](#) – gives an overview of all the other utilities and their operation. Current utilities are:
 - [xc16-bin2hex](#): Converts a linked object file into an Intel® hex file.
 - [xc16-nm Utility](#): Lists symbols from an object file.
 - [xc16-objdump Utility](#): Displays information about object files.
 - [xc16-ranlib Utility](#): Generates an index from the contents of an archive and stores it in the archive.
 - [xc16-strings Utility](#): Prints the printable character sequences.
 - [xc16-strip Utility](#): Discards all symbols from an object file.

Part 4 – Appendices

- [Appendix A. “Deprecated Features”](#) – describes features that are considered obsolete.
- [Appendix B. “Useful Tables”](#) – lists some useful tables: the ASCII character set and hexadecimal to decimal conversion.
- [Appendix C. “GNU Free Documentation License”](#) – details the license requirements for using the GNU language tools.

CONVENTIONS USED IN THIS GUIDE

The following conventions may appear in this documentation:

DOCUMENTATION CONVENTIONS

Description	Represents	Examples
Arial font:		
Italic characters	Referenced books	<i>MPLAB® X IDE User's Guide</i>
	Emphasized text	...is the <i>only</i> compiler...
Initial caps	A window	the Output window
	A dialog	the Settings dialog
	A menu selection	select Enable Programmer
Quotes	A field name in a window or dialog	"Save project before build"
Underlined, italic text with right angle bracket	A menu path	<u><i>File>Save</i></u>
Bold characters	A dialog button	Click OK
	A tab	Click the Power tab
Text in angle brackets < >	A key on the keyboard	Press <Enter>, <F1>
Courier font:		
Plain Courier	Sample source code	#define START
	Filenames	autoexec.bat
	File paths	c:\mcc18\h
	Keywords	_asm, _endasm, static
	Command-line options	-Opa+, -Opa-
	Bit values	0, 1
	Constants	0xFF, 'A'
Italic Courier	A variable argument	<i>file.o</i> , where <i>file</i> can be any valid filename
Square brackets []	Optional arguments	mpasmwin [options] <i>file</i> [options]
Curly brackets and pipe character: { }	Choice of mutually exclusive arguments; an OR selection	errorlevel {0 1}
Ellipses...	Replaces repeated text	var_name [, var_name...]
	Represents code supplied by user	void main (void) { ... }
Sidebar Text		
DD	Device Dependent. This feature is not supported on all devices. Devices supported will be listed in the title or text.	xmemory attribute

16-Bit Assembler, Linker and Utilities User's Guide

RECOMMENDED READING

This documentation describes how to use 16-bit language tools. Other useful documents are listed below. The following Microchip documents are available and recommended as supplemental reference resources.

Readme Files

For the latest information on Microchip tools, read the associated Readme files (HTML files) included with the software.

16-Bit Language Tools Getting Started (DS70094)

A guide to installing and working with the Microchip language tools for 16-bit devices. Examples using the 16-bit simulator SIM30 (a component of MPLAB SIM) are provided.

MPLAB® XC16 C Compiler User's Guide (DS50002071)

A guide to using the 16-bit C compiler. The 16-bit linker is used with this tool.

16-Bit Language Tools Libraries Reference Manual (DS50001456)

A descriptive listing of libraries available for Microchip 16-bit devices. This includes standard (including math) libraries and compiler built-in functions. DSP and 16-bit peripheral libraries are described in Release Notes provided with each peripheral library type.

Device-Specific Documentation

The Microchip website contains many documents that describe 16-bit device functions and features. Among these are:

- Individual and family data sheets
- Family reference manuals
- Programmer's reference manuals

THE MICROCHIP WEB SITE

Microchip provides online support via our web site at www.microchip.com. This web site is used as a means to make files and information easily available to customers. Accessible by using your favorite Internet browser, the web site contains the following information:

- **Product Support** – Data sheets and errata, application notes and sample programs, design resources, user's guides and hardware support documents, latest software releases and archived software
- **General Technical Support** – Frequently Asked Questions (FAQs), technical support requests, online discussion groups, Microchip consultant program member listing
- **Business of Microchip** – Product selector and ordering guides, latest Microchip press releases, listing of seminars and events, listings of Microchip sales offices, distributors and factory representatives

myMICROCHIP PERSONALIZED NOTIFICATION SERVICE

Microchip's personal notification service helps keep customers current on their Microchip products of interest. Subscribers will receive e-mail notification whenever there are changes, updates, revisions or errata related to a specified product family or development tool.

Please visit <http://www.microchip.com/pcn> to begin the registration process and select your preferences to receive personalized notifications. A FAQ and registration details are available on the page, which can be opened by selecting the link above.

When you are selecting your preferences, choosing "Development Systems" will populate the list with available development tools. The main categories of tools are listed below:

- **Compilers** – The latest information on Microchip C compilers, assemblers, linkers and other language tools. These include all MPLAB[®] C compilers; all MPLAB assemblers (including MPASM[™] assembler); all MPLAB linkers (including MPLINK[™] object linker); and all MPLAB librarians (including MPLIB[™] object librarian).
- **Emulators** – The latest information on Microchip in-circuit emulators. These include the MPLAB REAL ICE[™] and MPLAB ICE 2000 in-circuit emulators
- **In-Circuit Debuggers** – The latest information on Microchip in-circuit debuggers. These include the MPLAB ICD 2 and 3 in-circuit debuggers and PICKit[™] 2 and 3 debug express.
- **MPLAB[®] IDE** – The latest information on Microchip MPLAB IDE, the Windows[®] Integrated Development Environment for development systems tools. This list is focused on the MPLAB IDE, MPLAB IDE Project Manager, MPLAB Editor and MPLAB SIM simulator, as well as general editing and debugging features.
- **Programmers** – The latest information on Microchip programmers. These include the device (production) programmers MPLAB REAL ICE in-circuit emulator, MPLAB ICD 3 in-circuit debugger, MPLAB PM3, and PRO MATE[®] II and development (nonproduction) programmers MPLAB ICD 2 in-circuit debugger, PICSTART[®] Plus and PICKit 1, 2 and 3.
- **Starter/Demo Boards** – These include MPLAB Starter Kit boards, PICDEM[™] demo boards, and various other evaluation boards.

16-Bit Assembler, Linker and Utilities User's Guide

CUSTOMER SUPPORT

Users of Microchip products can receive assistance through several channels:

- Distributor or Representative
- Local Sales Office
- Field Application Engineer (FAE)
- Technical Support

Customers should contact their distributor, representative or field application engineer (FAE) for support. Local sales offices are also available to help customers. A listing of sales offices and locations is included in the back of this document.

Technical support is available through the web site at: <http://support.microchip.com>.

Send notification of documentation errors or comments to Microchip via e-mail to docerrors@microchip.com.



Part 1 – MPLAB XC16 Assembler

Chapter 1. Assembler Overview	15
Chapter 2. Assembler Command Line Options	19
Chapter 3. MPLAB XC16 Assembly Language	33
Chapter 4. Assembler Directives	53
Chapter 5. Assembler Listing File	87
Chapter 6. Assembler Errors/Warnings/Messages	91

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Chapter 1. Assembler Overview

1.1 Introduction

MPLAB XC16 Assembler produces relocatable machine code from symbolic assembly language for the dsPIC® DSC and PIC24 MCU families of devices. The assembler is an application that provides a platform for developing assembly language code. The assembler is a port of the GNU assembler from the Free Software Foundation.

Topics covered in this chapter are:

- [Feature Set](#)
- [Assembler Usage](#)
- [Input/Output Files](#)

1.2 Feature Set

Notable features of the assembler include:

- Support for the entire 16-bit instruction set
- Support for fixed-point and floating-point data
- Support for ELF and COFF object formats
- Available for Windows®, Linux® and Mac OS®
- Command Line Interface
- Rich Directive Set
- Flexible Macro Language
- Available for MPLAB® X IDE and MPLAB IDE v8

1.3 Assembler Usage

The MPLAB XC16 Assembler translates user assembly source files into relocatable object files. These object files can then be put into an archive (MPLAB XC16 Object Archiver/Librarian) or linked with other relocatable object files and archives to create an executable file (MPLAB XC16 Object Linker). See the “*MPLAB XC16 C Compiler User's Guide*” (DS50002071) for an overview of the tools process flow.

Typically the command-line driver, `xc16-gcc`, is used to invoke the assembler as it can be passed assembler source files as input; however, the options for the assembler are supplied here for instances where the assembler is being called directly, or when options need to be set in the assembler tab of the Build Options dialog (MPLAB IDE v8) or assembler category of the Project Properties window (MPLAB X IDE).

The assembler command line may contain options and file names. For details on command line option syntax, see [Section 2.2 “Command-Line Syntax”](#).

Note that the assembler will not produce any messages unless there are errors or warnings – there are no “assembly completed” messages. For more on messages, see [Chapter 6. “Assembler Errors/Warnings/Messages.”](#)

16-Bit Assembler, Linker and Utilities User's Guide

1.4 Input/Output Files

Standard assembler input and output files are listed below.

Extension	Description
Input	
.s	Source File
Output	
.o	Object File
.lst	Listing File

Unlike the MPASM™ assembler (for use with 8-bit PIC® MCUs), MPLAB XC16 Assembler does not generate error files, hex files, or symbol and debug files. The assembler is capable of creating a listing file and a relocatable object file (that may or may not contain debugging information). MPLAB XC16 Object Linker is used with the assembler to produce the final object files, map files and final executable file for debugging with MPLAB X IDE or MPLAB IDE v8 (see [Figure 1.2](#)).

1.4.1 Source File

The assembler accepts, as input, a source file that consists of 16-bit device instructions, assembler directives and comments. A sample source file is shown in [Example 1-1](#).

Note: Microchip Technology strongly suggests an .s extension for assembly source files. This will enable you to easily use the C compiler driver without having to specify the option to tell the driver that the file should be treated as an assembly file. See the “*MPLAB® XC16 C Compiler User's Guide*” (DS50002071) for more details on the C compiler driver.

Example 1-1: Sample Assembler Code

```
.title " Sample dsPIC Assembler Source Code"
.sbttl " For illustration only."

; dsPIC registers
.equ CORCONL, CORCON
.equ PSV,2

.section .const,psv
hello:
.ascii "Hello world!\n\0"

.text
.global __reset
__reset:
; set PSVPAG to page that contains 'hello'
mov    #psvpage(hello),w0
mov    w0,PSVPAG

; enable Program Space Visibility
bset.b CORCONL,#PSV

; make a pointer to 'hello'
mov    #psvoffset(hello),w0

.end
```

For more information, see [Section 2.2 “Command-Line Syntax”](#) and [Chapter 4. “Assembler Directives.”](#)

1.4.2 Object File

The assembler creates a relocatable object file. These object files do not yet have addresses resolved and must be linked before they can be used for executables.

By default, the name of the object file created is `a.out`. Specify the `-o` option (see [Chapter 2. “Assembler Command Line Options.”](#)) on the command line to override the default name.

By default, object files are created in the ELF format. To specify ELF or COFF format explicitly, use the `-omf` option on the command line, as shown:

```
xc16-as -omf=elf test.s
xc16-as -omf=coff test2.s
```

Alternatively, the environment variable `XC16_OMF` may be used to specify object file format for the 16-bit language tools.

1.4.3 Listing File

The assembler has the capability to produce listing files. For details on how to generate a listing file and the components of that file, see [Chapter 5. “Assembler Listing File.”](#)

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Chapter 2. Assembler Command Line Options

2.1 Introduction

MPLAB XC16 Assembler may be used on the command line interface as well as with MPLAB X IDE or MPLAB IDE v8. The following options may be used with any of these interfaces.

Topics covered in this chapter are:

- [Command-Line Syntax](#)
- [Options that Modify the Listing Output](#)
- [Options that Control Informational Output](#)
- [Options that Control Output File Creation](#)
- [Other Options](#)

2.2 Command-Line Syntax

The assembler command line may contain options and file names. Options may appear in any order and may be before, after or between file names. The order of file names determines the order of assembly.

```
xc16-as [options|sourcefiles]...
```

'--' (two hyphens) by itself names the standard input file explicitly as one of the files for the assembler to translate. Except for '--', any command line argument that begins with a hyphen ('-') is an option. Each option changes the behavior of the assembler, but no option changes the way another option works.

Some options require exactly one file name to follow them. The file name may either immediately follow the option's letter or it may be the next command line argument. For example, to specify an output file named `test.o`, either of the following options would be acceptable:

- `-o test.o`
- `-otest.o`

Note: Command line options are case sensitive.

2.3 Options that Modify the Listing Output

The following options are used to control the listing output. For debugging and general analysis of code operation, a listing file is helpful. Constructing one with useful information is accomplished using the options in this section.

- `-a[suboption] [=file]`
- `--listing-lhs-width #`
- `--listing-lhs-width2 #`
- `--listing-rhs-width #`
- `--listing-cont-lines #`

16-Bit Assembler, Linker and Utilities User's Guide

2.3.1 `-a[suboption] [=file]`

The `-a` option enables listing output. The `-a` option supports the following suboptions to further control what is included in the assembly listing:

<code>-ac</code>	Omit false conditionals
<code>-ad</code>	Omit debugging directives
<code>-ah</code>	Include high-level source
<code>-ai</code>	Include section information
<code>-al</code>	Include assembly
<code>-am</code>	Include macro expansions
<code>-an</code>	Omit forms processing
<code>-as</code>	Include symbols
<code>-a=file</code>	Output listing to specified file (must be in current directory).

If no suboptions are specified, the default suboptions used are `hls`; the `-a` option by itself requests high-level, assembly, and symbolic listing. You can use other letters to select specific options for the listing output.

The letters after the `-a` may be combined into one option. So, for example, instead of specifying `-al -an` on the command line, you could specify `-aln`. Most of the examples in the following sections combine the section's suboption with `-al`, because `-al` is required for an assembly listing.

2.3.1.1 `-ac`

`-ac` omits false conditionals from a listing. Any lines that are not assembled because of a false `.if` or `.ifdef` (or the `.else` of a true `.if` or `.ifdef`) will be omitted from the listing. [Example 2-1](#) shows a listing where the `-ac` option was not used.

[Example 2-2](#) shows a listing for the same source where the `-ac` option was used.

Example 2-1: Listing File Generated with `-al` Command Line Option

MPLAB ASM30 Listing: `example2.1.s` page 1

```
1          .data
2          .if 0
3              .if 1
4              .endif
5              .long 0
6              .if 0
7                  .long 0
8              .endif
9          .else
10             .if 1
11             .endif
12 0000 02 00 00 00          .long 2
13             .if 0
14                 .long 3
15             .else
16 0004 04 00 00 00          .long 4
17             .endif
18         .endif
19
20         .if 0
21             .long 5
22         .elseif 1
23             .if 0
24                 .long 6
```

Assembler Command Line Options

```
25 .elseif 1
26 0008 07 00 00 00 .long 7
27 .endif
28 .elseif 1
29 .long 8
30 .else
31 .long 9
32 .endif
```

Example 2-2: Listing File Generated with -alc Command Line Option

MPLAB ASM30 Listing: example2.2.s page 1

```
1 .data
2 .if 0
9 .else
10 .if 1
11 .endif
12 0000 02 00 00 00 .long 2
13 .if 0
15 .else
16 0004 04 00 00 00 .long 4
17 .endif
18 .endif
19
20 .if 0
22 .elseif 1
23 .if 0
25 .elseif 1
26 0008 07 00 00 00 .long 7
27 .endif
28 .elseif 1
30 .else
32 .endif
```

Note: Some lines have been omitted, due to the `-ac` option; i.e., lines 3-8, 14, 21, 24, 29 and 31.

2.3.1.2 -ad

`-ad` omits debugging directives from the listing. This is useful if a compiler that was given a debugging option generated the assembly source code. The compiler-generated debugging directives will not clutter the listing. [Example 2-3](#) shows a listing using both the `d` and `h` suboptions. Compared to using the `h` sub-option alone (see the next section), the listing is much cleaner.

16-Bit Assembler, Linker and Utilities User's Guide

Example 2-3: Listing File Generated with -alh Command Line Option

MPLAB ASM30 Listing: example2.3.s

page 1

```
1          .file "example2.3.c"
2          .text
3          .align 2
9          .global _main ; export
10         _main:
1:example2.3.c **** extern int ADD (int, int);
2:example2.3.c ****
3:example2.3.c **** int
4:example2.3.c **** main(void)
5:example2.3.c **** {
16          .set    __PA__,1
17 000000 00 00 FA      lnk    #0
18
6:example2.3.c **** return ADD(4, 5);
20 000002 51 00 20      mov    #5,w1
21 000004 40 00 20      mov    #4,w0
22 000006 00 00 02      call   _ADD
22          00 00 00
7:example2.3.c **** }
29
30 00000a 00 80 FA      ulnk
31 00000c 00 00 06      return
32          .set    __PA__,0
37
38          .end
```

2.3.1.3 -ah

-ah requests a high-level language listing. High-level listings require that the assembly source code is generated by a compiler, a debugging option like -g is given to the compiler, and assembly listings (-al) are requested. -al requests an output program assembly listing. [Example 2-4](#) shows a listing that was generated using the -alh command line option.

Assembler Command Line Options

Example 2-4: Listing File Generated with -alh Command Line Option

MPLAB ASM30 Listing: example2.4.s

page 1

```
1          .file "example2.4.c"
2          .text
3          .align 2
4          .def _main
5          .val _main
6          .scl 2
7          .type 044
8          .endef
9          .global _main ; export
10         _main:
11             .def .bf
12             .val .
13             .scl 101
14         1:example2.4.c **** extern int ADD (int, int);
15         2:example2.4.c ****
16         3:example2.4.c **** int
17         4:example2.4.c **** main(void)
18         5:example2.4.c **** {
19             .line 5
20             .endef
21             .set __PA__,1
22         17 000000 00 00 FA      lnk #0
23         18
24         6:example2.4.c **** return ADD(4, 5);
25         19             .ln 6
26         20 000002 51 00 20      mov #5,w1
27         21 000004 40 00 20      mov #4,w0
28         22 000006 00 00 02      call _ADD
29         22 00 00 00
30         7:example2.4.c **** }
31         23             .ln 7
32         24             .def .ef
33         25             .val .
34         26             .scl 101
35         27             .line 7
36         28             .endef
37         29
38         30 00000a 00 80 FA      ulnk
39         31 00000c 00 00 06      return
40         32             .set __PA__,0
41         33             .def _main
42         34             .val .
43         35             .scl -1
44         36             .endef
45         37
46         38             .end
```

2.3.1.4 -ai

-ai displays information on each of the code and data sections. This information contains details on the size of each of the sections and then a total usage of program and data memory. [Example 2-5](#) shows a listing where the -ai option was used.

16-Bit Assembler, Linker and Utilities User's Guide

Example 2-5: Listing File Generated with -ai Command Line Option

SECTION INFORMATION:

Section	Length (PC units)	Length (bytes) (dec)
-----	-----	-----
.text	0x16	0x21 (33)

TOTAL PROGRAM MEMORY USED (bytes): 0x21 (33)

Section	Length (bytes) (dec)
-----	-----
.data	0 (0)
.bss	0 (0)

TOTAL DATA MEMORY USED (bytes): 0 (0)

2.3.1.5 -al

`-al` requests an assembly listing. This sub-option may be used with other suboptions. See the other examples in this section.

2.3.1.6 -am

`-am` expands macros in a listing. [Example 2-6](#) shows a listing where the `-am` option was not used. [Example 2-7](#) shows a listing for the same source where the `-am` option was used.

Example 2-6: Listing File Generated with -al Command Line Option

MPLAB ASM30 Listing: example2.5.s page 1

```
1          .text
2          .macro div_s reg1, reg2
3              repeat #18-1
4                  div.sw \reg1,\reg2
5          .endm
6
7          .macro div_u reg1, reg2
8              repeat #18-1
9                  div.uw \reg1,\reg2
10         .endm
11
12 000000 40 01 20      mov #20, w0
13 000002 52 00 20      mov #5, w2
14 000004 11 00 09      div_u w0, w2
14          02 80 D8
15
16 000008 00 02 BE      mov.d w0, w4
17
18 00000a 40 01 20      mov #20, w0
19 00000c B3 FF 2F      mov #-5, w3
20 00000e 11 00 09      div_s w0, w3
20          03 00 D8
```


Assembler Command Line Options

Example 2-7: Listing File Generated with -alm Command Line Option

MPLAB ASM30 Listing: example2.6.s page 1

```
1 .text
2 .macro div_s reg1, reg2
3     repeat #18-1
4         div.sw \reg1,\reg2
5     .endm
6
7 .macro div_u reg1, reg2
8     repeat #18-1
9         div.uw \reg1,\reg2
10    .endm
11
12 000000 40 01 20    mov #20, w0
13 000002 52 00 20    mov #5, w2
14                    div_u w0, w2
14 000004 11 00 09    > repeat #18-1
14 000006 02 80 D8    > div.uw w0,w2
15
16 000008 00 02 BE    mov.d w0, w4
17
18 00000a 40 01 20    mov #20, w0
19 00000c B3 FF 2F    mov #-5, w3
20                    div_s w0, w3
20 00000e 11 00 09    > repeat #18-1
20 000010 03 00 D8    > div.sw w0,w3
```

Note: > signifies expanded macro instructions.

2.3.1.7 -an

-an turns off all forms processing that would be performed by the listing directives .psize, .eject, .title and .sbttl. [Example 2-8](#) shows a listing where the -an option was not used. [Example 2-9](#) shows a listing for the same source where the -an option was used.

Example 2-8: Listing File Generated with -al Command Line Option

MPLAB ASM30 Listing: example2.7.s page 1

User's Guide Example

Listing Options

```
1 .text
2 .title "User's Guide Example"
3 .sbttl " Listing Options"
4 .psize 10
5
6 000000 50 00 20    mov #5, w0
7 000002 61 00 20    mov #6, w1
```

MPLAB ASM30 Listing: example2.7.s page 2

User's Guide Example

Listing Options

```
8 000004 01 01 40    add w0, w1, w2
9 .eject
```

MPLAB ASM30 Listing: example2.7.s page 3

User's Guide Example

Listing Options

```
10
11 000006 24 00 20    mov #2, w4
```

16-Bit Assembler, Linker and Utilities User's Guide

```
12 000008 03 00 09          repeat #3
13 00000a 04 22 B8          mul.uu w4, w4, w4
14
15 00000c 16 00 20          mov #1, w6
16 00000e 64 33 DD          sl w6, #4, w6
MPLAB ASM30 Listing: example2.7.s          page 4
User's Guide Example
Listing Options
17
18 000010 06 20 E1          cp w4, w6
19 000012 00 00 32          bra z, done
20
21 000014 00 00 00          nop
22
23                          done:
MPLAB ASM30 Listing: example2.7.s          page 5
User's Guide Example
Listing Options
24
25                          .end
```

Example 2-9: Listing File Generated with -aln Command Line Option

```
1                          .text
2                          .title "User's Guide Example"
3                          .sbttl " Listing Options"
4                          .psize 10
5
6 000000 50 00 20          mov #5, w0
7 000002 61 00 20          mov #6, w1
8 000004 01 01 40          add w0, w1, w2
9                          .eject
10
11 000006 24 00 20          mov #2, w4
12 000008 03 00 09          repeat #3
13 00000a 04 22 B8          mul.uu w4, w4, w4
14
15 00000c 16 00 20          mov #1, w6
16 00000e 64 33 DD          sl w6, #4, w6
17
18 000010 06 20 E1          cp w4, w6
19 000012 00 00 32          bra z, done
20
21 000014 00 00 00          nop
22
23                          done:
24
25                          .end
```

Assembler Command Line Options

2.3.1.8 -as

`-as` requests a symbol table listing. [Example 2-10](#) shows a listing that was generated using the `-as` command line option. Note that both defined and undefined symbols are listed.

Example 2-10: Listing File Generated with `-as` Command Line Option

MPLAB ASM30 Listing: sample2b.s

```
DEFINED SYMBOLS
                                *ABS*:00000000 fake
                                sample2b.s:4 .text:00000000 __reset
                                sample2b.s:13 .text:0000001c L2
                                                .text:00000000 .text
                                                .data:00000000 .data
                                                .bss:00000000 .bss

UNDEFINED SYMBOLS
__i
__j
```

2.3.1.9 -a=*file*

`=file` defines the name of the output file. This file must be in the current directory.

2.3.2 --listing-lhs-width

The `--listing-lhs-width` option is used to set the width of the output data column of the listing file. By default, this is set to 3 for program memory and 4 for data memory. The following line is extracted from a listing. The output data column is in bold text.

```
6 000000 50 00 20 mov #5, w0
```

If the option `--listing-lhs-width 2` is used, then the same line will appear as follows in the listing:

```
6 000000 50 00      mov #5, w0
6          20
```

2.3.3 --listing-lhs-width2

The `--listing-lhs-width2` option is used to set the width of the continuation lines of the output data column of the listing file. By default, this is set to 3 for program memory and 4 for data memory. If the specified width is smaller than the first line, this option is ignored. The following lines are extracted from a listing. The output data column is in bold.

```
2 0000 50 6C 65 61 .ascii "Please pay inside."
2      73 65 20 70
2      61 79 20 69
2      6E 73 69 64
2      65 2E
```

If the option `--listing-lhs-width2 7` is used, then the same line will appear as follows in the listing:

```
2 0000 50 6C 65 61 .ascii "Please pay inside."
2      73 65 20 70 61 79 20
2      69 6E 73 69 64 65 2E
```

2.3.4 --listing-rhs-width

The `--listing-rhs-width` option is used to set the maximum width in characters of the lines from the source file. By default, this is set to 100. The following lines are extracted from a listing that was created without using the `--listing-rhs-width` option. The text in bold are the lines from the source file.

```
2 0000 54 68 69 73 .ascii "This line is long."
2          20 6C 69 6E
2          65 20 69 73
2          20 6C 6F 6E
2          67 65 72 20
```

If the option `--listing-rhs-width 20` is used, then the same line will appear as follows in the listing:

```
2 0000 54 68 69 73 .ascii "This line i
2          20 6C 69 6E
2          65 20 69 73
2          20 6C 6F 6E
2          67 65 72 20
```

The line is truncated (not wrapped) in the listing, but the data is still there.

2.3.5 --listing-cont-lines

The `--listing-cont-lines` option is used to set the maximum number of continuation lines used for the output data column of the listing. By default, this is 8. The following lines are extracted from a listing that was created without using the `--listing-cont-lines` option. The text in bold shows the continuation lines used for the output data column of the listing.

```
2 0000 54 68 69 73 .ascii "This is a long character sequence."
2          20 69 73 20
2          61 20 6C 6F
2          6E 67 20 63
2          68 61 72 61
2          63 74 65 72
2          20 73 65 71
2          75 65 6E 63
2          65 2E
```

Notice that the number of bytes displayed matches the number of bytes in the ASCII string; however, if the option `--listing-cont-lines 2` is used, then the output data will be truncated after 2 continuation lines as shown below.

```
2 0000 54 68 69 73 .ascii "This is a long character sequence."
2          20 69 73 20
2          61 20 6C 6F
```

2.4 Options that Control Informational Output

The options in this section control how information is output. Errors, warnings and messages concerning code translation and execution are controlled through several of the options in this section.

Any item in parenthesis shows the short method of specifying the option, e.g., `--no-warn` also may be specified as `-W`.

2.4.1 --fatal-warnings

Warnings are treated as if they were errors.

Assembler Command Line Options

2.4.2 `--no-warn (-W)`

Warnings are suppressed. If you use this option, no warnings are issued. This option only affects the warning messages. It does not change how your file is assembled. Errors are still reported.

2.4.3 `--warn`

Warnings are issued, if appropriate. This is the default behavior.

2.4.4 `-J`

No warnings are issued about signed overflow.

2.4.5 `--help`

The assembler will show a message regarding the command line usage and options. The assembler then exits.

2.4.6 `--target-help`

The assembler will show a message regarding the 16-bit device specific command line options. The assembler then exits.

2.4.7 `--version`

The assembler version number is displayed. The assembler then exits.

2.4.8 `--verbose (-v)`

The assembler version number is displayed. The assembler does not exit. If this is the only command line option used, then the assembler will print out the version and wait for entry of the assembly source through standard input. Use `<CTRL>-D` to send an EOF character to end assembly.

2.5 Options that Control Output File Creation

The options in this section control how the output file is created. For example, to change the name of the output object file, use `-o`.

Any item in parenthesis shows the short method of specifying the option, e.g., `--keep-locals` may be specified as `-L` also.

2.5.1 `-g`

Generate symbolic debugging information.

Note: For COFF, the option <code>-g</code> does not work with any section other than <code>.text</code> .
--

2.5.2 `--keep-locals (-L)`

Keep local symbols, i.e., labels beginning with `.L` (upper case only). Normally you do not see such labels when debugging, because they are intended for the use of programs (like compilers) that compose assembler programs. Normally both the assembler and linker discard such symbols. This option tells the assembler to retain those symbols in the object files.

2.5.3 `-o objfile`

Name the object file output *objfile*. In the absence of errors, there is always one object file output when you run the assembler. By default, it has the name `a.out`. Use this option (which takes exactly one filename) to give the object file a different name. Whatever the object file is called, the assembler overwrites any existing file with the same name.

2.5.4 `-omf = format`

Use this option to specify the object file format. Valid format names are ELF and COFF. Object file format names are not case sensitive.

2.5.5 `-R`

This option tells the assembler to write the object file as if all data-section data is located in the text section. The data section part of your object file is zero bytes long because all its bytes are located in the text section.

2.5.6 `--relax`

Turn relaxation on. Convert absolute calls and gotos to relative calls and branches when possible.

2.5.7 `--no-relax`

Turn relaxation off. This is the default behavior.

2.5.8 `-Z`

Generate object file even after errors. After an error message, the assembler normally produces no output. If for some reason, you are interested in object file output even after the assembler gives an error message, use the `-Z` option. If there are any errors, the assembler continues anyway, and writes an object file after a final warning message of the form "n errors, m warnings, generating bad object file".

2.5.9 `-MD file`

Write dependency information to *file*. The assembler can generate a dependency file. This file consists of a single rule suitable for describing the dependencies of the main source file. The rule is written to the file named in its argument. This feature can be used in the automatic updating of makefiles.

2.6 Other Options

The options in this section perform functions not defined in previous sections.

2.6.1 `--defsym sym=value`

Define symbol *sym* to given *value*.

2.6.2 `-I dir`

Use this option to add *dir* to the list of directories that the assembler searches for files specified in `.include` directives. You may use `-I` as many times as necessary to include a variety of paths. The current working directory is always searched first; after that, the assembler searches any `-I` directories in the same order as they were specified (left to right) on the command line.

2.6.3 `-p, --processor=PROC`

Specify the target processor, e.g.:

```
xc16-as -p30F2010 file.s
```

The assembler defines macros based on the target processor setting, which can be tested by conditional directives in source code. For example, include file `p30f2010.inc` contains the following:

```
.ifndef __30F2010
    .error "Include file does not match processor setting"
.endif
```

In addition to the target processor, a macro to identify the device family is also defined. For example:

```
.ifdef __dsPIC30F
    .print "dsPIC30F family selected"
.endif
```

Macros for the device families (see [Section 3.6.6 "Predefined Symbols"](#)) are defined based on target processor setting.

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Chapter 3. MPLAB XC16 Assembly Language

3.1 Introduction

The source language accepted by the macro assembler is described here. All opcode mnemonics and operand syntax are specific to the target device. The same assembler application is used for compiler-generated intermediate assembly and hand-written assembly source code.

Topics covered in this chapter are:

- [Internal Preprocessor](#)
- [Source Code Format](#)
- [Characters](#)
- [Constants](#)
- [Symbols](#)
- [Expressions](#)
- [Operators](#)
- [Special Operators](#)

3.2 Internal Preprocessor

The assembler has an internal preprocessor. The internal processor:

1. Adjusts and removes extra white space. It leaves one space or tab before the keywords on a line, and turns any other white space on the line into a single space.
2. Removes all comments, replacing them with a single space, or an appropriate number of new lines.
3. Converts character constants into the appropriate numeric value.

If you have a single character (e.g., 'b') in your source code, this will be changed to the appropriate numeric value. If you have a syntax error that occurs at the single character, the assembler will not display 'b', but instead display the first digit of the decimal equivalent.

For example, if you had `.global mybuf, 'b'` in your source code, the error message would say "Error: Rest of line ignored. First ignored character is '9'." Notice the error message says '9'. This is because the 'b' was converted to its decimal equivalent 98. The assembler is actually parsing `.global mybuf, 98`.

The internal processor does **not** do:

1. macro preprocessing
2. include file handling
3. anything else you may get from your C compiler's preprocessor

You can do include file preprocessing with the `.include` directive. See [Chapter 4. "Assembler Directives."](#)

You can use the C compiler driver to get other C-style preprocessing by giving the input file a `.S` suffix. See the "MPLAB® XC16 C Compiler User's Guide" (DS50002071) for more information.

16-Bit Assembler, Linker and Utilities User's Guide

If the first line of an input file is `#NO_APP` or if you use the `-f` option, white space and comments are not removed from the input file. Within an input file, you can ask for white space and comment removal in certain portions by putting a line that says `#APP` before the text that may contain white space or comments, and putting a line that says `#NO_APP` after this text. This feature is mainly intended to support assembly statements in compilers whose output is otherwise free of comments and white space.

Note: Excess white space, comments and character constants cannot be used in the portions of the input text that are not preprocessed.

3.3 Source Code Format

Assembly source code consists of statements and white spaces.

White space is one or more spaces or tabs. White space is used to separate pieces of a source line. White space should be used to make your code easier for people to read. Unless within character constants, any white space means the same as exactly one space.

Each *statement* has the following general format and is followed by a new line.

```
[label:] [mnemonic [operands] ]      [; comment]
```

OR

```
[label:] [directive [arguments] ]    [; comment]
```

- [Label](#)
- [Mnemonic](#)
- [Directive](#)
- [Operands](#)
- [Arguments](#)
- [Comments](#)

3.3.1 Label

A label is one or more characters chosen from the set composed of all letters, digits, the underline character (`_`), and the period (`.`). Labels may not begin with a decimal digit, except for the special case of a local symbol. (See [Section 3.6.2 “Local Symbols.”](#) for more information.) Case is significant. There is no length limit; all characters are significant.

Label definitions must be immediately followed by a colon. A space, a tab, an end of line, or assembler mnemonic or directive may follow the colon.

Label definitions may appear on a line by themselves and will reference the next address.

The value of a label after linking is the absolute address of a location in memory.

3.3.2 Mnemonic

Mnemonics tell the assembler which machine instructions to assemble. For example, addition (`ADD`), branches (`BRA`) or moves (`MOV`). Unlike labels that you create yourself, mnemonics are provided by the assembly language. Mnemonics are not case sensitive.

See the “*16-bit MCU and DSC Programmer's Reference Manual*” (DS70000157) for more details.

3.3.3 Directive

Assembler directives are commands that appear in the source code but are not translated directly into machine code. Directives are used to control the assembler, its input, output and data allocation. The first character of a directive is a dot (.). More details are provided in [Chapter 4. “Assembler Directives.”](#) on the available directives.

3.3.4 Operands

Each machine instruction takes 0 to 8 operands. See the “*16-bit MCU and DSC Programmer’s Reference Manual*” (DS70000157). Operands provide data and addressing information to the instruction. Operands must be separated from mnemonics by one or more spaces or tabs.

Commas should separate multiple operands. If commas do not separate operands, a warning will be displayed and the assembler will take its best guess on the separation of the operands. Operands consist of literals, file registers condition codes, destination select, and accumulator select.

3.3.4.1 Literals

Literal values are distinguished with a preceding pound sign (#). Literal values can be hexadecimal, octal, binary or decimal format. Hexadecimal numbers are distinguished by a leading 0x. Octal numbers are distinguished by a leading 0. Binary numbers are distinguished by a leading B. Decimal numbers require no special leading or trailing character.

Examples:

#0xe, #016, #0b1110 and #14 all represents the literal value 14.

#-5 represents the literal value -5.

#symbol represents the value of symbol.

3.3.4.2 File Registers

File registers represent on-chip general purpose and SFRs. File registers are distinguished from literal values because they do not have the preceding pound sign.

Each of the following examples tells the processor to move data located in the file register whose address is 14 to the working register w0:

```
mov 0xE, w0
mov 016, w0
mov 14, w0
.equ symbol, 14
mov symbol, w0
```

3.3.4.3 Registers

The following register names are built into the assembler:

w0, w1, w2, w3, w4, w5, w6, w7, w8, w9, w10, w11, w12, w13, w14, w15, W0, W1, W2, W3, W4, W5, W6, W7, W8, W9, W10, W11, W12, W13, W14, W15.

3.3.4.4 Condition Codes

Condition codes are used with BRA instructions. See the “*16-bit MCU and DSC Programmer’s Reference Manual*” (DS70000157) for more details.

```
bra C, label
```

3.3.4.5 Destination Select

The PIC18CXXX-compatible instructions accept `WREG` as an optional argument to specify whether the result should be placed into `WREG` (W0) or into the file register. See the “16-bit MCU and DSC Programmer’s Reference Manual” (DS70000157) for more details.

```
add sym, WREG
```

3.3.4.6 Accumulator Select

The DSP instructions take an accumulator select operand (A or B) to specify which accumulator to use.

```
ADD A
```

3.3.5 Arguments

Each directive takes 0 to 3 arguments. These arguments give additional information to the directive on how it should carry out the command. Arguments must be separated from directives by one or more spaces or tabs. Commas must separate multiple arguments. More details are provided in [Chapter 4, “Assembler Directives.”](#) on the available directives.

3.3.6 Comments

Comments can be represented in the assembler as single-line or multiple-line comments.

3.3.6.1 Single-Line Comment

This type of comment extends from the comment character to the end of the line. For a single line comment, use a semicolon (;).

Example:

```
mov w0, w1;The rest of this line is a comment.
```

3.3.6.2 Multiple-line Comment

This type of comment can span multiple lines. For a multiple-line comment, use `/* ... */`. Multiple-line comments cannot be nested.

Example:

```
/* All  
of these  
lines  
are  
comments */
```

3.4 Characters

The character set used is standard 7 bit ASCII. Alphabetic case is significant for identifiers, but not mnemonics and reserved words. Tabs are treated as equivalent to spaces.

3.4.1 Delimiters

All numbers and identifiers must be delimited by white space, non-alphanumeric characters or the end of a line.

3.4.2 Special Characters

There are a few characters that are special in certain contexts. Within a macro body, the character & is used for token concatenation. To use the bitwise & operator within a macro body, escape it by using && instead. In a macro argument list, the angle brackets < and > are used to quote macro arguments.

Other special characters are described below.

Table 3-1: Special Characters and Usage

Character	Character Description	Syntax Usage
.	period	begins a directive
;	semicolon	begins a single-line comment
/*	slash, asterisk	begins a multiple-line comment
*/	asterisk, slash	ends a multiple-line comment
:	colon	ends a label definition
#	pound	begins a literal value
'c'	character in single quotes	specifies a single character value
"string"	character string in double quotes	specifies a character string

3.5 Constants

A constant is a value written so that its value is known by inspection, without knowing any context. Examples are:

```
.byte 74, 0112, 0b01001010, 0x4A, 0x4a, 'J', '\J';All the same value
.ascii "Ring the bell\7";A string constant
.float 0f-31415926535897932384626433832795028841971.693993751E-40
```

3.5.1 Numeric Constants

The assembler distinguishes three kinds of numbers according to how they are stored in the machine. Integers are numbers that would fit into a `long` in the C language. Floating-point numbers are IEEE 754 floating-point numbers. Fixed-point numbers are in Q-15 fixed-point format.

3.5.1.1 Integers

A binary integer is '0b' or '0B' followed by zero or more of the binary digits '01'.

An octal integer is '0' followed by zero or more of the octal digits '01234567'.

A decimal integer starts with a non-zero digit followed by zero or more decimal digits '0123456789'.

A hexadecimal integer is '0x' or '0X' followed by one or more hexadecimal digits '0123456789abcdefABCDEF'.

To denote a negative integer, use the prefix operator '-'.

3.5.1.2 Floating-Point Numbers

A floating-point number is represented in IEEE 754 format. A floating-point number is written by writing (in order):

- an optional prefix, which consists of the digit '0', followed by the letter 'e', 'f' or 'd' in upper or lower case. Because floating point constants are used only with `.float` and `.double` directives, the precision of the binary representation is independent of the prefix.
- an optional sign: either '+' or '-'.
- an optional integer part: zero or more decimal digits.
- an optional fractional part: '.' followed by zero or more decimal digits.
- an optional exponent, consisting of:
 - an 'E' or 'e'.
 - an optional sign: either '+' or '-'.
 - one or more decimal digits.

At least one of the integer part or fractional part must be present. The floating-point number has the usual base-10 value.

Floating-point numbers are computed independently of any floating-point hardware in the computer running the assembler.

3.5.1.3 Fixed-Point Numbers

A fixed-point number is represented in Q-15 format. This means that 15 bits are used to represent the fractional portion of the number. The most significant bit is the sign bit, followed by an implied binary point, and 15 bits of magnitude, for example:

bit no.	15	.	14	13	12	...	1	0
value	$\pm 2^0$.	2^{-1}	2^{-2}	2^{-3}	...	2^{-14}	2^{-15}

The smallest number in this format is -1, represented by:

0x8000 (1.000 0000 0000 0000)

the largest number is nearly 1 (.99996948), represented by:

0x7FFF (0.111 1111 1111 1111)

A fixed-point number is written in the same format as a floating-point number, but its value is constrained to be in the range [-1.0, 1.0).

3.5.2 Character Constants

There are two types of character constants. A *character* stands for one character in one byte and its value may be used in numeric expressions. A *string* potentially can contain many bytes and its value may not be used in arithmetic expressions.

3.5.2.1 Characters

A single character may be written as a single quote immediately followed by that character, or as a single quote immediately followed by that character and another single quote. As an example, either 'a' or 'a'.

The assembler accepts escape characters to represent special control characters. As an example, '\n' represents a new-line character. All accepted escape characters are listed in the table below.

Table 3-2: Escape Characters

Escape Character	Description	Hex Value
\a	Bell (alert) character	07
\b	Backspace character	08
\f	Form-feed character	0C
\n	New-line character	0A
\r	Carriage return character	0D
\t	Horizontal tab character	09
\v	Vertical tab character	0B
\\	Backslash	5C
\?	Question mark character	3F
\"	Double quote character	22
\digit digit digit	Octal character code. The numeric code is 3 octal digits.	
\x hex-digits	Hex character code. All trailing hex digits are combined. Either upper or lower case x works.	

The value of a character constant in a numeric expression is the machine's byte-wide code for that character. The assembler assumes your character code is ASCII.

16-Bit Assembler, Linker and Utilities User's Guide

3.6 Symbols

A symbol is one or more characters chosen from the set composed of all letters, digits, the underline character (`_`), and the period (`.`). Symbols may not begin with a digit. The case of letters is significant (e.g., `foo` is a different symbol than `Foo`). There is no length limit and all characters are significant.

Each symbol has exactly one name. Each name in an assembly language program refers to exactly one symbol. You may use that symbol name any number of times in a program.

- [Reserved Names](#)
- [Local Symbols](#)
- [Giving Symbols Other Values](#)
- [The Special DOT Symbol](#)
- [Using Executable Symbols in a Data Context](#)
- [Predefined Symbols](#)

3.6.1 Reserved Names

The following symbol names (case-insensitive) are reserved for the assembler.

Do not use `.equ`, `.equiv` or `.set` (See [Chapter 4. "Assembler Directives."](#)) with these symbols.

Table 3-3: Symbol Names – Reserved

W0	W1	W2	W3	W4	W5	W6	W7
W8	W9	W10	W11	W12	W13	W14	W15
WREG	A	B	OV	C	Z	N	GE
LT	GT	LE	NOV	NC	NZ	NN	GEU
LTU	GTU	LEU	OA	OB	SA	SB	

3.6.2 Local Symbols

Local symbols are used when temporary scope for a label is needed. There are ten local symbol names, which can be reused throughout the program. They may be referred to by using the names `'0'`, `'1'`, ..., `'9'`. To define a local symbol, write a label of the form `'N:'` (where `N` represents any digit 0-9). To refer to the most recent previous definition of that symbol, write `'Nb'`, using the same digit as when you defined the label. To refer to the next definition of a local label, write `'Nf'`. The `'b'` stands for "backwards" and the `'f'` stands for "forwards". There is no restriction on how to use these labels; however, at any point in assembly, no more than 10 backward local labels and 10 forward local labels may be referred to.

Example 3-1:

```
print_string:
    mov     w0,w1
1:
    cp0.b   [w1]
    bra     z,9f
    mov.b   [w1++],w0
    call    print_char
    bra     1b
9:
    return
```


Local symbol names are only a notation device. They are immediately transformed into more conventional symbol names before the assembler uses them. The symbol names stored in the symbol table, appearing in error messages, and optionally emitted to the object file have the following parts:

Table 3-4: Symbol Parts

Parts	Description
L	All local labels begin with 'L'.
Digit	If the label is written '0:', then the digit is '0'. If the label is written '1', then the digit is '1'. And so on up through '9'.
CTRL-A	This unusual character is included so you do not accidentally invent a symbol of the same name. The character has ASCII value '\001'.
Ordinal number	This is a serial number to keep the labels distinct. The first '0:' gets the number '1'; the 15th '0:' gets the number '15'; and so on. Likewise for the other labels '1:' through '9:'. For instance, the first '1:' is named L1C-A1, the 44th '3:' is named L3C-A44.

Example 3-2:

```
00000100 <print_string>:
100: 80 00 78      mov.w      w0, w1

00000102 <L1:1>:
102: 11 04 e0      cp0.b      [w1]
104: 03 00 32      bra        Z, . + 0x8
106: 31 40 78      mov.b      [w1++], w0
108: 02 00 07      rcall      . + 0x6
10a: fb ff 37      bra        . + 0xFFFFFFFF8

0000010c <L9:1>:
10c: 00 00 06      return
```

3.6.3 Giving Symbols Other Values

A symbol can be given an arbitrary value by writing a symbol, followed by an equals sign '=', followed by an expression. This is equivalent to using the `.set` directive (see [Chapter 4. "Assembler Directives."](#)).

Example 3-3:

```
PSV = 4
```

3.6.4 The Special DOT Symbol

The special symbol '`.`' refers to the current address that is being assembled into. Thus, the expression:

```
melvin: .word . ; in a data section
```

defines `melvin` to contain its own data address. Assigning a value to `.` is treated the same as a `.org` directive. Thus the expression:

```
. = .+2
```

is the same as saying:

```
.org .+2
```

The symbol '\$' is accepted as a synonym for '`.`'.

When used in an executable section, '`.`' refers to a PC address. On the 16-bit device, the PC increments by 2 for each instruction word. Odd values are not permitted.

3.6.5 Using Executable Symbols in a Data Context

The 16-bit device modified-Harvard architecture includes separate address spaces for data storage and program storage. Most instructions and assembler directives imply a context which is compatible with symbols from one address space or the other. For example, the `CALL` instruction implies an executable context, so the assembler reports an error if a program tries to `CALL` a symbol located in a data section.

Likewise, instructions and directives that imply a data context cannot be used with symbols located in an executable section. Assembling the following code sequence will result in an error, as shown:

```
.text
msg: .asciz "Here is an important message"
    mov #msg,w0
:
:
Assembler messages:
Error: Cannot reference executable symbol (msg) in a data context
```

In this example the `mov` instruction implies a data context. Because symbol `msg` is located in an executable section, an error is reported. Possibly the programmer was trying to derive a pointer for use with the PSV window. The special operators described in [Section 3.9 "Special Operators."](#) can be used whenever an executable symbol must be referenced in a data context:

```
.text
msg: .asciz "Here is an important message"
    mov #psvoffset(msg),w0
```

Here the `psvoffset()` operator derives a 16-bit value which is suitable for use in a data context.

The next example shows how the special symbol `."` can be used with a data directive in an executable section:

```
.text
fred: .long paddr(.
```

Here the `paddr()` operator derives a 24-bit value which is suitable for use in a data context. The `.long` directive pads the value to 32 bits and encodes it into the `.text` section.

3.6.6 Predefined Symbols

The assembler predefines several symbols which can be tested by conditional directives in source code.

Table 3-5: Predefined Symbols

Symbol	Definition
Device Family Symbols	
__C30COFF	16-bit compiler COFF output
__C30ELF	16-bit compiler ELF output
__dsPIC30F	dsPIC30F target device family
__dsPIC33F	dsPIC33F target device family
__dsPIC33E	dsPIC33EP target device family
__PIC24F	PIC24FJ target device family
__PIC24FK	PIC24FK target device family
__PIC24H	PIC24H target device family
__PIC24E	PIC24EP target device family
__MCHP16	No target device family specified
Feature Symbols	
__HAS_DSP	Device has a DSP engine
__HAS_EEDATA	Device has EEDATA memory
__HAS_DMA	Device has DMA memory
__HAS_DMAV2	Device has DMA v2 support
__HAS_CODEGUARD	Device has Codeguard™ Security
__HAS_PMP	Device has Parallel Master Port (PMP)
__HAS_PMPV2	Device has PMP v2 support
__HAS_PMP_ENHANCED	Device has Enhanced PMP
__HAS_EDS	Device has EDS
__HAS_5VOLTS	Device is a 5-volt device

16-Bit Assembler, Linker and Utilities User's Guide

3.7 Expressions

An expression specifies an address or numeric value. White space may precede and/or follow an expression. The result of an expression must be an absolute number or an offset into a particular section. When an expression is not absolute and does not provide enough information for the assembler to know its section, the assembler terminates and generates an error message.

3.7.1 Empty Expressions

An empty expression has no value: it is just white space or null. Wherever an absolute expression is required, you may omit the expression, and the assembler assumes a value of (absolute) 0.

3.7.2 Integer Expressions

An integer expression is one or more arguments delimited by operators. Arguments are symbols, numbers or subexpressions. Subexpressions are a left parenthesis '(' followed by an integer expression, followed by a right parenthesis ')'; or a prefix operator followed by an argument.

Integer expressions involving symbols in program memory are evaluated in Program Counter (PC) units. On the 16-bit device, the PC increments by 2 for each instruction word.

Example 3-4: Branch After a Label

Branch to the next instruction after label `L` by specifying `L+2` as the destination.

```
bra L+2
```

3.8 Operators

Operators are arithmetic functions, like + or %. Prefix operators are followed by an argument. Infix operators appear between their arguments. Operators may be preceded and/or followed by white space.

Prefix operators have higher precedence than infix operators. Infix operators have an order of precedence dependent on their type.

3.8.1 Prefix Operators

The assembler has the following prefix operators. Each takes one argument, which must be absolute.

Table 3-6: Prefix Operators

Operator	Description	Example
-	Negation. Two's complement negation.	-1
~	Bit-wise not. One's complement.	~flags

3.8.2 Infix Operators

Infix operators take two arguments, one on either side. Operators have a precedence, by type, as shown in the table below; but, operations with equal precedence are performed left to right. Apart from + or –, both operators must be absolute, and the result is absolute.

Table 3-7: Infix Operators

Operator	Description	Example
Arithmetic		
*	Multiplication	5 * 4 (=20)
/	Division. Truncation is the same as the C operator '/'. /	23 / 4 (=5)
%	Remainder	30 % 4 (=2)
<<	Shift Left. Same as the C operator '<<'	2 << 1 (=4)
>>	Shift Right. Same as the C operator '>>'	2 >> 1 (=1)
Bit-Wise		
&	Bit-wise And	4 & 6 (=4)
^	Bit-wise Exclusive Or	4 ^ 6 (=2)
!	Bit-wise Or Not	0x1010 ! 0x5050 (=0xBFBF)
	Bit-wise Inclusive Or	2 4 (=6)
Simple Arithmetic		
+	Addition. If either argument is absolute, the result has the section of the other argument. You may not add together arguments from different sections.	4 + 10 (=14)
-	Subtraction. If the right argument is absolute, the result has the section of the left argument. If both arguments are in the same section, the result is absolute. You may not subtract arguments from different sections.	14 - 4 (=10)
Relational		
==	Equal to	.if (x == y)
!=	Not equal to (also <>)	.if (x != y)
<	Less than	.if (x < 5)
<=	Less than or equal to	.if (y <= 0)
>	Greater than	.if (x > a)
>=	Greater than or equal to	.if (x >= b)
Logical		
&&	Logical AND	.if ((x > 1) && (x < 10))
	Logical OR	.if ((y != x) (y < 100))

16-Bit Assembler, Linker and Utilities User's Guide

3.9 Special Operators

The assembler provides a set of special operators for each of the following actions:

- [Accessing Data in Program Memory](#)
- [Obtaining a Program Address of a Symbol or Constant](#)
- [Obtaining a Handle to a Program Address](#)
- [Obtaining the DMA Offset of a Symbol – PIC24H/dsPIC33F Devices Only](#)
- [Obtaining the Size of a Specific Section](#)
- [Obtaining the Starting Address of a Specific Section](#)
- [Accessing Functions in Boot or Secure Segments](#)

DD Table 3-8: Special Operators

Operators*	Description	Support
<code>tblpage (name)</code>	Get page for table read/write operations	All
<code>tbloffset (name)</code>	Get pointer for table read/write operations	All
<code>psvpage (name)</code>	Get page for PSV data window operations	All
<code>psvoffset (name)</code>	Get pointer for PSV data window operations	All
<code>paddr (label)</code>	Get 24-bit address of <i>label</i> in program memory	All
<code>handle (label)</code>	Get 16-bit reference to <i>label</i> in program memory	All
<code>dmepage (name)</code>	Get page suitable for DMA controller	24E/33E
<code>dmaoffset (name)</code>	Get offset of a symbol within DMA memory	24H/33F
<code>.sizeof. (name)</code>	Get size of section <i>name</i> in address units	All
<code>.startof. (name)</code>	Get starting address of section <i>name</i>	All
<code>boot (num)</code>	Get address of access slot <i>num</i> in the boot segment.	All
<code>secure (num)</code>	Get address of access slot <i>num</i> in the secure segment.	All
<code>edspage (name)</code>	Get page for EDS data window operations	All
<code>edsoffset (name)</code>	Get pointer for EDS data window operations	All

* You cannot use two special operators in an expression.

All = Support for all devices

24H = Support for PIC24H MCUs; 24E = Support for PIC24EP MCUs

33F = Support for dsPIC33F DSCs; 33E = Support for dsPIC33EP DSCs

3.9.1 Accessing Data in Program Memory

The 16-bit device modified-Harvard architecture is comprised of two separate address spaces: one for data storage and one for program storage. Data memory is 16 bits wide and is accessed with a 16-bit address; program memory is 24 bits wide and is accessed with a 24-bit address.

Normally, 16-bit instructions can read or write data values only from data memory, while program memory is reserved for instruction storage. This arrangement allows for very fast execution, since the two memory buses can work simultaneously and independently of each other. In other words, a 16-bit instruction can read, modify and write a location in data memory at the same time the next instruction is being fetched from program memory.

Occasionally, circumstances may arise when the programmer or application designer is willing to sacrifice some execution speed in return for the ability to read constant data directly from program memory. For example, certain DSP algorithms require large tables of coefficients that would otherwise consume the data memory needed to buffer real-time data. To accommodate these needs, the 16-bit device modified-Harvard architecture permits instructions to access data stored in program memory.

There are three methods available for accessing data in program memory:

- [Table Read/Write Instructions](#)
- [PSV Data Window](#)
- [EDS Data Window](#)

In any case, the programmer must compensate for the different address width between data memory and program memory. For example, a pointer is commonly used to access constant data tables, yet pointers for table read/write instructions can specify an address of only 16 bits. A pointer used to access the PSV data window can specify only 15 bits – the most significant bit must be set for an address in the data window range (0x8000 to 0xFFFF).

As explained in the “*16-bit MCU and DSC Programmer’s Reference Manual*” (DS70000157), SFRs can be used to specify the full flash address. For a PSV address, use `PSVPAG` (or `DSRPAG` for devices with EDS). For a table read/write address, use `TBLPAG`. For an EDS address, use `DSRPAG`.

3.9.1.1 Table Read/Write Instructions

The `tblpage()` and `tbloffset()` operators provided by the assembler can be used with table read/write instructions. These operators may be applied to any symbol (usually representing a table of constant data) in program memory.

Suppose a table of constant data is declared in program memory like this:

```
.text
fib_data:
.word 0, 1, 2, 3, 5, 8, 13
```

To access this table with table read/write instructions, use the `tblpage()` and `tbloffset()` operators as follows:

```
; Set TBLPAG to the page that contains the fib_data array.
mov    #tblpage(fib_data), w0
mov    w0, _TBLPAG
; Make a pointer to fib_data for table instructions
mov    #tbloffset(fib_data), w0
; Load the first data value
tblrdl [w0++], w1
```

16-Bit Assembler, Linker and Utilities User's Guide

The programmer must ensure that the constant data table does not exceed the program memory page size that is implied by the TBLPAG register. The maximum table size implied by the TBLPAG register is 64 Kbytes. If additional constant data storage is required, simply create additional tables each with its own symbol, and repeat the code sequence above to load the TBLPAG register and derive a pointer.

3.9.1.2 PSV Data Window

The `psvpage()` and `psvoffset()` operators can be used with the PSV data window. These operators may be applied to any symbol (usually representing a table of constant data) in program memory.

Suppose a table of constant data is declared in program memory like this:

```
.section *,psv
fib_data:
    .word 0, 1, 2, 3, 5, 8, 13
```

To access this table through the PSV data window, use the `psvpage()` and `psvoffset()` operators as follows:

```
; Enable Program Space Visibility (Note 1)
    bset.b CORCONL, #PSV

; Set PSVPAG (Note 2) to the page that contains the fib_data array.
    mov    #psvpage(fib_data), w0
    mov    w0, _PSVPAG

; Make a pointer to fib_data in the PSV data window
    mov    #psvoffset(fib_data), w0

; Load the first data value
    mov    [w0++], w1
```

- | |
|---|
| <p>Note 1: Some devices do not need PSV to be enabled. Please check the data sheet for your device.</p> <p>2: For devices with EDS, use <code>DSRPAG</code>. Please check the data sheet for your device.</p> |
|---|

The programmer must ensure that the constant data table does not exceed the program memory page size that is implied by the PSVPAG register (or the DSRPAG register for devices with EDS). The maximum table size implied by the PSVPAG or DSRPAG register is 32 Kbytes. If additional constant data storage is required, simply create additional tables each with its own symbol, and repeat the code sequence above to load the PSVPAG or DSRPAG register and derive a pointer.

3.9.1.3 EDS Data Window

The `edspage()` and `edsoffset()` operators can be used with the EDS data window. The EDS data window replaces the PSV data window in certain device families. However, these operators are supported on all devices.

The `edspage()` operator may be applied to any symbol in any on-chip memory space. The operator returns a 10-bit page value. Unlike `psvpage()`, a value of zero is never returned.

The `edsoffset()` operator may be applied to any symbol in any on-chip memory space. The operator returns a 16-bit data space pointer. Unlike `psvoffset()`, the value of this pointer may fall anywhere in the data address space (0x0 to 0xFFFF).

Suppose that a table of data is located in any on-chip memory space. To access this table through the EDS data window, use the `edspage()` and `edsoffset()` operators as follows:

```
; set DSRPAG to the page that contains the glob_data array
mov    #edspage(glob_data),w0
mov    w0, _DSRPAG
; make a pointer to glob_data
mov    #edsoffset(glob_data),w0
; Load the first data value
mov    [w0++], w1
```

In order to access multiple items from a data table, you must ensure that the table does not cross a page boundary. To prevent this, specify the `page` section directive when the data table is defined. If additional constant storage is required, simply create additional tables, each with its own symbol, and repeat the code sequence in Section 3.9.1.3 to load the DSRPAG register and derive a pointer.

3.9.2 Obtaining a Program Address of a Symbol or Constant

The `paddr()` operator can be used to obtain the program address of a constant or symbol. For example, if you wanted to set up an interrupt vector table without using the default naming conventions, you could use the `paddr()` operator.

```
.section ivt, code
goto reset
.pword paddr(iv1)
.pword paddr(iv2)
...
```

3.9.3 Obtaining a Handle to a Program Address

The `handle()` operator can be used to obtain the 16-bit reference to a label in program memory. If the final resolved PC address of the label fits in 16 bits, that value is returned by the `handle()` operator. If the final resolved address exceeds 16 bits, the address of a jump table entry is returned instead. The jump table entry is a `GOTO` instruction containing a 24-bit absolute address. The handle jump table is created by the linker and is always located in low program memory. Handles permit any location in program memory to be reached via a 16-bit address and are provided to facilitate the use of C function pointers.

The handle jump table is created by the linker and contains an entry for each unique label that is used with the `handle()` operator.



3.9.4 Obtaining the DMA Offset of a Symbol – PIC24H/dsPIC33F Devices Only

The `dmaoffset()` operator can be used to obtain the offset of a symbol within DMA memory. For example, to declare a buffer in DMA memory, and load its offset into a register, you could use:

```
.section *,bss,dma
buf: .space 256

.text
mov #dmaoffset(buf), W0
```

To construct a table of DMA offsets for several symbols, you could use:

```
.word dmaoffset(buf1)
.word dmaoffset(buf2)
.word dmaoffset(buf3)
...
```



3.9.5 Obtaining the DMA Offset of a Symbol – PIC24EP/dsPIC33EP Devices Only

The `dmaoffset()` and `dmapage()` operators can be used to obtain the offset of a symbol within DMA memory.

```
.word dmaoffset(buf1), dmapage(buf1)
.word dmaoffset(buf2), dmapage(buf2)
.word dmaoffset(buf3), dmapage(buf3)
...
```

3.9.6 Obtaining the Size of a Specific Section

The `.sizeof.(section_name)` operator can be used to obtain the size of a specific section after the link process has occurred. For example, to find the final size of the `.data` section, use:

```
mov #.sizeof(.data), w0
```

<p>Note: When the <code>.sizeof.(section_name)</code> operator is used on a section in program memory, the size returned is the size in PC units. The 16-bit device PC increments by 2 for each instruction word.</p>
--

3.9.7 Obtaining the Starting Address of a Specific Section

The `.startof.(section_name)` operator can be used to obtain the starting address of a specific section after the link process has occurred. For example, to obtain the starting address of the `.data` section, use:

```
mov #.startof(.data), w1
```

3.9.8 Accessing Functions in Boot or Secure Segments

Functions in the boot or secure segments without access entries can be referenced like any other function:

```
call func1           ; call func1
mov #handle(func1),w1 ; create 16 bit pointer to func1 (instr)
.word handle(func1)   ; create 16 bit pointer to func1 (data)
.pword func1          ; create 24 bit pointer to func1
```

In order to support the separate linking of boot and secure application segments, access entry points may be defined. Access entry points provide a way to transfer control across segments to a function that may not be defined at link time. For more information about access entry points, see [Section 4.6 “Directives that Define Symbols.”](#) and [Section 10.14 “Boot and Secure Segments.”](#)

The `boot()` and `secure()` operators can be used to reference boot or secure functions via access entry points. These operators can be applied in both instructions and data directives, and will return 16, 24, or 32 bits, depending on the context.

```
call boot(4)          ; call access entry 4 in the boot segment
rcall secure(4)        ; pc-relative call to secure access entry 4
mov #boot(4),w1        ; load 16 bit pointer to boot entry 4
.word secure(5)        ; create 16 bit pointer to secure entry 5
.pword secure(5)       ; create 24 bit pointer to secure entry 5
.long boot(6)          ; create 32 bit pointer to boot entry 6
goto boot(7)           ; jump to boot entry 7
bra secure(7)          ; unconditional branch to secure entry 7
bra cc, boot(8)        ; conditional branch to boot entry 8
```

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Chapter 4. Assembler Directives

4.1 Introduction

Directives are assembler commands that appear in the source code but are not usually translated directly into opcodes. They are used to control the assembler: its input, output, and data allocation. All 16-bit directives are preceded by a dot '.'.

Note 1: Directives are *not* instructions (i.e., `movlw`, `btfss`, `goto`, etc.). For instruction set information, consult your device data sheet.

2: Directives that are supported, but deprecated, are listed in [Appendix A. "Deprecated Features."](#)

Topics covered in this chapter are:

- [Directives that Define Sections](#)
- [Directives that Fill Program Memory](#)
- [Directives that Initialize Constants](#)
- [Directives that Declare Symbols](#)
- [Directives that Define Symbols](#)
- [Directives that Modify Section Alignment](#)
- [Directives that Format the Output Listing](#)
- [Directives that Control Conditional Assembly](#)
- [Directives for Substitution/Expansion](#)
- [Miscellaneous Directives](#)
- [Directives for Debug Information](#)

4.2 Directives that Define Sections

Sections are locatable blocks of code or data that will occupy contiguous locations in the 16-bit device memory. Three sections are pre-defined: `.text` for executable code, `.data` for initialized data, and `.bss` for uninitialized data. Other sections may be defined; the linker defines several that are useful for locating data in specific areas of 16-bit memory.

Section directives are:

- `.bss`
- `.data`
- `.memory name, size(nn) [, origin(aa)]`
- `.pushsection name [, attr1[,...,attrn]]`
- `.popsection`
- `.section name [, "flags"] (deprecated)`
- `.section name [, attr1[,...,attrn]]`
- `.text`

4.2.1 `.bss`

Assemble the following statements onto the end of the `.bss` (uninitialized data) section.

Note: You cannot reserve all of data memory for symbols; you will need room for the stack. See [Section 10.10 "Stack Allocation."](#)

Example

```
; The following symbols (B1 and B2) will be placed in
; the uninitialized data section.
.bss
B1:  .space 4    ; 4 bytes reserved for B1
B2:  .space 1    ; 1 byte reserved for B2
```

4.2.2 `.data`

Assemble the following statements onto the end of the `.data` (initialized data) section.

Note: You cannot reserve all of data memory for symbols; you will need room for the stack. See [Section 10.10 "Stack Allocation."](#)

Example

```
; The following symbols (D1 and D2) will be placed in
; the initialized data section.
.data
D1:  .long 0x12345678    ; 4 bytes
D2:  .byte 0xFF          ; 1 byte
```

The linker collects initial values for section `.data` (and other sections defined with the `data` attribute) and creates a data initialization template. This template can be processed during application start-up to transfer initial values into memory. For C applications, a library function is called for this purpose automatically. Assembly projects can utilize this library by linking with the `libpic30` library. For more information, see the discussion of Run-Time Library Support in [Section 10.8 "Initialized Data."](#)

4.2.3 .memory *name*, size(*nn*) [, origin(*aa*)]

Define an external memory region for allocation by the linker. Sections may be assigned to region *name* by use of the `memory` section attribute.

Example

```
; define an external memory region
.memory _memory1, size(8192), origin(0)

; allocate a section in external memory
.section mem1_sec1,memory(_memory1)
.global _mem1_array1
_mem1_array1:
.skip 50
```

4.2.4 .pushsection *name* [, *attr*₁[,...,*attr*_{*n*}]]

Push the current section description onto the section stack, and assemble the following code into a section named *name*. The syntax is identical to `.section`. Every `.pushsection` should have a matching `.popsection`.

4.2.5 .popsection

Replace the current section description with the top section on the section stack. This section is popped off the stack.

4.2.6 .section *name* [, “*flags*”] (deprecated)

See [Section A.3.1 “.section *name* \[, “*flags*”\].”](#)

4.2.7 .section *name* [, *attr*₁[,...,*attr*_{*n*}]]

Assemble the following code into a section named *name*. If the character `*` is specified for *name*, the assembler will generate a unique name for the section based on the input file name in the format `filename.scnn`, where *n* represents the number of auto-generated section names.

Sections named `*` can be used to conserve memory because the assembler will not add alignment padding to these sections. Sections that are not named `*` may be combined across several files, so the assembler must add padding in order to guarantee the requested alignment.

If the optional argument is not present, the section attributes depend on the section name. A table of reserved section names with implied attributes is given in [Reserved Section Names with Implied Attributes](#). If the section name matches a reserved name, the implied attributes will be assigned to that section. If the section name is not recognized as a reserved name, the default attribute will be `data` (initialized storage in data memory).

Implied attributes for reserved section names other than `[.text, .data, .bss]` are deprecated. A warning will be issued if implied attributes for these reserved section are used.

If the first optional argument is quoted, it is taken as one or more flags that describe the section attributes. Quoted section flags are deprecated. (See [Appendix A. “Deprecated Features.”](#)). A warning will be issued if quoted section flags are used.

If the first optional argument is not quoted, it is taken as the first element of an attribute list. Attributes may be specified in any order, and are case-insensitive. Two categories of section attributes exist: attributes that represent section types, and attributes that modify section types.

16-Bit Assembler, Linker and Utilities User's Guide

DD

4.2.7.1 Attributes that Represent Section Types

Attributes that represent section types are mutually exclusive. At most, one of the attributes listed below may be specified for a given section.

Table 4-1: Attributes That Represent Section Types

Attribute	Description	Support
auxflash	Executable code in auxiliary program memory	24EP/33EP (some)
auxpsv	Constant is in auxiliary program memory	24EP/33EP (some)
bss	Uninitialized storage in data memory	All
code	Executable code in program memory	All
data	Initialized storage in data memory	All
eedata	Non-volatile storage in data EEPROM	30/24FxxK
heap	Memory for dynamic allocation in C	All
memory	External or user-defined memory	All
packedflash	Use the upper byte of Flash via packed storage	All
persist	Persistent storage in data memory	All
psv	Constants in program memory	All
stack	Processor stack	All

All = Supported on all devices

24X = Supported on PIC24X MCUs, where X can be EP, FxxK, FJ, FV, HJ.

30 = Supported on dsPIC30F DSCs

33X = Supported on dsPIC33X DSCs, where X can be EP, EV, FJ.

DD

4.2.7.2 Attributes that Modify Section Types

Depending on the attribute, all or some section types may be modified by it, as below.

Table 4-2: Attributes That Modify Section Types

Attribute*	Description	Attribute applies to**
address(a)	locate at absolute address a	auxflash, bss, code, data, eedata, heap, memory, persist, psv, stack
near	locate in the first 8K of memory	bss, data, persist
xmemory	locate in X address space	bss, data, heap, persist
ymemory	locate in Y address space	bss, data, heap, persist
reverse(n)	align the ending address +1	auxflash, bss, data, eedata, memory, persist, psv
align(n)	align the starting address	auxflash, bss, code, data, eedata, heap, memory, persist, psv, stack
noload	allocate, do not load	auxflash, bss, code, data, eedata, memory, persist, psv
merge(n)	mergable elements of size n***	auxflash, code, data, eedata, psv
info	do not allocate or load	auxflash, bss, code, data
dma	locate in DMA space	bss, data, persist
boot	locate in boot segment	bss, code, eedata, psv
secure	locate in secure segment	bss, code, eedata, psv
eds	locate in extended data space	bss, data, persist
shared	use section outside of application	auxflash, bss, code, data, eedata, memory, psv, stack
preserved	preserve variables on restart	bss, data, memory, persist
update	initialize variables on restart	bss, data, memory, persist
priority(n)	group variable initializations together.	auxflash, bss, code, data
page	do not cross page boundary	auxflash, bss, code, data, persist, psv

* = Not all devices support all attributes.

** = See [Table 4-1](#) for descriptions and device support.

*** = This attribute could be used by a linker to merge identical constants across input files.
If n=0, the section contains null-terminated strings of variable length.

16-Bit Assembler, Linker and Utilities User's Guide

DD

Attributes that modify section types may be used in combination. For example, `xmemory, address (a)` is a valid attribute string, but `xmemory, address (a) , ymemory` is not.

Table 4-3: Combining Attributes that Modify Section Types

Attribute*	Attribute can be combined with*
address	near, xmemory, ymemory, noload, dma, boot, secure, eds, page
near	address, xmemory, ymemory, reverse, align, noload, merge
xmemory	address, near, reverse, align, noload, merge, eds, page
ymemory	address, near(30,33), reverse, align, noload, merge, eds, page
reverse	near, xmemory, ymemory(30/330, noload, merge, dma, boot, secure, eds, page
align	near, xmemory, ymemory(30/330, noload, merge, dma, boot, secure, eds, page
noload	address, near, xmemory, ymemory, reverse, align, dma, boot, secure, eds, page
merge	near, xmemory, ymemory, reverse, align, eds, page
info	N/A
dma	address, reverse, align, noload, eds, page
boot	address, reverse, align, noload, eds, page
secure	address, reverse, align, noload, eds, page
eds	address, xmemory, ymemory, reverse, align, noload, merge, dma, boot, secure, page
shared	address, near, xmemory, ymemory, reverse, align, noload, merge, dma, boot, secure, eds, preserved, update, priority, page
preserved	address, near, xmemory, ymemory, reverse, align, noload, merge, dma, eds, shared, priority, page
update	address, near, xmemory, ymemory, reverse, align, noload, merge, dma, eds, shared, priority, page
priority(n)	address, near, xmemory, ymemory, reverse, align, dma, boot, secure, eds, shared, preserved, update, page
page	address, xmemory, ymemory(30, 33), reverse, align, noload, merge, dma, boot, secure, eds

* = Not all devices support all attributes.

4.2.7.3 Reserved Section Names with Implied Attributes

The following section names are available for user applications and are recognized to have implied attributes:

Reserved Name	Implied Attribute(s)	Support
.text	code	All
.data	data	All
.bss	bss	All
.xbss	bss, xmemory	30/33
.xdata	data, xmemory	30/33
.nbss	bss, near	All
.ndata	data, near	All
.ndconst	data, near	All
.pbss	bss, persist	All
.dconst	data	All
.ybss	bss, ymemory	30/33
.ydata	data, ymemory	30/33
.const	psv	All
.eedata	eedata	30

All = Supported on all devices

30 = Supported on dsPIC30F DSCs

33 = Supported on dsPIC33x DSCs

Reserved section names may be used with explicit attributes. If the explicit attribute(s) conflict with any implied attribute(s), an error will be reported.

Implied attributes for reserved section names other than [.text, .data, .bss] are deprecated. A warning will be issued if these names are used without explicit attributes.

4.2.7.4 Section Directive Examples

```
.section foo                                ;foo is initialized data memory.

.section bar,bss,xmemory,align(256) ;bar is uninitialized
                                   ;X data memory, aligned.
.section *,data,near                ;section is near
                                   ;initialized data memory.
.section buf1,bss,address(0x800)    ;buf1 is uninitialized
                                   ;data memory at 0x800.
.section tab1,psv,address(0x10000)  ;tab1 is psv constants
                                   ;at 0x10000.
```

4.2.7.5 Section Directive Examples - Boot/Secure Segments

Program Memory

Attributes can be used to declare protected functions in secure segments:

```
.section *,code,boot
.global func1
func1:
    return

.section *,code,secure
.global func2
func2:
    return
```

16-Bit Assembler, Linker and Utilities User's Guide

A secure function is defined by the combination of `.section` and `.global` directives, and a label. It is recommended that each secure function be defined in a separate section. If the function will be assigned an access entry point, separate sections are required.

An optional argument to `boot` or `secure` can be used to specify a protected access entry point:

```
.section *,code,boot(3)
.global func3
func3:
    return

.section *,code,secure(4)
.global func4
func4:
    return
```

The optional argument is valid only in code sections. Integers that represent access entry slots must be in the range 0..15 or 17..31. In addition to an entry slot number, the value `unused` may be used to specify an entry for all unused slots in the access entry area:

```
.section *,code,boot(unused)
.global func_default
func_default:
    return
```

An interrupt service routine may be specified with the value `isr`:

```
.section *,code,boot(isr)
.global func_isr
func_isr:
    retfie
```

A section identified with `boot(isr)` or `secure(isr)` will be assigned to access entry slot 16, which is reserved for interrupt functions.

Data Memory

The `boot` and `secure` attributes can be used to define protected variables in boot RAM or secure RAM:

```
.section *,bss,boot
.global boot_dat
boot_dat:
    .space 32

.section *,bss,secure
.global secure_dat
secure_dat:
    .space 32
```

There is no support for initialized data in protected RAM segments. Therefore `boot` or `secure` cannot be used in combination with attribute `data`. A diagnostic will be reported if initial values are specified in a section that has been designated `boot` or `secure`.

Constants in Non-Volatile Memory

Constants in non-volatile memory can be protected by using the `boot` or `secure` attribute in combination with `psv` or `eedata`:

```
.section *,psv,boot
.global key1
key1:
.ascii "abcdefg"
.section *,eedata,boot
.global key2
key2:
.ascii "hijklm"
```

4.2.8 .text

Assemble the following statements onto the end of the `.text` (executable code) section.

Example

```
; The following code will be placed in the executable
; code section.
.text
.global __reset
__reset:
    mov BAR, w1
    mov FOO, w0
LOOP:
    cp0.b [w0]
    bra Z, DONE
    mov.b [w0++], [w1++]
    bra LOOP
DONE:
.end
```

4.3 Directives that Fill Program Memory

These directives are only allowed in a code (executable) section. If they are not in a code section, a warning is generated and the rest of the line is ignored.

Fill directives are:

- `.fillupper [value]`
- `.fillvalue [value]`
- `.pfillvalue [value]`
- [Section Example](#)

4.3.1 `.fillupper [value]`

Define the upper byte (bits 16-23) to be used when this byte is skipped due to alignment or data defining directives. If *value* is not specified, it is reset to the default 0x00. Directives that may cause an upper byte to be filled are: `.align`, `.ascii`, `.asciz`, `.byte`, `.double`, `.fill`, `.fixed`, `.float`, `.hword`, `.int`, `.long`, `.skip`, `.space`, `.string` and `.word`. The value is persistent for a given code section, throughout the entire source file, and may be changed to another value by issuing subsequent `.fillupper` directives.

Example

See the [Section Example](#) table that follows.

4.3.2 `.fillvalue [value]`

Define the byte value to be used as fill in a code section when the lower word (bits 0-15) is skipped due to alignment or data defining directives. If *value* is not specified, the default value of 0x0000 is used. Directives that may cause the lower word to be filled are: `.align`, `.fill`, `.skip`, `.org` and `.space`. The value is persistent for a given code section, throughout the entire source file, and may be changed to another value by issuing subsequent `.fillvalue` directives.

Example

See the [Section Example](#) table that follows.

4.3.3 `.pfillvalue [value]`

Define the byte value to be used as fill in a code section when memory (bits 0-23) is skipped due to an alignment or data defining `p` directive. If *value* is not specified, it is reset to its default 0x000000. Directives that may cause a program word to be filled are: `.palign`, `.pfill`, `.pskip`, `.porg`, and `.pspace`. The value is persistent for a given code section, throughout the entire source file, and may be changed to another value by issuing subsequent `.pfillvalue` directives.

Example

See the [Section Example](#) table that follows.

4.3.4 Section Example

			.section .myconst, code
			.fillvalue 0x12
			.fillupper 0x34
			.pfillvalue 0x56
0x12	0x12	0x34	.fill 4
0x12	0x12		
		0x34	.align 2 ;Align to next p-word
0x56	0x56	0x56	.pfill 8
0x56	0x56	0x56	
0x56	0x56		
		0x56	.palign 2 ;Align to next p-word
			.fillvalue ;Reset fillvalue
			.pfillvalue ;Reset pfillvalue
0x00	0x00	0x34	.fill 4
0x00	0x00		
		0x34	.align 2 ;Align to next p-word
0x00	0x00	0x00	.pfill 8
0x00	0x00	0x00	
0x00	0x00		
		0x00	.palign 2 ;Align to next p-word

4.4 Directives that Initialize Constants

Constant initialization directives are:

- `.ascii "string1" | <##>1 [, ..., "stringn" | <##>n]`
- `.pascii "string1" | <##>1 [, ..., "stringn" | <##>n]`
- `.pascii "string1"`
- `.asciz "string1" | <##>1 [, ..., "stringn" | <##>n]`
- `.pasciz "string1" | <##>1 [, ..., "stringn" | <##>n]`
- `.pasciz "string2"`
- `.byte expr1 [, ..., exprn]`
- `.pbyte expr1 [, ..., exprn]`
- `.double value1 [, ..., valuen]`
- `.fixed value1 [, ..., valuen]`
- `.float value1 [, ..., valuen]`
- `.single value1 [, ..., valuen]`
- `.hword expr1 [, ..., exprn]`
- `.int expr1 [, ..., exprn]`
- `.long expr1 [, ..., exprn]`
- `.short expr1 [, ..., exprn]`
- `.string "str"`
- `.pstring "str"`
- `.pstring "string2"`
- `.word expr1 [, ..., exprn]`
- `.pword expr1 [, ..., exprn]`

4.4.1 `.ascii "string1" | <##>1 [, ..., "stringn" | <##>n]`

Assembles each string (with no automatic trailing zero byte) or `<##>` into successive bytes in the current section. `<##>` is a way of specifying a character by its ASCII code. For example, given that the ASCII code for a new line character is `0xa`, the following two lines are equivalent:

```
.ascii "hello\n","line 2\n"
.ascii "hello",<0xa>,"line 2",<0xa>
```

Note: If the `##` is not a number, 0 will be assembled. If the `##` is greater than 255, then the value will be truncated to a byte.

If in a code (executable) section, the upper program memory byte will be filled with the last `.fillupper` value specified or the NOP opcode (0x00) if no `.fillupper` has been specified.

4.4.2 `.pascii "string1" | <##>1 [, ..., "stringn" | <##>n]`

Assembles each string (with no automatic trailing zero byte) or `<##>` into successive bytes into program memory, including the upper byte. `<##>` is a way of specifying a character by its ASCII code. For example, given that the ASCII code for a new line character is `0xa`, the following two lines are equivalent:

```
.pascii "hello\n","line 2\n"
.pascii "hello",<0xa>,"line 2",<0xa>
```

Note: If the `##` is not a number, 0 will be assembled. If the `##` is greater than 255, then the value will be truncated to a byte.

4.4.3 .pascii “string₁”

Stores a sequence of ASCII characters (with no automatic trailing zero byte) into program memory, including the upper byte.

4.4.4 .asciz “string₁” | <##>₁ [, ..., “string_n” | <##>_n]

Assembles each string with an automatic trailing zero byte or <##> into successive bytes in the current section.

Note: If the ## is not a number, 0 will be assembled. If the ## is greater than 255, then the value will be truncated to a byte.

If in a code (executable) section, the upper program memory byte will be filled with the last `.fillupper` value specified or the NOP opcode (0x00) if no `.fillupper` has been specified.

4.4.5 .pasciz “string₁” | <##>₁ [, ..., “string_n” | <##>_n]

Assembles each string with an automatic trailing zero byte or <##> into program memory, including the upper byte.

Note: If the ## is not a number, 0 will be assembled. If the ## is greater than 255, then the value will be truncated to a byte.

4.4.6 .pasciz “string₂”

Stores a sequence of ASCII characters (with an automatic trailing zero byte) into program memory, including the upper byte.

4.4.7 .byte expr₁ [, ..., expr_n]

Assembles one or more successive bytes in the current section.

If in a code (executable) section, the upper program memory byte will be filled with the last `.fillupper` value specified or the NOP opcode (0x00) if no `.fillupper` has been specified.

4.4.8 .pbyte expr₁ [, ..., expr_n]

Assembles one or more successive bytes in the current section. This directive will allow you to create data in the upper byte of program memory.

This directive is only allowed in a code section. If not in a code section, a warning is generated and the rest of the line is ignored.

4.4.9 .double value₁ [, ..., value_n]

Assembles one or more double-precision (64-bit) floating-point constants into consecutive addresses in little-endian format.

If in a code (executable) section, the upper program memory byte will be filled with the last `.fillupper` value specified or the NOP opcode (0x00) if no `.fillupper` has been specified.

Floating point numbers are in IEEE format (see [Section 3.5.1.2 “Floating-Point Numbers.”](#)).

The following statements are equivalent:

```
.double 12345.67
.double 1.234567e4
.double 1.234567e04
.double 1.234567e+04
.double 1.234567E4
.double 1.234567E04
.double 1.234567E+04
```

It is also possible to specify the hexadecimal encoding of a floating point constant. The following statements are equivalent and encode the value 12345.67 as a 64-bit double-precision number:

```
.double 0e:40C81CD5C28F5C29
.double 0f:40C81CD5C28F5C29
.double 0d:40C81CD5C28F5C29
```

4.4.10 **.fixed value₁[, ..., value_n]**

Assembles one or more 2-byte fixed-point constants (range $-1.0 \leq f < 1.0$) into consecutive addresses in little-endian format. Fixed-point numbers are in Q-15 format (see [Section 3.5.1.3 "Fixed-Point Numbers."](#)).

4.4.11 **.float value₁[, ..., value_n]**

Assembles one or more single-precision (32-bit) floating-point constants into consecutive addresses in little-endian format.

If in a code (executable) section, the upper program memory byte will be filled with the last `.fillupper` value specified or the NOP opcode (0x00) if no `.fillupper` has been specified.

Floating point numbers are in IEEE format (see [Section 3.5.1.2 "Floating-Point Numbers."](#)).

The following statements are equivalent:

```
.float 12345.67
.float 1.234567e4
.float 1.234567e04
.float 1.234567e+04
.float 1.234567E4
.float 1.234567E04
.float 1.234567E+04
```

It is also possible to specify the hexadecimal encoding of a floating-point constant. The following statements are equivalent and encode the value 12345.67 as a 32-bit double-precision number:

```
.float 0e:4640E6AE
.float 0f:4640E6AE
.float 0d:4640E6AE
```

4.4.12 **.single value₁[, ..., value_n]**

Assembles one or more single-precision (32-bit), floating-point constants into consecutive addresses in little-endian format.

If in a code (executable) section, the upper program memory byte will be filled with the last `.fillupper` value specified or the NOP opcode (0x00) if no `.fillupper` has been specified.

Floating point numbers are in IEEE format.

4.4.13 **.hword *expr*₁[, ..., *expr*_{*n*}]**

Assembles one or more 2-byte numbers into consecutive addresses in little-endian format.

4.4.14 **.int *expr*₁[, ..., *expr*_{*n*}]**

Assembles one or more 2-byte numbers into consecutive addresses in little-endian format.

4.4.15 **.long *expr*₁[, ..., *expr*_{*n*}]**

Assembles one or more 4-byte numbers into consecutive addresses in little-endian format.

4.4.16 **.short *expr*₁[, ..., *expr*_{*n*}]**

Same as `.word`.

4.4.17 **.string “*str*”**

Same as `.asciz`.

4.4.18 **.pstring “*str*”**

Same as `.pasciz`.

4.4.19 **.pstring “*string*₂”**

Same as `.pasciz "string2"`.

4.4.20 **.word *expr*₁[, ..., *expr*_{*n*}]**

Assembles one or more 2-byte numbers into consecutive addresses in little-endian format.

4.4.21 **.pword *expr*₁[, ..., *expr*_{*n*}]**

Assembles one or more 3-byte numbers into consecutive addresses in the current section.

This directive is only allowed in a code section. If not in a code section, a warning is generated and the rest of the line is ignored.

4.5 Directives that Declare Symbols

Declare symbol directives are:

- `.bss symbol, length [, algn]`
- `.comm symbol, length [, algn]`
- `.extern symbol`
- `.global symbol .globl symbol`
- `.lcomm symbol, length`
- `.weak symbol`

4.5.1 `.bss symbol, length [, algn]`

Reserve *length* (an absolute expression) bytes for a local symbol. The addresses are allocated in the bss section, so that at run-time the bytes start off zeroed. *symbol* is declared local so it is not visible to other objects. If *algn* is specified, it is the address alignment required for *symbol*. The bss location counter is advanced until it is a multiple of the requested alignment. The requested alignment must be a power of 2.

4.5.2 `.comm symbol, length [, algn]`

Declares a common symbol named *symbol*. When linking, a common symbol in one object file may be merged with a defined or common symbol of the same name in another object file. If the linker does not see a definition for that symbol, then it will allocate *length* bytes of uninitialized memory. If the linker sees multiple common symbols with the same name, and they do not all have the same size, the linker will allocate space using the largest size.

If *algn* is specified, it is the address alignment required for *symbol*. The requested alignment must be a power of two. *algn* is supported when the object file format is ELF; otherwise, it is ignored.

4.5.3 `.extern symbol`

Declares a symbol name that may be used in the current module, but it is defined as global in a different module.

4.5.4 `.global symbol` `.globl symbol`

Declares a symbol *symbol* that is defined in the current module and is available to other modules.

4.5.5 `.lcomm symbol, length`

Reserve *length* bytes for a local common denoted by *symbol*. The section and value of *symbol* are those of the new local common. The addresses are allocated in the bss section, so that at run-time, the bytes start off zeroed. *symbol* is not declared global so it is normally not visible to the linker.

4.5.6 `.weak symbol`

Marks the symbol named *symbol* as weak. When a weak-defined symbol is linked with a normal-defined symbol, the normal-defined symbol is used with no error. When a weak-undefined symbol is linked and the symbol is not defined, the value of the weak symbol becomes zero with no error.

4.6 Directives that Define Symbols

Define symbol directives are:

- `.equ symbol, expression`
- `.equiv symbol, expression`
- `.set symbol, expression`

4.6.1 `.equ symbol, expression`

Set the value of *symbol* to *expression*. You may set a symbol any number of times in assembly. If you set a global symbol, the value stored in the object file is the last value equated to it.

4.6.2 `.equiv symbol, expression`

Like `.equ`, except the assembler will signal an error if symbol is already defined.

4.6.3 `.set symbol, expression`

Same as `.equ`.

4.7 Directives that Modify Section Alignment

There are two ways to modify section alignment: implicitly and explicitly. Implicit alignment occurs first.

- [Implicit Alignment in Program Memory](#)
- [Explicit Section Alignment Directives](#)

4.7.1 Implicit Alignment in Program Memory

In addition to directives that explicitly align the location counter (such as `.align`, `.palign`, `.org`, `.porg`, etc.), many statements cause an implicit alignment to occur under certain conditions. Implicit alignment occurs when padding is inserted so that the next statement begins at a valid address. Padding uses the current `.fillvalue` and `.fillupper` values if specified; otherwise the value zero is used.

In data memory, a valid address is available for each byte. Since no data directives specify memory in quantities of less than one byte, implicit alignment is not required in data memory.

In program memory, a valid address is available for each instruction word (3 bytes). Since data directives can specify individual bytes, implicit alignment to the next valid address is sometimes required.

The following conditions cause implicit alignment in program memory:

1. Labels must be aligned to a valid address.

For example, the following source code:

```
.text
.pbyte 0x11
L1:
.pbyte 0x22
.pbyte 0x33,0x44
```

generates implicit alignment as shown:

```
Disassembly of section .text:
00000000 <.text>:
0:  11 00 00      nop
00000002 <L1>:
2:  22 33 44      .pword 0x443322
```

Note: Two bytes of padding were inserted so that label <code>L1</code> would be aligned to a valid address.
--

2. Instructions must be aligned to a valid address.

For example, the following source code:

```
.text
.pbyte 0x11
mov     w2,w3
```

generates implicit alignment as shown:

Disassembly of section .text:

```
00000000 <.text>:
0:   11 00 00          nop
2:   82 01 78          mov.w     w2, w3
```

Note: Two bytes of padding were inserted so that the `mov` instruction would be aligned to a valid address.

3. Transitions between p-type data directives (`.pbyte`, `.pspace`, etc.) and normal data directives (`.byte`, `.space`, etc.), in either direction, are aligned to a valid address.

For example, the following source code:

```
.text
.byte 0x11
.pbyte 0x22
.pbyte 0x33,0x44
```

generates implicit alignment as shown:

Disassembly of section .text:

```
00000000 <.text>:
0:   11 00 00          nop
2:   22 33 44          .pword 0x443322
```

Note: Two bytes of padding were inserted so that the transition from normal to p-type directive would be aligned to a valid address.

4.7.2 Explicit Section Alignment Directives

Directives that explicitly modify section alignment are:

- `.align align[, fill[, max-skip]]`
- `.palign align[, fill[, max-skip]]`
- `.fill repeat[, size[, fill]]`
- `.pfill repeat[, size[, fill]]`
- `.org new-lc[, fill]`
- `.porg new-lc[, fill]`
- `.skip size[, fill] .space size[, fill]`
- `.pskip size[, fill] .pspace size[, fill]`
- `.struct expression`

4.7.3 `.align align[, fill[, max-skip]]`

Pad the location counter (in the current subsection) to a particular storage boundary.

align is the address alignment required. The location counter is advanced until it is a multiple of the requested alignment. If the location counter is already a multiple of the requested alignment, no change is needed or made. In a code section, an alignment of 2 is required to align to the next instruction word. The requested alignment must be a power of 2.

fill is optional. If not specified:

- In a data section, a value of 0x00 is used to fill the skipped bytes.
- In a code section, the last specified `.fillvalue` is used to fill the lower two bytes of program memory and the last specified `.fillupper` is used to fill the upper program memory byte.

max-skip is optional. If specified, it is the maximum number of bytes that should be skipped by this directive. If doing the alignment would require skipping more bytes than the specified maximum, then the alignment is not done at all.

Alignment within a section is required for modulo addressing. It is worth noting that the overall section alignment reflects the greatest alignment of any `.align` directives that are included. Further, the assembler must pad out the section length to match its alignment. This is done in order to preserve the requested alignment in case the section is combined with other sections of the same name during the link. To avoid unnecessary padding of aligned sections, use the section name `*`, which identifies a unique section that will never be combined.

4.7.4 .palign *align*[, *fill*[, *max-skip*]]

Pad the location counter (in the current subsection) to a particular storage boundary. This directive is only allowed in a code section. If not in a code section, a warning is generated and the rest of the line is ignored.

align is the address alignment required. The location counter is advanced until it is a multiple of the requested alignment. If the location counter is already a multiple of the requested alignment, no change is needed. In a code section, an alignment of 2 is required to align to the next instruction word. The requested alignment must be a power of 2.

fill is optional. If not specified, the last `.pfillvalue` specified is used to fill the skipped bytes. All three bytes of the program memory word are filled.

max-skip is optional. If specified, it is the maximum number of bytes (including the upper byte) that should be skipped by this directive. If doing the alignment would require skipping more bytes than the specified maximum, then the alignment is not done at all.

4.7.5 .fill *repeat*[, *size*[, *fill*]]

Reserve *repeat* copies of *size* bytes. *repeat* may be zero or more. *size* may be zero or more, but if it is more than 8, then it is deemed to have the value 8. The content of each *repeat* bytes is taken from an 8-byte number. The highest order 4 bytes are zero. The lowest order 4 bytes are value rendered in the little-endian byte-order. Each *size* bytes in a repetition is taken from the lowest order *size* bytes of this number.

size is optional and defaults to one, if omitted.

fill is optional. If not specified:

- In a data section, a value of 0x00 is used to fill the skipped bytes.
- In a code section, the last specified `.fillvalue` is used to fill the lower two bytes of program memory and the last specified `.fillupper` is used to fill the upper program memory byte.

4.7.6 .pfill *repeat*[, *size*[, *fill*]]

Reserve *repeat* copies of *size* bytes including the upper byte. *repeat* may be zero or more. *size* may be zero or more, but if it is more than 8, then it is deemed to have the value 8. The content of each *repeat* byte is taken from an 8-byte number. The highest order 4 bytes are zero. The lowest order 4 bytes are value rendered in the little-endian byte-order. Each *size* byte in a repetition is taken from the lowest order *size* bytes of this number.

This directive is only allowed in a code section. If not in a code section, a warning is generated and the rest of the line is ignored.

size is optional and defaults to one, if omitted. Size is the number of bytes to reserve (including the upper byte).

fill is optional. If not specified, it defaults to the last `.pfillvalue` specified. All three bytes of each instruction word are filled.

4.7.7 .org *new-lc* [, *fill*]

Advance the location counter of the current section to *new-lc*. In program memory, *new-lc* is specified in PC units. On the 16-bit device, the PC increments by 2 for each instruction word. Odd values are not permitted.

Note: A location counter is not an absolute address but the offset from the start of the section in which the .org occurs.

The bytes between the current location counter and the new location counter are filled with *fill*. *new-lc* is an absolute expression. You cannot .org backwards. You cannot use .org to cross sections.

The new location counter is relative to the current module and is not an absolute address.

fill is optional. If not specified:

- In a data section, a value of 0x00 is used to fill the skipped bytes.
- In a code section, the last specified .fillvalue is used to fill the lower two bytes of program memory and the last specified .fillupper is used to fill the upper program memory byte.

4.7.8 .porg *new-lc* [, *fill*]

Advance the location counter of the current section to *new-lc*. In program memory, *new-lc* is specified in PC units. On the 16-bit device, the PC increments by 2 for each instruction word. Odd values are not permitted.

Note: A location counter is not an absolute address but the offset from the start of the section in which the .porg occurs.

The bytes between the current location counter and the new location counter are filled with *fill*. *new-lc* is an absolute expression. You cannot .porg backwards. You cannot use .porg to cross sections.

The new location counter is relative to the current module and is not an absolute address.

This directive is only allowed in a code section. If not in a code section, a warning is generated and the rest of the line is ignored.

fill is optional. If not specified, it defaults to the last .pfillvalue specified. All three bytes of each instruction word are filled.

4.7.9 .skip *size* [, *fill*] .space *size* [, *fill*]

Reserve *size* bytes. Each byte is filled with the value *fill*.

fill is optional. If the value specified for *fill* is larger than a byte, a warning is displayed and the value is truncated to a byte. If not specified:

- In a data section, a value of 0x00 is used to fill the skipped bytes.
- In a code section, the last specified .fillvalue is used to fill the lower two bytes of program memory and the last specified .fillupper is used to fill the upper program memory byte.

4.7.10 **.pskip *size*, *fill*** **.pspace *size*, *fill***

Reserve *size* bytes (including the upper byte). Each byte is filled with the value *fill*.

This directive is only allowed in a code section. If not in a code section, a warning is generated and the rest of the line is ignored.

The new location counter is relative to the current module and is not an absolute address.

fill is optional. If the value specified for *fill* is larger than a byte, a warning is displayed and the value is truncated to a byte. If not specified, it defaults to the last *.pfillvalue* specified. All three bytes of each instruction word are filled.

4.7.11 **.struct *expression***

Switch to the absolute section, and set the section offset to *expression*, which must be an absolute expression. You might use this as follows:

```
.struct 0
field1:
    .struct field1 + 4
field2:
    .struct field2 + 4
field3:
```

This would define the symbol *field1* to have the value 0, the symbol *field2* to have the value 4, and the symbol *field3* to have the value 8. Assembly would be left in the absolute section, and you would need to use a *.section* directive of some sort to change to some other section before further assembly.

4.8 Directives that Format the Output Listing

Output listing format directives are:

- `.eject`
- `.list`
- `.nolist`
- `.psize lines[, columns]`
- `.sbttl "subheading"`
- `.title "heading"`

4.8.1 `.eject`

Force a page break at this point when generating assembly listings.

4.8.2 `.list`

Controls (in conjunction with `.nolist`) whether assembly listings are generated. This directive increments an internal counter (which is one initially). Assembly listings are generated if this counter is greater than zero.

Only functional when listings are enabled with the `-a` command line option and forms processing has not been disabled with the `-an` command line option.

4.8.3 `.nolist`

Controls (in conjunction with `.list`) whether assembly listings are generated. This directive decrements an internal counter (which is one initially). Assembly listings are generated if this counter is greater than zero.

Only functional when listings are enabled with the `-a` command line option and forms processing has not been disabled with the `-an` command line option.

4.8.4 `.psize lines[, columns]`

Declares the number of lines, and optionally, the number of columns to use for each page when generating listings.

Only functional when listings are enabled with the `-a` command line option and forms processing has not been disabled with the `-an` command line option.

4.8.5 `.sbttl "subheading"`

Use subheading as a subtitle (third line, immediately after the title line) when generating assembly listings. This directive affects subsequent pages, as well as the current page, if it appears within ten lines of the top.

4.8.6 `.title "heading"`

Use heading as the title (second line, immediately after the source file name and page number) when generating assembly listings.

4.9 Directives that Control Conditional Assembly

Conditional assembly directives are:

- `.else`
- `.elseif expr`
- `.endif`
- `.err`
- `.error "string"`
- `.if expr`
- `.ifdecl symbol`
- `.ifndecl symbol .ifnotdecl symbol`
- `.ifdef symbol`
- `.ifndef symbol .ifnotdef symbol`

4.9.1 `.else`

Used in conjunction with the `.if` directive to provide an alternative path of assembly code should the `.if` evaluate to false.

4.9.2 `.elseif expr`

Used in conjunction with the `.if` directive to provide an alternative path of assembly code should the `.if` evaluate to false, and a second condition exists.

4.9.3 `.endif`

Marks the end of a block of code that is only assembled conditionally.

4.9.4 `.err`

If the assembler sees an `.err` directive, it will print an error message, and unless the `-Z` option was used, it will not generate an object file. This can be used to signal an error in conditionally compiled code.

4.9.5 `.error "string"`

Similar to `.err`, except that the specified string is printed.

4.9.6 `.if expr`

Marks the beginning of a section of code that is only considered part of the source program being assembled if the argument `expr` is non-zero. The end of the conditional section of code must be marked by an `.endif`; optionally, you may include code for the alternative condition, flagged by `.else`.

4.9.7 `.ifdecl symbol`

Assembles the following section of code if the specified symbol has been declared.

4.9.8 `.ifndecl symbol` `.ifnotdecl symbol`

Assembles the following section of code if the specified symbol has **not** been declared.

4.9.9 `.ifdef symbol`

Assembles the following section of code if the specified symbol has been defined (i.e., assigned a value).

4.9.10 *.ifndef symbol* *.ifnotdef symbol*

Assembles the following section of code if the specified symbol has **not** been defined (i.e., not assigned a value).

4.10 Directives for Substitution/Expansion

Substitution/expansion directives are:

- `.exitm`
- `.irp symbol, value1 [, ..., valuen]endr`
- `.irpc symbol, valueendr`
- `.macro symbol arg1[=default] [, ..., argn [=default]]endm`
- `.purgem "name"`
- `.rept countendr`

4.10.1 `.exitm`

Exit early from the current marco definition. See `.macro` directive.

4.10.2 `.irp symbol, value1 [, ..., valuen]endr`

Evaluate a sequence of statements assigning different values to *symbol*. The sequence of statements starts at the `.irp` directive, and is terminated by a `.endr` directive. For each *value*, *symbol* is set to *value*, and the sequence of statements is assembled. If no *value* is listed, the sequence of statements is assembled once, with *symbol* set to the null string. To refer to *symbol* within the sequence of statements, use `\symbol`.

For example, assembling:

```
.irp reg,0,1,2,3  
push w\reg  
.endr
```

is equivalent to assembling:

```
push w0  
push w1  
push w2  
push w3
```

4.10.3 `.irpc symbol, value`

...
`.endr`

Evaluate a sequence of statements assigning different values to *symbol*. The sequence of statements starts at the `.irpc` directive and is terminated by a `.endr` directive. For each character in *value*, *symbol* is set to the character, and the sequence of statements is assembled. If no *value* is listed, the sequence of statements is assembled once, with *symbol* set to the null string. To refer to *symbol* within the sequence of statements, use `\symbol`.

For example, assembling:

```
irpc reg,0123
push w\reg
.endr
```

is equivalent to assembling:

```
push w0
push w1
push w2
push w3
```

4.10.4 `.macro symbol arg1[=default] [, ..., argn [=default]]`

...
`.endm`

Define macros that generate assembly output. A macro accepts optional arguments and can call other macros or even itself recursively.

If a macro definition requires arguments, specify their names after the macro name, separated by commas or spaces. To refer to arguments within the macro block, use `\arg` or `&arg&`. The second form can be used to combine an argument with additional characters to create a symbol name.

For example, assembling:

```
.macro display_int sym
    mov \sym,w0
    rcall display
.endm
```

```
display_int result
```

is equivalent to assembling:

```
mov result,w0
rcall display
```

In the next example, a macro is used to define HI- and LO-word constants for a 32-bit integer.

```
.macro LCONST name,value
.equ    \name,\value
.equ    &name&LO,(\value) & 0xFFFF
.equ    &name&HI,((\value)>>16) & 0xFFFF
.endm
```

```
LCONST seconds_per_day 60*60*24
```

```
mov    #seconds_per_dayLO,w0
mov    #seconds_per_dayHI,w1
```


`xc16-as` maintains a counter of how many macros have been executed in the pseudo-variable `\@`. This value can be copied to the assembly output, but only within a macro definition. In the following example, a recursive macro is used to allocate an arbitrary number of labeled buffers.

```
.macro make_buffers num,size
BUF\@: .space \size
      .if (\num - 1)
make_buffers (\num - 1),\size
      .endif
      .endm

      .bss
make_buffers 4,16      ; create BUF0..BUF3, 16 bytes each
```

4.10.5 `.purgem "name"`

Undefine the macro *name*, so that later uses of the string will not be expanded. See `.macro` directive.

4.10.6 `.rept count`

```
...
.endr
```

Repeat the sequence of lines between the `.rept` directive and the next `.endr` directive count times.

For example, assembling

```
.rept 3
.long 0
.endr
```

is equivalent to assembling

```
.long 0
.long 0
.long 0
```

4.11 Miscellaneous Directives

Miscellaneous directives are:

- `.abort`
- `.appline line-number .ln line-number`
- `.end`
- `.fail expression`
- `.ident "comment"`
- `.incbin "file" [,skip[,count]]`
- `.include "file"`
- `.loc file-number, line-number`
- `.pincbin "file" [,skip[,count]]`
- `.print "string"`
- `.version "string"`

4.11.1 `.abort`

Prints out the message ".abort detected. Abandoning ship." and exits the program.

4.11.2 `.appline line-number` `.ln line-number`

Change the logical line number. The next line has that logical line number.

4.11.3 `.end`

End program

4.11.4 `.fail expression`

Generates an error or a warning. If the value of the *expression* is 500 or more, `as` will print a warning message. If the value is less than 500, `xc16-as` will print an error message. The message will include the value of *expression*. This can occasionally be useful inside complex nested macros or conditional assembly.

4.11.5 `.ident "comment"`

Appends *comment* to the section named `.comment`. This section is created if it does not exist. The 16-bit linker will ignore this section when allocating program and data memory, but will combine all `.comment` sections together, in link order.

4.11.6 `.incbin "file" [,skip[,count]]`

The `.incbin` directive includes *file* verbatim at the current location. The file is assumed to contain binary data. The search paths used can be specified with the `-I` command-line option (see [Chapter 2. "Assembler Command Line Options"](#)). Quotation marks are required around *file*.

The *skip* argument skips a number of bytes from the start of the file. The *count* argument indicates the maximum number of bytes to read. Note that the data is not aligned in any way, so it is the user's responsibility to make sure that proper alignment is provided both before and after the `.incbin` directive.

When used in an executable section, `.incbin` fills only the lower 16 bits of each program word.

4.11.7 .include “file”

Provides a way to include supporting files at specified points in your source code. The code is assembled as if it followed the point of the `.include`. When the end of the included file is reached, assembly of the original file continues at the statement following the `.include`.

4.11.8 .loc file-number, line-number

`.loc` is essentially the same as `.ln`. The assembler expects that this directive occurs in the `.text` section. `file-number` is ignored.

4.11.9 .pincbin “file”[,skip[,count]]

The `.pincbin` directive includes `file` verbatim at the current location. The file is assumed to contain binary data. The search paths used can be specified with the `-I` command-line option (see [Chapter 2. “Assembler Command Line Options”](#)). Quotation marks are required around `file`.

The `skip` argument skips a number of bytes from the start of the file. The `count` argument indicates the maximum number of bytes to read. Note that the data is not aligned in any way, so it is the user's responsibility to make sure that proper alignment is provided both before and after the `.pincbin` directive.

`.pincbin` is supported only in executable sections, and fills all 24 bits of each program word.

4.11.10 .print “string”

Prints `string` on the standard output during assembly.

4.11.11 .version “string”

This directive creates a `.note` section and places into it an ELF formatted note of type `NT_VERSION`. The note's name is set to `string`. `.version` is supported when the output file format is ELF; otherwise, it is ignored.

4.12 Directives for Debug Information

Debug information directives are:

- `.def name`
- `.dim`
- `.endef`
- `.file "string"`
- `.line line-number`
- `.scl class`
- `.size expression`
- `.size name, expression`
- `.sleb128 expr1 [, ..., exprn]`
- `.tag structname`
- `.type value`
- `.type name, description`
- `.uleb128 expr1 [, ..., exprn]`
- `.val addr`

4.12.1 `.def name`

Begin defining debugging information for a symbol *name*; the definition extends until the `.endef` directive is encountered.

4.12.2 `.dim`

Generated by compilers to include auxiliary debugging information in the symbol table. Only permitted inside `.def/.endef` pairs.

4.12.3 `.endef`

Flags the end of a symbol definition begun with `.def`.

4.12.4 `.file "string"`

Tells the assembler that it is about to start a new logical file. This information is placed into the object file.

4.12.5 `.line line-number`

Generated by compilers to include auxiliary symbol information for debugging. Only permitted inside `.def/.endef` pairs.

4.12.6 `.scl class`

Set the storage class value for a symbol. May only be used within `.def/.endef` pairs.

4.12.7 `.size expression`

Generated by compilers to include auxiliary debugging information in the symbol table. Only permitted inside `.def/.endef` pairs.

4.12.8 `.size name, expression`

Generated by compilers to include auxiliary information for debugging. This variation of `.size` is supported when the output file format is in Executable and Linking Format (ELF).

4.12.9 **.sleb128** *expr*₁ [, ..., *expr*_{*n*}]

Signed little endian base 128. Compact variable length representation of numbers used by the DWARF symbolic debugging format.

4.12.10 **.tag** *structname*

Generated by compilers to include auxiliary debugging information in the symbol table. Only permitted inside `.def/.endef` pairs. Tags are used to link structure definitions in the symbol table with instances of those structures.

4.12.11 **.type** *value*

Records the integer value as the type attribute of a symbol table entry. Only permitted within `.def/.endef` pairs.

4.12.12 **.type** *name, description*

Sets the type of symbol name to be either a function symbol or an object symbol. This variation of `.type` is supported when the output file format is ELF. For example,

```
.text
.type foo,@function
foo:
    return

.data
.type dat,@object
dat:    .word 0x1234
```

4.12.13 **.uleb128** *expr*₁ [, ..., *expr*_{*n*}]

Unsigned little endian base 128. Compact variable length representation of numbers used by the DWARF symbolic debugging format.

4.12.14 **.val** *addr*

Records the address *addr* as the value attribute of a symbol table entry. Only permitted within `.def/.endef` pairs.

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Chapter 5. Assembler Listing File

5.1 Introduction

The assembler has the capability to produce listing files. These listing files are not absolute listing files, and the addresses that appear in the listing are relative to the start of sections.

Topics covered in this chapter are:

- [Generation](#)
- [Contents](#)

5.2 Generation

To generate a listing file whether in MPLAB X IDE, in MPLAB IDE v8 or on the command line, you will need to specify an option described in [Section 2.3 “Options that Modify the Listing Output.”](#) By default, a listing file is written to a `.lst` file.

16-Bit Assembler, Linker and Utilities User's Guide

5.3 Contents

The listing files produced by the assembler are composed of the several elements. [Example 5-1](#) shows a sample listing file.

The example listing file contains these elements:

- **Header** – contains the name of the assembler, the name of the file being assembled, and a page number. This is not shown if the `-an` option is specified.
- **Title Line** – contains the title specified by the `.title` directive. This is not shown if the `-an` option is specified.
- **Subtitle** – contains the subtitle specified by the `.sbttl` directive. This is not shown if the `-an` option is specified.
- **High-level source** if the `-ah` option is given to the assembler. The format for high-level source is:

```
<line #>:<filename>          **** <source>
```

For example:

```
1:hello.c          **** #include <stdio.h>
```

- **Assembler source** if the `-al` option is given to the assembler. The format for assembler source is:

```
<line #> <addr> <encoded bytes> <source>
```

For example:

```
245 000004 00 0F 78          mov      w0,[w14]
```

Notes:

- 1: Line numbers may be repeated.
- 2: Addresses are relative to sections in this module and are not absolute.
- 3: Instructions are encoded in “little endian” order.

- **Symbol table** if the `-as` option is given to the assembler. Both, a list of defined and undefined symbols will be given. The defined symbols will have the format:

```
DEFINED SYMBOLS
<filename>:<line #> <section>:<addr> <symbol>
```

For example:

```
DEFINED SYMBOLS
foo.s:229      .text:00000000 _main
```

The undefined symbols will have the format:

```
UNDEFINED SYMBOLS
<symbol>
```

For example:

```
UNDEFINED SYMBOLS
printf
```


Example 5-1: Sample Assembler Listing File

MPLAB XC16_ASM_Listing: example1.1.s page 1
Sample dsPIC Assembler Source Code
For illustration only.

```
1
2                                .title " Sample dsPIC Assembler Source Code"
3                                .sbttl " For illustration only."
4
5                                ; dsPIC registers
6                                .equ CORCONL, CORCON
7                                .equ PSV,2
8
9                                .section .const,psv
10                               hello:
11 0000 48 65 6C 6C                .ascii "Hello world!\n\0"
11      6F 20 77 6F
11      72 6C 64 21
11      0A 00
12
13                                .text
14                                .global __reset
15                               __reset:
16                                ; set PSVPAG to page that contains 'hello'
17 000000 00 00 20                mov     #psvpage(hello),w0
18 000002 00 00 88                mov     w0,PSVPAG
19
20                                ; enable Program Space Visibility
21 000004 00 40 A8                bset.b  CORCONL,#PSV
22
23                                ; make a pointer to 'hello'
24 000006 00 00 20                mov     #psvoffset(hello),w0
25
26                                .end
```

MPLAB XC16_ASM_Listing: example1.1.s page 2
Sample dsPIC Assembler Source Code
For illustration only.

DEFINED SYMBOLS

	ABS:00000000	fake
example1.1.s:10	.const:00000000	hello
example1.1.s:15	.text:00000000	__reset
	.text:00000000	.text
	.data:00000000	.data
	.bss:00000000	.bss
	.const:00000000	.const

UNDEFINED SYMBOLS
CORCON
PSVPAG

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Chapter 6. Assembler Errors/Warnings/Messages

6.1 Introduction

MPLAB XC16 Assembler may generate errors, warnings and messages. To select the messages that are generated, see [Section 2.4 “Options that Control Informational Output.”](#)

The following topics are covered in this appendix:

- [Fatal Errors](#)
- [Errors](#)
- [Warnings](#)
- [Messages](#)

For information on assembler **limitations** and **known problems**, see the Readme file.

6.2 Fatal Errors

The following errors indicate that an internal error has occurred in the assembler. Please contact Microchip Technology for support if any of the following errors are generated:

- A dummy instruction cannot be used!
- bad floating-point constant: exponent overflow, probably assembling junk
- bad floating-point constant: unknown error code=error_code
- C_EFCN symbol out of scope
- Can't continue
- Can't extend frag num. chars
- Can't open a bfd on stdout name
- Case value val unexpected at line _line_ of file “_file_”
- emulations not handled in this configuration
- error constructing pop_table_name pseudo-op table: err_txt
- expr.c(operand): bad atof_generic return val val
- failed sanity check.
- filename:line_num: bad return from bfd_install_relocation: val
- filename:line_num: bad return from bfd_install_relocation
- Inserting “name” into symbol table failed: error_string
- pic30_get_g_or_h_mode_value called with an invalid operand type
- pic30_get_p_or_q_mode_value called with an invalid operand type
- pic30_insert_dsp_writeback called with an invalid operand type
- pic30_insert_dsp_x_prefetch_operation called with an invalid offset
- pic30_insert_dsp_x_prefetch_operation called with an invalid operand type
- pic30_insert_dsp_y_prefetch_operation called with an invalid offset
- pic30_insert_dsp_y_prefetch_operation called with an invalid operand type
- invalid segment “name”; segment “name” assumed
- label “temp\$” redefined
- macros nested too deeply

- missing emulation mode name
- multiple emulation names specified
- Relocation type not supported by object file format
- reloc type not supported by object file format
- rva not supported
- rva without symbol
- unrecognized emulation name 'em'
- Unsupported BFD relocation size in bytes

6.3 Errors

Symbol [A](#) [B](#) [C](#) [D](#) [E](#) [F](#) [I](#) [L](#) [M](#) [N](#) [O](#) [P](#) [R](#) [S](#) [T](#) [U](#) [W](#)

Symbol

.abort detected. Abandoning ship.

User error invoked with the `.abort` directive.

.else without matching .if - ignored.

An `.else` directive was seen without a preceding `.if` directive.

“.elseif” after “.else” - ignored

An `.elseif` directive specified after a `.else` directive. Modify your code so that the `.elseif` directive comes before the `.else` directive.

“.elseif” without matching “.if” - ignored.

An `.elseif` directive was seen without a preceding `.if` directive.

“.endif” without “.if”

An `.endif` directive was seen without a preceding `.if` directive.

.err encountered.

A user error invoked with the `.err` directive.

sign not valid in data allocation directive.

The `#` sign cannot be used within a data allocation directive (`.byte`, `.word`, `.pword`, `.long`, etc.)

warnings, treating warnings as errors.

The `--fatal-warnings` command line option was specified on the command line and warnings existed.

A

absolute address can not be specified for section '.const'.

Section `.const` is a C compiler resource. Although it is permissible for an application to allocate constants in section `.const` explicitly, it is not permissible to assign an absolute address for this section.

Absolute address must be greater than or equal to 0.

A negative absolute address was specified as the target for the `DO` or `BRA` instruction. The assembler does not know anything about negative addresses.

Alignment in CODE section must be at least 2 units.

The alignment value for the `.align` directive must be at least 2 units. Either no alignment was specified or an alignment less than 2 was specified. Modify the `.align` directive to have an alignment of at least 2.

Attributes for section 'name' conflict with implied attributes.

Certain section names have implied attributes. In this case, the attributes specified in a `.section` directive conflict with its implied attributes. See [Section 4.2 “Directives that Define Sections.”](#) for more information.

B**backw. ref to unknown label “#.”, 0 assumed.**

A backwards reference was made to a local label that was not seen. See [Section 3.6.1 “Reserved Names.”](#) for more information on local labels.

bad defsym; format is --defsym name=value.

The format for the command line option `--defsym` is incorrect. Most likely, you are missing the `=` between the name and the value.

Bad expression.

The assembler did not recognize the expression. See [Chapter 2. “Assembler Command Line Options”](#) and [Chapter 3. “MPLAB XC16 Assembly Language”](#) for more details on assembler syntax.

bignum invalid; zero assumed.

The big number specified in the expression is not valid.

Byte operations expect an offset between -512 and 511.

The offset specified in `[Wn+offset]` or `[Wn-offset]` exceeded the maximum or minimum value allowed for byte instructions.

C**Cannot call a symbol (name) that is not located in an executable section.**

Attempted to `CALL` a symbol that is not located in a `CODE` section.

Cannot create floating-point number.

Could not create a floating-point number because of exponent overflow or because of a floating-point exception that prohibits the assembler from encoding the floating-point number.

Cannot redefine executable symbol ‘s’.

A statement label or an executable section cannot be redefined with a `.set` or `.equ` directive.

Cannot reference executable symbol (name) in a data context.

An attempt was made to use a symbol in an executable section as a data address. To reference an executable symbol in a data context, the `psvoffset()` or `tbloffset()` operator is required.

Cannot use a constant as the argument of dmaoffset.

An attempt was made to use a constant as the argument to a `dmaoffset`.

Can not use dmaoffset on a symbol (name) that is not located in a dma section.

For some devices, the `dmaoffset()` operator can only be used on symbols that are located in dma memory.

16-Bit Assembler, Linker and Utilities User's Guide

Cannot use operator on a symbol (name) that is not located in an executable or read-only section.

The following operators can be applied to symbols in executable or read-only sections only: `tbloffset()`, `psvoffset()`, `tblpage()`, `psvpage()`, `handle()`, `paddr()`.

Cannot use operator on a symbol (name) that is not located in a code, psv or eedata section.

You cannot use one of the special operators (`tbloffset`, `tblpage`, `psvoffset`, `psvpage`, `handle` or `paddr`) on a symbol that is not located in a code, psv or eedata section.

Cannot use operator with this directive.

An attempt was made to use a special operator (`tbloffset`, `tblpage`, `psvoffset`, `psvpage`, `handle` or `paddr`) with a data allocation directive that does not allocate enough bytes to store the requested data.

Cannot write to output file.

For some reason, the output file could not be written to. Ensure that you have write permission to the file and that there is enough disk space.

Can't open file_name for reading.

The specified input source file could not be opened. Ensure that the file exists and that you have permission to access the file.

D

directive directive not supported in pic30 target.

The pic30 target does not support this directive. This directive is available in other versions of the assembler, but the pic30 target does not support it for one reason or another. Please check [Chapter 4. "Assembler Directives"](#) for a complete list of supported directives.

duplicate "else" - ignored.

Two `.else` directives were specified for the same `.if` directive.

E

end of file inside conditional.

The file ends without terminating the current conditional. Add a `.endif` to your code.

end of macro inside conditional.

A conditional is unterminated inside a macro. The `.endif` directive to end the current conditional was not specified before seeing the `.endm` directive.

Expected comma after symbol-name: rest of line ignored.

Missing comma from the `.comm` directive after the symbol name.

Expected constant expression for fill argument.

The fill argument for the `.fill`, `.pfill`, `.skip`, `.pskip`, `.space` or `.pspace` directive must be a constant value. Attempted to use a symbol. Replace symbol with a constant value.

Expected constant expression for new-lc argument.

The new location counter argument for the `.org` directive must be a constant value. Attempted to use a symbol. Replace symbol with a constant value.

Expected constant expression for repeat argument.

The repeat argument for the `.fill`, `.pfill`, `.skip`, `.pskip`, `.space` or `.pspace` directive must be a constant value. Attempted to use a symbol. Replace symbol with a constant value.

Expected constant expression for size argument.

The size argument for the `.fill` or `.pfill` directive must be a constant value. Attempted to use a symbol. Replace symbol with a constant value.

Expression too complex.

An expression is too complex for the assembler to process.

F**floating point number invalid; zero assumed.**

The floating-point number specified in the expression is not valid.

I**Ignoring attempt to re-define symbol 'symbol'.**

The symbol that you are attempting to define with `.comm` or `.lcomm` has already been defined and is not a common symbol.

Invalid expression (expr) contained inside of the brackets.

Assembler did not recognize the expression between the brackets.

invalid identifier for ".ifdef".

The identifier specified after the `.ifdef` must be a symbol. See [Section 3.6.1 "Reserved Names."](#) and [Section 4.9 "Directives that Control Conditional Assembly."](#) for more details.

Invalid mnemonic: 'token'.

The token being parsed is not a valid mnemonic for the instruction set.

invalid listing option 'optarg'.

The sub-option specified is not valid. Acceptable suboptions are `c`, `d`, `h`, `l`, `m`, `n`, `v` and `=`.

Invalid operands specified ('insn'). Check operand #n.

The operands specified were invalid. The assembler was able to match `n-1` operands successfully. Although there is no assurance that operand `#n` is the culprit, it is a general idea of where you should begin looking.

Invalid operand syntax ('insn').

This message usually comes hand-in-hand with one of the previous operand syntax errors.

Invalid post increment value. Must be +/- 2, 4 or 6.

Assembler saw `[Wn] +=value`, where `value` is expected to be a +/- 2, 4 or 6. `value` was not correct. Specify a `value` of +/- 2, 4 or 6.

Invalid post decrement value. Must be +/- 2, 4 or 6.

Assembler saw `[Wn] -=value`, where `value` is expected to be a +/- 2, 4 or 6. `value` was not correct. Specify a `value` of +/- 2, 4 or 6.

Invalid register in operand expression.

Assembler was attempting to find either pre- or post-increment or decrement. The operand did not contain a register. Specify one of the registers `w0-w16` or `W0-W16`.

16-Bit Assembler, Linker and Utilities User's Guide

Invalid register in expression reg.

Assembler saw [junk] or [junk] +=n or [junk] -=n. Was expecting a register between the brackets. Specify one of the registers w0-w16 or W0-W16 between the brackets.

Invalid use of ++ in operand expression.

Assembler was attempting to find either pre- or post-increment. The operand specified was neither pre-increment [++Wn] nor post-increment [Wn++]. Make sure that you are not using the old syntax of [Wn]++.

Invalid use of -- in operand expression.

Assembler was attempting to find either pre- or post-decrement. The operand specified was neither pre-decrement [--Wn] nor post-decrement [Wn--]. Make sure that you are not using the old syntax of [Wn]--.

Invalid value (#) for relocation name.

The final value of the relocation is not a valid value for the operand associated with the given relocation.

'name' is not a valid attribute name.

While processing a .section directive, the assembler found an identifier that is not a valid section attribute.

L

Length of .comm "sym" is already #. Not changed to #.

An attempt was made to redefine the length of a common symbol.

M

misplaced)

Missing parenthesis when expanding a macro. The syntax \(...) will literally substitute the text between the parenthesis into the macro. The trailing parenthesis was missing from this syntax.

Missing model parameter.

Missing symbol in the .irp or .irpc directive.

Missing right bracket.

The assembler did not see the terminating bracket ']'.

Missing size expression.

The .lcomm directive is missing the length expression.

Missing ')' after formals.

Missing trailing parenthesis when listing the macro formals inside of parenthesis.

Missing ')' assumed.

Expected a terminating parenthesis ')' while parsing the expression. Did not see one where expected so assumes where you wanted the trailing parenthesis.

Missing ']' assumed.

Expected a terminating brace ']' while parsing the expression. Did not see one where expected so assumes where you wanted the trailing brace.

Mnemonic not found.

The assembler was expecting to parse an instruction and could not find a mnemonic.

N

Negative of non-absolute symbol name.

Attempted to take the negative of a symbol name that is non-absolute. For example, `.word -sym`, where `sym` is external.

New line in title.

The `.title` heading is missing a terminating quote.

non-constant expression in “.elseif” statement.

The argument of the `.elseif` directive must be a constant value able to be resolved on the first pass of the directive. Ensure that any `.equ` of a symbol used in this argument is located before the directive. See [Section 4.9 “Directives that Control Conditional Assembly.”](#) for more details.

non-constant expression in “.if” statement.

The argument of the `.if` directive must be a constant value able to be resolved on the first pass of the directive. Ensure that any `.equ` of a symbol used in this argument is located before the directive. See [Section 4.9 “Directives that Control Conditional Assembly.”](#) for more details.

Number of operands exceeds maximum number of 8.

Too many operands were specified in the instruction. The largest number of operands accepted by any of the 16-bit device instructions is 8.

O

Only support plus register displacement (i.e., [Wb+Wn]).

Assembler found [Wb-Wn]. The syntax only supports a plus register displacement.

Operands share encoding bits. The operands must encode identically.

Two operands are register with displacement addressing mode [Wb+Wn]. The two operands share encoding bits so the Wn portion must match or be able to be switched to match the Wb of the other operand.

operation combines symbols in different segments.

The left-hand side of the expression and the right-hand side of the expression are located in two different sections. The assembler does not know how to handle this expression.

operator modifier must be preceded by a #.

The modifier (`tbloffset`, `tblpage`, `psvoffset`, `psvpage`, `handle`) was specified inside of an instruction, but was not preceded by a `#`. Include the `#` to represent that this is a literal.

P

paddr modifier not allowed in instruction.

The `paddr` operator was specified in an instruction. This operator can only be specified in a `.pword` or `.long` directive as those are the only two locations that are wide enough to store all 24 bits of the program address.

PC relative expression is not a valid GOTO target.

The assembler does not support expressions which modify the PC of a GOTO destination such as `“ . + 4 ”` or `“ sym + 100 ”`.

R

Register expected as first operand of expression expr.

Assembler found [junk+anything] or [junk-anything]. The only valid expression contained in brackets with a + or a - requires that the first operand be a register.

Register or constant literal expected as second operand of expression expr.

Assembler found [Wn+junk] or [Wn-junk]. The only valid operand for this format is a register with plus or minus literal offset or a register with displacement.

Requested alignment 'n' is greater than alignment of absolute section 'name'

When the `address()` attribute is used to specify an absolute address for a section, it constrains the ability of the assembler to align objects within the section. The alignment specified in a `.align` or `.palign` directive must not be greater than the alignment implied by the section address.

S

section alignment must be a power of two.

The argument to an `align()` or `reverse()` section attribute was invalid.

section address 0xn timer exceeds near data range.

section address must be even.

section address must be in range [0..0x7ffffe].

The argument to an `address()` section attribute was invalid.

Symbol 'name' can not be both weak and common.

Both the `.weak` directive and `.comm` directive were used on the same symbol within the same source file.

syntax error in .startof. or .sizeof.

The assembler found either `.startof.` or `.sizeof.`, but did not find the beginning parenthesis '(' or ending parenthesis ')'. See [Section 3.9.6 "Obtaining the Size of a Specific Section."](#) and [Section 3.9.7 "Obtaining the Starting Address of a Specific Section."](#) for details on the `.startof.` and `.sizeof.` operators.

T

This expression is not a valid GOTO target.

The assembler does not support expressions that include unresolved symbols as a GOTO destination.

Too few operands ('insn').

Too few operands were specified for this instruction.

Too many operands ('insn').

Too many operands were specified for this instruction.

U

unexpected end of file in irp or irpc.

The end of the file was seen before the terminating `.endr` directive.

unexpected end of file in macro definition.

The end of the file was seen before the terminating `.endm` directive.

Unknown pseudo-op: 'directive'.

The assembler does not recognize the specified directive. Check to see that you have spelled the directive correctly.

Note: the assembler expects that anything that is preceded by a dot (.) is a directive.

W**WAR hazard detected.**

The assembler found a Write After Read hazard in the instruction. A WAR hazard occurs when a common W register is used for both the source and destination given that the source register uses pre/post-increment/decrement.

Word operations expect even offset.

An attempt was made to specify [Wn+offset] or [Wn-offset] where offset is even with a word instruction.

Word operations expect an even offset between -1024 and 1022.

The offset specified in [Wn+offset] or [Wn-offset] was even, but exceeded the maximum or minimum value allowed for word instructions.

6.4 Warnings

The assembler generates warnings when an assumption is made so that the assembler could continue assembling a flawed program. Warnings should not be ignored. Each warning should be specifically looked at and corrected to ensure that the assembler understands what was intended. Warning messages can sometimes point out bugs in your program.

Symbol

.def pseudo-op used inside of .def/.endef: ignored.

The specified directive is not allowed within a `.def/.endef` pair. `.def/.endef` directives are used for specifying debugging information and normally are only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, note the following:

1. you want to use the `.line` directive to specify the line number information for the symbol, and
2. you cannot nest `.def/.endef` directives.

.dim pseudo-op used outside of .def/.endef: ignored.

The specified directive is only allowed within a `.def/.endef` pair. These directives are used to specify debugging information and normally are only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, you must first specify a `.def` directive before specifying this directive.

.endef pseudo-op used outside of .def/.endef: ignored.

The specified directive is only allowed within a `.def/.endef` pair. These directives are used to specify debugging information and normally are only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, you must first specify a `.def` directive before specifying this directive.

.fill size clamped to 8.

The size argument (second argument) of the `.fill` directive specified was greater than eight. The maximum size allowed is eight.

.fillupper expects a constant positive byte value. 0xXX assumed.

The `.fillupper` directive was specified with an argument that is not a constant positive byte value. The last `.fillupper` value that was specified will be used.

.fillupper not specified in a code section. .fillupper ignored.

The specified directive must be specified in a code section. The assembler has seen this directive in a data section. This warning probably indicates that you forgot to change sections to a code section.

.fillvalue expects a constant positive byte value. 0xXX assumed.

The `.fillvalue` directive was specified with an argument that is not a constant positive byte value. The last `.fillvalue` value that was specified will be used.

.fillvalue not specified in a code section. .fillvalue ignored.

The specified directive must be specified in a code section. The assembler has seen this directive in a data section. This warning probably indicates that you forgot to change sections to a code section.

.In pseudo-op inside .def/.endef: ignored.

The specified directive is not allowed within a `.def/.endef` pair. `.def/.endef` directives are used for specifying debugging information and normally are only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, note the following:

1. you want to use the `.line` directive to specify the line number information for the symbol, and
2. you cannot nest `.def/.endef` directives.

.loc outside of .text.

The `.loc` directive must be specified in a `.text` section. The assembler has seen this directive in a non-`.text` section. The directive has no effect.

.loc pseudo-op inside .def/.endef: ignored.

The specified directive is not allowed within a `.def/.endef` pair. `.def/.endef` directives are used for specifying debugging information and normally are only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, note the following:

1. you want to use the `.line` directive to specify the line number information for the symbol, and
2. you cannot nest `.def/.endef` directives.

.palign not specified in a code section. .palign ignored.

The specified directive must be specified in a code section. The assembler has seen this directive in a data section. This warning probably indicates that you forgot to change sections to a code section.

.pbyte not specified in a code section. .pbyte ignored.

The specified directive must be specified in a code section. The assembler has seen this directive in a data section. This warning probably indicates that you forgot to change sections to a code section.

.pfill not specified in a code section. .pfill ignored.

The specified directive must be specified in a code section. The assembler has seen this directive in a data section. This warning probably indicates that you forgot to change sections to a code section.

.pfill size clamped to 8.

The size argument (second argument) of the `.fill` directive specified was greater than eight. The maximum size allowed is eight.

.pfillvalue expects a constant positive byte value. 0xFF assumed.

The `.pfillvalue` directive was specified with an argument that is not a constant positive byte value. The last `.pfillvalue` value that was specified will be used as if this directive did not exist.

.pfillvalue not specified in a code section. .pfillvalue ignored.

The specified directive must be specified in a code section. The assembler has seen this directive in a data section. This warning probably indicates that you forgot to change sections to a code section.

.pword not specified in a code section. .pword ignored.

The specified directive must be specified in a code section. The assembler has seen this directive in a data section. This warning probably indicates that you forgot to change sections to a code section.

.size pseudo-op used outside of .def/.endef ignored.

The specified directive is only allowed within a `.def/ .endef` pair. These directives are used to specify debugging information and normally are only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, you must first specify a `.def` directive before specifying this directive.

.scl pseudo-op used outside of .def/.endef ignored.

The specified directive is only allowed within a `.def/ .endef` pair. These directives are used to specify debugging information and normally are only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, you must first specify a `.def` directive before specifying this directive.

.tag pseudo-op used outside of .def/.endef ignored.

The specified directive is only allowed within a `.def/ .endef` pair. These directives are used to specify debugging information and normally are only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, you must first specify a `.def` directive before specifying this directive.

.type pseudo-op used outside of .def/.endef ignored.

The specified directive is only allowed within a `.def/ .endef` pair. These directives are used to specify debugging information and normally are only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, you must first specify a `.def` directive before specifying this directive.

.val pseudo-op used outside of .def/.endef ignored.

The specified directive is only allowed within a `.def/ .endef` pair. These directives are used to specify debugging information and normally are only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, you must first specify a `.def` directive before specifying this directive.

A

Alignment too large: 2¹⁵ assumed.

An alignment greater than 2¹⁵ was requested. 2¹⁵ is the largest alignment request that can be made.

B

badly formed .dim directive ignored

The arguments for the `.dim` directive were unable to be parsed. This directive is used to specify debugging information and normally is only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, the arguments for the `.dim` directive are constant integers separated by a comma.

D

Directive not specified in a code section. Directive ignored.

The directive on the indicated line must be specified in a code section. The assembler has seen this directive in a data section. This warning probably indicates that you forgot to change sections to a code section.

E

error setting flags for “section_name”: error_message.

If this warning is displayed, then the GNU code has changed as the if statement always evaluates false.

Expecting even address. Address will be rounded.

The absolute address specified for a CALL or GOTO instruction was odd. The address is rounded up. You will want to ensure that this is the intended result.

Expecting even offset. Offset will be rounded.

The PC-relative instruction at this line contained an odd offset. The offset is rounded up to ensure that the PC-relative instruction is working with even addresses.

I

Ignoring changed section attributes for section_name.

This section's attributes have already been set, and the new attributes do not match those previously set.

Ignoring fill value in absolute section.

A fill argument cannot be specified for either the `.org` or `.org` directive when the current section is absolute.

Implied attributes for section 'name' are deprecated.

Certain section names have implied attributes. In this case, a section was defined without listing its implied attributes. For clarity and future compatibility, section attributes should be listed explicitly. See [Section 4.2 “Directives that Define Sections.”](#) for more information.

L

Line numbers must be positive integers.

The line number argument of the `.ln` or `.loc` directive was less than or equal to zero after specifying debugging information for a function. These directives are used to specify debugging information and normally are only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, note that function symbols can only exist on positive line numbers.

M

Macro 'name' has a previous definition.

A macro has been redefined without removing the previous definition with the `.purgem` directive.

mismatched .eb

The assembler has seen a `.eb` directive without first seeing a matching `.bb` directive. The `.bb` and `.eb` directives are the begin block and end block directives and must always be specified in pairs.

O

Overflow/underflow for .long may lose significant bits.

A constant value specified in a `.long` directive is too large and will lose significant bits when encoded.

Q

Quoted section flags are deprecated, use attributes instead.

Previous versions of the assembler recommended the use of single character section flags. For clarity and future compatibility, attribute names should be used instead.

R

Repeat argument < 0. .fill ignored.

The repeat argument (first argument) of the `.fill` directive specified was less than zero. The repeat argument must be an integer that is greater than or equal to zero.

Repeat argument < 0. .pfill ignored.

The repeat argument (first argument) of the `.pfill` directive specified was less than zero. The repeat argument must be an integer that is greater than or equal to zero.

S

Size argument < 0. .fill ignored.

The size argument (second argument) of the `.fill` directive specified was less than zero. The size argument must be an integer that is between zero and eight, inclusive. If the size argument is greater than eight, it is deemed to have a value of eight.

Size argument < 0. .pfill ignored

The size argument (second argument) of the `.pfill` directive specified was less than zero. The size argument must be an integer that is between zero and eight, inclusive. If the size argument is greater than eight, it is deemed to have a value of eight.

‘symbol_name’ symbol without preceding function.

A `.bf` directive was seen without the preceding debugging information for the function symbol. This directive is used to specify debugging information and normally is only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, you must first `.def` the function symbol and give it a `.type` of function (`C_FCN = 101`).

T

tag not found for .tag symbol_name.

This warning should not be seen unless the assembler was unable to create the given symbol name. Check your code for errors. If you still receive this warning, contact technical support.

U

unexpected storage class sclass.

The assembler is processing the `.endef` directive and has either seen a storage class that it does not recognize or has not seen a storage class. This directive is used to specify debugging information and normally is only generated by the compiler. If you are attempting to specify debugging information for your assembly language program, you must specify a storage class using the `.scl` directive, and that storage class cannot be one of the following:

1. Undefined static (`C_USTATIC = 14`)
2. External definition (`C_EXTDEF = 5`)
3. Undefined label (`C_ULABEL = 7`)
4. Dummy entry (end of block) (`C_LASTENT = 20`)
5. Line # reformatted as symbol table entry (`C_LINE = 104`)
6. Duplicate tag (`C_ALIAS = 105`)
7. External symbol in dmert public library (`C_HIDDEN = 106`)
8. Weak symbol - GNU extension to COFF (`C_WEAKEXT = 127`)

unknown section attribute ‘flag’.

The `.section` directive does not recognize the specified section flag. Please see [Section 4.2 “Directives that Define Sections.”](#), for the supported section flags.

unsupported section attribute ‘i’.

The `.section` directive does not support the “i” section flag for COFF. Please see [Section 4.2 “Directives that Define Sections.”](#), for the supported section flags.

unsupported section attribute ‘l’.

The `.section` directive does not support the “l” section flag for COFF. Please see [Section 4.2 “Directives that Define Sections.”](#), for the supported section flags.

unsupported section attribute ‘o’.

The `.section` directive does not support the “o” section flag for COFF. Please see [Section 4.2 “Directives that Define Sections.”](#), for the supported section flags.

V

Value get truncated to use.

The fill value specified for either the `.skip`, `.pskip`, `.space`, `.pspace`, `.org` or `.porg` directive was larger than a single byte. The value has been truncated to a byte.

6.5 Messages

The assembler generates messages when a non-critical assumption is made so that the assembler could continue assembling a flawed program. Messages may be ignored. However, messages can sometimes point out bugs in your program.



Part 2 – MPLAB XC16 Object Linker

Chapter 7. Linker Overview	109
Chapter 8. Linker Command Line Options	113
Chapter 9. Linker Scripts.....	135
Chapter 10. Linker Processing	173
Chapter 11. Linker Examples	205
Chapter 12. Linker Map File	215
Chapter 13. Linker Errors/Warnings	223

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Chapter 7. Linker Overview

7.1 Introduction

MPLAB XC16 Object Linker produces binary code from relocatable object code, and any archive/library files, for the dsPIC® DSC and PIC24 MCU families of devices. The 16-bit linker is an application that provides a platform for developing executable code. The linker is a part of the GNU linker from the Free Software Foundation.

Topics covered in this chapter are:

- [Feature Set](#)
- [Linker Usage](#)
- [Input/Output Files](#)

7.2 Feature Set

Notable features of the linker include:

- Automatic or user-defined stack allocation
- Supports 16-bit Program Space Visibility (PSV) window
- Available for Windows, Linux and Mac OS
- Command Line Interface
- Linker scripts for all 16-bit devices
- Available for MPLAB® X IDE and MPLAB IDE v8

7.3 Linker Usage

The MPLAB XC16 Object Linker translates object files from the MPLAB XC16 assembler, and archive/library files from the MPLAB XC16 archiver/librarian, into an executable file. See the “*MPLAB XC16 C Compiler User's Guide*” (DS50002071) for an overview of the tools process flow.

In most instances it will not be necessary to invoke the linker directly, as the compiler driver, `xc16-gcc`, will automatically execute the linker with all necessary arguments. Using the linker directly is not simple, and should be attempted only by those with a sound knowledge of the compiler and linking in general. The compiler often makes assumptions about the way in which the program will be linked. If the linker sections are not linked correctly, code failure may result.

16-Bit Assembler, Linker and Utilities User's Guide

7.4 Input/Output Files

Linker input and output files are listed below.

Table 7-1: Linker Files

Extension	Description
Input	
.o	object file
.a	library file
.gld	linker script file
Output	
.exe, .out	binary file
.map	map file

Unlike the MPLINK linker, the MPLAB XC16 linker does not generate absolute listing files. The 16-bit linker is capable of creating a map file and a binary file (that may or may not contain debugging information).

7.4.1 Object Files

Relocatable code produced from source files. The linker accepts ELF format object files by default. To specify ELF or COFF object format explicitly, use the `-omf` command line option, as shown:

```
xc16-ld -omf=elf ...
```

Alternatively, the environment variable `XC16_OMF` may be used to specify object file format for the 16-bit language tools.

7.4.2 Library Files

A collection of object files grouped together for convenience.

7.4.3 Linker Script File

Linker scripts, or command files:

- Instruct the linker where to locate sections
- Specify memory ranges for a given part
- Can be customized to locate user-defined sections at specific addresses

For more on linker script files, see [Chapter 9. "Linker Scripts."](#)

Example 7-1: Linker Script

```
OUTPUT_FORMAT("coff-pic30")
OUTPUT_ARCH("pic30")

MEMORY
{
    data (a!xr) : ORIGIN = 0x800, LENGTH = 1024
    program (xr) : ORIGIN = 0, LENGTH = (8K * 2)
}

SECTIONS
{
    .text :
    {
        *(.vector);
        *(.handle);
        *(.text);
    } >program

    .bss (NOLOAD) :
    {
        *(.bss);
    } >data

    .data :
    {
        *(.data);
    } >data
} /* SECTIONS */

WREG0 = 0x00;
WREG1 = 0x02;
```

7.4.4 Linker Output File

By default, the name of the linker output binary file is `a.out`. You can override the default name by specifying the `-o` option on the command line. The format of the binary file is an executable ELF file by default. To specify a ELF or COFF executable file, use the `-omf` option as shown in [Section 7.4.1 “Object Files.”](#)

7.4.5 Map File

The linker has the capability to produce map files. For details on how to generate a map file and the components of that file, see [Chapter 12. “Linker Map File.”](#)

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Chapter 8. Linker Command Line Options

8.1 Introduction

MPLAB XC16 Object Linker may be used on the command line interface as well as with an IDE.

8.2 Highlights

Topics covered in this chapter are:

- [Syntax](#)
- [Options that Control Output File Creation](#)
- [Options that Control Run-time Initialization](#)
- [Options that Control Informational Output](#)
- [Options that Modify the Link Map Output](#)
- [Options that Specify CodeGuard™ Security Features](#)
- [Options that Control the Preprocessor](#)

8.3 Syntax

The linker supports many command line options, but in actual practice few of them are used in any particular context.

```
xc16-ld [options] file...
```

Note: Command line options are case sensitive.

For example, `xc16-ld` links object files and archives to produce a binary file. To link a file `hello.o`:

```
xc16-ld -o output hello.o -lpic30
```

This tells `xc16-ld` to produce a file called `output` as the result of linking the file `hello.o` with the archive `libpic30.a`.

When linking a C application, there are typically several archives (also known as “libraries”) which are included in the link command. The list of archives may be specified within `--start-group`, `--end-group` options to help resolve circular references:

```
xc16-ld -o output hello.o --start-group -lpic30 -lm -lc --end-group
```

The command line options to `xc16-ld` may be specified in any order, and may be repeated at will. Repeating most options with a different argument will either have no further effect, or override prior occurrences (i.e., those farther to the left on the command line) of that option. Options that may be meaningfully specified more than once are noted in the descriptions below.

Non-option arguments are object files that are to be linked together. They may follow, precede or be mixed in with command line options, except that an object file argument may not be placed between an option and its argument.

16-Bit Assembler, Linker and Utilities User's Guide

Usually the linker is invoked with at least one object file, but you can specify other forms of binary input files using `-l` (lowercase L) and the script command language. If no binary input files are specified, the linker does not produce any output, and issues the message 'No input files'.

If the linker cannot recognize the format of an object file, it will assume that it is a linker script. A script specified in this way augments the main linker script used for the link (either the default linker script or the one specified by using `-T`). This feature permits the linker to link against a file that appears to be an object or an archive; but, actually, merely defines some symbol values, or uses `INPUT` or `GROUP` to load other objects.

For options with names that are a single letter, option arguments must either follow the option letter without intervening white space, or be given as separate arguments immediately following the option that requires them.

For options with names that are multiple letters, either one dash or two can precede the option name; for example, `-trace-symbol` and `--trace-symbol` are equivalent. There is one exception to this rule. Multiple-letter options that begin with the letter `o` can only be preceded by two dashes.

Arguments to multiple-letter options must either be separated from the option name by an equal sign, or be given as separate arguments immediately following the option that requires them. For example, `--trace-symbol srec` and `--trace-symbol=srec` are equivalent. Unique abbreviations of the names of multiple-letter options are accepted.

8.4 Options that Control Output File Creation

Output file creation options are:

Option	'No' Option
<code>--add-flags-code=,--add-flags-data=,--add-flags-const=</code>	
<code>--application-id=name</code>	
<code>--architecture arch (-A arch)</code>	
<code>-(archives -), --start-group archives, --end-group</code>	
<code>--coresident</code>	
<code>-d, -dc, -dp</code>	
<code>--defsym sym=expr</code>	
<code>--discard-all (-x)</code>	
<code>--discard-locals (-X)</code>	
<code>--fill=option</code>	
<code>--fill-upper value</code>	
<code>--force-exe-suffix</code>	
<code>--force-link</code>	<code>--no-force-link</code>
<code>--gc-sections</code>	
<code>--isr</code>	<code>--no-isr</code>
<code>--ivt</code>	<code>--no-ivt</code>
<code>-legacy-libc</code>	
<code>--library libname (-l libname)</code>	
<code>--library-path <dir> (-L <dir>)</code>	
<code>--msecondary-id, --msecondary-id-location</code>	
<code>-mreserve</code>	
	<code>--no-keep-memory</code>
	<code>--noinhibit-exec</code>
<code>-omf=format</code>	
<code>--output file (-o file)</code>	
<code>--pad-flash=size</code>	
<code>--preserve=executable</code>	
<code>--preserve-all</code>	
<code>-p,--processor PROC</code>	
<code>--relocatable (-r, -i, -Ur)</code>	
<code>--reserve-const=size</code>	
<code>--retain-symbols-file file</code>	
<code>--script file (-T file)</code>	
<code>--select-objects</code>	<code>--no-select-objects</code>
<code>--smart-io</code>	<code>--no-smart-io</code>
<code>--strip-all (-s)</code>	
<code>--strip-debug (-S)</code>	
<code>-Tbss address</code>	
<code>-Tdata address</code>	
<code>-Ttext address</code>	
<code>--undefined symbol (-u symbol)</code>	<code>--no-undefined</code>

16-Bit Assembler, Linker and Utilities User's Guide

Option	'No' Option
<code>--wrap symbol</code>	

8.4.1 `--add-flags-code=`, `--add-flags-data=`, `--add-flags-const=`

Add section attributes to all code, data, or PSV (.const) type sections.

Each option can be repeated or have a comma-separated list of additions. The flag names are the same as you would find in the assembler documentation, with one exception: 'name()' is used to prepend a name to the currently-provided section name.

```
--add-flags-code=attr1,attr2,attr3
--add-flags-code=attr1 --add-flags-code=attr2 --add-flags-code=attr3
--add-flags-code=name(fred)
```

`=name(x)` will prepend the current section name with `x`. This allows 'sequentially' mapped sections to be mapped without having to modify the linker script, i.e., to add 'boot' to all text type sections:

```
--add-flags-code=name(remapped) --add-flags-code=boot
```

Any incompatible section will be flagged.

8.4.2 `--application-id=name`

Used for co-residency. Specify an application ID for the application code being compiled. This will define extra symbols using the application ID symbol name. For example:

```
xc16-gcc -DVERSION=1 foo.c -o foo.exe -Wl,--application-id=foo
```

where symbol `_version` would be mapped to `_foo_version`.

8.4.3 `--architecture arch (-A arch)`

Set architecture.

The architecture argument identifies the particular architecture in the 16-bit devices, enabling some safeguards and modifying the archive-library search path.

8.4.4 `-(archives -)`, `--start-group archives`, `--end-group`

Start and end a group.

The archives should be a list of archive files. They may be either explicit file names, or `-l` options. The specified archives are searched repeatedly until no new undefined references are created. Normally, an archive is searched only once in the order that it is specified on the command line. If a symbol in that archive is needed to resolve an undefined symbol referred to by an object in an archive that appears later on the command line, the linker would not be able to resolve that reference. By grouping the archives, they will all be searched repeatedly until all possible references are resolved. Using this option has a significant performance cost. It is best to use it only when there are unavoidable circular references between two or more archives.

8.4.5 `--coresident`

Co-residency Linking.

Inform the linker that a coresident link is being performed and to omit the Reset vector from a link. Used with `--no-isr`.

Example:

```
$ xc16-gcc example.c -mcpu=30f6014 -T p30F6014.gld -o
example.exe -Wl,--no-isr,--coresident
```

8.4.6 -d, -dc, -dp

Force common symbols to be defined.

Assign space to common symbols even when a relocatable output file is specified (i.e., with `-r`).

8.4.7 --defsym sym=expr

Define a symbol.

Create a global symbol in the output file that contains the absolute address given by *expr*. You may use this option as many times as necessary to define multiple symbols in the command line. A limited form of arithmetic is supported for the *expr* in this context: you may give a hexadecimal constant or the name of an existing symbol, or use `+` and `-` to add or subtract hexadecimal constants or symbols.

Note: There should be no white space between <i>sym</i> , the equals sign (" <code>=</code> ") and <i>expr</i> .

8.4.8 --discard-all (-x)

Discard all local symbols.

8.4.9 --discard-locals (-X)

Discard temporary local symbols.

8.4.10 --fill=option

Fill unused program memory. The format is:

```
--fill=[wn:]expression[@address[:end_address] | unused]
```

address and *end_address* will specify the range of program memory addresses to fill. If *end_address* is not provided, then the *expression* will be written to the specific memory location at address *address*. The optional literal value *unused* may be specified to indicate that all unused memory will be filled. If none of the location parameters are provided, all unused memory will be filled. *expression* will describe how to fill the specified memory. The following options are available:

Single value

```
xc16-ld --fill=0x12345678@unused
```

Range of values

```
xc16-ld --fill=1,2,3,4,097@0x9d000650:0x9d000750
```

Incrementing value

```
xc16-ld --fill=7+=911@unused
```

By default, the linker will fill using data that is instruction-word length. For 16-bit devices, the default fill width is 24 bits. However, you may specify the value width using `[wn:]`, where *w* is the fill value's width and *n* belongs to [1, 3].

Multiple fill options may be specified on the command line; the linker will always process fill options at specific locations first.

8.4.11 `--fill-upper value`

Set fill value for upper byte of data.

Use *value* as the upper byte (bits 16-23) when encoding data into program memory. This option affects the encoding of sections created with the `psv` or `eedata` attribute, and also the data initialization template if the `--no-pack-data` option is enabled. If this option is not specified, a default value of 0 will be used.

8.4.12 `--force-exe-suffix`

Force generation of file with `.exe` suffix.

8.4.13 `--force-link`

Force linking of objects that may not be compatible.

If a target processor has been specified with the `-p`, `--processor` option, the linker will compare it to information contained in the objects combined during the link. If a possible conflict is detected, an error (i.e., in the case of a possible instruction set incompatibility) or a warning (i.e., in the case of possible register incompatibility) will be reported. Specify this option to override such errors or warnings.

8.4.14 `--no-force-link`

Do not force linking of objects that may not be compatible (**this is the default**).

8.4.15 `--gc-sections`

Remove unused (dead) functions from code at link time.

Support is for ELF projects only. In order to make the best use of this feature, add the `-ffunction-sections` option to the compiler command line.

8.4.16 `--isr`

Create an interrupt function for unused vectors (**this is the default**).

If a function named `__DefaultInterrupt` is defined by an application, the linker will insert its address into unused slots in the primary and alternate vector tables. If this function is not defined, create a function that consists of a single reset instruction and insert the address of this function.

8.4.17 `--no-isr`

Do not create an interrupt function for vectors unused by the application.

Do not create a default interrupt function if an application does not provide one.

The unused vector slots will remain unfilled and can be defined in a future link (as in co-resident applications).

8.4.18 `--ivt`

The linker is instructed to generate an IVT or AIVT, unless one is explicitly created in the linker script or by other means.

8.4.19 `--no-ivt`

The linker is instructed *not* to generate an IVT or AIVT, unless one is explicitly created in the linker script or by other means.

8.4.20 `-legacy-libc`

Use legacy include files and libraries (those distributed with v3.24 and before).

The content of include file and libraries changed in v3.25 to be compatible with the HI-TECH C compiler.

8.4.21 `--library libname (-l libname)`

Search for library *libname*.

Add archive file *libname* to the list of files to link. This option may be used any number of times. `xc16-ld` will search its path-list for occurrences of `liblibname.a` for every *libname* specified. The linker will search an archive only once, at the location where it is specified on the command line. If the archive defines a symbol that was undefined in some object that appeared before the archive on the command line, the linker will include the appropriate file(s) from the archive. However, an undefined symbol in an object appearing later on the command line will not cause the linker to search the archive again. See the `-r` option for a way to force the linker to search archives multiple times. You may list the same archive multiple times on the command line.

If the format of the archive file is not recognized, the linker will ignore it. Therefore, a version mismatch between libraries and the linker may result in “undefined symbol” errors.

If file `liblibname.a` is not found, the linker will search for an `omf`-specific version of the library with name `liblibname-coff.a` or `liblibname-elf.a`.

8.4.22 `--library-path <dir> (-L <dir>)`

Add *<dir>* to library search path.

Add path *<dir>* to the list of paths that `xc16-ld` will search for archive libraries and `xc16-ld` control scripts. You may use this option any number of times. The directories are searched in the order in which they are specified on the command line. All `-L` options apply to all `-l` options, regardless of the order in which the options appear. The library paths can also be specified in a link script with the `SEARCH_DIR` command. Directories specified this way are searched at the point in which the linker script appears in the command line.

8.4.23 `--msecondary-id=id, --msecondary-id-location=loc`

`--msecondary-id=` is used to assign a unique ID to a secondary executable. This can be used when there are multiple secondary applications in a single project and the main core may dynamically select which image is programmed. Giving these applications a unique ID will allow the IDE to select the correct image during debugging. By default the linker will allocate this location at the end of memory; using `--msecondary-id-location=` allows selection of this location.

8.4.24 -mreserve

The link shell will process this option by creating a temporary linker script that will define a section to reserve a certain range of memory:

`[start_address:end_address].`

Syntax:

`-mreserve=memory_region@start_address:end_address`

The section created will have the name pattern:

`reserve_memory_region_start_address`

The address of the section will be: `start_address`

The length of the section will be: `end_address - start_address + 1`

Example:

`-mreserve=data@0x1000:0x1050`

A temporary linker script with the following content will be created:

```
SECTIONS {
reserve_data_0x1000 0x1000: {
SHORT(0x0);
. = 0x51;
}
}
```

Multiple `-mreserve` options can be passed to the linker to reserve multiple ranges.

8.4.25 --no-keep-memory

Use less memory and more disk I/O.

`xc16-ld` normally optimizes for speed over memory usage by caching the symbol tables of input files in memory. This option tells `xc16-ld` to instead optimize for memory usage by rereading the symbol tables, as necessary. This may be required if `xc16-ld` runs out of memory space while linking a large executable.

8.4.26 --noinhibit-exec

Create an output file even if errors occur.

Retain the executable output file whenever it is still usable. Normally, when the linker encounters an error during the link process, it will exit without writing an output file.

8.4.27 -omf=format

`xc16-ld` produces ELF format output binary files by default. Use this option to specify ELF or COFF format explicitly. Alternatively, the environment variable `XC16_OMF` may be used to specify object file format for the 16-bit language tools.

Note: The input and output file formats must match. The `-omf` option can be used to specify both input and output file formats.

8.4.28 --output file (-o file)

Set output file name.

Use `file` as the name for the program produced by `xc16-ld`; if this option is not specified, the name `a.out` is used by default.

8.4.29 `--pad-flash=size`

Ensure that the linked output section is padded to a size byte boundary.
Used with co-resident applications.

8.4.30 `--preserve=executable`

Use a previously-compiled executable to identify where to allocate preserved variables.

8.4.31 `--preserve-all`

Preserve all variables unless explicitly marked with update.

8.4.32 `-p,--processor PROC`

Specify the target processor (e.g., 30F2010).

Specify a target processor for the link. This information will be used to detect possible incompatibility between objects during the link. See `--force-link` for more information.

8.4.33 `--relocatable (-r, -i, -Ur)`

Generate relocatable output.

That is, generate an output file that can, in turn, serve as input to `xc16-ld`. This is often called partial linking. If this option is not specified, an absolute file is produced.

8.4.34 `--reserve-const=size`

Reserve the specified amount of const data (size). If no value is specified, the maximum is reserved.

Used with co-resident applications.

8.4.35 `--retain-symbols-file file`

Keep only symbols listed in *file*.

Retain only the symbols listed in the file *file*, discarding all others. *file* is simply a flat file, with one symbol name per line. This option is especially useful in environments where a large global symbol table is accumulated gradually, to conserve run-time memory. `--retain-symbols-file` does not discard undefined symbols, or symbols needed for relocations. You may only specify `--retain-symbols-file` once in the command line. It overrides `-s` and `-S`.

8.4.36 `--script file (-T file)`

Read linker script.

Read link commands from the file *file*. These commands replace the default link script of `xc16-ld` (rather than adding to it), so *file* must specify everything necessary to describe the target format. If *file* does not exist, `xc16-ld` looks for it in the directories specified by any preceding `-L` options. Multiple `-T` options accumulate.

8.4.37 `--select-objects`

Select library objects based on options (**this is the default**).

Some compiler options, such as `-mlarge-arrays`, must be set consistently across all objects in an application. In order to maintain full compatibility, pre-compiled libraries must contain multiple versions of each object. Library objects are selected based on a signature which is created by the compiler and reflects the options used to create the object. Objects from older libraries that lack a signature are considered to be compatible if the restrictive compiler options have not been set.

8.4.38 `--no-select-objects`

Don't select library objects based on options.

This option causes the linker to load the first instance of a library object, regardless of the options signature. This option can be used to force library compatibility with restrictive compiler options, even if the library lacks a signature.

8.4.39 `--smart-io`

Merge I/O library functions when possible (**this is the default**).

Several I/O functions in the standard C library exist in multiple versions. For example, there are separate output conversion functions for integers, short doubles and long doubles. If this option is enabled, the linker will merge function calls to reduce memory usage whenever possible. Library function merging will not result in a loss of functionality.

8.4.40 `--no-smart-io`

Don't merge I/O library functions.

Do not attempt to conserve memory by merging I/O library function calls. In some instances, the use of this option will increase memory usage.

8.4.41 `--strip-all (-s)`

Strip all symbols.

Omit all symbol information from the output file.

8.4.42 `--strip-debug (-S)`

Strip debugging symbols.

Omit debugger symbol information (but not all symbols) from the output file.

8.4.43 `-Tbss address`

Set address of `.bss` section.

Use *address* as the starting address for the `bss` segment of the output file. *address* must be a single hexadecimal integer. For compatibility with other linkers, you may omit the leading `0x` usually associated with hexadecimal values.

Normally the address of this section is specified in a linker script.

8.4.44 `-Tdata address`

Set address of `.data` section.

Use *address* as the starting address for the `data` segment of the output file. *address* must be a single hexadecimal integer. For compatibility with other linkers, you may omit the leading `0x` usually associated with hexadecimal values.

Normally the address of this section is specified in a linker script.

8.4.45 `-Ttext address`

Set address of `.text` section.

Use *address* as the starting address for the `text` segment of the output file. *address* must be a single hexadecimal integer. For compatibility with other linkers, you may omit the leading `0x` usually associated with hexadecimal values.

Normally the address of this section is specified in a linker script.

8.4.46 --undefined symbol (-u symbol)

Start with undefined reference to *symbol*.

Force *symbol* to be entered into the output file as an undefined symbol. Doing this may, for example, trigger linking of additional modules from standard libraries. `-u` may be repeated with different option arguments to enter additional undefined symbols.

8.4.47 --no-undefined

Allow no undefined symbols.

8.4.48 --wrap symbol

Use wrapper functions for *symbol*.

Any undefined reference to *symbol* will be resolved to `_wrap__symbol`. Any undefined reference to `_real__symbol` will be resolved to *symbol*. This can be used to provide a wrapper for a system function. The wrapper function should be called `_wrap__symbol`. If it wishes to call the system function, it should call `_real__symbol`.

Here is an example:

```
#include <stdio.h>

#ifdef WRAP
#define WRAPIT2(X) _wrap__ ## X
#define WRAPIT(X) WRAPIT2(X)
#define REAL2(X) _real__ ## X
#define REAL(X) REAL2(X)
#else
#define WRAPIT(X) X
#define REAL(X) X
#endif

int REAL(foo)(int x) {

    /* this is the real function */
    /* we if we want to wrap it, then this is called _real__foo (or
    __real__foo from the linkers perspective
    */
    return x +1;
}

#ifdef WRAP
/* wrap it and call the real one */
int WRAPIT(foo)(int x) {
    return REAL(foo)(x)+1;
}
#endif

main() {
    fprintf(stderr,"foo returns %d\n", foo(1));
}
```

The following will wrap `foo` and make it call `__wrapfoo`; `__real__foo` will call the real `foo` function. The output will be 3.

```
xc16-gcc wrap.c -save-temps -DWRAP -Wl,--wrap,_foo
```

16-Bit Assembler, Linker and Utilities User's Guide

The following will not wrap `foo`, and make the 'real' `foo` be simply `foo`. The output will be 2.

```
xc16-gcc wrap.c -save-temps
```

The `--wrap` option can be used to intercept a precompiled library call. Be aware that this option requires an assembly level symbol name. For example, in order to wrap the C symbol `foo`, you must specify `--wrap _foo`, since all C symbols will be given a leading underscore by the C compiler.

8.5 Options that Control Run-time Initialization

Run-time initialization options are:

Option	'No' Option
<code>--data-init</code>	<code>--no-data-init</code>
<code>--handles</code>	<code>--no-handles</code>
<code>--heap size</code>	
<code>--local-stack</code>	<code>--no-local-stack</code>
<code>--pack-data</code>	<code>--no-pack-data</code>
<code>--stack size</code>	
<code>--stackguard size</code>	

8.5.1 `--data-init`

Support initialized data (**this is the default**).

Create a special output section named `.dinit` as a template for the run-time initialization of data. The C start-up module in `libpic30.a` interprets this template and copies initial data values into initialized data sections. Other data sections (such as `.bss`) are cleared before the `main()` function is called. Note that the persistent data section (`.pbss`) is not affected by this option.

8.5.2 `--no-data-init`

Don't support initialized data.

Suppress the template which is normally created to support run-time initialization of data. When this option is specified, the linker will select a shorter form of the C start-up module in `libpic30.a`. If the application includes data sections which require initialization, a warning message will be generated and the initial data values discarded. Storage for the data sections will be allocated as usual.

8.5.3 `--handles`

Support far code pointers (**this is the default**).

Create a special output section named `.handles` as a jump table for accessing far code pointers. Entries in the jump table are used only when the address of a code pointer exceeds 16 bits. The jump table must be loaded in the lowest range of program memory (as defined in the linker scripts).

8.5.4 `--no-handles`

Don't support far code pointers.

Suppress the handle jump table which is normally created to access far code pointers. The programmer is responsible for making certain that all code pointers can be reached with a 16 bit address. If this option is specified and the address of a code pointer exceeds 16 bits, an error is reported.

8.5.5 `--heap size`

Set heap to *size* bytes.

Allocate a run-time heap of *size* bytes for use by C programs. The heap is allocated from unused data memory. If sufficient memory is unavailable, an error is reported.

8.5.6 `--local-stack`

Prevent allocating the stack in extended data space memory (**this is the default**).

8.5.7 `--no-local-stack`

Allow allocating the stack in extended data space memory.

8.5.8 `--pack-data`

Pack initial data values (**this is the default**).

Fill the upper byte of each instruction word in the data initialization template with data. This option conserves program memory and causes the template to appear as random, and possibly invalid instructions, if viewed in the disassembler.

8.5.9 `--no-pack-data`

Don't pack initial data values.

Fill the upper byte of each instruction word in the data initialization template with 0x0 or another value specified with `--fill-upper`. This option consumes additional program memory and causes the template to appear as `NOP` instructions if viewed in the disassembler (and will be executed as such by the 16-bit device).

8.5.10 `--stack size`

Set minimum stack to *size* bytes (default=16).

By default, the linker allocates all unused data memory for the run-time stack. Alternatively, the programmer may allocate the stack by defining a section with the `stack` attribute. Use this option to ensure that at least a minimum-sized stack is available. The actual stack size is reported in the link map output file. If the minimum size is not available, an error is reported. The default minimum stack size does not include a stack guardband, as described in the next section.

8.5.11 `--stackguard size`

Set stack guardband to *size* bytes (default=16).

By default a portion of the physical stack is reserved for a guardband.

The stack guardband ensures that enough stack space is available to process a stack overflow exception. The default value (16 bytes) was chosen to handle the worst-case scenario, and guarantees that an exception handler can be invoked. This option can be used to reserve additional stack space for exception processing, or to reduce the guardband size, freeing up additional memory for the stack.

8.6 Options that Control Informational Output

Information output options are:

Option	'No' Option
<code>--check-sections</code>	<code>--no-check-sections</code>
<code>--help</code>	
<code>--memory-usage</code>	
	<code>--no-psrd-psrd-check</code>
<code>--report-mem</code>	
<code>--trace (-t)</code>	
<code>--trace-symbol symbol (-y symbol)</code>	
<code>-V</code>	
<code>--verbose</code>	
<code>--version (-v)</code>	
<code>--warn-common</code>	
	<code>--no-warn-mismatch</code>
<code>--warn-once</code>	
<code>--warn-section-align</code>	

8.6.1 `--check-sections`

Check section addresses for overlaps (**this is the default**).

8.6.2 `--no-check-sections`

Do not check section addresses for overlaps.

8.6.3 `--help`

Print option help.

Print a summary of the command line options on the standard output and exit.

8.6.4 `--memory-usage`

Specify FLASH and Data memory usage. Useful for co-resident applications.

Enables the writing of two data tables, one for FLASH and one for data memory. Each table is `NULL` terminated. Each ROW contains a pair of values in FLASH, the first is the start address of the consumed memory and the second is the last address.

The FLASH table can be accessed via the global symbol `__ROM_USAGE` and the RAM table can be accessed with the global symbol `__RAM_USAGE`. The symbols can be accessed using `tblrd` instructions or any other access method that can read the upper byte of FLASH.

8.6.5 `--no-psrd-psrd-check`

This is a linker option that can be used to disable the automatic check for PSRD PSRD violations (back-to-back data flash reads). In general, it is not recommended that this option be used.

8.6.6 `--no-warn-mismatch`

Do not warn about mismatched input files.

Normally `xc16-ld` will give an error if you try to link together input files that are mismatched for some reason, perhaps because they have been compiled for different processors or for different endiannesses. This option tells `xc16-ld` that it should silently permit such possible errors. This option should only be used with care in cases when you have taken some special action that ensures that the linker errors are inappropriate.

Note: This option does not apply to library files specified with <code>-l</code> .

8.6.7 `--report-mem`

Print a memory usage report.

Print a summary of memory usage to standard output during the link. This report also appears in the link map.

8.6.8 `--trace (-t)`

Trace file.

Print the names of the input files as `xc16-ld` processes them.

8.6.9 `--trace-symbol symbol (-y symbol)`

Trace mentions of *symbol*.

Print the name of each linked file in which *symbol* appears. This option may be given any number of times. On many systems, it is necessary to prep-end an underscore to the *symbol*. This option is useful when you have an undefined symbol in your link but do not know where the reference is coming from.

8.6.10 `-v`

Print version and other information.

8.6.11 `--verbose`

Output lots of information during link.

Display the version number for `xc16-ld`. Display the input files that can and cannot be opened. Display the linker script if using a default built-in script.

8.6.12 `--version (-v)`

Print version information.

8.6.13 --warn-common

Warn about duplicate common symbols.

Warn when a common symbol is combined with another common symbol or with a symbol definition. Unix linkers allow this somewhat sloppy practice, but linkers on some other operating systems do not. This option allows you to find potential problems from combining global symbols. Unfortunately, some C libraries use this practice, so you may get some warnings about symbols in the libraries as well as in your programs.

There are three kinds of global symbols, illustrated here with C examples:

A definition, which goes in the initialized data section of the output file.

```
int i = 1;
```

An undefined reference, which does not allocate space. There must be either a definition or a common symbol for the variable somewhere.

```
extern int i;
```

A common symbol. If there are only (one or more) common symbols for a variable, it goes in the uninitialized data area of the output file.

```
int i;
```

The linker merges multiple common symbols for the same variable into a single symbol. If they are of different sizes, it picks the largest size. The linker turns a common symbol into a declaration if there is a definition of the same variable.

The `--warn-common` option can produce five kinds of warnings. Each warning consists of a pair of lines: the first describes the symbol just encountered, and the second describes the previous symbol encountered with the same name. One or both of the two symbols will be a common symbol.

Turning a common symbol into a reference, because there is already a definition for the symbol.

```
file(section): warning: common of 'symbol' overridden by definition
file(section): warning: defined here
```

Turning a common symbol into a reference, because a later definition for the symbol is encountered. This is the same as the previous case, except that the symbols are encountered in a different order.

```
file(section): warning: definition of 'symbol' overriding common
file(section): warning: common is here
```

Merging a common symbol with a previous same-sized common symbol.

```
file(section): warning: multiple common of 'symbol'
file(section): warning: previous common is here
```

Merging a common symbol with a previous larger common symbol.

```
file(section): warning: common of 'symbol' overridden by larger common
file(section): warning: larger common is here
```

Merging a common symbol with a previous smaller common symbol. This is the same as the previous case, except that the symbols are encountered in a different order.

```
file(section): warning: common of 'symbol' overriding smaller common
file(section): warning: smaller common is here
```

8.6.14 --warn-once

Warn only once per undefined symbol.

Only warn once for each undefined symbol, rather than once per module that refers to it.

8.6.15 `--warn-section-align`

Warn if start of section changes due to alignment.

Warn if the address of an output section is changed because of alignment. This means a gap has been introduced into the (normally sequential) allocation of memory.

Typically, an input section will set the alignment. The address will only be changed if it is not explicitly specified; that is, if the `SECTIONS` command does not specify a start address for the section.

8.7 Options that Modify the Link Map Output

Link map output modifying options are:

```
--cref
--print-map (-M)
-Map file
```

8.7.1 `--cref`

Output cross-reference table.

If a linker map file is being generated, the cross-reference table is printed to the map file. Otherwise, it is printed on the standard output. The format of the table is intentionally simple, so that a script may easily process it, if necessary. The symbols are printed out, sorted by name. For each symbol, a list of file names is given. If the symbol is defined, the first file listed is the location of the definition. The remaining files contain references to the symbol.

8.7.2 `--print-map (-M)`

Print map file on standard output.

Print a link map to the standard output. A link map provides information about the link, including the following:

Where object files and symbols are mapped into memory.

How common symbols are allocated.

All archive members included in the link, with a mention of the symbol which caused the archive member to be brought in.

8.7.3 `-Map file`

Write a map file.

Print a link map to the *file*. See the description of the `--print-map (-M)` option.

8.8 Options that Specify CodeGuard™ Security Features

Three linker options are related to CodeGuard Security:

- boot *LIST* – Specify options for the boot segment
- secure *LIST* – Specify options for the secure segment
- general *LIST* – Specify options for the general segment

LIST may include a single segment option or several segment options separated by colons. Multiple instances of *boot*, *secure*, or *general* options are accepted and will be combined. An optional equals sign (=) may precede *LIST*.

8.8.1 CodeGuard Security Segment Options

The following segment options correspond to specific CodeGuard Security settings as described in the CodeGuard Security documentation. The linker will validate that any CodeGuard Security option(s) specified are supported by the target device. An error will be reported if the target device does not support a particular option. Valid options settings will be encoded as configuration words for the target device.

For MPLAB X IDE, these options will appear in the Project Properties window under xc16-ld options. For MPLAB IDE v8, these options will appear in the Build Options dialog as a sub-category of the **XC16 LINK** tab. They will be passed to the linker via command line.

Table 8-1: CodeGuard™ Security Segment Options

Option	Segment(s) Supported		
	boot	secure	general
no_ram **	X	X	
small_ram	X	X	
medium_ram	X	X	
large_ram	X	X	
no_flash **	X	X	
small_flash_std	X	X	
medium_flash_std	X	X	
large_flash_std	X	X	
small_flash_high	X	X	
medium_flash_high	X	X	
large_flash_high	X	X	
no_eeprom **	X	X	
eeprom	X		
small_eeprom		X	
medium_eeprom		X	
large_eeprom		X	
no_write_protect **	X	X	X
write_protect	X	X	X
no_code_protect **			X
code_protect_std			X
code_protect_high			X

** default setting

Example 8-1: CodeGuard Security segment options

```
--boot small_flash_std
--boot=small_ram:medium_flash_std:eeprom

--secure no_ram:small_flash_std
--secure=medium_ram:large_flash_high

--general write_protect
--general=no_write_protect:code_protect_high
```

8.8.2 User-Defined Segment Options

The following segment options are supported for any device. They enable the programmer to take advantage of special language features created for CodeGuard Security, including separately linked application segments and access entry branch tables. These options do not require CodeGuard Security support in hardware and will not be encoded as configuration word settings.

Note: User-defined segment options should not be combined with CodeGuard Security options. They are intended for debugging and/or special bootloader applications.

Table 8-2: User-Defined Segment Options

Option	Segment(s) Supported		
	boot	secure	general
ram_size=nn	X	X	
flash_size=nn	X	X	

nn is a positive integer in decimal or hex format

Example 8-2: User-Defined segment options

```
--boot flash_size=128
--boot=ram_size=64:flash_size=256

--secure flash_size=256
--secure=ram_size=64:flash_size=256
```

8.9 Options that Control the Preprocessor

Linker scripts are passed to the C preprocessor before actual linking begins. This provides an opportunity to substitute macro definitions and to include conditional blocks of code. The C preprocessor is well-known by programmers and documentation is widely available.

Linker preprocessor options are listed in the sections below.

```
-D<macro> [=value]  
--no-cpp  
--save-gld
```

8.9.1 -D<macro> [=value]

Define a macro (with optional value) to the preprocessor.

Macros can be used to substitute literal values into a script, such as for the origin or length of memory regions. They can also be used to select conditional blocks of code using directives such as `#ifdef`, `#endif`.

8.9.2 --no-cpp

Do not preprocess linker scripts.

Linker script preprocessing is enabled by default. This option can be used to disable preprocessing.

Care should be used in selecting this option. If a linker script requires preprocessing (such as for conditional blocks of text), using this option will cause a processing error.

8.9.3 --save-gld

Save preprocessed linker scripts.

By default the result of preprocessing is a temporary file. This option can be used to save the preprocessed linker script. A filename is automatically generated based on the linker script filename.

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Chapter 9. Linker Scripts

9.1 Introduction

Linker scripts are used to control MPLAB XC16 Object Linker functions. You can customize your linker script for specialized control of the linker.

9.2 Highlights

Topics covered in this chapter are:

- [Overview of Linker Scripts](#)
- [Command Line Information](#)
- [Contents of a Linker Script](#)
- [Creating a Custom Linker Script](#)
- [Linker Script Command Language](#)
- [Expressions in Linker Scripts](#)

9.3 Overview of Linker Scripts

Linker scripts control all aspects of the link process, including:

- allocation of data memory and program memory
- mapping of sections from input files into the output file
- construction of special data structures (such as interrupt vector tables)
- assignment of absolute SFR addresses for the target device

9.3.1 Contents

Linker scripts are text files that contain a series of commands. Each command is either a keyword, possibly followed by arguments, or an assignment to a symbol. Comments may be included just as in C, delimited by `/*` and `*/`. As in C, comments are syntactically equivalent to white space. Unlike C, white space is significant and is often not permitted between syntax elements.

9.3.2 File Names and Locations

The 16-bit Language Tools include a set of standard linker scripts: device-specific linker scripts (e.g., `p30f3014.gld`) and one generic linker script (`p30sim.gld`).

Standard linker script files are provided for each device and are located under:

`Install_Dir/support/DeviceFamily/gld`

where `Install_Dir` is the installation directory for the MPLAB XC16 C compiler and `DeviceFamily` is the name of the device family (e.g., `dsPIC33EP`) or generic.

16-Bit Assembler, Linker and Utilities User's Guide

9.4 Command Line Information

Linker scripts are specified on the command line using either the `-T` option or the `--script` option (see [Section 8.4 “Options that Control Output File Creation.”](#)):

```
xc16-ld -o output.cof output.o --script
    ..\support\dsPIC30F\gld\p30f3014.gld
```

If the linker is invoked through `xc16-gcc`, add the `-Wl,` prefix to allow the option to be passed to the linker:

```
xc16-gcc -o output.cof output.s -Wl,--script,
    ..\support\dsPIC30F\gld\p30f3014.gld
```

If no linker script is specified, the linker will use an internal version known as the default linker script. The default linker script has memory range information and SFR definitions that are appropriate for the command line simulator (mdb). The default linker script can be examined by invoking the linker with the `--verbose` option:

```
xc16-ld --verbose
```

<p>Note: The default linker script is functionally equivalent to the generic linker script <code>p30sim.gld</code>.</p>
--

Linker scripts are located by using the library search path which, by default, includes the standard directories provided with the install.

9.5 Contents of a Linker Script

In the next several sections, a device-specific linker script for the dsPIC30F3014 will be examined. The linker script contains the following categories of information:

- [Processor and Startup Modules](#)
- [Memory Region Information](#)
- [Base Memory Addresses](#)
- [Input/Output Section Map](#)
- [Interrupt Vector Tables](#)
- [SFR Addresses](#)

9.5.1 Processor and Startup Modules

The first several lines of a linker script define the processor and startup modules:

```
/*
** Linker Script for 30f3014
*/

OUTPUT_ARCH("30F3014")
CRT0_STARTUP(crt0_standard.o)
CRT1_STARTUP(crt1_standard.o)

OPTIONAL(-lp30F3014)
```

The `OUTPUT_ARCH` command specifies the target processor. The `CRTn_STARTUP` commands specify two C run-time startup modules to be loaded from archives. The linker will select one of these based on whether data initialization has been enabled. The `OPTIONAL` command specifies a device-specific library that should be opened if available. If the library file cannot be found, the link will continue without error unless there are unresolved references in the application.

9.5.2 Memory Region Information

The next section of a linker script defines the various memory regions for the target device using the `MEMORY` command.

For the dsPIC30F3014, several memory regions are defined:

```
/*
** Memory Regions
*/
MEMORY
{
    data          : ORIGIN = 0x800,      LENGTH = 2048
    program       : ORIGIN = 0x100,      LENGTH = ((8K * 2) - 0x100)
    reset         : ORIGIN = 0,          LENGTH = (4)
    ivt           : ORIGIN = 0x04,       LENGTH = (62 * 2)
    aivt          : ORIGIN = 0x84,       LENGTH = (62 * 2)
    __FOSC        : ORIGIN = 0xF80000,   LENGTH = (2)
    __FWDTC       : ORIGIN = 0xF80002,   LENGTH = (2)
    __FBORPOR     : ORIGIN = 0xF80004,   LENGTH = (2)
    __CONFIG4     : ORIGIN = 0xF80006,   LENGTH = (2)
    __CONFIG5     : ORIGIN = 0xF80008,   LENGTH = (2)
    __FGS         : ORIGIN = 0xF8000A,   LENGTH = (2)
    __FUID0       : ORIGIN = 0x8005C0,   LENGTH = (2)
    __FUID1       : ORIGIN = 0x8005C2,   LENGTH = (2)
    __FUID2       : ORIGIN = 0x8005C4,   LENGTH = (2)
    __FUID3       : ORIGIN = 0x8005C6,   LENGTH = (2)
    eedata        : ORIGIN = 0x7FFC00,   LENGTH = (1024)
}
```

Each memory region is range-checked as sections are added during the link process. If any region overflows, a link error is reported.

`MEMORY` regions are:

- [Data Region](#)
- [Program Region](#)
- [Reset, Ivt and Aivt Regions](#)
- [Fuse Configuration Regions](#)
- [Unit ID Regions](#)
- [EEDATA Memory Region](#)

9.5.2.1 Data Region

```
data          : ORIGIN = 0x800,      LENGTH = 2048
```

The data region corresponds to the RAM memory of the dsPIC30F3014 device, and is used for both initialized and uninitialized variables. The starting address of region `data` is `0x800`. This is the first usable location in RAM, after the space reserved for memory-mapped SFRs.

16-Bit Assembler, Linker and Utilities User's Guide

9.5.2.2 Program Region

```
program      : ORIGIN = 0x100, LENGTH = ((8K * 2) - 0x100)
```

The program region corresponds to the Flash memory of the dsPIC30F3014 device that is available for user code, library code and constants. The starting address of region program is 0x100. This is the first location in Flash that is available for general use. Addresses below 0x100 are reserved for the Reset instruction and the two vector tables.

The length specification of the program region deserves particular emphasis. The (8K * 2) portion indicates that the dsPIC30F3014 has 8K instruction words of Flash memory, and that each instruction word is 2 address units wide. The - 0x100 portion reflects the fact that some of the Flash is reserved for the Reset instruction and vector tables.

Note: Instruction words in the 16-bit devices are 24 bits, or 3 bytes, wide. However, the PC increments by 2 for each instruction word for compatibility with data memory. Address and lengths in program memory are expressed in PC units.

9.5.2.3 Reset, Ivt and Aivt Regions

```
reset        : ORIGIN = 0,          LENGTH = (4)
```

The Reset region corresponds to the 16-bit Reset instruction at address 0 in program memory. The Reset region is 4 address units, or 2 instruction words, long. This region always contains a GOTO instruction that is executed upon device reset. The GOTO instruction is encoded by data commands in the section map (see Section 9.5.4.1).

```
ivt          : ORIGIN = 0x04,      LENGTH = (62 * 2)
aivt         : ORIGIN = 0x84,      LENGTH = (62 * 2)
```

The `ivt` and `aivt` regions correspond to the interrupt vector table and alternate interrupt vector table, respectively. Each interrupt vector table contains 62 entries, each is 2 address units in length. Each entry represents a word of program memory, which contains a 24-bit address. The linker initializes the vector tables with appropriate data, according to standard naming conventions.

Regions `reset`, `ivt` and `aivt` comprise the low address portion of Flash memory that is not available for user programs.

9.5.2.4 Fuse Configuration Regions

```
__FOSC       : ORIGIN = 0xF80000, LENGTH = (2)
__FWDTC      : ORIGIN = 0xF80002, LENGTH = (2)
__FBORPOR    : ORIGIN = 0xF80004, LENGTH = (2)
__CONFIG4    : ORIGIN = 0xF80006, LENGTH = (2)
__CONFIG5    : ORIGIN = 0xF80008, LENGTH = (2)
__FGS        : ORIGIN = 0xF8000A, LENGTH = (2)
```

These regions correspond to the dsPIC30F3014 configuration registers.

Each fuse configuration region is exactly one instruction word long. If sections are defined in the application source code with the standard naming convention, the section contents will be written into the appropriate configuration register(s). Otherwise, the registers are left uninitialized. If more than one value is defined for any configuration region, a link error will be reported.

9.5.2.5 Unit ID Regions

```
__FUID0      : ORIGIN = 0x8005C0, LENGTH = (2)
__FUID1      : ORIGIN = 0x8005C2, LENGTH = (2)
__FUID2      : ORIGIN = 0x8005C4, LENGTH = (2)
__FUID3      : ORIGIN = 0x8005C6, LENGTH = (2)
```

The unit ID regions correspond to locations in program memory that may be programmed with application-specific information.

9.5.2.6 EEDATA Memory Region

```
eedata       : ORIGIN = 0x7FFC00, LENGTH = (1024)
```

The `eedata` region corresponds to non-volatile data flash memory located in high memory. Although located in program memory space, the data flash is organized like data memory. The total length is 1024 bytes.

9.5.3 Base Memory Addresses

This portion of the linker script defines the base addresses of several output sections in the application. Each base address is defined as a symbol with the following syntax:

```
name = value;
```

The symbols are used to specify load addresses in the section map. For the dsPIC30F3014, several base memory addresses are defined. Not all of these symbols are referenced in the section map; some are included for informational purposes.

```
/*
** Base Memory Addresses - Program Memory
*/
__RESET_BASE = 0;          /* Reset Instruction */
__IVT_BASE   = 0x04;       /* Interrupt Vector Table */
__AIVT_BASE  = 0x84;       /* Alternate Interrupt Vector Table */
__CODE_BASE  = 0x100;      /* Handles, User Code, Library Code */

/*
** Base Memory Addresses - Data Memory
*/
__SFR_BASE   = 0;          /* Memory-mapped SFRs */
__DATA_BASE  = 0x800;      /* X and General Purpose Data Memory */
__YDATA_BASE = 0x0C00;     /* Y Data Memory for DSP Instructions */
```

9.5.4 Input/Output Section Map

The section map is the heart of the linker script. It defines how input sections are mapped to output sections. Note that input sections are portions of an application that are defined in source code, while output sections are created by the linker. Generally, several input sections may be combined into a single output section.

For example, suppose that an application is comprised of five different functions, and each function is defined in a separate source file. Together, these source files will produce five input sections. The linker will combine these input sections into a single output section. Only the output section has an absolute address.

If any input or output sections are empty, there is no penalty or storage cost for the linked application. Most applications will use only a few of the many sections that appear in the section map.

- [Output Section .reset](#)
- [Output Section .text](#)
- [User-Defined Section in Program Memory](#)
- [Output Sections in Configuration Memory](#)
- [User-Defined Section in Data Flash Memory](#)
- [In-Circuit Debugger Memory](#)
- [User-Defined Section in Data Memory](#)

9.5.4.1 Output Section `.reset`

Section `.reset` contains a `GOTO` instruction, created at link time, from output section data commands:

```
/*
** Reset Instruction
*/
.reset __RESET_BASE :
{
    SHORT(ABSOLUTE(__reset));
    SHORT(0x04);
    SHORT((ABSOLUTE(__reset) >> 16) & 0x7F);
    SHORT(0);
} >reset
```

Each `SHORT()` data command causes a 2 byte value to be included. There are two expressions which include the symbol `__reset`, which by convention is the first function invoked after a device reset. Each expression calculates a portion of the address of the Reset function. These declarations encode a 24-bit `GOTO` instruction, which is two instruction words long.

The `ABSOLUTE()` function specifies the final value of a program symbol after linking. If this function were omitted, a relative (before-linking) value of the program symbol would be used.

The `>reset` portion of this definition indicates that this section should be allocated in the Reset memory region.

9.5.4.2 Output Section `.text`

Section `.text` collects executable code from all of the application's input files.

```
/*
** User Code and Library Code
*/
.text :
{
    *(.init);
    *(.user_init);
    keep(*(.handle));
    keep(*(.isr));
    *(.libc) *(.libm) *(.libdsp); /* keep together in this order */
    *(.lib*);
} >program
```

Several different input sections are collected into one output section. This was done to ensure the order in which the input sections are loaded.

Table 9-1: Section Types and Names

Section Type	Section Name	Description
input	<code>.init</code>	Contains the startup code that is executed immediately after device reset. It is positioned first so that its address may be readily available.
input	<code>.user_init</code>	Contains a call table for user initialization functions.
input	<code>.handle</code>	Used for function pointers and is loaded first at low addresses. <code>keep</code> is required to prevent <code>-gcc-sections</code> from deleting this code.
input	<code>.isr</code>	Used for interrupt service functions. Again, <code>keep</code> is used to preserve the code.
library	<code>.libc</code> <code>.libm</code> <code>.libdsp</code>	These sections must be grouped together to ensure locality of reference.
library	<code>.lib*</code>	Collects other libraries, such as the peripheral libraries (which are allocated in section <code>.libperi</code>).

The input section `.text` is not explicitly mapped so that the linker may distribute code around PSV sections in order to more successfully satisfy PSV address requirements.

9.5.4.3 User-Defined Section in Program Memory

A stub is included for user-defined output sections in program memory. This stub may be edited as needed to support the application requirements. Once a standard linker script has been modified, it is called a “custom linker script.” In practice, it is often simpler to use section attributes in source code to locate user-defined sections in program memory. See [Chapter 11. “Linker Examples.”](#) for more information.

```
/*
** User-Defined Section in Program Memory
**
** note: can specify an address using
**       the following syntax:
**
**       usercode 0x1234 :
**           {
**               *(usercode);
**           } >program
**/
usercode :
{
    *(usercode);
} >program
```

An exact, absolute starting address can be specified, if necessary. If the address is greater than the current location counter, the intervening memory space will be skipped and filled with zeros. If the address is less than the current location counter, a section overlap will occur. Whenever two output sections occupy the same address range, a link error will be reported. Overlapping sections in program memory can not be supported.

Note: Each memory region has its own location counter.

9.5.4.4 User-Defined Constants in Program Memory

A comment block is included that describes how to define sections that will be accessed via the PSV window or the EDS window. Such sections are defined with the `psv` attribute. The syntax used to represent a PSV section address is different from other type sections. In particular, the Load Memory Address (LMA) should be defined, not the Virtual Memory Address (VMA). The LMA is unique and describes where the section is located in program memory. The VMA describes a location in the data window that may be shared by multiple pages of program memory, and is therefore not unique.

```
/*
** User-Defined Constants in Program Memory
**
** For PSV-type sections, the Load Memory Address (LMA)
** should be specified as follows:
**
**      userconst : AT(0x1234)
**      {
**          *(userconst);
**      } >program
**
** Note that mapping PSV sections in linker scripts
** is not generally recommended.
**
** Because of page alignment restrictions, memory is
** often used more efficiently when PSV sections
** do not appear in the linker script.
**
** For more information on memory allocation,
** please refer to chapter 10, "Linker Processing"
** in the Assembler, Linker manual (DS50001317).
**
*/
```

As noted, defining PSV-type sections in the linker script is not generally recommended. This is because sections that appear in the linker script are allocated sequentially, and PSV sections have significant page alignment restrictions. For more information on memory allocation and PSV sections, see [Chapter 10. "Linker Processing."](#)

16-Bit Assembler, Linker and Utilities User's Guide

9.5.4.5 Output Sections in Configuration Memory

Several sections are defined that match the Fuse Configuration memory regions:

```
/*
** Configuration Fuses
*/
__FOSC :
{ *(__FOSC.sec)      } >__FOSC
__FWDT :
{ *(__FWDT.sec)      } >__FWDT
__FBORPOR :
{ *(__FBORPOR.sec)   } >__FBORPOR
__CONFIG4 :
{ *(__CONFIG4.sec)   } >__CONFIG4
__CONFIG5 :
{ *(__CONFIG5.sec)   } >__CONFIG5
__FGS :
{ *(__FGS.sec)       } >__FGS
__FICD :
{ *(__FICD.sec)      } >__FICD
__FUID0 :
{ *(__FUID0.sec)     } >__FUID0
__FUID1 :
{ *(__FUID1.sec)     } >__FUID1
__FUID2 :
{ *(__FUID2.sec)     } >__FUID2
__FUID3 :
{ *(__FUID3.sec)     } >__FUID3
```

The Configuration Fuse sections are supported by macros defined in the 16-bit device-specific include files in `support/inc` and the C header files in `support/h`.

For example, to disable the Watchdog Timer in assembly language:

```
.include "p30f6014.inc"
config __FWDT, WDT_OFF
```

The equivalent operation in C would be:

```
#include "p30f6014.h"
_FWDT(WDT_OFF);
```

Configuration macros have the effect of changing the current section. In C, the macro should be used outside of any function. In assembly language, the macro should be followed by a `.section` directive.

9.5.4.6 User-Defined Section in Data Flash Memory

A stub is included for user-defined output sections in EEDATA memory. This stub may be edited as needed to support the application requirements. Once a standard linker script has been modified, it is called a “custom linker script.” In practice, it is often simpler to use section attributes in source code to locate user-defined sections in data flash memory. See [Chapter 11. “Linker Examples.”](#) for more information.

```
/*
** User-Defined Section in Data Flash Memory
**
** note: can specify an address using
**       the following syntax:
**
**       eedata 0x7FF100 :
**       {
**           *(eedata);
**       } >eedata
**/
eedata :
{
    *(eedata);
} >eedata
```

An exact, absolute starting address can be specified, if necessary. If the address is greater than the current location counter, the intervening memory will be skipped and filled with zeros. If the address is less than the current location counter, a section overlap will occur. Whenever two output sections occupy the same address range, a link error will be reported. Overlapping sections in EEDATA memory can not be supported.

Note: Each memory region has its own location counter.

9.5.4.7 In-Circuit Debugger Memory

An in-circuit debugger/emulator requires a portion of data memory for its variables and stack. Since the debugger is linked separately and in advance of user applications, the block of memory must be located at a fixed address and dedicated for use by the debugger.

```
/*
** ICD Debug Exec
**
** This section provides optional storage for
** the in-circuit debugger. Define a global symbol
** named __ICD2RAM to enable the debugger. This section
** must be loaded at data address 0x800.
**/
.icd __DATA_BASE (NOLOAD):
{
    . += (DEFINED (__ICD2RAM) ? 0x50 : 0 );
} > data
```

Section `.icd` is designed to optionally reserve memory for the in-circuit debugger/emulator. If global symbol `__ICD2RAM` is defined at link time, 0x50 bytes of memory at address 0x800 will be reserved. The `(NOLOAD)` attribute indicates that no initial values need to be loaded for this section. The name for this symbol was created when there was only one in-circuit debugger, the MPLAB ICD 2.

9.5.4.8 User-Defined Section in Data Memory

A stub is included for user-defined output sections in data memory. This stub may be edited as needed to support the application requirements. Once a standard linker script has been modified, it is called a “custom linker script.” In practice, it is often simpler to use section attributes in source code to locate user-defined sections in data memory. See [Chapter 11. “Linker Examples.”](#) for more information.

```
/*
** User-Defined Section in Data Memory
**
** note: can specify an address using
**       the following syntax:
**
**       userdata 0x1234 :
**       {
**           *(userdata);
**       } >data
**/
userdata :
{
    *(userdata);
} >data
```

An exact, absolute starting address can be specified, if necessary. If the address is greater than the current location counter, the intervening memory space will be skipped and filled with zeros. If the address is less than the current location counter, a section overlap will occur. Whenever two output sections occupy the same address range, a link error will be reported. Overlapping sections in data memory cannot be supported.

9.5.5 Interrupt Vector Tables

The primary and alternate interrupt vector tables are defined in a second section map, near the end of the standard linker script:

```
/*
** Section Map for Interrupt Vector Tables
*/
SECTIONS
{

/*
** Primary Interrupt Vector Table
*/
.ivt __IVT_BASE :
{
    LONG(DEFINED(__ReservedTrap0) ? ABSOLUTE(__ReservedTrap0) :
        ABSOLUTE(__DefaultInterrupt));
    LONG(DEFINED(__OscillatorFail) ? ABSOLUTE(__OscillatorFail) :
        ABSOLUTE(__DefaultInterrupt));
    LONG(DEFINED(__AddressError) ? ABSOLUTE(__AddressError) :
        ABSOLUTE(__DefaultInterrupt));
    :
    :
    LONG(DEFINED(__Interrupt53) ? ABSOLUTE(__Interrupt53) :
        ABSOLUTE(__DefaultInterrupt));
} >ivt
```

The vector table is defined as a series of `LONG()` data commands. Each vector table entry is 4 bytes in length (3 bytes for a program memory address plus an unused phantom byte). The data commands include an expression using the `DEFINED()` function and the `?` operator. A typical entry may be interpreted as follows:

If symbol “__OscillatorFail” is defined, insert the absolute address of that symbol. Otherwise, insert the absolute address of symbol “__DefaultInterrupt”.

By convention, a function that will be installed as the second interrupt vector should have the name `__OscillatorFail`. If such a function is included in the link, its address is loaded into the entry. If the function is not included, the address of the default interrupt handler is loaded instead. If the application has not provided a default interrupt handler (i.e., a function with the name `__DefaultInterrupt`), the linker will generate one automatically. The simplest default interrupt handler is a Reset instruction.

Note: The programmer must insure that functions installed in interrupt vector tables conform to the architectural requirements of interrupt service routines.

16-Bit Assembler, Linker and Utilities User's Guide

The contents of the alternate interrupt vector table are defined as follows:

```
/*
** Alternate Interrupt Vector Table
*/
.aivt __AIVT_BASE :
{
    LONG(DEFINED(__AltReservedTrap0) ? ABSOLUTE(__AltReservedTrap0)
:
    (DEFINED(__ReservedTrap0) ? ABSOLUTE(__ReservedTrap0)
:
    ABSOLUTE(__DefaultInterrupt)));
    LONG(DEFINED(__AltOscillatorFail) ? ABSOLUTE(__AltOscillatorFail)
:
    (DEFINED(__OscillatorFail) ? ABSOLUTE(__OscillatorFail)
:
    ABSOLUTE(__DefaultInterrupt)));
    LONG(DEFINED(__AltAddressError) ? ABSOLUTE(__AltAddressError)
:
    (DEFINED(__AddressError) ? ABSOLUTE(__AddressError)
:
    ABSOLUTE(__DefaultInterrupt)));
:
:
    LONG(DEFINED(__AltInterrupt53) ? ABSOLUTE(__AltInterrupt53)
:
    (DEFINED(__Interrupt53) ? ABSOLUTE(__Interrupt53)
:
    ABSOLUTE(__DefaultInterrupt)));
} >aivt
```

The syntax of the alternate interrupt vector table is similar to the primary, except for an additional expression that causes each alternate table entry to default to the corresponding primary table entry.

9.5.6 SFR Addresses

Absolute addresses for the SFRs are defined as a series of symbol definitions:

```
**=====
=
**
**          dsPIC Core Register Definitions
**
**=====*
```

```
/
WREG0 = 0x0000;
_WREG0 = 0x0000;
WREG1 = 0x0002;
_WREG1 = 0x0002;
:
:
```

Note: If identifiers in a C or assembly program are defined with the same names as SFRs, multiple definition linker errors will result.

Two versions of each SFR address are included, with and without a leading underscore. This is to enable both C and assembly language programmers to refer to the SFR using the same name. By convention, the C compiler adds a leading underscore to every identifier.

9.6 Creating a Custom Linker Script

The standard 16-bit linker scripts are general purpose and will satisfy the demands of most applications. However, occasions may arise where a custom linker script is required.

To create a custom linker script, start with a copy of the standard linker script that is appropriate for the target device. For example, to customize a linker script for the dsPIC30F3014 device, start with a copy of `p30f3014.gld`.

Customizing a standard linker script will usually involve editing sections or commands that are already present. For example, stubs for user-defined sections in both data memory and program memory are included. These stubs may be renamed and/or customized with absolute addresses if required.

It is recommended that unused sections be retained in a custom linker script, since unused sections will not impact application memory usage. If a section must be removed for a custom script, C style comments can be used to disable it.

9.7 Linker Script Command Language

Linker scripts are text files that contain a series of commands. Each command is either a keyword (possibly followed by arguments) or an assignment to a symbol. Multiple commands may be separated using semicolons. White space is generally ignored.

Strings such as file or format names can normally be entered directly. If the file name contains a character, such as a comma, which would otherwise serve to separate file names, the file name may be specified in double quotes. There is no way to use a double quote character in a file name.

Comments may be included just as in C, delimited by `/*` and `*/`. As in C, comments are syntactically equivalent to white space.

- [Basic Linker Script Concepts](#)
- [Commands Dealing with Files](#)
- [Assigning Values to Symbols](#)
- [MEMORY Command](#)
- [SECTIONS Command](#)
- [Other Linker Script Commands](#)

9.7.1 Basic Linker Script Concepts

The linker combines input files into a single output file. The output file and each input file are in a special data format known as an object file format. Each file is called an object file. Each object file has, among other things, a list of sections. A section in an input file is called an input section; similarly, a section in the output file is an output section.

Each section in an object file has a name and a size. Most sections also have an associated block of data, known as the section contents. A section may be marked as loadable, which means that the contents should be loaded into memory when the output file is run. A section with no contents may be allocatable (which means that an area in memory should be set aside), but nothing in particular should be loaded there (in some cases, this memory must be zeroed out).

Every loadable or allocatable output section has two addresses. The first is the VMA, or virtual memory address. This is the address the section will have when the output file is run. The second is the LMA, or load memory address. This is the address at which the section will be loaded. In most cases, the two addresses will be the same. An example of when they might be different is when a section is intended for use in the PSV window. In this case, the program memory address would be the LMA, and the data memory address would be the VMA.

The sections in an object file can be viewed by using the `xc16-objdump` program with the `-h` option.

Every object file also has a list of symbols, known as the symbol table. A symbol may be defined or undefined. Each symbol has a name, and each defined symbol has an address, among other information. If a C or C++ program is compiled into an object file, a defined symbol will be created for every defined function and global or static variable. Every undefined function or global variable which is referenced in the input file will become an undefined symbol.

Symbols in an object file can be viewed by using the `xc16-nm` program, or by using the `xc16-objdump` program with the `-t` option.

9.7.2 Commands Dealing with Files

Several linker script commands deal with files.

CRT0_STARTUP(object file)

This command identifies which primary startup module should be loaded from the compiler libraries. The primary startup module defines reserved symbol `__resetPRI` and is responsible for initializing the C runtime environment. Multiple versions of this module exist in order to support architectural differences between device families. Although the linker expects to find this command in every linker script, a default startup module will be selected if the command is missing (as might be the case with custom linker scripts in legacy projects.)

CRT1_STARTUP(object file)

This command identifies which alternate startup module should be loaded from the compiler libraries. The alternate startup module defines reserved symbol `__resetALT` and is responsible for initializing the C runtime environment without data initialization. Multiple versions of this module exist in order to support architectural differences between device families. Although the linker expects to find this command in every linker script, a default startup module will be selected if the command is missing (as might be the case with custom linker scripts in legacy projects.)

INCLUDE filename

Include the linker script filename at this point. The file will be searched for in the current directory, and in any directory specified with the `-L` option. Calls to `INCLUDE` may be nested up to 10 levels deep.

INPUT(file, file, ...)

INPUT(file file ...)

The `INPUT` command directs the linker to include the named files in the link, as though they were named on the command line. The linker will first try to open the file in the current directory. If it is not found, the linker will search through the archive library search path. See the description of `-L` in [Section 8.4.22 “--library-path <dir> \(-L <dir>\)”](#).

If `INPUT (-lfile)` is used, `xc16-ld` will transform the name to `libfile.a`, as with the command line argument `-l`.

When the `INPUT` command appears in an implicit linker script, the files will be included in the link at the point at which the linker script file is included. This can affect archive searching.

```
GROUP(file, file, ...)  
GROUP(file file ...)
```

The `GROUP` command is like `INPUT`, except that the named files should all be archives, and they are searched repeatedly until no new undefined references are created. See the description of `- (` in [Section 8.4.4 “-\(archives -\), --start-group archives, --end-group.”](#).

```
OPTIONAL(file, file, ...)  
OPTIONAL(file file ...)
```

The `OPTIONAL` command is analogous to the `INPUT` command, except that the named files are not required for the link to succeed. This is particularly useful for specifying archives (or libraries) that may or may not be installed with the compiler.

```
OUTPUT(filename)
```

The `OUTPUT` command names the output file. Using `OUTPUT(filename)` in the linker script is exactly like using `-o filename` on the command line (see [Section 8.4.28 “--output file \(-o file\).”](#)). If both are used, the command line option takes precedence.

```
SEARCH_DIR(path)
```

The `SEARCH_DIR` command adds `path` to the list of paths where the linker looks for archive libraries. Using `SEARCH_DIR(path)` is exactly like using `-L path` on the command line (see [Section 8.4.22 “--library-path <dir> \(-L <dir>\).”](#)). If both are used, then the linker will search both paths. Paths specified using the command line option are searched first.

```
STARTUP(filename)
```

The `STARTUP` command is just like the `INPUT` command, except that `filename` will become the first input file to be linked, as though it were specified first on the command line.

9.7.3 Assigning Values to Symbols

A value may be assigned to a symbol in a linker script. This will define the symbol as a global symbol.

- [Simple Assignments](#)
- [PROVIDE Command](#)

9.7.3.1 Simple Assignments

A symbol may be assigned using any of the C assignment operators:

```
symbol = expression ;
symbol += expression ;
symbol -= expression ;
symbol *= expression ;
symbol /= expression ;
symbol <<= expression ;
symbol >>= expression ;
symbol &= expression ;
symbol |= expression ;
```

The first case will define symbol to the value of expression. In the other cases, symbol must already be defined, and the value will be adjusted accordingly.

The special symbol name '.' indicates the location counter. This symbol may only be used within a `SECTIONS` command.

The semicolon after expression is required.

Expressions are defined in [Section 9.8 "Expressions in Linker Scripts."](#)

Symbol assignments may appear as commands in their own right, or as statements within a `SECTIONS` command, or as part of an output section description in a `SECTIONS` command.

The section of the symbol will be set from the section of the expression; for more information, see [Section 9.8.6 "The Section of an Expression."](#)

Here is an example showing the three different places that symbol assignments may be used:

```
floating_point = 0;
SECTIONS
{
    .text :
    {
        *(.text)
        _etext = .;
    }
    _bdata = (. + 3) & ~ 4;
    .data : { *(.data) }
}
```

In this example, the symbol `floating_point` will be defined as zero. The symbol `_etext` will be defined as the address following the last `.text` input section. The symbol `_bdata` will be defined as the address following the `.text` output section aligned upward to a 4-byte boundary.

9.7.3.2 PROVIDE Command

In some cases, it is desirable for a linker script to define a symbol only if it is referenced and is not defined by any object included in the link. For example, traditional linkers defined the symbol `etext`. However, ANSI C requires that `etext` may be used as a function name without encountering an error. The `PROVIDE` keyword may be used to define a symbol, such as `etext`, only if it is referenced but not defined. The syntax is `PROVIDE(symbol = expression)`.

Here is an example of using `PROVIDE` to define `etext`:

```
SECTIONS
{
    .text :
    {
        *(.text)
        _etext = .;
        PROVIDE(etext = .);
    }
}
```

In this example, if the program defines `_etext` (with a leading underscore), the linker will give a multiple definition error. If, on the other hand, the program defines `etext` (with no leading underscore), the linker will silently use the definition in the program. If the program references `etext` but does not define it, the linker will use the definition in the linker script.

9.7.4 MEMORY Command

The linker's default configuration permits allocation of all available memory. This can be overridden by using the `MEMORY` command.

The `MEMORY` command describes the location and size of blocks of memory in the target. It can be used to describe which memory regions may be used by the linker and which memory regions it must avoid. Sections may then be assigned to particular memory regions. The linker will set section addresses based on the memory regions and will warn about regions that become too full. The linker will not shuffle sections around to fit into the available regions.

The syntax of the `MEMORY` command is:

```
MEMORY
{
    name [(attr)] : ORIGIN = origin, LENGTH = len
    ...
}
```

The name is a name used in the linker script to refer to the region. The region name has no meaning outside of the linker script. Region names are stored in a separate name space, and will not conflict with symbol names, file names or section names. Each memory region must have a distinct name.

The `attr` string is an optional list of attributes associated with the memory region. Historically it was used to determine where unmapped sections should be located by the sequential memory allocator. This capability is no longer used because unmapped sections are now located by the best-fit allocator (for more information see [Section 10.5 "Linker Allocation."](#)).

The origin is an expression for the start address of the memory region. The expression must evaluate to a constant before memory allocation is performed, which means that section relative symbols may not be used. The keyword `ORIGIN` may be abbreviated to `org` or `o` (but not, for example, `ORG`).

The `len` is an expression for the size in bytes of the memory region. As with the origin expression, the expression must evaluate to a constant before memory allocation is performed. The keyword `LENGTH` may be abbreviated to `len` or `l`.

<p>Note: It is possible to use a preprocessor macro instead of a literal value for the origin and/or length of a memory region.</p>
--

Once a memory region is defined, the linker can be directed to place specific output sections into that memory region by using the `>region` output section attribute. For example, to specify a memory region named `mem`, use `>mem` in the output section definition. If no address was specified for the output section, the linker will set the address to the next available address within the memory region. If the combined output sections directed to a memory region are too large for the region, the linker will issue an error message.

9.7.5 SECTIONS Command

The `SECTIONS` command tells the linker how to map input sections into output sections and how to place the output sections in memory.

The format of the `SECTIONS` command is:

```
SECTIONS
{
    sections-command
    sections-command
    ...
}
```

Each `SECTIONS` command may be one of the following:

- an `ENTRY` command (see [Section 9.7.6 “Other Linker Script Commands.”](#))
- a symbol assignment (see [Section 9.7.3 “Assigning Values to Symbols.”](#))
- an output section description
- an overlay description

The `ENTRY` command and symbol assignments are permitted inside the `SECTIONS` command for convenience in using the location counter in those commands. This can also make the linker script easier to understand because those commands can be used at meaningful points in the layout of the output file.

Output section descriptions and overlay descriptions are described below.

If a `SECTIONS` command does not appear in the linker script, the linker will place each input section into an identically named output section in the order that the sections are first encountered in the input files. If all input sections are present in the first file, for example, the order of sections in the output file will match the order in the first input file. The first section will be at address zero.

- [Input Section Description](#)
- [Input Section Wildcard Patterns](#)
- [Input Section for Common Symbols](#)
- [Input Section Example](#)
- [Output Section Description](#)
- [Output Section Address](#)
- [Output Section Data](#)
- [Output Section Discarding](#)
- [Output Section Attributes](#)
- [Output Section LMA](#)
- [Output Section Region](#)
- [Output Section Fill](#)
- [Overlay Description](#)

9.7.5.1 Input Section Description

The most common output section command is an input section description.

The input section description is the most basic linker script operation. Output sections tell the linker how to lay out the program in memory. Input section descriptions tell the linker how to map the input files into the memory layout.

An input section description consists of a file name optionally followed by a list of section names in parentheses.

The file name and the section name may be wildcard patterns, which are described further below.

The most common input section description is to include all input sections with a particular name in the output section. For example, to include all input `.text` sections, one would write:

```
*(.text)
```

Here the `*` is a wildcard which matches any file name. To exclude a list of files from matching the file name wildcard, `EXCLUDE_FILE` may be used to match all files except the ones specified in the `EXCLUDE_FILE` list. For example:

```
*(EXCLUDE_FILE (*crtend.o *otherfile.o) .ctors)
```

will cause all `.ctors` sections from all files except `crtend.o` and `otherfile.o` to be included.

There are two ways to include more than one section:

```
*(.text .rdata)
*(.text) *(.rdata)
```

The difference between these is the order in which the `.text` and `.rdata` input sections will appear in the output section. In the first example, they will be intermingled. In the second example, all `.text` input sections will appear first, followed by all `.rdata` input sections.

A file name can be specified to include sections from a particular file. This would be useful if one of the files contain special data that needs to be at a particular location in memory. For example:

```
data.o(.data)
```

If a file name is specified without a list of sections, then all sections in the input file will be included in the output section. This is not commonly done, but it may be useful on occasion. For example:

```
data.o
```

When a file name is specified which does not contain any wild card characters, the linker will first see if the file name was also specified on the linker command line or in an `INPUT` command. If not, the linker will attempt to open the file as an input file, as though it appeared on the command line. This differs from an `INPUT` command because the linker will not search for the file in the archive search path.

9.7.5.2 Input Section Wildcard Patterns

In an input section description, either the file name or the section name or both may be wildcard patterns.

The file name of `*` seen in many examples is a simple wildcard pattern for the file name.

The wildcard patterns are like those used by the UNIX shell.

<code>*</code>	matches any number of characters
<code>?</code>	matches any single character
<code>[chars]</code>	matches a single instance of any of the <i>chars</i> ; the <code>-</code> character may be used to specify a range of characters, as in <code>[a-z]</code> to match any lower case letter
<code>\</code>	quotes the following character

When a file name is matched with a wildcard, the wildcard characters will not match a `/` character (used to separate directory names on UNIX). A pattern consisting of a single `*` character is an exception; it will always match any file name, whether it contains a `/` or not. In a section name, the wildcard characters will match a `/` character.

File name wildcard patterns only match files which are explicitly specified on the command line or in an `INPUT` command. The linker does not search directories to expand wild cards.

If a file name matches more than one wildcard pattern, or if a file name appears explicitly and is also matched by a wildcard pattern, the linker will use the first match in the linker script. For example, this sequence of input section descriptions is probably in error, because the `data.o` rule will not be used:

```
.data : { *(.data) }  
.data1 : { data.o(.data) }
```

Normally, the linker will place files and sections matched by wild cards in the order in which they are seen during the link. This can be changed by using the `SORT` keyword, which appears before a wildcard pattern in parentheses (e.g., `SORT(.text*)`). When the `SORT` keyword is used, the linker will sort the files or sections into ascending order by name before placing them in the output file.

To verify where the input sections are going, use the `-M` linker option to generate a map file. The map file shows precisely how input sections are mapped to output sections.

This example shows how wildcard patterns might be used to partition files. This linker script directs the linker to place all `.text` sections in `.text` and all `.bss` sections in `.bss`. The linker will place the `.data` section from all files beginning with an upper case character in `.DATA`; for all other files, the linker will place the `.data` section in `.data`.

```
SECTIONS {  
  .text : { *(.text) }  
  .DATA : { [A-Z]*(.data) }  
  .data : { *(.data) }  
  .bss : { *(.bss) }  
}
```

9.7.5.3 Input Section for Common Symbols

A special notation is needed for common symbols, because common symbols do not have a particular input section. The linker treats common symbols as though they are in an input section named `COMMON`.

File names may be used with the `COMMON` section just as with any other input sections. This will place common symbols from a particular input file in one section, while common symbols from other input files are placed in another section.

In most cases, common symbols in input files will be placed in the `.bss` section in the output file. For example:

```
.bss { *(.bss) *(COMMON) }
```

If not otherwise specified, common symbols will be assigned to section `.bss`.

9.7.5.4 Input Section Example

The following example is a complete linker script. It tells the linker to read all of the sections from file `all.o` and place them at the start of output section `outputa` which starts at location `0x10000`. All of section `.input1` from file `foo.o` follows immediately, in the same output section. All of section `.input2` from `foo.o` goes into output section `outputb`, followed by section `.input1` from `foo1.o`. All of the remaining `.input1` and `.input2` sections from any files are written to output section `outputc`.

```
SECTIONS {
    outputa 0x10000 :
    {
        all.o
        foo.o (.input1)
    }
    outputb :
    {
        foo.o (.input2)
        foo1.o (.input1)
    }
    outputc :
    {
        *(.input1)
        *(.input2)
    }
}
```

9.7.5.5 Output Section Description

The full description of an output section looks like this:

```
name [address] [(type)] : [AT(lma)]
{
    output-section-command
    output-section-command
    ...
} [>region] [AT>lma_region] [=fillexp]
```

Most output sections do not use most of the optional section attributes.

The white space around *name* and *address* is required. The colon and the curly braces are also required. The line breaks and other white space are optional.

A section name may consist of any sequence of characters, but a name which contains any unusual characters such as commas must be quoted.

Each output-section-command may be one of the following:

- a symbol assignment (see [Section 9.7.3 “Assigning Values to Symbols.”](#))
- an input section description (see [Section 9.7.5.1 “Input Section Description.”](#))
- data values to include directly (see [Section 9.7.5.7 “Output Section Data.”](#))

9.7.5.6 Output Section Address

The *address* is an expression for the VMA (the virtual memory address) of the output section. If *address* is not provided, the linker will set it based on *region* if present, or otherwise based on the current value of the location counter.

If *address* is provided, the address of the output section will be set to precisely that. If neither *address* nor *region* is provided, then the address of the output section will be set to the current value of the location counter aligned to the alignment requirements of the output section. The alignment requirement of the output section is the strictest alignment of any input section contained within the output section.

For example,

```
.text . : { *(.text) }
```

and

```
.text : { *(.text) }
```

are subtly different. The first will set the address of the `.text` output section to the current value of the location counter. The second will set it to the current value of the location counter aligned to the strictest alignment of a `.text` input section.

The address may be an arbitrary expression (see [Section 9.8 “Expressions in Linker Scripts.”](#)). For example, to align the section on a 0x10 byte boundary, so that the lowest four bits of the section address are zero, the command could look like this:

```
.text ALIGN(0x10) : { *(.text) }
```

This works because `ALIGN` returns the current location counter aligned upward to the specified value.

Specifying *address* for a section will change the value of the location counter.

9.7.5.7 Output Section Data

Explicit bytes of data may be inserted into an output section by using `BYTE`, `SHORT`, or `LONG` as an output section command. Each keyword is followed by an expression in parentheses providing the value to store. The value of the expression is stored at the current value of the location counter.

The `BYTE`, `SHORT`, or `LONG` commands store one, two, or four bytes (respectively). For example, this command will store the four byte value of the symbol `addr`:

```
LONG(addr)
```

After storing the bytes, the location counter is incremented by the number of bytes stored. When using data commands in a program memory section, it is important to note that the linker considers program memory to be 32-bits wide, even though only 24 bits are physically implemented. Therefore, the most significant 8 bits of a `LONG` data value are not loaded into device memory.

Data commands only work inside a section description and not between them, so the following will produce an error from the linker:

```
SECTIONS { .text : { *(.text) } LONG(1) .data : { *(.data) } }
```

whereas this will work:

```
SECTIONS { .text : { *(.text) ; LONG(1) } .data : { *(.data) } }
```

The `FILL` command may be used to set the fill pattern for the current section. It is followed by an expression in parentheses. Any otherwise unspecified regions of memory within the section (for example, gaps left due to the required alignment of input sections) are filled with the two least significant bytes of the expression, repeated as necessary. A `FILL` statement covers memory locations after the point at which it occurs in the section definition; by including more than one `FILL` statement, different fill patterns may be used in different parts of an output section.

This example shows how to fill unspecified regions of memory with the value `0x9090`:

```
FILL(0x9090)
```

The `FILL` command is similar to the `=fillexp` output section attribute (see [Section 9.7.5.9 "Output Section Attributes."](#)), but it only affects the part of the section following the `FILL` command, rather than the entire section. If both are used, the `FILL` command takes precedence.

9.7.5.8 Output Section Discarding

The linker will not create an output section which does not have any contents. This is for convenience when referring to input sections that may or may not be present in any of the input files. For example:

```
.foo { *(.foo) }
```

will only create a `.foo` section in the output file if there is a `.foo` section in at least one input file.

If anything other than an input section description is used as an output section command, such as a symbol assignment, then the output section will always be created, even if there are no matching input sections.

The special output section name `/DISCARD/` may be used to discard input sections. Any input sections which are assigned to an output section named `/DISCARD/` are not included in the output file.

9.7.5.9 Output Section Attributes

To review, the full description of an output section is:

```
name [address] [(type)] : [AT(lma)]
{
    output-section-command
    output-section-command
    ...
} [>region] [AT>lma_region] [:phdr :phdr ...] [=fillexp]
```

name, *address* and *output-section-command* have already been described. In the following sections, the remaining section attributes will be described.

9.7.5.10 Output Section Type

Each output section may have a type. The type is a keyword in parentheses. The following types are defined:

NOLOAD

The section should be marked as not loadable, so that it will not be loaded into memory when the program is run.

DSECT, COPY, INFO, OVERLAY

These type names are supported for backward compatibility, and are rarely used. They all have the same effect: the section should be marked as not allocatable, so that no memory is allocated for the section when the program is run.

The linker normally sets the attributes of an output section based on the input sections which map into it. This can be overridden by using the section type. For example, in the script sample below, the `ROM` section is addressed at memory location 0 and does not need to be loaded when the program is run. The contents of the `ROM` section will appear in the linker output file as usual.

```
SECTIONS {
    ROM 0 (NOLOAD) : { ... }
    ...
}
```

16-Bit Assembler, Linker and Utilities User's Guide

9.7.5.11 Output Section LMA

Every section has a virtual address (VMA) and a load address (LMA). The address expression which may appear in an output section description sets the VMA.

The linker will normally set the LMA equal to the VMA. This can be changed by using the `AT` keyword. The expression `lma` that follows the `AT` keyword specifies the load address of the section. Alternatively, with `AT>lma_region` expression, a memory region may be specified for the section's load address (see [Section 9.7.4 "MEMORY Command."](#)).

This feature is designed to make it easy to build a ROM image. For example, the following linker script creates three output sections: one called `.text`, which starts at `0x1000`, one called `.mdata`, which is loaded at the end of the `.text` section even though its VMA is `0x2000`, and one called `.bss` to hold uninitialized data at address `0x3000`. The symbol `_data` is defined with the value `0x2000`, which shows that the location counter holds the VMA value, not the LMA value.

```
SECTIONS
{
    .text 0x1000 : { *(.text) _etext = . ; }
    .mdata 0x2000 :
        AT ( ADDR (.text) + SIZEOF (.text) )
        { _data = . ; *(.data); _edata = . ; }
    .bss 0x3000 :
        { _bstart = . ; *(.bss) *(COMMON) ; _bend = . ; }
}
```

The run-time initialization code for use with a program generated with this linker script would include a function to copy the initialized data from the ROM image to its run-time address. The initialization function could take advantage of the symbols defined by the linker script.

It would rarely be necessary to write such a function, however. The 16-bit linker includes automatic support for the initialization of BSS-type and data-type sections. Instead of mapping a data section into both program memory and data memory (as this example implies), the linker creates a special template in program memory which includes all of the relevant information. See [Section 10.8 "Initialized Data."](#) for details.

9.7.5.12 Output Section Region

A section can be assigned to a previously defined region of memory by using `>region`. See [Section 9.7.4 "MEMORY Command."](#)

Here is a simple example:

```
MEMORY { rom : ORIGIN = 0x1000, LENGTH = 0x1000 }
SECTIONS { ROM : { *(.text) } >rom }
```

9.7.5.13 Output Section Fill

A fill pattern can be set for an entire section by using `=fillexp`. `fillexp` as an expression. Any otherwise unspecified regions of memory within the output section (for example, gaps left due to the required alignment of input sections) will be filled with the two least significant bytes of the value, repeated as necessary.

The fill value can also be changed with a `FILL` command in the output section commands; see [Section 9.7.5.7 "Output Section Data."](#)

Here is a simple example:

```
SECTIONS { .text : { *(.text) } =0x9090 }
```

9.7.5.14 Overlay Description

An overlay description provides an easy way to describe sections which are to be loaded as part of a single memory image but are to be run at the same memory address. At run time, some sort of overlay manager will copy the overlaid sections in and out of the run-time memory address as required, perhaps by simply manipulating addressing bits.

This approach is not suitable for defining sections that will be used with the PSV window, because the `OVERLAY` command does not permit individual load addresses to be specified for each section. Instead, the 16-bit linker provides automatic support for read-only sections in the PSV window. See [Section 10.9 “Read-only Data.”](#) for details.

Overlays are described using the `OVERLAY` command. The `OVERLAY` command is used within a `SECTIONS` command, like an output section description. The full syntax of the `OVERLAY` command is as follows:

```
OVERLAY [start] : [NOCROSSREFS] [AT ( ldaddr )]  
{  
    secname1  
    {  
        output-section-command  
        output-section-command  
        ...  
    } [:phdr...] [=fill]  
    secname2  
    {  
        output-section-command  
        output-section-command  
        ...  
    } [:phdr...] [=fill]  
    ...  
} [>region] [:phdr...] [=fill]
```

Everything is optional except `OVERLAY` (a keyword), and each section must have a name (*secname1* and *secname2* above). The section definitions within the `OVERLAY` construct are identical to those within the general `SECTIONS` construct, except that no addresses and no memory regions may be defined for sections within an `OVERLAY`.

The sections are all defined with the same starting address. The load addresses of the sections are arranged such that they are consecutive in memory starting at the load address used for the `OVERLAY` as a whole (as with normal section definitions, the load address is optional, and defaults to the start address; the start address is also optional, and defaults to the current value of the location counter).

If the `NOCROSSREFS` keyword is used, and there are any references among the sections, the linker will report an error. Since the sections all run at the same address, it normally does not make sense for one section to refer directly to another.

For each section within the `OVERLAY`, the linker automatically defines two symbols. The symbol `__load_start_secname` is defined as the starting load address of the section. The symbol `__load_stop_secname` is defined as the final load address of the section. Any characters within *secname* which are not legal within C identifiers are removed. C (or assembler) code may use these symbols to move the overlaid sections around as necessary.

16-Bit Assembler, Linker and Utilities User's Guide

At the end of the overlay, the value of the location counter is set to the start address of the overlay plus the size of the largest section.

Here is an example. Remember that this would appear inside a `SECTIONS` construct.

```
OVERLAY 0x1000 : AT (0x4000)
{
    .text0 { o1/*.o(.text) }
    .text1 { o2/*.o(.text) }
}
```

This will define both `.text0` and `.text1` to start at address `0x1000`. `.text0` will be loaded at address `0x4000`, and `.text1` will be loaded immediately after `.text0`. The following symbols will be defined: `__load_start_text0`, `__load_stop_text0`, `__load_start_text1`, `__load_stop_text1`.

C code to copy overlay `.text1` into the overlay area might look like the following:

```
extern char __load_start_text1, __load_stop_text1;
memcpy ((char *) 0x1000, &__load_start_text1,
        &__load_stop_text1 - &__load_start_text1);
```

The `OVERLAY` command is a convenience, since everything it does can be done using the more basic commands. The previous example could have been written identically as follows.

```
.text0 0x1000 : AT (0x4000) { o1/*.o(.text) }
__load_start_text0 = LOADADDR (.text0);
__load_stop_text0 = LOADADDR (.text0) + SIZEOF (.text0);
.text1 0x1000 : AT (0x4000 + SIZEOF (.text0)) { o2/*.o(.text) }
__load_start_text1 = LOADADDR (.text1);
__load_stop_text1 = LOADADDR (.text1) + SIZEOF (.text1);
. = 0x1000 + MAX (SIZEOF (.text0), SIZEOF (.text1));
```

9.7.6 Other Linker Script Commands

There are several other linker script commands, which are described briefly:

ENTRY (*symbol*)

Specify *symbol* as the first instruction to execute in the program. The linker will record the address of this symbol in the output object file header. This does not affect the Reset instruction at address zero, which must be generated in some other way. By convention, the 16-bit linker scripts construct a `GOTO __reset` instruction at address zero.

EXTERN(*symbol symbol ...*)

Force *symbol* to be entered in the output file as an undefined symbol. Doing this may, for example, trigger linking of additional modules from standard libraries. Several symbols may be listed for each `EXTERN`, and `EXTERN` may appear multiple times. This command has the same effect as the `-u` command line option.

FORCE_COMMON_ALLOCATION

This command has the same effect as the `-d` command line option: to make 16-bit linker assign space to common symbols even if a relocatable output file is specified (`-r`).

NOCROSSREFS(*section section ...*)

This command may be used to tell 16-bit linker to issue an error about any references among certain output sections. In certain types of programs, when one section is loaded into memory, another section will not be. Any direct references between the two sections would be errors.

The `NOCROSSREFS` command takes a list of output section names. If the linker detects any cross references between the sections, it reports an error and returns a non-zero exit status. The `NOCROSSREFS` command uses output section names, not input section names.

OUTPUT_ARCH(*processor_name*)

Specify a target processor for the link. This command has the same effect as the `-p,--processor` command line option. If both are specified, the command line option takes precedence. The processor name should appear in quotes; for example "30F6014", "24FJ128GA010", or "33FJ128GP706".

OUTPUT_FORMAT(*format_name*)

The `OUTPUT_FORMAT` command names the object file format to use for the output file.

TARGET(*bfdname*)

The `TARGET` command names the object file format to use when reading input files. It affects subsequent `INPUT` and `GROUP` commands.

9.8 Expressions in Linker Scripts

The syntax for expressions in the linker script language is identical to that of C expressions. All expressions are evaluated as 32-bit integers.

You can use and set symbol values in expressions.

The linker defines several special purpose built-in functions for use in expressions.

- [Constants](#)
- [Symbol Names](#)
- [The Location Counter](#)
- [Operators](#)
- [Evaluation](#)
- [The Section of an Expression](#)
- [Built-in Functions](#)

9.8.1 Constants

All constants are integers.

As in C, the linker considers an integer beginning with 0 to be octal, and an integer beginning with 0x or 0X to be hexadecimal. The linker considers other integers to be decimal.

In addition, you can use the suffixes K and M to scale a constant by 1024 or 1024*1024 respectively. For example, the following all refer to the same quantity:

```
_fourk_1 = 4K;  
_fourk_2 = 4096;  
_fourk_3 = 0x1000;
```

9.8.2 Symbol Names

Unless quoted, symbol names start with a letter, underscore, or period and may include letters, digits, underscores, periods and hyphens. Unquoted symbol names must not conflict with any keywords. You can specify a symbol which contains odd characters or has the same name as a keyword by surrounding the symbol name in double quotes:

```
"SECTION" = 9;  
"with a space" = "also with a space" + 10;
```

Since symbols can contain many non-alphabetic characters, it is safest to delimit symbols with spaces. For example, A-B is one symbol, whereas A - B is an expression involving subtraction.

9.8.3 The Location Counter

The special linker variable dot '.' always contains the current output location counter. Since the '.' always refers to a location in an output section, it may only appear in an expression within a `SECTIONS` command. The '.' symbol may appear anywhere that an ordinary symbol is allowed in an expression.

Assigning a value to '.' will cause the location counter to be moved. This may be used to create holes in the output section. The location counter may never be moved backwards.

```
SECTIONS
{
    output :
    {
        file1(.text)
        . = . + 1000;
        file2(.text)
        . += 1000;
        file3(.text)
    } = 0x1234;
}
```

In the previous example, the `.text` section from `file1` is located at the beginning of the output section `output`. It is followed by a 1000 byte gap. Then the `.text` section from `file2` appears, also with a 1000 byte gap following before the `.text` section from `file3`. The notation `= 0x1234` specifies what data to write in the gaps.

'.' actually refers to the byte offset from the start of the current containing object. Normally this is the `SECTIONS` statement, whose start address is 0, hence '.' can be used as an absolute address. If '.' is used inside a section description, however, it refers to the byte offset from the start of that section, not an absolute address, as shown in the following script:

```
SECTIONS
{
    . = 0x100
    .text: {
        *(.text)
        . = 0x200
    }
    . = 0x500
    .data: {
        *(.data)
        . += 0x600
    }
}
```

The `.text` section will be assigned a starting address of 0x100 and a size of exactly 0x200 bytes, even if there is not enough data in the `.text` input sections to fill this area. (If there is too much data, an error will be produced because this would be an attempt to move '.' backwards). The `.data` section will start at 0x500 and it will have an extra 0x600 bytes worth of space after the end of the values from the `.data` input sections and before the end of the `.data` output section itself.

16-Bit Assembler, Linker and Utilities User's Guide

9.8.4 Operators

The linker recognizes the standard C set of arithmetic operators, with the following standard bindings and precedence levels:

Table 9-2: Precedence of Operators

Precedence	Associativity	Operators	Description
1 (highest)	left	! - ~	Prefix operators
2	left	* / %	multiply, divide, modulo
3	left	+ -	add, subtract
4	left	>> <<	bit shift right, left
5	left	== != > < <= >=	Relational
6	left	&	bitwise and
7	left		bitwise or
8	left	&&	logical and
9	left		logical or
10	right	? :	Conditional
11 (lowest)	right	&= += -= *= /=	Symbol assignments

9.8.5 Evaluation

The linker evaluates expressions lazily. It only computes the value of an expression when absolutely necessary.

The linker needs some information, such as the value of the start address of the first section, and the origins and lengths of memory regions, in order to do any linking at all. These values are computed as soon as possible when the linker reads in the linker script.

However, other values (such as symbol values) are not known or needed until after storage allocation. Such values are evaluated later, when other information (such as the sizes of output sections) is available for use in the symbol assignment expression.

The sizes of sections cannot be known until after allocation, so assignments dependent upon these are not performed until after allocation.

Some expressions, such as those depending upon the location counter '.', must be evaluated during section allocation.

If the result of an expression is required, but the value is not available, then an error results. For example, a script like the following:

```
SECTIONS
{
    .text 9+this_isnt_constant :
    { *(.text) }
}
```

will cause the error message "non-constant expression for initial address".

9.8.6 The Section of an Expression

When the linker evaluates an expression, the result is either absolute or relative to some section. A relative expression is expressed as a fixed offset from the base of a section.

The position of the expression within the linker script determines whether it is absolute or relative. An expression which appears within an output section definition is relative to the base of the output section. An expression which appears elsewhere will be absolute.

A symbol set to a relative expression will be relocatable if you request relocatable output using the `-r` option. That means that a further link operation may change the value of the symbol. The symbol's section will be the section of the relative expression.

A symbol set to an absolute expression will retain the same value through any further link operation. The symbol will be absolute, and will not have any particular associated section.

You can use the built-in function `ABSOLUTE` to force an expression to be absolute when it would otherwise be relative. For example, to create an absolute symbol set to the address of the end of the output section `.data`:

```
SECTIONS
{
    .data : { *(.data) _edata = ABSOLUTE(.); }
}
```

If `ABSOLUTE` were not used, `_edata` would be relative to the `.data` section.

9.8.7 Built-in Functions

The linker script language includes a number of built-in functions for use in linker script expressions.

```
ABSOLUTE(exp)
ADDR(section)
ALIGN(exp)
ASSERT(exp, message)
BLOCK(exp)
DEFINED(symbol)
LOADADDR(section)
MAX(exp1, exp2)
MIN(exp1, exp2)
NEXT(exp)
SIZEOF(section)
```

9.8.7.1 ABSOLUTE(exp)

Return the absolute (non-relocatable, as opposed to non-negative) value of the expression `exp`. Primarily useful to assign an absolute value to a symbol within a section definition, where symbol values are normally section relative (see [Section 9.8.6 “The Section of an Expression.”](#)).

9.8.7.2 ADDR(*section*)

Return the absolute address (the VMA) of the named section. Your script must previously have defined the location of that section. In the following example, `symbol_1` and `symbol_2` are assigned identical values:

```
SECTIONS { ...
    .output1 :
    {
        start_of_output_1 = ABSOLUTE(.);
        ...
    }
    .output :
    {
        symbol_1 = ADDR(.output1);
        symbol_2 = start_of_output_1;
    }
    ...
}
```

9.8.7.3 ALIGN(*exp*)

Return the location counter (.) aligned to the next *exp* boundary. *exp* must be an expression whose value is a power of two. This is equivalent to:

$$(. + exp - 1) \& \sim(exp - 1)$$

ALIGN doesn't change the value of the location counter; it just does arithmetic on it. Here is an example which aligns the output `.data` section to the next 0x2000 byte boundary after the preceding section and sets a variable within the section to the next 0x8000 boundary after the input sections:

```
SECTIONS { ...
    .data ALIGN(0x2000): {
        *(.data)
        variable = ALIGN(0x8000);
    }
    ...
}
```

The first use of ALIGN in this example specifies the location of a section because it is used as the optional address attribute of a section definition, see [Section 9.7.5 "SECTIONS Command."](#) The second use of ALIGN is used to define the value of a symbol.

The built-in function NEXT is closely related to ALIGN.

9.8.7.4 ASSERT(*exp*, *message*)

Ensure that *exp* is non-zero. If it is zero, then exit the linker with an error code, and print *message*. E.g.,

```
__CHECK = ASSERT(1, "OK");
```

9.8.7.5 BLOCK(*exp*)

This is a synonym for ALIGN, for compatibility with older linker scripts. It is most often seen when setting the address of an output section.

9.8.7.6 `DEFINED(symbol)`

Return 1 if `symbol` is in the linker global symbol table and is defined; otherwise return 0. You can use this function to provide default values for symbols. For example, the following script fragment shows how to set a global symbol `begin` to the first location in the `.text` section, but if a symbol called `begin` already existed, its value is preserved:

```
SECTIONS { ...
    .text : {
        begin = DEFINED(begin) ? begin : . ;
        ...
    }
    ...
}
```

9.8.7.7 `LOADADDR(section)`

Return the absolute LMA of the named section. This is normally the same as `ADDR`, but it may be different if the `AT` attribute is used in the output section definition, see [Section 9.7.5 “SECTIONS Command.”](#)

9.8.7.8 `MAX(exp1, exp2)`

Returns the maximum of `exp1` and `exp2`.

9.8.7.9 `MIN(exp1, exp2)`

Returns the minimum of `exp1` and `exp2`.

9.8.7.10 `NEXT(exp)`

Return the next unallocated address that is a multiple of `exp`. This function is equivalent to `ALIGN(exp)`.

9.8.7.11 `SIZEOF(section)`

Return the size in bytes of the named section, if that section has been allocated. If the section has not been allocated when this is evaluated, the linker will report an error. In the following example, `symbol_1` and `symbol_2` are assigned identical values:

```
SECTIONS{ ...
    .output {
        .start = . ;
        ...
        .end = . ;
    }
    symbol_1 = .end - .start ;
    symbol_2 = SIZEOF(.output);
    ...
}
```

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Chapter 10. Linker Processing

10.1 Introduction

How the MPLAB XC16 Object Linker builds an application from input files and the linker script is discussed here.

10.2 Highlights

Topics covered in this chapter are:

- [Overview of Linker Processing](#)
- [Memory Addressing](#)
- [Linker Allocation](#)
- [Global and Weak Symbols](#)
- [Handles](#)
- [Initialized Data](#)
- [Read-only Data](#)
- [Stack Allocation](#)
- [Heap Allocation](#)
- [Interrupt Vector Tables](#)
- [Optimizing Memory Usage](#)
- [Boot and Secure Segments](#)
- [Co-resident Application Linking](#)
- [Notable Symbols](#)

10.3 Overview of Linker Processing

A linker combines one or more object files, with optional archive files, into a single executable output file. The object files contain relocatable sections of code and data which the linker will allocate into target memory. The entire process is controlled by a linker script, also known as a link command file. A linker script is required for every link.

The link process may be broken down into 6 steps:

1. [Loading Input Files](#)
2. [Allocating Memory](#)
3. [Resolving Symbols](#)
4. [Creating Special Sections](#)
5. [Computing Absolute Addresses](#)
6. [Building the Output File](#)

10.3.1 Loading Input Files

The initial task of the linker is to interpret link command options and load input files. If a linker script is specified, that file is opened and interpreted. Otherwise an internal default linker script is used. In either case, the linker script provides a description of the target device, including specific memory region information and SFR addresses. See [Chapter 9. “Linker Scripts”](#) for more details.

Next the linker opens all of the input object files. Each input file is checked to make sure the object format is compatible. If the object format is not compatible, an error is generated. The contents of each input file are then loaded into internal data structures. Typically each input file will contain multiple sections of code or data. Each section contains a list of relocation entries which associate locations in a section's raw data with relocatable symbols.

10.3.2 Allocating Memory

After all of the input files have been loaded, the linker allocates memory. This is accomplished by assigning each input section to an output section. The relation between input and output sections is defined by a section map in the linker script. An output section may or may not have the same name as an input section. Each output section is then assigned to a memory region in the target device.

<p>Note: Input sections are derived from source code by the compiler or the assembler. Output sections are created by the linker.</p>
--

If an input section is not explicitly assigned to an output section, the linker will allocate the unassigned section according to section attributes. For more information about linker allocation, see [Section 10.5 “Linker Allocation.”](#)

10.3.3 Resolving Symbols

Once memory has been allocated, the linker begins the process of resolving symbols. Symbols defined in each input section have offsets that are relative to the beginning of the section. The linker converts these values into output section offsets.

Next, the linker attempts to match all external symbol references with a corresponding symbol definition. Multiple definitions of the same external symbol result in an error. If an external symbol is not found, an attempt is made to locate the symbol definition in an archive file. If the symbol definition is found in an archive, the corresponding archive module is loaded.

Modules loaded from archives may contain additional symbol references, so the process continues until all external symbol references have matching definitions. External symbols that are defined as “weak” receive special processing, as explained in [Section 10.6 “Global and Weak Symbols.”](#) If any external symbol reference remains undefined, an error is generated.

References to redundant functions in archive files will be merged in order to conserve memory. For example, both integer and floating-point versions of the standard C formatted I/O functions are included in `libc.a`. The 16-bit compiler will generate references to the appropriate function, based on a static analysis of format strings. When multiple object files are combined by the linker, both versions of a particular I/O function may be referenced. In such cases the integer functions are redundant, since they represent a subset of the floating-point functionality. The linker will detect this situation, and merge the I/O functions together to conserve memory. This optimization may be disabled with the `--no-smart-io` option.

10.3.4 Creating Special Sections

After the symbols have been resolved, the linker constructs any special input or output sections that are required. For example, the compiler or assembler may have created function pointers using the `handle()` operator. The linker then builds a special input section named `.handle` to implement a jump table. For more information about handles, see [Section 10.7 “Handles.”](#)

The linker also constructs a special input section named `.dinit` to support initialized data. Section `.dinit` is an initialization template that is interpreted by the C run-time library. For more information about initialized data, see [Section 10.8 “Initialized Data.”](#)

If the application has not defined a default interrupt handler, the linker will create one automatically in a special input section named `.isr`. Unused slots in the interrupt vector tables are populated with the address of this function. For more information on the default interrupt handler, see [Section 10.12 “Interrupt Vector Tables.”](#)

10.3.5 Computing Absolute Addresses

After the special sections have been created, the final sizes of all output sections are known. The linker then computes absolute addresses for all output sections and external symbols. Each output section is checked to make sure it falls within its assigned memory regions. If any section falls outside of its memory region, an error is generated. Any symbols defined in the linker script are also computed.

Boundaries of the stack and heap are calculated, based on the extent of unused data memory. If insufficient memory is available, an error is generated. For more information about the stack and heap, see [Section 10.10 “Stack Allocation.”](#) and [Section 10.11 “Heap Allocation.”](#)

10.3.6 Building the Output File

Finally, the linker builds the output file. Relocation entries in each section are patched using absolute addresses. If the address computed for a symbol does not fit in the relocation entry, a link error results. This can occur, for example, if a function pointer is referenced without the `handle()` operator and its address is too large to fit in 16 bits.

A link map is also generated if requested with the appropriate option. The link map includes a memory usage report, which shows the starting address and length of all sections in data memory and program memory. For more information about the link map, see [Section 9.5.4 “Input/Output Section Map.”](#)

16-Bit Assembler, Linker and Utilities User's Guide

10.4 Memory Addressing

The 16-bit devices use a modified Harvard architecture with separate data and program memory spaces. Data memory is both byte-oriented (8 bits wide) and word-oriented (16 bits wide). Bytes are assigned sequential addresses, starting with 0, 1, 2, 3 and so on. Words are assigned sequential even addresses, starting with 0, 2, 4, 6 and so on.

Program memory is word-oriented, where each instruction word is 24 bits wide. Instruction words are assigned sequential even addresses, starting with 0, 2, 4, 6 and so on. The PC indicates the next instruction to be executed, and increments by 2 for each instruction word. Individual bytes in a program memory word are not addressable.

While a traditional Harvard architecture does not permit access to data stored in program memory, the 16-bit architecture provides three ways to accomplish this task:

- [Table Access Instructions](#)
- [Program Space Visibility \(PSV\) Window](#)
- [Extended Data Space \(EDS\) Window](#)

10.4.1 Table Access Instructions

The table access instructions `tblrdl`, `tblrdh`, `tblwtl` and `tblwth` can be used to access data stored in program memory. Data is addressed through a 16-bit data register pointer in combination with the 8-bit TBLPAG register. The special operators `tbloffset()` and `tblpage()` facilitate table access in assembly language. See the 16-bit assembler documentation, "Table Read/Write Instructions", for more information.

The linker resolves symbolic references to labels in program memory for use with the table access instructions. Although data in program memory can be specified one byte at a time, only the least-significant byte in each instruction word has a unique address. For example, consider the following assembly source code example:

```
.section prog,code
L1: .pbyte 1
L2: .pbyte 2
L3: .pbyte 3
L4: .pbyte 4
    .pbyte 5
    .pbyte 6
    .pbyte 7,8,9
```

In this example, the `code` section attribute designates a section to be allocated in program memory, and the `.pbyte` directives define individual byte constants. Since labels must resolve to a valid PC address, the assembler adds padding after each of the first three constants. Subsequent constants do not require padding. The following assembly listing excerpt illustrates the organization of these constants in program memory:

```
1                                     .section prog,code
2 000000 01 00 00 L1:.pbyte 1
3 000002 02 00 00 L2:.pbyte 2
4 000004 03 00 00 L3:.pbyte 3
5 000006 04      L4:.pbyte 4
6           05      .pbyte 5
7           06      .pbyte 6
8 000008 07 08 09 .pbyte 7,8,9
```

Constants 1, 2, 3 are padded out to a full instruction word and have unique PC addresses. Constants 4, 5, 6 are packed into a single instruction word and share the same address.

10.4.2 Program Space Visibility (PSV) Window

The Program Space Visibility window can be used to access data stored in the least significant 16 bits of program memory. When PSV is enabled, the upper 32K of data memory space (0x8000-0xFFFF) functions as a window into program memory. Data is addressed through a 16-bit data register pointer in combination with the 8-bit PSVPAG register. The special operators `psvoffset()` and `psvpage()` are provided to facilitate PSV access in assembly language. Built-in functions `__builtin_psvoffset()` and `__builtin_psvpage()` are provided to facilitate PSV access in C.

The linker supports PSV window operations through the use of read-only data sections. For a detailed discussion of read-only sections, see [Section 10.9 “Read-only Data.”](#)

10.4.3 Extended Data Space (EDS) Window

Some device families support a new data memory architecture called Extended Data Space (EDS). EDS extends the functionality of the PSV window to access additional pages of RAM as well as memory-mapped peripherals. On an EDS device, the PSVPAG register has been replaced by two registers:

- DSRPAG for reading from Flash, RAM, etc.
- DSWPAG for writing to RAM

The operation of the EDS window is analogous to the PSV window. When the page registers are set appropriately, a portion of program memory (or extended data memory) can be accessed in the data address range 0x8000 to 0xFFFF. Unlike the PSV window, the EDS window is always enabled. Another difference is that certain page number ranges imply different address spaces:

EDS Page Range	Description
0x001 to 0x1FF	read/write access to RAM
0x200 to 0x2FF	read-only access to lower 16 bits of program memory
0x300 to 0x3FF	read-only access to upper 8 bits of program memory

Note: EDS page 0 is undefined. Application code should not attempt to access the EDS window with a page value of zero. Such access is prohibited and a hardware exception will occur.

The special operators `edsoffset()` and `edspage()` are provided to facilitate EDS access from assembly language. Built-in functions `__builtin_edsoffset()` and `__builtin_edspage()` are provided to facilitate EDS access from C.

The EDS special operators may be used to access any object in on-chip memory, including local RAM (i.e., RAM located within the first 32K of data address space). Consequently, `edsoffset()` may return a pointer in the range 0x0 to 0xFFFF. `edspage()` will return a page value in the range 0x001 to 0x2FF. Page values greater than 0x300 are not currently supported.

10.5 Linker Allocation

Linker allocation is controlled by the linker script, and proceeds in three steps:

1. [Mapping Input Sections to Output Sections](#)
2. [Assigning Output Sections to Regions](#)
3. [Allocating Unmapped Sections](#)

Steps 1 and 2 are performed by a sequential memory allocator. Input sections which appear in the linker script are assigned to specific memory regions in the target devices. Addresses within a memory region are allocated sequentially, beginning with the lowest address and growing upwards.

Step 3 is performed by a best-fit memory allocator. Input sections which do not appear in the linker script are assigned to memory regions according to their attributes. The best-fit allocator makes efficient use of any remaining memory, including gaps between output sections that may have been left by the sequential allocator.

If memory has been reserved for the `boot` and/or `secure` segments, it will be allocated by the best-fit allocator in step 3. The sequential allocator will avoid these segments, so sections designated with the `boot` or `secure` attributes should not appear in the linker script.

10.5.1 Mapping Input Sections to Output Sections

Input sections are grouped and mapped into output sections, according to the section map. When an output section contains several different input sections, the exact ordering of input sections may be important. For example, consider the following output section definition:

```
/*
** User Code and Library Code
*/
.text :
{
    *(.init);
    *(.user_init);
    *(.handle);
    *(.libc) *(.libm) *(.libdsp); /* keep together in this order */
    *(.lib*);
} >program
```

Here the output section named `.text` is defined. Notice that the contents of this section are specified within curly braces `{}`. After the closing brace, `>program` indicates that this output section should be assigned to memory region `program`.

The contents of output section `.text` may be interpreted as follows:

- Input sections named `.init` are collected and mapped into the output section. By convention, there is only one `.init` section, and it contains the startup code for an application. It appears first in the output section (i.e., at the lowest address) so that its address is readily available if necessary.
- Input sections named `.user_init` are collected and mapped into the output section. These sections are created by the compiler and refer to functions that have been decorated with the `user_init` attribute. Their position within the output section is not critical, but since they are associated with section `.init`, they are located immediately after.
- All input sections named `.handle` are collected and mapped into the output section. `.handle` sections occupy a relatively low address range, which is a requirement for code handles.

- Input sections named `.libc`, `.libm` and `.libdsp` are collected and mapped into the output section. Grouping these sections ensures locality of reference for the run-time library functions, so that PC-relative instructions can be used for maximum efficiency.
- Input sections which match the wildcard pattern `.lib*` are collected and mapped into the output section. This includes libraries such as the peripheral libraries (which are allocated in section `.libperi`).

10.5.2 Assigning Output Sections to Regions

Once the sizes of all output sections are known, they are assigned to memory regions. Normally a region is specified in the output section definition. If a region is not specified, the first defined memory region will be used.

Memory regions are filled sequentially, from lower to higher addresses, in the same order that sections appear in the section map. Memory reserved for `boot` or `secure` segments will be avoided, as well as sections that have been marked with the `address` attribute in source code. A location counter, unique to each region, keeps track of the next available memory location. There are two conditions which may cause gaps in the allocation of memory within a region:

1. The section map specifies an absolute address for an output section, or
2. The output section has a particular alignment requirement.

In either case, any intervening memory between the current location counter and the absolute (or aligned) address is skipped. Once a range of memory has been skipped, it is available for use by the best-fit allocator. The exact address of all items allocated in memory may be determined from the link map file.

Section alignment requirements typically arise in DSP programming. To utilize modulo addressing, it is necessary to align a block of memory to a particular storage boundary. This can be accomplished with the `aligned` attribute in C, or with the `.align` directive in assembly language. The section containing an aligned memory block must also be aligned, to the same (or greater) power of 2. If two or more input sections have different alignment requirements, the largest alignment is used for the output section.

Another restriction on memory allocation is associated with read-only data sections. Read-only data sections are identified with the `psv` section attribute and are dedicated for use in the PSV window or the Extended Data Space (EDS) window. The C compiler creates a read-only data section named `.const` to store constants when the `--mconst-in-code` option is selected.

To allow efficient access of constant tables in the PSV or EDS window, the linker ensures that a read-only section will not cross a page boundary. Therefore a single setting of the page register can be used to access the entire section. If necessary, output sections in program memory will be re-sorted after the sequential allocation pass to accommodate this restriction. If an absolute address has been specified in the linker script for a particular section, it will not be moved. In general, fully relocatable sections provide the most flexibility for efficient memory allocation.

Note: Sections with specific alignment requirements, such as `psv` sections or sections intended for modulo addressing, may be allocated most efficiently by the best-fit allocator. For best-fit allocation, these sections should not appear in the linker script.

10.5.3 Allocating Unmapped Sections

After all sections that appear in the section map are allocated, any remaining sections are considered to be unmapped. Unmapped sections are allocated according to section attributes. The linker uses a best-fit memory allocator to determine the most efficient arrangement in memory. The primary emphasis of the best-fit allocator is the reduction or elimination of memory gaps due to address alignment restrictions.

Since data memory is limited on many 16-bit devices, and several architectural features imply address alignment restrictions, efficient allocation of data memory is particularly important. By convention, data memory sections are not explicitly mapped in linker scripts, thus providing maximum flexibility for the best-fit memory allocator.

Section attributes affect memory allocation as described below. For a general discussion of section attributes, see [Section 4.2 “Directives that Define Sections.”](#)

code

The code attribute specifies that a section should be allocated in program memory, as defined by region `program` in the linker script. The following attributes may be used in conjunction with code and will further specify the allocation:

- `address()` specifies an absolute address
- `align()` specifies alignment of the section starting address
- `boot` specifies the boot segment
- `secure` specifies the secure segment

data

The data attribute specifies that a section should be allocated as initialized storage in data memory, as defined by region `data` in the linker script. The following attributes may be used in conjunction with data and will further specify the allocation:

`address()` specifies an absolute address

`near` specifies the first 8K of data memory

DD `xmemory` specifies X address space, which includes all of region data below the address `__YDATA_BASE` as defined in the linker script (dsPIC30F/33F DSCs only)

DD `ymemory` specifies Y address space, which includes all of region data above the address `__YDATA_BASE` as defined in the linker script (dsPIC30F/33F DSCs only)

`align()` specifies alignment of the section starting address

`reverse()` specifies alignment of the section ending address + 1

DD `dma` specifies dma address space, which includes the portion of region data between addresses `__DMA_BASE` and `__DMA_END` as defined in the linker script (for PIC24H MCUs and dsPIC33F DSCs only).

bss

The `bss` attribute specifies that a section should be allocated as uninitialized storage in data memory, as defined by region `data` in the linker script. The following attributes may be used in conjunction with `bss` and will further specify the allocation:

`address()` specifies an absolute address

`near` specifies the first 8K of data memory

DD `xmemory` specifies X address space, which includes all of region data below the address `__YDATA_BASE` as defined in the linker script (dsPIC30F/33F DSCs only)

DD `ymemory` specifies Y address space, which includes all of region data above the address `__YDATA_BASE` as defined in the linker script (dsPIC30F/33F DSCs only)

`align()` specifies alignment of the section starting address

`reverse()` specifies alignment of the section ending address + 1

DD `dma` specifies dma address space, which includes the portion of region data between addresses `__DMA_BASE` and `__DMA_END` as defined in the linker script (for PIC24H MCUs and dsPIC33F DSCs only).

`boot` specifies the boot segment

`secure` specifies the secure segment

persist

The `persist` attribute specifies that a section should be allocated as persistent storage in data memory, as defined by region `data` in the linker script. Persistent storage is not cleared or initialized by the C run-time library. The following attributes may be used in conjunction with `persist` and will further specify the allocation:

`address()` specifies an absolute address

`near` specifies the first 8K of data memory

DD `xmemory` specifies X address space, which includes all of region data below the address `__YDATA_BASE` as defined in the linker script (dsPIC30F/33F DSCs only)

DD `ymemory` specifies Y address space, which includes all of region data above the address `__YDATA_BASE` as defined in the linker script (dsPIC30F/33F DSCs only)

`align()` specifies alignment of the section starting address

`reverse()` specifies alignment of the section ending address + 1

DD `dma` specifies dma address space, which includes the portion of region data between addresses `__DMA_BASE` and `__DMA_END` as defined in the linker script (for PIC24H MCUs and dsPIC33F DSCs only).

16-Bit Assembler, Linker and Utilities User's Guide

psv

The psv attribute specifies that a section should be allocated in program memory, as defined by region `program` in the linker script. psv sections are intended for use with the Program Space Visibility window or the Extended Data Space (EDS) window, and will be located so that the entire contents may be accessed using a single setting of the page register. This allocation rule implies that the total size of a psv section can not exceed 32K. The following attributes may be used in conjunction with psv and will further specify the allocation:

- `address()` specifies an absolute address
- `align()` specifies alignment of the section starting address
- `reverse()` specifies alignment of the section ending address + 1
- `boot` specifies the boot segment
- `secure` specifies the secure segment

memory

The memory attribute specifies that a section should be allocated in external or user-defined memory. The following attributes may be used in conjunction with memory and will further specify the allocation:

- `address()` specifies an absolute address
- `align()` specifies alignment of the section starting address
- `reverse()` specifies alignment of the section ending address + 1
- `noload` specifies that the section should not be loaded with the primary application

Note: Sections allocated in external or user-defined memory cannot be accessed by the PSV window or the EDS window.

DD eedata – dsPIC30F DSCs only

The eedata attribute specifies that a section should be allocated in data EEPROM memory, as defined by region `eedata` in the linker script. The following attributes may be used in conjunction with eedata and will further specify the allocation:

- `address()` specifies an absolute address
- `align()` specifies alignment of the section starting address
- `reverse()` specifies alignment of the section ending address + 1
- `boot` specifies the boot segment
- `secure` specifies the secure segment

heap

The heap attribute specifies that a section should be designated for use by the C run-time library for dynamic memory allocation. The heap must always be allocated in local data memory (address range 0x0 to 0x7FFE). The following attributes may be used in conjunction with heap and will further specify the allocation:

- `address()` specifies an absolute address

DD `xmemory` specifies X address space, which includes all of region data below the address `__YDATA_BASE` as defined in the linker script (dsPIC30F/33F DSCs only)

DD `ymemory` specifies Y address space, which includes all of region data above the address `__YDATA_BASE` as defined in the linker script (dsPIC30F/33F DSCs only)

- `align()` specifies alignment of the section starting address

stack

The stack attribute specifies that a section should be designated for use as the processor stack. On most devices, the stack must always be allocated in local data memory (address range 0x0 to 0x7FFE). On some devices, the stack may be located anywhere in EDS page 1 (address range 0x0 to 0xFFFE). The following attributes may be used in conjunction with stack and will further specify the allocation:

`address()` specifies an absolute address

`align()` specifies alignment of the section starting address

10.6 Global and Weak Symbols

When a symbol reference appears in an object file without a corresponding definition, the symbol is declared external. By default, external symbols have global binding and are referred to as global symbols. External symbols may be explicitly declared with weak binding, using the `__weak__` attribute in C or the `.weak` directive in assembly language.

As the name implies, global symbols are visible to all input files involved in the link. There must be one (and only one) definition for every global symbol referenced. If a global definition is not found among the input files, archives will be searched and the first archive module found that contains the needed definition will be loaded. If no definition is found for a global symbol a link error is reported.

Weak symbols share the same name space as global symbols, but are handled differently. Multiple definitions of a weak symbol are permitted. If a weak definition is not found among the input files, archives are not searched and a value of 0 is assumed for all references to the weak symbol. A global symbol definition of the same name will take precedence over a weak definition (or the lack of one). In essence, weak symbols are considered optional and may be replaced by global symbols, or ignored entirely.

10.7 Handles

The modified Harvard architecture of dsPIC30F devices supports two memory spaces of unequal size. Data memory space can be fully addressed with 16 bits while program memory space requires 24 bits. Since the native integer data type (register width) is only 16 bits, there is an inherent difficulty in the allocation and manipulation of function pointers that require a full 24 bits. Reserving a pair of 16-bit registers to represent every function pointer is inefficient in terms of code space and execution speed, since many programs will fit in 64K words of program space or less. However, the linker must accommodate function pointers throughout the full 24-bit range of addressable program memory.

Note: Future versions of the compiler may define function pointers to be 24 bits or larger. In such cases, handles will not be used.

In order to ensure a valid 16-bit pointer for any function in the full program memory address space, the 16-bit assembler and linker support the `handle()` operator. The C compiler uses this operator whenever a function address is taken. Assembly programmers can use this operator three different ways:

```
mov    #handle(func),w0 ; handle() used in an instruction
.word  handle(func)      ; handle() used with a data word directive
.pword handle(func)      ; handle() used with a instruction word
                        ;directive
```

The linker searches all input files for handle operators and constructs a jump table in a section named `.handle`. For each function that is referenced by one or more handle operators, a single entry is made in the jump table. Each entry is a `GOTO` instruction. Note that `GOTO` is capable of reaching any function in the full 24-bit address space. Section `.handle` is allocated low in program memory, well within the range of a 16-bit pointer.

When the output file is built, the absolute addresses of all functions are known. Each handle relocation entry is filled with an absolute address. If the address of the target function fits in 16 bits, it is inserted directly into the object code. If the absolute address of the target function exceeds 16 bits, the address of the corresponding entry in the jump table is used instead. Only functions located beyond the range of 16-bit addressing suffer any performance penalty with this technique. However, there is a code space penalty for each unused entry in the jump table.

In order to conserve program memory, the handle jump table can be suppressed for certain devices, or whenever the application programmer is sure that all function pointers will fit in 16 bits. One way is to specify the `--no-handles` link option on the command line or in the IDE. Another way is to define a symbol named `__NO_HANDLES` in the linker script:

```
__NO_HANDLES = 1;
```

Linker scripts for 16-bit devices with 32K instruction words or less all contain the `__NO_HANDLES` definition to suppress the handle jump table.

Note: If the handle jump table is suppressed, and the target address of a function pointer does not fit in 16 bits, a “relocation truncated” link error will be generated.

10.8 Initialized Data

The linker provides automatic support for initialized variables in data memory. Variables are allocated in sections. Each data section is declared with a flag that indicates whether it is initialized, or not initialized.

To control the initialization of the various data sections, the linker constructs a data initialization template. The template is allocated in program memory, and is processed at start-up by the run-time library. When the application main program takes control, all variables in data memory have been initialized.

- [Standard Data Section Names](#)
- [Data Initialization Template](#)
- [Run-Time Library Support](#)

10.8.1 Standard Data Section Names

Traditionally, linkers based on the GNU technology support three sections in the linked binary file:

Table 10-1: Traditional Section Names

Section Name	Description	Attribute
.text	executable code	code
.data	data memory that receives initial values	data
.bss	data memory that is not initialized	bss

The name “bss” dates back several decades, and means memory “Block Started by Symbol.” By convention, bss memory is filled with zeros during program start-up.

The traditional section names are considered to have implied attributes as listed in Table 10-1. The `code` attribute indicates that the section contains executable code and should be loaded in program memory. The `bss` attribute indicates that the section contains data storage that is not initialized, but will be filled with zeros at program start-up. The `data` attribute indicates that the section contains data storage that receives initial values at start-up.

Assembly applications may define additional sections with explicit attributes using the section directive described in [Section 4.2 “Directives that Define Sections.”](#) For C applications, the 16-bit compiler will automatically define sections to contain variables and functions as needed. For more information on the attributes of variables and functions that may result in automatic section definition, see the “*MPLAB XC16 C Compiler User's Guide*” (DS50002071).

Note: Whenever a section directive is used, all declarations that follow are assembled into the named section. This continues until another section directive appears, or the end of file. For more information on defining sections and section attributes, see [Section 4.2 “Directives that Define Sections.”](#)

10.8.2 Data Initialization Template

As noted in [Section 10.8.1 “Standard Data Section Names.”](#), the 16-bit Language Tools support BSS-type sections (memory that is not initialized) as well as data-type sections (memory that receives initial values). The data-type sections receive initial values at start-up and the BSS-type sections are filled with zeros.

A generic data initialization template is used that supports any number of arbitrary BSS-type sections or data-type sections. The data initialization template is created by the linker and is loaded into an output section named `.dinit` in program memory. Start-up code in the run-time library interprets the template and initializes data memory accordingly.

The data initialization template contains one record for each output section in data memory. The template is terminated by a null instruction word. The format of a data initialization record is:

```
/* data init record */
struct data_record {
    char      *dst;          /* destination address */
    unsigned int len;        /* length in bytes */
    unsigned int format:7;   /* format code */
    unsigned int page:9;     /* destination page */
    char      dat[0];        /* variable length data */
};
```

The first element of the record is a pointer to the section in data memory. The second and third elements are the section length and format code, respectively. The fourth element is the page value of the section. On EDS devices, the page value will be in the range 0x001 to 0x1FF. On all other devices, the page value will be zero. The last element is an optional array of data bytes. For BSS-type sections, no data bytes are required.

The format code has three possible values.

Table 10-2: Format Code Values

Format Code	Description
0	Fill the output section with zeros
1	Copy 2 bytes of data from each instruction word in the data array
2	Copy 3 bytes of data from each instruction word in the data array

By default, data records are created using format 2. Format 2 conserves program memory by using the entire 24-bit instruction word to store initial values. Note that this format causes the encoded instruction words to appear as random and possibly invalid instructions if viewed in the disassembler.

Format 1 data records may be created by specifying the `--no-pack-data` option. Format 1 uses only the lower 16 bits of each 24-bit instruction word to store initial values. The upper byte of each instruction word is filled with 0x0 by default and causes the template to appear as NOP instructions if viewed in the disassembler (and will be executed as such by the 16-bit device). A different value may be specified for the upper byte of the data template with the `--fill-data` option.

10.8.3 Run-Time Library Support

In order to initialize variables in data memory, the data initialization template must be processed at start-up, before the application's main function takes control. For C programs, this task is performed by C start-up modules in the runtime library. Assembly language programs can also use the C start-up modules by linking with `libpic30-coff.a` or `libpic30-elf.a`.

Multiple versions of the start-up modules are contained within the runtime library. The linker will select a startup module based on commands in the linker script. For example:

```
CRT0_STARTUP(crt0_standard.o)
CRT1_STARTUP(crt1_standard.o)
```

For each device, two start-up modules are specified: a primary module (CRT0) and an alternate module (CRT1).

To utilize a start-up module, the application must allow the run-time library to take control at device Reset. This happens automatically for C programs. The application's `main()` function is invoked after the start-up module has completed its work. Assembly language programs should use the following naming conventions to specify which routine takes control at device Reset.

Table 10-3: Main Entry Points

Main Entry Name	Description
<code>__reset</code>	Takes control immediately after device Reset
<code>_main</code>	Takes control after the start-up module completes its work

Note that the first entry name (`__reset`) includes two leading underscore characters. The second entry name (`_main`) includes only one leading underscore character. The linker scripts construct a `GOTO __reset` instruction at location 0 in program memory, which transfers control upon device Reset.

The primary start-up module is linked by default and performs the following:

1. The stack pointer (W15) and stack pointer limit register (SPLIM) are initialized, using values provided by the linker or a custom linker script. For more information, see [Section 10.10 "Stack Allocation."](#)
2. If a `.const` section is defined, it is mapped into the PSV window by initializing the PSVPAG and CORCON registers. On devices which support EDS the DSRPAG register will be initialized. Note that a `.const` section is defined when the "Constants in code space" option is selected in MPLAB IDE, or the `-mconst-in-code` option is specified on the compiler command line.
3. The data initialization template in section `.dinit` is read, causing all uninitialized sections to be cleared, and all initialized sections to be initialized with values read from program memory.
4. If the application has defined `user_init` functions, section `.user_init` is called.
5. The function `main` is called with no parameters.
6. If `main` returns, the processor will reset.

The alternate start-up module is linked when the `--no-data-init` option is specified. It performs the same operations, except for step (3), which is omitted. The alternate start-up module is much smaller than the primary module, and can be selected to conserve program memory if data initialization is not required.

Source code for both modules is provided in the `src` directory of the MPLAB XC16 C compiler installation directory. The start-up modules may be modified if necessary. For example, if an application requires `main` to be called with parameters, a conditional assembly directive may be switched to provide this support.

10.9 Read-only Data

Read-only data sections are located in program memory, but are defined and accessed just like data memory. They are useful for storing constant tables that are too large for available data memory. The C compiler creates a read-only section named `.const` when the `-mconst-in-code` option is specified.

Access to read-only data sections is provided by means of the PSV window, or the EDS window. In either case, a reference to the read-only data is resolved to a data address within the PSV or EDS window.

C programmers can use the `space` attribute to allocate variables in read-only data sections. Access to such variables can be managed automatically by the compiler, or by explicit application code. For additional information on using read-only variables in C, refer to “*MPLAB® XC16 C Compiler User's Guide*” (DS50002071), Section 4.14 “Program Space Visibility (PSV) Usage” and Section 6.2 “Managed PSV Pointers”.

The `psv` section attribute is used to designate read-only data sections in assembly language. The contents of read-only data sections may be specified with data directives, as shown in the following assembly source example:

```
.section ronly,psv
L1: .byte 1
L2: .byte 2
```

In this example, section `ronly` will be allocated in program memory. Both byte constants will be located in the same program memory word, followed by a pad byte. Unlike other sections in program memory, read-only sections are byte addressable. Each label is resolved to a unique address that lies with the PSV or EDS address range.

The linker allocates read-only sections such that they do not cross a page boundary. Therefore, a single setting of the page register will access the entire section. A maximum length restriction is implied; the linker will issue an error message if any read-only data section exceeds 32 Kbytes. Only the least significant 16 bits of each instruction word are available for data storage (bits 16-23). The upper byte of each program word is filled with 0x0 or another value specified with the `--fill-upper` option. None of the p-variant assembler directives (including `.pbyte` and `.pword`) are permitted in read-only data sections.

16-Bit Assembler, Linker and Utilities User's Guide

The following examples illustrate how bytes in read-only sections may be accessed:

```
; example 1
mov    #psvpage(L1),w0
mov     w0,PSVPAG           ; set page register
mov     #psvoffset(L1),w0
mov     #psvoffset(L2),w1
mov.b   [w0],w2             ; load the byte at L1
mov.b   [w1],w3             ; load the byte at L2
```

```
; example 2
mov     #edspage(L1),w0
mov     w0,DSRPAG          ; set page register
mov     #edsoffset(L1),w0
mov     #edsoffset(L2),w1
mov.b   [w0],w2            ; load the byte at L1
mov.b   [w1],w3            ; load the byte at L2
```

User-defined read-only sections do not require a custom linker script. Based on the `psv` section attribute, the linker will locate the section in program memory and map its labels into the PSV or EDS window. If the programmer wishes to declare a read-only section in a custom linker script, the following syntax may be used:

```
/*
** User-Defined Constants in Program Memory
**
** This section is identified as a read-only section
** by use of the psv section attribute. It will be
** loaded into program memory and mapped into data
** memory using the PSV or EDS window.
**/
userconstants ADDR : AT (LOADADDR)
{
    *(userconstants);
} >program
```

In this example, `LOADADDR` specifies the load address in program memory.

It is not generally recommended to define read-only data sections in the linker script. This is because sections that appear in the linker script are allocated sequentially, and read-only data sections have significant page alignment restrictions. Because of these alignment restrictions, sequential allocation can fragment memory and result in less efficient memory utilization.

Likewise, it is not recommended to specify an absolute address for read-only data sections using attributes in source code. Absolute sections also fragment memory and can result in less efficient memory utilization.

10.10 Stack Allocation

The 16-bit device dedicates register W15 for use as a software stack pointer. All processor stack operations, including function calls, interrupts and exceptions, use the software stack. Upon Power-on or Reset, register W15 is initialized to point to a region of memory reserved for the stack. The stack grows upward, towards higher memory addresses.

The 16-bit device also supports stack overflow detection. If the stack limit register SPLIM is initialized, the device will test for overflow on all stack operations. If an overflow should occur, the processor will initiate a stack error exception. By default, this will result in a processor Reset. Applications may also install a stack error exception handler by defining an interrupt function named `__StackError`. See [Section 10.12 “Interrupt Vector Tables.”](#) for details.

By default, 16-bit linker allocates the largest stack possible from unused data memory. Therefore, care should be taken when assigning symbols to data sections so as to leave room for the stack (see [Section 4.2 “Directives that Define Sections.”](#)).

The location and size of the stack is reported in the link map output file, under the heading Dynamic Memory Usage. Applications can ensure that at least a minimum sized stack is available by using the `--stack` command option. For example:

```
xc16-ld -o t.exe t1.o --stack=0x100
```

While performing automatic stack allocation, 16-bit linker increases the minimum required size by a small amount to accommodate the processing of stack overflow exceptions. The stack limit register SPLIM is initialized to point just below this extra space, which acts as a stack overflow guardband. If not enough memory is available for the minimum size stack plus guardband, the linker will report an error.

The default stack guardband size is 16 bytes. Applications can specify a different size by using the `--stackguard` command option. For example:

```
xc16-ld -o t.exe t1.o --stackguard=32
```

As an alternative to automatic stack allocation, the stack may be allocated directly with a user-defined section in assembly language. For example:

```
.section my_stack, stack
.space 0x100
```

When the stack is allocated in this way, the usable stack space will be slightly less than 0x100 bytes, since a portion of the user-defined section will be reserved for the stack guardband.

An appropriate worst case stack usage requirement can be determined using the Stack Usage Guidance analysis tool, which can provide a compile-time analysis static of the end executable. For more on this tool, see the “*MPLAB XC16 C Compiler User’s Guide*” (DS50002071).

Regardless of how the stack is allocated (automatically or by user-defined section) the linker creates two symbols for use by the startup module. `__SP_init` defines the initial value for the stack pointer (W15), and `__SPLIM_init` defines the initial value for the stack limit register (SPLIM).

The start-up module uses these symbols to initialize the stack pointer and stack pointer limit register. Normally the start-up module is provided by `libpic30.a`. In special cases, the application may provide its own start-up code. The following stack initialization sequence may be used:

```
mov     #__SP_init,w15    ; initialize w15
mov     #__SPLIM_init,w0  ;
mov     w0,__SPLIM        ; initialize SPLIM
```

16-Bit Assembler, Linker and Utilities User's Guide

10.11 Heap Allocation

The 16-bit compiler standard C library, `libc.a`, supports dynamic memory allocation functions such as `malloc()` and `free()`. Applications which utilize these functions must instruct the linker to reserve a portion of 16-bit data memory for this purpose. The reserved memory is called a heap.

Applications can specify the heap size by using the `--heap` command option. For example:

```
xc16-ld -o t.exe t1.o --heap=0x100
```

While performing automatic heap allocation, the linker allocates the heap from unused data memory. The heap size is always specified by the programmer. In contrast, the linker sets the stack size to a maximum value, utilizing all remaining data memory.

As an alternative to automatic heap allocation, the heap may be allocated directly with a user-defined section in assembly source code. For example:

```
.section my_heap, heap  
.space 0x100
```

The location and size of the heap are reported in the link map output file, under the heading Dynamic Memory Usage. If the requested size is not available, the linker reports an error.

DD 10.12 Interrupt Vector Tables

dsPIC30F/33F DSC and PIC24F/H MCU devices have two interrupt vector tables - a primary and an alternate table, each containing exception vectors, as well as a `RESET` instruction at location zero. By convention, the linker initializes the `RESET` instruction and interrupt vector tables automatically, using information provided in the standard linker scripts.

The 16-bit compiler provides a special syntax for writing interrupt handlers. See the “*MPLAB® XC16 C Compiler User's Guide*” (DS50002071) for more information.

Assembly language programmers can install interrupt handlers simply by following the standard naming conventions. Interrupt handlers declared with the standard names and defined as globals are automatically installed into the vector tables.

By convention, the entry point named `__reset` takes control at device Reset. All applications written in assembly language must include a Reset function with this name. For C programs, the Reset function is provided in `libpic30`, which initializes the C run-time environment.

Note: Applications may provide a default interrupt handler, which will be installed into any unused vector table entries. In assembly language, the name of the default interrupt handler is `__DefaultInterrupt`. In C the name is `_DefaultInterrupt`.

If the application does not provide a default interrupt handler, the linker will create one in section `.isr` that contains a `reset` instruction. Creation of a default interrupt handler by the linker may be suppressed with the `--no-isr` option. In that case unused slots in the interrupt vector tables will be filled with zeros.

10.12.1 Interrupt Handler Example

The following example provides a Reset function and a default interrupt handler in assembly language. The default interrupt handler uses persistent data storage to keep a count of unexpected interrupts and/or error traps.

```
.include "p30f6014.inc"
.text

.global __reset
__reset:
    ;; takes control at device reset/power-on
    mov    #__SP_init,w15        ; initialize stack pointer
    mov    #__SPLIM_init,w0      ; and stack limit register
    mov    w0,SPLIM              ;

    btst   RCON,#POR             ; was this a power-on reset?
    bra    z,start              ; branch if not

    clr    FaultCount            ; else clear fault counter
    bclr   RCON,#POR             ; and power-on bit
start:
    goto   main                  ; start application

.global __T1Interrupt
__T1Interrupt:
    ;; services timer 1 interrupts
    bclr   IFS0,#T1IF           ; clear the interrupt flag
    retfie                          ; and return from interrupt

.global __DefaultInterrupt
__DefaultInterrupt:
    ;; services all other interrupts & traps
    inc    FaultCount            ; increment the fault counter
    reset                          ; and reset the device

    .section .pbss,persist      ; persistent data storage
    .global FaultCount          ; is not affected by reset
FaultCount:
    .space 2                    ; count of unexpected interrupts
```

The standard naming conventions for interrupt handlers are described in the sections below.

Note: The compiler requires only one leading underscore before any of the interrupt handler names. The assembler requires two leading underscores before any of the interrupt handler names. The compiler format is shown in tables in the following sections.

10.12.2 Interrupt Tables Location

For tables of interrupt vectors by device family:

- In MPLAB X IDE, for newer versions of the compiler, open the Dashboard window and click on the **Compiler Help** button.
- On the command-line, see the docs subdirectory of the MPLAB XC16 C compiler install directory (e.g., C:\Program Files (x86)\Microchip\xc16\v1.25). Open the XC16MasterIndex file and click on the "Interrupt Vector Tables Reference" link.

16-Bit Assembler, Linker and Utilities User's Guide

10.13 Optimizing Memory Usage

For memory intensive applications, it is often necessary to optimize memory usage by reducing or eliminating any unused gaps. The linker will optimize memory allocation automatically in most cases. However, certain constructs in source code and/or linker scripts may introduce gaps and should be avoided.

Memory gaps generally fall into the following categories:

- [Gaps Between Variables of Different Types](#)
- [Gaps Between Aligned Variables](#)
- [Gaps Between Input Sections](#)
- [Gaps Between Output Sections](#)

10.13.1 Gaps Between Variables of Different Types

Gaps may be inserted between variables of different types to satisfy address alignment requirements. For example, the following sequence of C statements will result in a gap:

```
char c1;
int i;
char c2;
int j;
```

Because the processor requires integers to be aligned on a 16-bit boundary, a padding byte was inserted after variables `c1` and `c2`. To eliminate this padding, variables of the same type should be defined together, as shown:

```
char c1,c2;
int i,j;
```

Gaps between variables are not visible to the linker, and are not reported in the link map. To detect these gaps, an assembly listing file must be created. The following procedure can be used:

1. If the source file is written in C, specify the `-save-temps` command line option to the compiler. This will cause an assembly version of the source file to be saved in `filename.s`.

```
xc16-gcc test.c -save-temps
```

2. Specify the `-ai` listing option to the assembler. This will cause a table of section information to be generated.

```
xc16-as test.s -ai
```

SECTION INFORMATION:

Section	Length (PC units)	Length (bytes) (dec)
-----	-----	-----
.text	0	0 (0)

TOTAL PROGRAM MEMORY USED (bytes): 0 (0)

Section	Alignment Gaps	Length (bytes) (dec)
-----	-----	-----
.data	0	0 (0)
.bss	0	0 (0)
.nbss	0x2	0x8 (8)

TOTAL DATA MEMORY USED (bytes): 0x8 (8)

In this example, 2 bytes of unused memory were inserted into section `.nbss`. Gaps between ordinary C variables will not exceed 1 byte per variable.

10.13.2 Gaps Between Aligned Variables

Variables may be defined in C with the `aligned` attribute in order to specify special alignment requirements for modulo addressing or other purposes. Use of the `aligned` attribute will cause the variable to be allocated in a unique section. Since a unique section is never combined with other input sections, no alignment padding is necessary and the linker will allocate memory for the aligned variable in the most efficient way possible.

For example, the following sequence of C statements will not result in an alignment gap, because variable `buf` is allocated in a unique section automatically:

```
char c1,c2;
int i,j;
int __attribute__((aligned(256))) buf[128];
```

When allocating space for aligned variables in assembly language, the source code must also specify a section name. Unless the aligned variable is defined in a unique section, alignment padding may be inserted. For example, the following sequence of assembly statements would result in a large alignment gap, and should be avoided:

```
.section my_vars,bss
.global _var1,_var2,_buf
_var1: .space 2
_var2: .space 2
; location counter is now 4
.align 256
_buf: .space 256
; location counter is now 512
```

Re-ordering the statements so that `_buf` is defined first will not eliminate the gap. A named input section will be padded so that its length is a multiple of the requested alignment. This is necessary in order to guarantee correct alignment when multiple input sections with the same name are combined by the linker. Therefore reordering statements would cause the gap to move, but would not eliminate the gap.

Aligned variables in assembly must be defined in a unique section in order to avoid alignment padding. It is not sufficient to specify a section name that is used only once, because the assembler does not know if that section will be combined with others by the linker. Instead, the special section name `*` should be used. As explained in [Section 4.2 “Directives that Define Sections.”](#), the section name `*` instructs the assembler to create a unique section that will not be combined with other sections.

To avoid alignment gaps, the previous example could be written as:

```
.section my_vars,bss
.global _var1,_var2
_var1: .space 2
_var2: .space 2

.section *,bss
.global _buf
.align 256
_buf: .space 256
```

The alignment requirement for `_buf` could also be specified in the `.section` directive, as shown:

```
.section *,bss,align(256)
.global _buf
_buf: .space 256
```

10.13.3 Gaps Between Input Sections

Gaps between input sections are similar to gaps between aligned variables, except that the padding is inserted by the linker, not the assembler. This type of gap can occur when variables with different alignment requirements are defined in separate source files.

A necessary condition for the insertion of alignment gaps by the linker is explicit mapping of input sections in the linker script. For example, older versions of the 16-bit compiler (prior to version 1.30) included the following definition:

```
/*
** Initialized Data and Constants
*/
.data :
{
    *(.data);
    *(.dconst);
} >data
```

This example maps all input sections named `.data` and all input sections named `.dconst` into a single output section. The various input sections will be combined sequentially. If the alignment requirement of any section exceeds that of the previous section, the linker will insert padding as needed and report an alignment gap in the link map:

Data Memory Usage

section	address	alignment gaps	total length (dec)
-----	-----	-----	-----
.data	0x800	0x10	0x90 (144)

Total data memory used (bytes): 0x90 (144) <1%

The remedy for this type of gap is to simply eliminate the mapping of input sections in linker scripts. Unmapped sections are allocated individually by the linker, so that no special alignment padding is necessary. Newer versions of the 16-bit compiler (version 1.30 and later) do not explicitly map any input sections in data memory for this reason.

10.13.4 Gaps Between Output Sections

Gaps between output sections can occur when the alignment requirements differ and multiple sections are allocated sequentially into the same memory region.

A necessary condition for the insertion of alignment gaps between output sections is explicit mapping of output sections in the linker script. For example, older versions of the 16-bit compiler (prior to version 1.30) included the following definitions:

```
/*
** Persistent Data
*/
.pbss (NOLOAD) :
{
    *(.pbss);
} >data

/*
** Static Data
*/
.bss (NOLOAD) :
{
    *(.bss);
} >data
```

This example creates two output sections (`.pbss` and `.bss`) and maps them into memory region `data`. Because the output sections are allocated sequentially, any difference in alignment requirements will result in gap.

In some instances the linker will make use of this gap, depending on the availability, size, and alignment requirements of any unmapped sections. In general it is preferable to eliminate the explicit mapping of output sections in linker scripts. When all output sections are unmapped, the linker is free to perform a best-fit allocation based on section attributes.

One consequence of best-fit allocation is that gaps between output sections may appear in unexpected places. The linker tries to use small memory blocks first, and will locate sections to leave the largest unused portions. When memory is segmented, such as by the introduction of an absolute section, the arrangement in memory may change dramatically. This should not be a problem unless the programmer expects a certain area of memory to remain unused. In such cases the programmer should reserve memory explicitly, using an array definition in source code, or by editing the linker script.

Explicit mapping of output sections in linker scripts is recommended only when the proximity or relative ordering of sections is important, and can't be satisfied using the section attributes described in [Section 4.2 "Directives that Define Sections."](#)

10.14 Boot and Secure Segments

The linker supports boot, secure, and general segments as described in the “*CodeGuard™ Security Reference Manual*” (DS70000180). The security model which includes segment sizes and configuration options may be specified in multiple ways. The linker allocates memory according to this security model and supports independent linking of application segments.

10.14.1 Specifying the Security Model

The application security model (including the sizes of various secure segments in FLASH, RAM, and EEDATA) can be specified in two ways:

1. In source code using macros currently defined for the FBS, FSS, FGS configuration words. See processor-specific include files for details and examples.
2. Using linker command options (see [Section 8.8 “Options that Specify CodeGuard™ Security Features.”](#)).

If both methods are used to provide conflicting information, the linker will issue a diagnostic. Likewise, a diagnostic will be issued if a security model is specified that can not be supported by the target device. The security model will be encoded by the linker into the executable file as contents for the FBS, FSS, and FGS configuration words.

A summary of CodeGuard Security options and segment sizes is written to the link map file. For example:

```
Selected CodeGuard Options:
  FBS:BSS:STRD_SMALL_BOOT_CODE
  FSS:SSS:STRD_SMALL_SEC_CODE
CodeGuard FLASH Memory:
  boot      0x100 to 0x3fe
  secure    0x400 to 0x1ffe
  general   0x2000 to 0x17ffe
CodeGuard RAM Memory:
  general   0x800 to 0x279f
  secure    (none)
  boot      (none)
```

10.14.2 User-Defined Boot and Secure Segments

User-defined `boot` and `secure` segments are supported in program memory and data memory. This allows an application to take advantage of the CodeGuard Security language extensions on any device, not just CodeGuard Security-enabled devices. User-defined segments are specified with the `ram_size` and `flash_size` options (see [Section 8.8 “Options that Specify CodeGuard™ Security Features.”](#)).

A summary of user-defined boot and secure segments is written to the link map file. For example:

```
User-Defined CodeGuard Segments
  boot RAM:           0x20 bytes
  secure RAM:         0x80 bytes
CodeGuard FLASH Memory:
  boot   (none)
  secure (none)
  general 0x100 to 0x17ffe
CodeGuard RAM Memory:
  general 0x800 to 0x26ff
  secure 0x2700 to 0x277f
  boot   0x2780 to 0x279f
```

User-defined segment options should not be combined with CodeGuard Security options. They are intended for debugging and/or special bootloader applications. User-defined segment options are not encoded in the FBS, FSS, FGS configuration words.

10.14.3 Boot and Secure Segment Allocation

The linker will collect input sections designated as `boot` or `secure` and allocate them according to the security model. Diagnostics will be issued for errors such as overflow of a secure segment, or requests for a type of protected memory that does not match the security model.

The linker reserves memory for boot and secure segments by adjusting boundaries of the following memory regions: program, data, and eedata. Therefore the name, origin, and length of these regions expressed in the linker script should reflect the original values, not values adjusted for boot and secure segments.

Note: Only sections explicitly designated as `boot` or `secure` will be allocated in the `boot` and `secure` segments. For independently linked applications, `boot` and `secure` functions must not call any library functions, or have any section dependencies that are not explicitly designated as `boot` or `secure`.

If access entry points have been defined, the linker will construct branch tables as needed for the boot or secure segment. Branch tables fill the entire access area (32 instruction words), regardless of how many access entry slots are actually used. This ensures that secure segment object code can be reached only by access entry point. Unused slots in the branch table will be filled with the default entry if one has been specified.

Execution flow may reach access entry points in several different ways, using a combination of machine instructions and data directives. Each access entry consists of a single, unconditional branch instruction, which targets the actual object code for a secure function.

10.14.4 Resolving Symbols

Symbol references within CodeGuard Security segments, and between CodeGuard Security segments, will be processed normally. If access entry points have been specified in a code address reference or in a function call reference, they will be resolved to specific offsets in the access entry tables. This mechanism allows the linker to resolve references to `boot` or `secure` functions that are defined only in terms of their access entry slot number, and is the key to supporting independently-linked applications.

Interrupt service routines designated for the `boot` or `secure` segments will be installed as a vector in slot 16 of the appropriate segment. Unused slots in the access entry tables are resolved to the `unused` function handler if one has been defined.

Note: The linker implements the security model in terms of memory allocation, but does not enforce a security policy. For example, references to a function defined in a secure segment from a lower privileged segment are permitted. Therefore it is possible to successfully link an application that fails at runtime due to CodeGuard Security hardware protection. This should be a relatively uncommon occurrence, since in practice strict CodeGuard Security protection implies independently-linked application segments.

10.14.5 Example of Simple Bootloader Application

A simple bootloader might look like this:

```
#include <xc.h>
#include "bootloader.h"

volatile int safe_to_continue = 0;

void __attribute__((interrupt)) _T1Interrupt(void) {
    safe_to_continue = 1;
    IFS0bits.T1IF = 0;
}

/*
 * will be filled in by target application
 */
void (*volatile startup_location)(void)
__attribute__((section("startup"), shared));

main() {

    IPC0bits.T1IP = 6;
    IEC0bits.T1IE=1;
    PR1 = 1000;
    T1CONbits.TON = 1;
    while (safe_to_continue == 0);
    T1CONbits.TON = 0;

    // presumably we have no communications request, and can proceed
    if (startup_location) startup_location();
}
```


The variable `startup_location` is shared and defined in the bootloader. The end application may redefine the contents of the variable, but not the address. The target application will provide an over-ridden definition of this value to be the address of its startup location. For a C application, this will be the function `_reset`.

```
#include <xc.h>
#include <stdio.h>
#include "bootloader.h"

extern void _reset(void);

void (*volatile startup_location)(void)
__attribute__((section("startup"),shared)) = &_reset;

main() {
    fprintf(stderr, "Hello world\n");
}
```

The rest of the bootloader application could be filled-in to accept some communication and re-program the target. When the whole system starts, the NULL initialization of `startup_location` may be over-written by the target application.

10.15 Co-resident Application Linking

Co-resident applications are programs that share the same physical memory space on an MCU or DSC. These applications are linked together in such a way that they can share the device memory resource.

10.15.1 Associated Options

There are several optional controls:

- There is only one reset vector and one set of fuses. Only one application can control these locations. The linker command line option `--coresident` asks the linker not to fill these locations, and once they are filled in it would be a link error to try and define them again. See [Section 8.4.5 “--coresident.”](#)
- There is still a restriction that there be only one constant section for the const-in-code memory model, which must be shared amongst all linked applications. It is possible to reserve some space for future growth by using the option `--reserve-const=size`. This option will set aside up to size total bytes of space to be used for the `.const` section in a future link. Remember the total size of the `.const` section cannot be larger than 32K bytes. This option must be used if const memory space is being used in a co-resident application domain. See [Section 8.4.34 “--reserve-const=size.”](#)
- Padding general FLASH areas. There are already methods to add padding to prelinked, input sections. However, this could be wasteful if all we need to do is ensure that the contiguous allocation ends on a particular boundary. The option `--pad-flash=size` will ensure that the linked output section is padded to a size byte boundary. See [Section 8.4.29 “--pad-flash=size.”](#)
- The linker option `--no-isr` has been extended to prevent the linker from completing the vector table. With this option specified, the linked executable will contain entries only for vector slots that are used in the application. The others will remain unfilled and can be defined in a future link. It is safer, in a final application, to have all vector slots completely defined. Without the option `--no-isr`, the linker will fill unused slots to call the `_DefaultInterrupt(void)` handler function, which it will define if none exists. See [Section 8.4.17 “--no-isr.”](#)
- The linker option `--no-isr` instructs the linker not to generate an IVT or AIVT, unless one is explicitly created in the linker script or by other means. See [Section 8.4.19 “--no-ivt.”](#)
- The command line option `--application-id=name` causes the linker to create alias names for each external symbol. The provided name is prepended to the normal symbol name; name should be C appropriate. See [Section 8.4.2 “--application-id=name.”](#)
- The command line option `--memory-usage` causes the linker to create information about the static memory usage for the linked application space. This information is represented in the executable as a null terminated sequence of pairs of start and end addresses for Flash usage and RAM usage. Flash usage information is placed in a section named `.flash_usage`, RAM usage is in `.ram_usage`. See [Section 8.6.4 “--memory-usage.”](#) and [Section 10.16 “Notable Symbols.”](#)

10.15.2 Associated Attributes

An attribute, `shared`, can be applied to a function or data in C or a section in assembly to indicate that the entry may be used outside of the application. A data item will be initialized at startup of any application in the co-resident set.

10.15.3 Co-resident Usage Restrictions

When using co-resident applications:

- A function may be shared between co-resident applications, but calling such a function should only be attempted if you are sure that all data is initialized, i.e., the called function should only reference shared or constant data.
- Only one co-resident application can fill in each vector slot.
- Each co-resident application will share the same const-in-code Flash PSV page.
- A co-resident application should have a stack provided. For a non-coresident application, the language tool will select the largest block of free space to be the stack. In a co-resident application, this should be defined by the programmer. Fortunately this is easily done with a simple assembly file. The following will reserve 1024 bytes for the stack:

```
.section *,stack  
.space 1024
```

16-Bit Assembler, Linker and Utilities User's Guide

10.16 Notable Symbols

The following symbols are defined by the linker and may be useful in code development.

__DATA_LENGTH **__CODE_LENGTH**

Description:	Symbols that represent the maximum length of their respective data sections.
Include:	libpic30.h
Prototype:	<pre>extern int __DATA_LENGTH; extern int __CODE_LENGTH;</pre>
Remarks:	These symbols are defined in the default linker scripts. They are treated like assembler equates but can be used from C.
Default Behavior:	The address of the symbol (its value in equate terms) represents the maximum length of the data section.
Example:	

_PROGRAM_END

Description:	A symbol defined in program memory to mark the highest address used by a CODE or PSV section.
Include:	libpic30.h
Prototype:	<pre>__attribute__((space(prog))) int _PROGRAM_END</pre>
Remarks:	In C, the symbol should be referenced with the address operator (&), as in a built-in function call that accepts the address of an object in program memory. Also, this symbol can be used by applications as an end point for checksum calculations. In assembly language, it should be referenced with an extra underbar character in the prefix.
Default Behavior:	The highest address used by a CODE or PSV section.
Examples:	<p>C code:</p> <pre>__builtin_tblpage(&_PROGRAM_END) __builtin_tbloffset(&_PROGRAM_END) _prog_addressT big_addr; _init_prog_address(big_addr, &_PROGRAM_END)</pre> <p>Assembly code:</p> <pre>mov #tblpage(__PROGRAM_END),w0 mov #tbloffset(__PROGRAM_END),w1 .pword __PROGRAM_END .long __PROGRAM_END</pre>

Chapter 11. Linker Examples

11.1 Introduction

The 16-bit devices include many architectural features that require special handling by the linker. The 16-bit compiler and assembler each provide a syntax that can be used to designate certain elements of an application for special handling. In C, a rich set of attributes are available to modify variable and function definitions (see the “*MPLAB XC16 C Compiler User's Guide*” - DS50002071). In assembly language, variables and functions are abstracted into memory sections, which become inputs to the linker. The assembler provides another set of attributes that are available to modify section definitions (see [Section 4.7 “Directives that Modify Section Alignment”](#)).

This chapter includes a number of 16-bit specific linker examples and shows the equivalent syntax in C and assembly language.

11.2 Highlights

Topics covered in this chapter are:

- [Memory Addresses and Relocatable Code](#)
- [Locating a Variable at a Specific Address](#)
- [Locating a Function at a Specific Address](#)
- [Using More than 32K of Constants](#)
- [Locating a Constant at a Specific Address in Program Memory](#)
- [Locating and Accessing Data in EEPROM Memory](#)
- [Creating an Incrementing Modulo Buffer in X Memory](#)
- [Creating a Decrementing Modulo Buffer in Y Memory](#)
- [Locating the Stack at a Specific Address](#)
- [Locating and Reserving Program Memory](#)

16-Bit Assembler, Linker and Utilities User's Guide

11.3 Memory Addresses and Relocatable Code

For most applications it is preferable to write fully relocatable source code, thus allowing the linker to determine the exact addresses in memory where functions and variables are placed. The final address of external symbols in data memory and program memory can be determined from the link map output, as shown in this excerpt:

```
...
External Symbols in Data Memory (by address):

0x0802      __curbrk
0x0804      _Stdin
0x082c      _Stdout
0x0854      _Stderr
0x087c      _Files
0x088c      _Aldata
0x0890      _Size_block

...
External Symbols in Data Memory (by name):

0x0802      __curbrk
0x088c      _Aldata
0x087c      _Files
0x0890      _Size_block
0x0854      _Stderr
0x0804      _Stdin
0x082c      _Stdout

...
```

In some cases it is necessary for the programmer to specify the address where a certain variable or function should be located. Traditionally this is done by creating a user-defined section and writing a custom linker script. The 16-bit assembler and compiler provide a set of attributes that can be used to specify absolute addresses and memory spaces directly in source code. When these attributes are used, custom linker scripts are not required.

<p>Note: By specifying an absolute address, the programmer assumes the responsibility to ensure the specified address is reasonable and available. If the specified address is out of range, or conflicts with a statically allocated resource, a link error will occur.</p>

11.4 Locating a Variable at a Specific Address

In this example, array `buf1` is located at a specific address in data memory. The address of `buf1` can be confirmed by executing the program in the simulator, or by examining the link map.

```
#include "stdio.h"
int __attribute__((address(0x900))) buf1[128];
void main()
{
    printf("0x900 = 0x%x\n", &buf1);
}
```

The equivalent array definition in assembly language appears below. The `.align` directive is optional and represents the default alignment in data memory. Use of `*` as a section name causes the assembler to generate a unique name based on the source file name.

```
        .section *,address(0x900),bss,near
        .global  _buf1
        .align   2
_buf1:  .space   256
```

11.5 Locating a Function at a Specific Address

In this example, function `func` is located at a specific address. Two built-in compiler functions are used to calculate the program memory address, which is not otherwise available in C.

```
#include "stdio.h"
void __attribute__((address(0x2000))) func()
{}
void main()
{
    long addr;

    addr = ((long) __builtin_tblpage(func) << 16)
           + __builtin_tbloffset(func);
    printf("0x2000 = 0x%lx\n", addr);
}
```

The equivalent function definition in assembly language appears below. The `.align` directive is optional and represents the default alignment in program memory. Use of `*` as a section name causes the assembler to generate a unique name based on the source file name.

```
        .section *,address(0x2000),code
        .global  _func
        .align   2
_func:  return
```

11.6 Using More than 32K of Constants

By default, the compiler collects const-qualified variables and string literals into a compiler managed section named `.const`. This section is allocated in program memory, and is mapped into data memory by means of the Program Space Visibility (PSV) window, or the Extended Data Space (EDS) window. Variables may be explicitly assigned to this section with the `space(auto_psv)` attribute.

Because `.const` is a PSV-type section, it is limited to 32K of total constants. To use more constants, variables may be assigned to other sections with the `space(psv)` attribute. This attribute causes the variable to be allocated in a program memory section that is designated for use with the PSV or EDS window.

For example:

```
const int __attribute__((space(psv))) table1[] =
    { 1, 2, 3, /* and so on */ };
```

`space(psv)` specifies the allocation of the variable, but it does not describe how the variable will be accessed. In order to access variables in `space(psv)`, the PSV or EDS page register must be managed so that the correct range of program memory is visible. Two options for managing the page register are available: compiler-managed access, or user-managed access.

11.6.1 Compiler-Managed Access

With this option, the compiler generates additional instruction as needed to save, set, and restore the PSV or EDS window page register. To specify compiler-managed access, add the `__psv__` access qualifier to the variable definition. For example:

```
__psv__ const int __attribute__((space(psv))) table1[] =
    { 1, 2, 3, /* and so on */ };
```

The `__psv__` access qualifier works with any variable allocated in `space(psv)`. It can be used on any 16-bit device, and directs the compiler to generate code automatically for managing the PSV or EDS window page register.

11.6.2 User-Managed Access

User-managed access means that the programmer must write explicit code to save, set, and restore the PSV or EDS window page register. In certain situations, this could result in faster execution speed.

In the following example, the constant `status_string` is located in the compiler-managed PSV section, while the constant `gamma_factor` is located in a separate PSV section.

Note: To modify this example to run on a device which supports the EDS window, replace references to PSVPAG with DSRPAG.

The compiler will initialize the page register only for the compiler-managed PSV section on startup. To properly access `gamma_factor`, you must manually manage the page register. Namely, save the current page value, set the page register to access `gamma_factor`, and restore the original page value after. To determine the correct page value for a constant stored in program memory, use the `__builtin_psvpage()` helper function.

When the page register has been modified to access `gamma_factor`, be careful not to access constants stored in the compiler-managed PSV section, such as string constants used with `printf()`. Any attempts to access constants stored in the compiler-managed PSV section with an incorrect page value will fail.

Note: On devices with less than 16K instruction words, there is only one page and manual management of the page register is not required.

```
#include "stdio.h"
#include "p30fxxx.h"

const char __attribute__((space(auto_psv))) status_string[2][10] =
{"System OK", "Key Made"};
const int __attribute__((space(psv))) gamma_factor[3] = {13, 23, 7};

int main(void)
{
    unsigned psv_shadow;
    unsigned key, seed = 17231;

    /* print the first status string */
    printf ("%s\n", status_string[0]);

    /* save the PSVPAG */
    psv_shadow = PSVPAG;

    /* set the PSVPAG for accessing gamma_factor[] */
    PSVPAG = __builtin_psvpage (gamma_factor);

    /* build the key from gamma_factor */
    key = (seed + gamma_factor[0] + gamma_factor[1]) / gamma_factor[2];

    /* restore the PSVPAG for the compiler-managed PSVPAG */
    PSVPAG = psv_shadow;

    /* print the second status message */
    printf ("%s \n", status_string[1]);
}
```

16-Bit Assembler, Linker and Utilities User's Guide

11.7 Locating a Constant at a Specific Address in Program Memory

In this example, the constant `table` is located at a specific address in program memory. When a constant is specifically placed at an address in program memory, it must be placed in its own PSV section using the `space(psv)` attribute. If a device has only one PSV page (16K instruction words or less), the `(psv)` section and `(auto_psv)` section will share the same PSV page by default.

Note: It is not possible to place a constant at a specific address in Program Memory using the `space(auto_psv)` attribute. Only the `space(psv)` attribute may be used to perform this task.

The `__builtin_tbladdress()` helper function can be used to find the address of a constant stored in program memory. The `__psv__` access qualifier is used to specify compiler-managed access.

```
#include "stdio.h"
#include "p30fxxxx.h"

__psv__ const unsigned __attribute__((space(psv),
    address (0x2000))) table[10] = {0, 1, 2, 3, 4, 5, 6, 7, 8, 9};

int main(void)
{
    unsigned sum=0, u;
    long addr;

    /* compute the address of table and print it */
    addr = __builtin_tbladdress(table);

    /* print the address of table */
    printf ("table[] is stored at address 0x%x\n", addr);

    /* sum the values in table[] */
    for (u=0; u<10; u++) {
        sum += table[u];
    }

    /* print the sum */
    printf ("sum is %d\n", sum);
}
```

The equivalent constant definition for the array `table` in assembly language appears below. The `.align` directive is optional and represents the default alignment in program memory. Use of `*` as a section name causes the assembler to generate a unique name based on the source file name.

```
.section *,address(0x2000),psv
.global _table
.align 2
_table:
    .word 0,1,2,3,4,5,6,7,8,9
```

In order to allocate `table` in data memory, the `space(psv)` attribute could be changed to `space(data)`. In this case, the specified address would be a data memory address. In the absence of a `space` attribute, the keyword `const` directs the C compiler to allocate the variable in the same space as other compiler constants. Constants are allocated in program memory by default, or in data memory if the constants-in-data memory model is selected.

11.8 Locating and Accessing Data in EEPROM Memory

In this example, two arrays are defined in data EEPROM. `Table1` is aligned to a 32-bit address, so it will be eligible for erasing or programming using the row programming algorithm. `Table2` is defined with standard alignment, so it must be erased or programmed one word at a time. The macro `_EEDATA` is used to place a variable in the Data EEPROM section of memory and align the variable to the specified byte boundary. This macro is defined in the processor header files for devices which contain data flash. This example is targeted for the dsPIC30F6014 processor, and includes the processor header file `p30f6014.h`.

The compiler and linker treat Data EEPROM like any other custom-defined (`psv`) section. The `__psv__` access qualifier is used to instruct the compiler to generate the necessary instructions to manage the PSV or EDS page register automatically.

```
/* load SFR definitions and macros */
#include "p30f6014.h"

/* load standard I/O definitions */
#include "stdio.h"

__psv__ unsigned int _EEDATA(32) Table1[16];

__psv__ unsigned int _EEDATA(2) Table2[4]= {0x1234, 0x5678, 0x9ABC,
0xDEF0};

unsigned int i, temp_data[4];
__psv__ unsigned int *ee_rd_ptr;

int main( void )
{
    /* initialize EEPROM read pointer */
    ee_rd_ptr = &Table2[0];

    /* read integer data from EEPROM */
    temp_data[0] = *ee_rd_ptr++;
    temp_data[1] = *ee_rd_ptr++;
    temp_data[2] = *ee_rd_ptr++;
    temp_data[3] = *ee_rd_ptr;

    /* display it */
    for ( i = 0; i < 4; i++)
        printf(" %x", temp_data[i]);
    printf("\n");
}
```

16-Bit Assembler, Linker and Utilities User's Guide

The equivalent array definitions for `Table1` and `Table2` in assembly language appear below. Use of `*` as a section name causes the assembler to generate a unique name based on the source file name.

```
.global _Table1
.section *,eedata
.align 32
_Table1:
.space 32

.global _Table2
.section *,eedata
.align 2
_Table2:
.word 0x1234
.word 0x5678
.word 0x9ABC
.word 0xDEF0
```

11.9 Creating an Incrementing Modulo Buffer in X Memory

An incrementing modulo buffer for use in assembly language can be easily defined in C. In this example, the macro `_XBSS` is used to define an array whose memory alignment is the smallest power of two that is greater than or equal to its size. `_XBSS` is defined in the processor header file, which in this example is `p30f6014.h`.

```
#include "p30f6014.h"
#include "stdio.h"

int _XBSS(128) xbuf[50];

void main()
{
    printf("Should be zero: %x\n", (int) &xbuf % 128);
}
```

The equivalent definition in assembly language appears below. The section alignment could have specified with a separate `.align` directive. By using `*` as a section name, the linker is afforded maximum flexibility to allocate memory.

```
.global _xbuf
.section *,xmemory,bss,align(128)
_xbuf: .space 100
```

11.10 Creating a Decrementing Modulo Buffer in Y Memory

A decrementing modulo buffer for use in assembly language can be easily defined in C. In this case, the ending address +1 of the array must be aligned. There is not a suitable predefined macro in the processor header files for this purpose, so variable attributes are specified directly. The `far` attribute is recommended because Y memory does not fall within the near space on all devices, and the compiler uses a small-data memory model by default.

```
#include "stdio.h"

int __attribute__((space(ymemory), far, reverse(128))) ybuf[50];

void main()
{
    printf("Should be zero: %x\n",
        ((int) &ybuf + sizeof(ybuf)) % 128);
}
```

Notes:

- Note 1:** The `reverse()` attribute can be used with constants stored in program memory only if they are located in a PSV section, not the compiler-managed `auto_psv` section.
- 2:** The `reverse()` attribute can be used with constants stored in Data EEPROM memory.

The equivalent definition in assembly language appears below. Reverse section alignment can only be specified as an argument to the `.section` directive.

```
.global _ybuf
.section *,ymemory,reverse(128)
_ybuf: .space 100
```

11.11 Locating the Stack at a Specific Address

By default, the linker allocates a maximum-size stack using the largest unused block of data memory. In cases where it is necessary for the programmer to specify the location and size of the stack explicitly, the stack may be defined in assembly language, using the `stack` attribute:

```
.section my_stack, stack, address(0x1800) .space 0x100
```

When the stack is allocated in this way, the usable stack space will be slightly less than 0x100 bytes, since a portion of the user-defined section will be reserved for the stack guardband.

16-Bit Assembler, Linker and Utilities User's Guide

11.12 Locating and Reserving Program Memory

In this example, a block of program memory is reserved for a special purpose, such as a bootloader. An arbitrary sized function is allocated in the block, with the remaining space reserved for expansion or other purposes.

The following output section definition is added to a custom linker script:

```
BOOT_START = 0xA200;
BOOT_LEN = 0x400;

my_boot BOOT_START :
{
    *(my_boot);
    . = BOOT_LEN; /* advance dot to the maximum length */
} > program
```

Note the “dot assignment” (.=) that appears inside the section definition after the input sections. Dot is a special variable that represents the location counter, or next fill point, in the current section. It is an offset relative to the start of the section. The statement in effect says “no matter how big the input sections are, make sure the output section is full size.”

The following C function will be allocated in the reserved block:

```
void __attribute__((section("my_boot"))) func1()
{
    /* etc. */
}
```

The equivalent assembly language would be:

```
.section my_boot,code
.global _func1
_func1:
    ; and so on..
    return
```

If the bootloader is allocated at the start of program memory, a custom linker script is not be required. Instead, the function could be defined with attribute boot. For example:

```
void __attribute__((boot)) func1()
{
    /* and so on.. */
}
```

The equivalent definition in assembly language:

```
.section *,code,boot
.global _func1
_func1:
    ; and so on..
    return
```

In this case, program memory will be automatically reserved by specifying a CodeGuard Security™ boot segment in FBS configuration word settings, or by specifying a user-defined boot segment with linker command option. See [Section 10.14 “Boot and Secure Segments”](#) for more information.

Chapter 12. Linker Map File

12.1 Introduction

The linker has the capability to produce map files. These map files list archive files included, memory usage, external symbols, linker script information and memory maps.

Topics covered in this chapter are:

- [Generation](#)
- [Contents](#)

12.2 Generation

To generate a map file whether in MPLAB X IDE, in MPLAB IDE v8 or on the command line, you will need to specify an option described in [Section 8.7 "Options that Modify the Link Map Output."](#) By default, a map file is written to a `.map` file.

12.3 Contents

The map files produced by the linker consist of the following items.

Table 12-1: Linker Map File Contents

Item	Description
Tool Name and Command	Path and executable name of the linker, as well as command line options used.
Archive Members	The name of any members from archive files that are included in the link
Memory Usage Report	The starting address and length of all output sections in program memory, data memory and dynamic memory
External Symbol Table	All external symbols in data and program memory
Memory Configuration	All of the memory regions defined for the link
Linker Script and Memory Map	Modules, sections and symbols that are included in the link as specified in the linker script

16-Bit Assembler, Linker and Utilities User's Guide

Example 12-1: Map File

The following is an example of a linker map file for a PIC24FJ MCU project. Note that the `.debug_aranges` register list has been shortened to save space.

Microchip Technology Inc, v1.34 (9)

Tool Name:

c:\program files (x86)\microchip\xc16\v1.34\bin\bin\..\bin\elf-ld.exe

Command:

```
-p24FJ128GA010 \  
-o \  
dist/default/production/PIC24FJ_Count.X.production.elf \  
-Lc:/program files (x86)/microchip/xcl6/v1.34/bin/bin/.. \  
-Lc:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../lib \  
-Lc:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../support/PIC24E/gld \  
-Lc:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../support/PIC24F/gld \  
-Lc:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../support/PIC24H/gld \  
-Lc:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../support/dsPIC30F/gld \  
-Lc:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../support/dsPIC33C/gld \  
-Lc:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../support/dsPIC33E/gld \  
-Lc:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../support/dsPIC33F/gld \  
build/default/production/counter.o \  
build/default/production/timer.o \  
--defsym=__MPLAB_BUILD=1 \  
-Tp24FJ128GA010.gld.00 \  
--stack=16 \  
--check-sections \  
--data-init \  
--pack-data \  
--handles \  
--isr \  
--no-gc-sections \  
--fill-upper=0 \  
--stackguard=16 \  
--no-force-link \  
--smart-io \  
-Map=pic24fj_count.map \  
--report-mem \  
--memorysummary \  
dist/default/production/memoryfile.xml \  
-start-group \  
--library=lega-pic30-elf \  
--library=m-elf \  
--library=lega-c-elf \  
-end-group
```

Optional library libpPIC24Fxxx.a not found

Archive member included because of file (symbol)

c:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../lib

\liblega-pic30-elf.a(crt0_standard.o)

(__resetPRI)

c:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../lib

\liblega-pic30-elf.a(data_init_standard.o)

c:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../lib
\liblega-pic30-elf.a(crt0_standard.o) (__data_init_standard)

Linker Map File

```
c:/program files (x86)/microchip/xc16/v1.34/bin/bin/../../lib
  \liblega-pic30-elf.a(crt_start_mode_normal.Leo)
    (__crt_start_mode_normal)
```

xc16-ld 1.34 (9)

"program" Memory [Origin = 0x200, Length = 0x155fc]

section	address	length (PC units)	length (bytes)	(dec)
-----	-----	-----	-----	-----
.text	0x200	0xce	0x135	(309)
.text	0x2ce	0x56	0x81	(129)
.dinit	0x324	0x8	0xc	(12)

Total "program" memory used (bytes): 0x1c2 (450) <1%

"data" Memory [Origin = 0x800, Length = 0x2000]

section	address	alignment	gaps	total length	(dec)
-----	-----	-----	-----	-----	-----
.nbss	0x800		0	0x2	(2)

Total "data" memory used (bytes): 0x2 (2) <1%

Dynamic Memory Usage

region	address	maximum length	(dec)
-----	-----	-----	-----
heap	0	0	(0)
stack	0x802	0x1ffe	(8190)

Maximum dynamic memory (bytes): 0x1ffe (8190)

External Symbols in Data Memory (by address):

0x0800 _counter

External Symbols in Data Memory (by name):

0x0800 _counter

External Symbols in Program Memory (by address):

0x000200 __resetPRI
0x00023c __psv_init
0x000252 __data_init
0x000252 __data_init_standard
0x0002c8 __DefaultInterrupt
0x0002ce _TimerInit
0x0002ea _TimerIsOverflowEvent
0x000302 _main

External Symbols in Program Memory (by name):

0x0002ce _TimerInit
0x0002ea _TimerIsOverflowEvent

16-Bit Assembler, Linker and Utilities User's Guide

0x0002c8	__DefaultInterrupt
0x000252	__data_init
0x000252	__data_init_standard
0x00023c	__psv_init
0x000200	__resetPRI
0x000302	__main

Memory Configuration

Name	Origin	Length	Attributes
data	0x000800	0x002000	a !xr
reset	0x000000	0x000004	
ivt	0x000004	0x0000fc	
__reserved	0x000100	0x000004	
aivt	0x000104	0x0000fc	
program	0x000200	0x0155fc	xr
CONFIG2	0x0157fc	0x000002	
CONFIG1	0x0157fe	0x000002	
default	0x000000	0xffffffff	

Linker script and memory map

```
LOAD build/default/production/counter.o
LOAD build/default/production/timer.o
        0x0001          __MPLAB_BUILD = 0x1
LOAD pPIC24Fxxx
LOAD c:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../lib\libfx-elf.a
        0x157fc          __CONFIG2 = 0x157fc
        0x157fe          __CONFIG1 = 0x157fe
        0x0200          __CODE_BASE = 0x200
        0x155fc          __CODE_LENGTH = 0x155fc
        0x0004          __IVT_BASE = 0x4
        0x0104          __AIVT_BASE = 0x104
        0x0800          __DATA_BASE = 0x800
        0x2000          __DATA_LENGTH = 0x2000

.reset          0x000000          0x4
                0x000000          0x2 SHORT 0x200 ABSOLUTE (__reset)
                0x000001          0x2 SHORT 0x4
                0x000002          0x2 SHORT 0x0 ((ABSOLUTE (__reset) >> 0x10) &
0x7f)
                0x000003          0x2 SHORT 0x0

.text           0x000200          0xce
*(.init)
.init          0x000200          0x4c c:/program files (x86)/microchip/xcl6/v1.34
                /bin/bin/../../lib
                \liblega-pic30-elf.a(crt0_standard.o)
                0x000200          __resetPRI
                0x000200          __reset
                0x00023c          __psv_init
.init          0x00024c          0x7c c:/program files (x86)/microchip/xcl6/v1.34
                /bin/bin/../../lib
                \liblega-pic30-elf.a(data_init_standard.o)
                0x000252          __data_init_standard
                0x000252          __data_init

*(.user_init)
*(.handle)
*(.isr*)
.isr           0x0002c8          0x4 default_isr
```

Linker Map File

```

                                0x0002c8                __DefaultInterrupt

*(.libc)
*(.libm)
*(.libdsp)
*(.lib*)
.libpic30_crt_start_mode
                                0x0002cc                0x2  c:/program files (x86)/microchip/xcl6/v1.34
                                                                /bin/bin/../../lib\liblega-pic30-elf.a
                                                                (crt_start_mode_normal.Leo)
                                0x0002cc                __crt_start_mode
                                0x0002cc                __crt_start_mode_normal

usercode
*(usercode)

__CONFIG2                      0x0157fc                0x2
*(__CONFIG2.sec*)
__CONFIG2.sec                  0x0157fc                0x2  build/default/production/counter.o

__CONFIG1                      0x0157fe                0x2
*(__CONFIG1.sec*)
__CONFIG1.sec                  0x0157fe                0x2  build/default/production/counter.o

.comment
*(.comment)

.debug_info                    0x000000                0x70b
*(.debug_info)
.debug_info                     0x000000                0x246  build/default/production/counter.o
.debug_info                     0x000246                0x4c5  build/default/production/timer.o
*(.gnu.linkonce.wi.*)

.debug_abbrev                   0x000000                0x1d3
*(.debug_abbrev)
.debug_abbrev                   0x000000                0xc2  build/default/production/counter.o
.debug_abbrev                   0x0000c2                0x111  build/default/production/timer.o

.debug_line                     0x000000                0x1e9
*(.debug_line)
.debug_line                     0x000000                0xf4  build/default/production/counter.o
.debug_line                     0x0000f4                0xf5  build/default/production/timer.o

.debug_frame                   0x000000                0x7c
*(.debug_frame)
.debug_frame                    0x000000                0x30  build/default/production/counter.o
.debug_frame                    0x000030                0x4c  build/default/production/timer.o

.debug_str                     0x000000                0x12
*(.debug_str)
.debug_str                      0x000000                0x12  build/default/production/timer.o

.debug_loc
*(.debug_loc)

.debug_macinfo
*(.debug_macinfo)

.debug_pubnames
                                0x000000                0x7a
*(.debug_pubnames)
```

16-Bit Assembler, Linker and Utilities User's Guide

```
.debug_pubnames
0x000000      0x41 build/default/production/counter.o
.debug_pubnames
0x000041      0x39 build/default/production/timer.o

.debug_ranges
*(.debug_ranges)

.debug_aranges 0x000000      0x30
*(.debug_aranges)
.debug_aranges
0x000000      0x18 build/default/production/counter.o
.debug_aranges
0x000018      0x18 build/default/production/timer.o
0x0000      WREG0 = 0x0
0x0000      _WREG0 = 0x0
0x0002      WREG1 = 0x2
0x0002      _WREG1 = 0x2
0x0004      WREG2 = 0x4
0x0004      _WREG2 = 0x4
0x0006      WREG3 = 0x6
0x0006      _WREG3 = 0x6
0x0008      WREG4 = 0x8
0x0008      _WREG4 = 0x8
0x000a      WREG5 = 0xa
0x000a      _WREG5 = 0xa
:
0x0240      SPI1 = 0x240
0x0240      _SPI1 = 0x240
0x0260      SPI2 = 0x260
0x0260      _SPI2 = 0x260
0x0220      UART1 = 0x220
0x0220      _UART1 = 0x220
0x0230      UART2 = 0x230
0x0230      _UART2 = 0x230

START GROUP
LOAD c:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../lib\liblega-pic30-elf.a
LOAD c:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../lib\libm-elf.a
LOAD c:/program files (x86)/microchip/xcl6/v1.34/bin/bin/../../lib\liblega-c-elf.a
END GROUP
OUTPUT(dist/default/production/PIC24FJ_Count.X.production.elf elf32-pic30)
LOAD default_isr
LOAD data_init

.debug_pubtypes
0x000030      0x78
.debug_pubtypes
0x000030      0x1f build/default/production/counter.o
.debug_pubtypes
0x00004f      0x59 build/default/production/timer.o

__c30_signature
0x0000a8      0x18
__c30_signature
0x0000a8      0x6 build/default/production/counter.o
__c30_signature
0x0000ae      0x6 build/default/production/timer.o
__c30_signature
```

Linker Map File

	0x0000b4	0x6 c:/program files (x86)/microchip/xcl6/v1.34 /bin/bin/../../../../lib \liblega-pic30-elf.a(crt0_standard.o)
__c30_signature	0x0000ba	0x6 c:/program files (x86)/microchip/xcl6/v1.34 /bin/bin/../../../../lib \liblega-pic30-elf.a(data_init_standard.o)
.nbss	0x0800	0x2
.nbss	0x0800	0x2 build/default/production/counter.o
	0x0800	_counter
.text	0x0002ce	0x56
.text	0x0002ce	0x34 build/default/production/timer.o
	0x0002ce	_TimerInit
	0x0002ea	_TimerIsOverflowEvent
.text	0x000302	0x22 build/default/production/counter.o
	0x000302	_main
.dinit	0x000324	0x8
.dinit	0x000324	0x8 data_init

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Chapter 13. Linker Errors/Warnings

13.1 Introduction

MPLAB XC16 Object Linker generates errors and warnings. A descriptive list of these outputs is shown here.

For information on linker **limitations** and **known problems**, see the Readme file.

13.2 Highlights

The following topics covered in this appendix:

- [Errors](#)
- [Warnings](#)

13.3 Errors

Symbols

% by zero

Modulo by zero is not computable.

/ by zero

Division by zero is not computable.

A

A heap is required, but has not been specified.

A heap must be specified when using Standard C input/output functions.

Address 0x8 of *filename* section .reset is not within region reset.

This error indicates a problem with the linker script. Normally section `.reset` is created by the linker script and includes a single GOTO instruction. If a linker script is included in the link as an input file, it will augment the built-in script instead of replacing it. Then section `.reset` will be created twice, resulting in an overflow. To correct this error, specify `--script` or `-T` on the link command before the linker script file name.

Address *addr* of *filename* section *secname* is not within region *region*.

Section `secname` has overflowed the memory region to which it was assigned.

C

Cannot access symbol (*name*) with file register addressing. Value must be less than 8192.

`name` is not located in near address space. A read or write of `name` could not be resolved with the small data memory model.

Cannot access symbol (*name*) at an odd address.

Instructions that operate on word-sized data require operands to be allocated at even addresses.

16-Bit Assembler, Linker and Utilities User's Guide

cannot move location counter backwards (from *address1* to *address2*).

The location counter can be advanced but it cannot be moved backwards. An operation is attempting to move it from *address1* backwards to *address2*.

cannot open linker script file *name*.

Unable to open the specified linker script file. Check the file name and/or the path.

cannot open *name*:

Cannot open the input file *name*. Check for correct spelling, extension or path.

cannot PROVIDE assignment to location counter.

The PROVIDE keyword may not be used to make an assignment to the location counter.

Cannot use relocation type *reloc* on a symbol (name) that is located in an executable section.

An attempt was made to use a symbol in an executable section as a data address. To reference an executable symbol in a data context, the `psvoffset()` or `tbloffset()` operator is required.

Could not allocate data memory.

The linker could not find a way to allocate all of the sections that have been assigned to region 'data'.

Could not allocate program memory.

The linker could not find a way to allocate all of the sections that have been assigned to region 'program'.

Could not allocate eedata memory.

The linker could not find a way to allocate all of the sections that have been assigned to region 'eedata'.

Could not allocate section '*name*', because '*ymemory,near*' is not a valid combination on this device.

The linker could not allocate section name because the combination of section attributes [*ymemory,near*] is not valid on the current device.

Could not allocate section *secname* at address *addr*.

An address has been specified for *secname* that conflicts with another section or the limit of memory.

Could not allocate section '*section name*' it is illegal to use the last word of program memory

Using the last word of program memory is illegal and a link error will be generated if you attempt to place any code there.

D

Data region overlaps PSV window (%d bytes).

The data region address range must be less than the start address for the PSV window. This error occurs when the C compiler's "constants in code" option is selected and more than 32K of data memory is required for program variables.

--data-init and --no-data-init options can not be used together.

--data-init creates a special output section named `.dinit` as a template for the run-time initialization of data, --no-data-init does not. Only one option can be used.

__DMA_BASE is needed, but not defined (check linker script?)

__DMA_END is needed, but not defined (check linker script?)

The symbols `__DMA_BASE` and `__DMA_END` must be defined in order to allocate variables or sections in dma memory. By convention these symbols are defined in the linker script for a particular device, if that device supports dma memory.

E

EOF in comment.

An end-of-file marker (EOF) was found in a comment.

F

op forward reference of section *secname*.

The section name being used in the operation has not been defined yet.

G

--gc-sections and -r may not be used together.

Do not use `--gc-sections` option which enables garbage collection of unused input sections with the `-r` option which generates relocatable output.

H

--handles and --no-handles options cannot be used together.

`--handles` supports far code pointers; `--no-handles` does not. Only one option can be used.

I

includes nested too deeply.

`include` statements should be nested no deeper than 10 levels.

Illegal value for DO instruction offset (-2, -1 or 0).

These values are not permitted.

invalid assignment to location counter.

The operation is not a valid assignment to the location counter.

invalid hex number '*num*'.

A hexadecimal number can only use the digits 0-9 and A-F (or a-f). The number is identified as a hex value by using 0x as the prefix.

invalid syntax in flags.

The region attribute flags must be `w`, `x`, `a`, `r`, `i` and/or `l`. ('!' is used to invert the sense of any following attributes.) Any other letters or symbols will produce the invalid syntax error.

M

macros nested too deeply.

Macros should be nested no deeper than 10 levels.

missing argument to -m.

The emulation option (`-m`) requires a name for the emulation linker.

N

Near data space has overflowed by *num* bytes.

Near data space must fit within the lowest 8K address range. It includes the sections `.nbss` for static or non-initialized variables, and `.ndata` for initialized variables.

no input files.

The 16-bit linker requires at least one object file.

non constant address expression for section *secname*.

The address for the specified section must be a constant expression.

nonconstant expression for *name*.

name must be a constant expression.

non constant address expression specified. Section will be allocated at the current address in the current region.

If a load address is specified for a section in the linker script using the AT (symbol) expression and “symbol” is not defined, a warning will be generated and the section will be allocated at the current address in the current region.

Not enough contiguous memory for section *secname*.

The linker attempted to reallocate program memory to prevent a read-only section from crossing a PSV page boundary, but a memory solution could not be found.

Not enough memory for heap (*num* bytes available).

There was not enough memory free to allocate the heap.

Not enough memory for stack (*num* bytes available).

There was not enough memory free to allocate the minimum-sized stack.

O

object *name* was created for the *processor* which is not instruction set compatible with the target *processor*.

An object file to be linked was created for a different processor family than the link target, and the instruction sets are not compatible.

Odd values are not permitted for a new location counter.

When a `.org` or `.porg` directive is used in a code section, the new location counter must be even. This error also occurs if an odd value is assigned to the special DOT variable.

P

--pack-data and --no-pack-data options cannot be used together.

`--pack-data` fills the upper byte of each instruction word in the data initialization template with data. `--no-pack-data` does not. Only one option can be used.

PSV section *secname* exceeds 32 Kbytes (actual size = *num*).

The constant data table may not exceed the program memory page size that is implied by the PSVPAG register which is 32 Kbytes.

R

region *region* is full (*filename* section *secname*).

The memory region *region* is full, but section *secname* has been assigned to it.

--relax and -r may not be used together.

The option `--relax` which turns relaxation on may not be used with the `-r` option which generates relocatable output.

relocation truncated to fit: PC RELATIVE BRANCH *name*.

The relative displacement to function *name* is greater than 32K instruction words. A function call to *name* could not be resolved with the small code memory model.

relocation truncated to fit: *relocation_type* *name*.

The relocated value of *name* is too large for its intended use.

S

section *.handle* must be allocated low in program memory.

A custom linker script has organized memory such that section *.handle* is not located within the first 32K words of program memory.

section *secname1* [*startaddr1*—*startaddr2*] overlaps section *secname2* [*startaddr1*—*startaddr2*]\n”),

There is not enough region memory to place both of the specified sections or they have been assigned to addresses that result in an overlap.

-shared not supported.

The option `-shared` is not supported by the 16-bit linker.

Symbol (*name*) is not located in an executable section.

An attempt was made to call or branch to a symbol in a bss, data or readonly section.

syntax error.

An incorrectly formed expression or other syntax error was encountered in a linker script.

U

undefined symbol ‘*__reset*’ referenced in expression.

The library `-lpic30` is required, or some other input file that contains a start-up function. This error may result from a version or architecture mismatch between the linker and library files.

undefined symbol ‘*symbol*’ referenced in expression.

The specified symbol has not been defined.

undefined reference to ‘*_Ctype*’.

undefined reference to ‘*_Tolotab*’.

undefined reference to ‘*_Touptab*’.

These errors indicate a version mismatch between include files and library files, or between library files and precompiled object files. Make sure that all object files to be linked have been compiled with the same version of the 16-bit compiler. If you are using a precompiled object or library file from another vendor, request an update that is compatible with the latest version of the compiler.

undefined reference to '*symbol*.'

The specified symbol has not been defined. Either an input file has been omitted, a library file is incomplete, a library file requires a symbol from an earlier library, or a circular reference exists between libraries. Circular references can be resolved with the `--start-group`, `--end-group` options.

unrecognized emulation mode: *target***Supported emulations:**

The specified target is not an emulation mode supported by the linker. The list of supported emulations follows the error message.

unrecognized -a option '*argument*.'

The `-a` option is not supported by 16-bit devices; so it is ignored.

unrecognized -assert option '*option*.'

The `-assert` option is not supported by 16-bit devices; so it is ignored.

unrecognized option '*option*'.

The specified option is not a recognized linker option. Check the option and its usage information with the `--help` option.

***op* uses undefined section *secname*.**

The section referred to in the operation is not defined.

X

X data space has overflowed by *num* bytes.

The address range for X data space must be less than the start of Y data space. The start of Y data space is determined by the processor used.

Y

`__YDATA_BASE` is needed, but not defined.

By convention, the starting address of Y data memory for a particular device is defined in linker scripts using this name. The linker needed this information to allocate a section with `xmemory` or `ymemory` attribute, but could not find it.

13.4 Warnings

A

Addresses specified for READONLY section *name* are not valid for PSV window.

The application has specified absolute addresses for a read-only section that are not consistent with the PSV window. If two addresses have been specified, the least-significant 15 bits should be identical. Also, the most significant bit of the virtual address should be set.

C

cannot find entry symbol *symbol* defaulting to *value*.

The linker can't find the entry symbol, so it will use the first address in the text section. This message may occur if the `-e` option incorrectly contains an equal sign ('=') in the option (i.e., `-e=0x200`).

common of '*name*' overridden by definition defined here.

The specified variable name has been declared in more than one file with one instance being declared as common. The definition will override the common symbol.

common of '*name*' overridden by larger common larger common is here.

The specified variable name has been declared in more than one file with different values. The smaller value will be overridden with the larger value.

common of '*name*' overriding smaller common smaller common is here.

The specified variable name has been declared in more than one file with different values. The first one encountered was smaller and will be overridden with the larger value.

D

data initialization has been turned off, therefore section *secname* will not be initialized.

The specified section requires initialization, but data initialization has been turned off; so, the initial data values are discarded. Storage for the data sections will be allocated as usual.

data memory region not specified. Using default upper limit of *addr*.

The linker has allocated a maximum-size stack. Since the data memory region was not specified, a default upper limit was used.

definition of '*name*' overriding common common is here.

The specified variable name has been declared in more than one file with one instance being declared as common. The definition will override the common symbol.

H

--heap option overrides HEAPSIZE symbol.

The `--heap` option has been specified and the `HEAPSIZE` symbol has been defined but they have different values so the `--heap` value will be used.

I

initial values were specified for a non-loadable data section (*name*). These values will be ignored.

By definition, a persistent data section implies data that is not initialized; therefore the values are discarded. Storage for the section will be allocated as usual.

M

**multiple common of '*name*'
previous common is here.**

The specified variable name has been declared in more than one file.

N

no memory region specified for section '*secname*'.

Section `secname` has been assigned to a default memory region, but other non-default regions are also defined.

O

object *name* was created for the *processor* and references register *name*.

An object file to be linked was created for a different processor family than the link target, and references an SFR that may not be compatible.

P

program memory region not specified. Using default upper limit of *addr*.

The linker has reallocated program memory to prevent a read-only section from crossing a PSV page boundary. Since the program memory region was not specified, a default upper limit was used.

R

READONLY section *secname* at *addr* crosses a PSVPAG boundary.

Address `addr` has been specified for a read-only section, causing it to cross a PSV page boundary. To allow efficient access of constant tables in the PSV window, it is recommended that the section should not cross a PSVPAG boundary.

'-retain-symbols-file' overrides '-s' and '-S'

If the strip all symbols option (`-s`) or the strip debug symbols option (`-S`) is used with `--retain-symbols-file FILE` only the symbols specified in the file will be kept.

S

--stack option overrides STACKSIZE symbol.

The `--stack` option has been specified and the `STACKSIZE` symbol has been defined but they have different values so the `--stack` value will be used.

T

target processor '*name*' does not match linker script.

The link target processor specified on the command line does not match the linker script `OUTPUT_ARCH` command. The processor name specified on the command line takes precedence.



Part 3 – 16-Bit Utilities (including the Archiver/Librarian)

Chapter 14. MPLAB XC16 Object Archiver/Librarian	233
Chapter 15. Other Utilities	239

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Chapter 14. MPLAB XC16 Object Archiver/Librarian

14.1 Introduction

The MPLAB XC16 Object Archiver/Librarian creates, modifies and extracts files from archives. This tool is one of several utilities (`xc16-ar`). An “archive” is a single file holding a collection of other files in a structure that makes it possible to retrieve the original individual files (called “members” of the archive).

The original files’ contents, mode (permissions), timestamp, owner and group are preserved in the archive, and can be restored on extraction.

The 16-bit archiver/librarian can maintain archives whose members have names of any length; however, if an `f` modifier is used, the file names will be truncated to 15 characters.

The archiver is considered a binary utility because archives of this sort are most often used as “libraries” holding commonly needed subroutines.

The archiver creates an index to the symbols defined in relocatable object modules in the archive when you specify the modifier `s`. Once created, this index is updated in the archive whenever the archiver makes a change to its contents (save for the `q` update operation). An archive with such an index speeds up linking to the library and allows routines in the library to call each other without regard to their placement in the archive.

You may use `xc16-nm -s` or `xc16-nm --print-armap` to list this index table. If an archive lacks the table, another form of the 16-bit archiver/librarian called `xc16-ranlib` can be used to add only the table.

The 16-bit archiver/librarian is designed to be compatible with two different facilities. You can control its activity using command line options or, if you specify the single command line option `-M`, you can control it with a script supplied via standard input.

14.2 Highlights

The following topics are covered in this chapter:

- [Archiver/Librarian and Other Development Tools](#)
- [Feature Set](#)
- [Input/Output Files](#)
- [Syntax](#)
- [Options](#)
- [Scripts](#)

16-Bit Assembler, Linker and Utilities User's Guide

14.3 Archiver/Librarian and Other Development Tools

The 16-bit librarian creates an archive file from object files created by the 16-bit assembler. Archive files may then be linked by the 16-bit linker with other relocatable object files to create an executable file. See the “*MPLAB XC16 C Compiler User's Guide*” (DS50002071) for an overview of the tools process flow.

14.4 Feature Set

Notable features of the librarian include:

- Available for Windows
- Command Line Interface

14.5 Input/Output Files

The 16-bit archiver/librarian generates archive files (.a). An archive file is a single file holding a collection of other files in a structure that makes it possible to retrieve the original individual files.

By default, object files are processed in the ELF format. To specify ELF or COFF format explicitly, use the `-omf` option on the command line, as shown:

```
xc16-ar -omf=coff [options...]  
xc16-ar -omf=elf [options...]
```

Alternatively, the environment variable `XC16_OMF` may be used to specify object file format for the 16-bit language tools.

14.6 Syntax

```
xc16-ar [-]P[MOD [RELPOS] [COUNT]] ARCHIVE [MEMBER...]  
xc16-ar -M [ <mri-script ]
```

14.7 Options

When you use the 16-bit archiver/librarian with command line options, the archiver insists on at least two arguments to execute: one key letter specifying the operation (optionally accompanied by other key letters specifying modifiers), and the archive name.

```
xc16-ar [-]P[MOD [RELPOS][COUNT]] ARCHIVE [MEMBER...]
```

Note: Command line options are case sensitive.

Most operations can also accept further *MEMBER* arguments, specifying archive members. Without specifying members, the entire archive is used.

The 16-bit archiver/librarian allows you to mix the operation code *P* and modifier flags *MOD* in any order, within the first command line argument. If you wish, you may begin the first command line argument with a dash.

The *P* keyletter specifies what operation to execute; it may be any of the following, but you must specify only one of them.

Table 14-1: Operation to Execute

Option	Function
d	Delete modules from the archive. Specify the names of modules to be deleted as <i>MEMBER...</i> ; the archive is untouched if you specify no files to delete. If you specify the <i>v</i> modifier, the 16-bit archiver/librarian lists each module as it is deleted.
m	Use this operation to move members in an archive. The ordering of members in an archive can make a difference in how programs are linked using the library, if a symbol is defined in more than one member. If no modifiers are used with <i>m</i> , any members you name in the <i>MEMBER</i> arguments are moved to the end of the archive; you can use the <i>a</i> , <i>b</i> or <i>i</i> modifiers to move them to a specified place instead.
p	Print the specified members of the archive, to the standard output file. If the <i>v</i> modifier is specified, show the member name before copying its contents to standard output. If you specify no <i>MEMBER</i> arguments, all the files in the archive are printed.
q	Append the files <i>MEMBER...</i> into <i>ARCHIVE</i> .
r	Insert the files <i>MEMBER...</i> into <i>ARCHIVE</i> (with replacement). If one of the files named in <i>MEMBER...</i> does not exist, the archiver displays an error message, and leaves undisturbed any existing members of the archive matching that name. By default, new members are added at the end of the file; but you may use one of the modifiers <i>a</i> , <i>b</i> or <i>i</i> to request placement relative to some existing member. The modifier <i>v</i> used with this operation elicits a line of output for each file inserted, along with one of the letters <i>a</i> or <i>r</i> to indicate whether the file was appended (no old member deleted) or replaced.
t	Display a table listing the contents of <i>ARCHIVE</i> , or those of the files listed in <i>MEMBER...</i> , that are present in the archive. Normally only the member name is shown; if you also want to see the modes (permissions), timestamp, owner, group and size, you can request that by also specifying the <i>v</i> modifier. If you do not specify a <i>MEMBER</i> , all files in the archive are listed. For example, if there is more than one file with the same name (<i>file</i>) in an archive (<i>b.a</i>), then <code>xc16-ar t b.a file</code> lists only the first instance; to see them all, you must ask for a complete listing in <code>xc16-ar t b.a</code> .
x	Extract members (named <i>MEMBER</i>) from the archive. You can use the <i>v</i> modifier with this operation, to request that the archiver list each name as it extracts it. If you do not specify a <i>MEMBER</i> , all files in the archive are extracted.

16-Bit Assembler, Linker and Utilities User's Guide

A number of modifiers (MOD) may immediately follow the P keyletter to specify variations on an operation's behavior.

Table 14-2: Modifiers

Option	Function
a	Add new files after an existing member of the archive. If you use the modifier <code>a</code> , the name of an existing archive member must be present as the <code>RELPOS</code> argument, before the <code>ARCHIVE</code> specification.
b	Add new files before an existing member of the archive. If you use the modifier <code>b</code> , the name of an existing archive member must be present as the <code>RELPOS</code> argument, before the <code>ARCHIVE</code> specification. (Same as <code>i</code> .)
c	Create the archive. The specified <code>ARCHIVE</code> is always created if it did not exist, when you requested an update. But a warning is issued unless you specify in advance that you expect to create it, by using this modifier.
f	Truncate names in the archive. The 16-bit archiver/librarian will normally permit file names of any length. This will cause it to create archives that are not compatible with the native archiver program on some systems. If this is a concern, the <code>f</code> modifier may be used to truncate file names when putting them in the archive.
i	Insert new files before an existing member of the archive. If you use the modifier <code>i</code> , the name of an existing archive member must be present as the <code>RELPOS</code> argument, before the <code>ARCHIVE</code> specification. (Same as <code>b</code> .)
l	This modifier is accepted but not used.
N	Uses the <code>COUNT</code> parameter. This is used if there are multiple entries in the archive with the same name. Extract or delete instance <code>COUNT</code> of the given name from the archive.
o	Preserve the original dates of members when extracting them. If you do not specify this modifier, files extracted from the archive are stamped with the time of extraction.
P	Use the full path name when matching names in the archive. The 16-bit archiver/librarian cannot create an archive with a full path name (such archives are not POSIX compliant), but other archive creators can. This option will cause the archiver to match file names using a complete path name, which can be convenient when extracting a single file from an archive created by another tool.
s	Write an object-file index into the archive, or update an existing one, even if no other change is made to the archive. You may use this modifier flag either with any operation, or alone. Running <code>xc16-ar s</code> on an archive is equivalent to running <code>xc16-ranlib</code> on it.
S	Do not generate an archive symbol table. This can speed up building a large library in several steps. The resulting archive cannot be used with the linker. In order to build a symbol table, you must omit the <code>S</code> modifier on the last execution of the archiver, or you must run <code>ranlib</code> on the archive.
u	Normally, <code>xc16-ar r...</code> inserts all files listed into the archive. If you would like to insert only those of the files you list that are newer than existing members of the same names, use this modifier. The <code>u</code> modifier is allowed only for the operation <code>r</code> (replace). In particular, the combination <code>qu</code> is not allowed, since checking the timestamps would lose any speed advantage from the operation <code>q</code> .
v	This modifier requests the verbose version of an operation. Many operations display additional information, such as, file names processed when the modifier <code>v</code> is appended.
V	This modifier shows the version number of the 16-bit archiver/librarian.

14.8 Scripts

If you use the single command line option `-M` with the archiver, you can control its operation with a rudimentary command language.

```
xc16-ar -M [ <SCRIPT >
```

Note: Command line options are case sensitive.

This form of the 16-bit archiver/librarian operates interactively if standard input is coming directly from a terminal. During interactive use, the archiver prompts for input (the prompt is `AR >`) and continues executing even after errors. If you redirect standard input to a script file, no prompts are issued and the 16-bit archiver/librarian abandons execution (with a nonzero exit code) on any error.

The archiver command language is **not** designed to be equivalent to the command line options; in fact, it provides somewhat less control over archives. The only purpose of the command language is to ease the transition to the 16-bit archiver/librarian for developers who already have scripts written for the MRI “librarian” program.

The syntax for the 16-bit archiver/librarian command language is straightforward:

- commands are recognized in upper or lower case; for example, `LIST` is the same as `list`. In the following descriptions, commands are shown in upper case for clarity.
- a single command may appear on each line; it is the first word on the line.
- empty lines are allowed, and have no effect.
- comments are allowed; text after either of the characters “*” or “;” is ignored.
- Whenever you use a list of names as part of the argument to an `xc16-ar` command, you can separate the individual names with either commas or blanks. Commas are shown in the explanations below, for clarity.
- “+” is used as a line continuation character; if “+” appears at the end of a line, the text on the following line is considered part of the current command.

Table 14-3 shows the commands you can use in archiver scripts, or when using the archiver interactively. Three of them have special significance.

Table 14-3: Archiver Scripts Commands

Option	Function
OPEN or CREATE	Specify a “current archive”, which is a temporary file required for most of the other commands.
SAVE	Commits the changes so far specified by the script. Prior to <code>SAVE</code> , commands affect only the temporary copy of the current archive.
ADDLIB ARCHIVE ADDLIB ARCHIVE (MODULE, MODULE, ...MODULE)	Add all the contents of <code>ARCHIVE</code> (or, if specified, each named <code>MODULE</code> from <code>ARCHIVE</code>) to the current archive. Requires prior use of <code>OPEN</code> or <code>CREATE</code> .
ADDMOD MEMBER, MEMBER, ... MEMBER	Add each named <code>MEMBER</code> as a module in the current archive. Requires prior use of <code>OPEN</code> or <code>CREATE</code> .
CLEAR	Discard the contents of the current archive, canceling the effect of any operations since the last <code>SAVE</code> . May be executed (with no effect) even if no current archive is specified.

16-Bit Assembler, Linker and Utilities User's Guide

Table 14-3: Archiver Scripts Commands (Continued)

Option	Function
CREATE ARCHIVE	Creates an archive, and makes it the current archive (required for many other commands). The new archive is created with a temporary name; it is not actually saved as ARCHIVE until you use SAVE. You can overwrite existing archives; similarly, the contents of any existing file named ARCHIVE will not be destroyed until SAVE.
DELETE MODULE, MODULE, ... MODULE	Delete each listed MODULE from the current archive; equivalent to <code>xc16-ar -d ARCHIVE MODULE ... MODULE</code> . Requires prior use of OPEN or CREATE.
DIRECTORY ARCHIVE (MODULE, ... MODULE) [OUTPUTFILE]	List each named MODULE present in ARCHIVE. The separate command VERBOSE specifies the form of the output: when verbose output is off, output is like that of <code>xc16-ar -t ARCHIVE MODULE....</code> When verbose output is on, the listing is like <code>xc16-ar -tv ARCHIVE MODULE....</code> Output normally goes to the standard output stream; however, if you specify OUTPUTFILE as a final argument, the 16-bit archiver/librarian directs the output to that file.
END	Exit from the archiver with a 0 exit code to indicate successful completion. This command does not save the output file; if you have changed the current archive since the last SAVE command, those changes are lost.
EXTRACT MODULE, MODULE, ... MODULE	Extract each named MODULE from the current archive, writing them into the current directory as separate files. Equivalent to <code>xc16-ar -x ARCHIVE MODULE....</code> Requires prior use of OPEN or CREATE.
LIST	Display full contents of the current archive, in “verbose” style regardless of the state of VERBOSE. The effect is like <code>xc16-ar tv ARCHIVE</code> . (This single command is a 16-bit archiver/librarian enhancement, rather than present for MRI compatibility.) Requires prior use of OPEN or CREATE.
OPEN ARCHIVE	Opens an existing archive for use as the current archive (required for many other commands). Any changes as the result of subsequent commands will not actually affect ARCHIVE until you next use SAVE.
REPLACE MODULE, MODULE, ... MODULE	In the current archive, replace each existing MODULE (named in the REPLACE arguments) from files in the current working directory. To execute this command without errors, both the file, and the module in the current archive, must exist. Requires prior use of OPEN or CREATE.
VERBOSE	Toggle an internal flag governing the output from DIRECTORY. When the flag is on, DIRECTORY output matches output from <code>xc16-ar -tv ...</code>
SAVE	Commits your changes to the current archive and actually saves it as a file with the name specified in the last CREATE or OPEN command. Requires prior use of OPEN or CREATE.

Chapter 15. Other Utilities

15.1 Introduction

This chapter discusses general information about other utilities for PIC24 MCUs and dsPIC DSCs.

15.2 Highlights

In addition to the archiver/librarian, other utilities are tools available for use with the assembler and/or linker.

Table 15-1: Available Utilities

Utility	Description
xc16-bin2hex*	Converts a linked object file into an Intel® hex file.
xc16-nm	Lists symbols from an object file.
xc16-objdump	Displays information about object files.
xc16-ranlib	Generates an index from the contents of an archive and stores it in the archive.
xc16-strings	Prints the printable character sequences.
xc16-strip	Discards all symbols from an object file.

*See the “MPLAB XC16 C Compiler User's Guide” (DS50002071) for an overview of the tools process flow.

15.3 xc16-bin2hex Utility

15.3.1 Introduction

The binary-to-hexadecimal (`xc16-bin2hex`) utility converts binary files (from the 16-bit linker) to Intel hex format files, suitable for loading into device programmers.

15.3.2 Highlights

The following topics are covered in this section:

- [Input/Output Files](#)
- [Syntax](#)
- [Options](#)

15.3.3 Input/Output Files

- Input: ELF or COFF formatted binary object files
- Output: Intel hex files

By default, object files are processed in the ELF format. To specify ELF or COFF format explicitly, use the `-omf` option on the command line, as shown:

```
xc16-bin2hex -omf=coff file1.out
xc16-bin2hex -omf=elf file2.out
```

Alternatively, the environment variable `XC16_OMF` may be used to specify object file format for the dsPIC30F language tools.

Because the Intel hex file format is byte-oriented, and the 16-bit PC is not, program memory sections require special treatment. Each 24-bit program word is extended to 32 bits by inserting a so-called “phantom byte”. Each program memory address is multiplied by 2 to yield a byte address.

For example, a section that is located at 0x100 in program memory will be represented in the hex file as 0x200. Consider the following assembly language source:

```
; file test.s
.section foo,code,address(0x100)
.pword 0x112233
```

The following commands will assemble the source file and create an Intel hex file:

```
xc16-as -o test.o test.s
xc16-bin2hex test.o
```

The file “test.hex” will be produced, with the following contents:

```
:020000040000fa
:040200003322110096
:00000001FF
```

Notice that the data record (line 2) has a load address of 0200, while the source code specified address 0x100. Note also that the data is represented in “little-endian” format, meaning the least significant byte appears first. The phantom byte appears last, just before the checksum.

15.3.4 Syntax

Command line syntax is:

```
xc16-bin2hex object_file [-v] [-a] [-u] [-omf=format]
```

Example 15.1: hello.cof

Convert the absolute COFF executable file `hello.cof` to `hello.hex`

```
xc16-bin2hex hello.cof
```


15.3.5 Options

The following options are supported.

Table 15-2: xc16-bin2hex Options

Option	Function
<i>object_file -a</i>	Sort the contents of the object file in ascending address order. For a summary of the object file contents, add the <i>-v</i> option (<i>-va</i>).
<i>-omf=format</i>	Specify object file format. The following formats are supported: ELF, COFF. Format names are case-insensitive. ELF is the default.
<i>-u</i>	Use upper-case hexadecimal digits
<i>-v</i>	Print a table of diagnostic information to standard output in the format shown in Example 15-2 .

Example 15-2: -va Option Output

writing hello.hex

section	PC address	byte address	length (w/pad)	actual length	(dec)
-----	-----	-----	-----	-----	-----
.reset	0	0	0x8	0x6	(6)
.ivt	0x4	0x8	0xf8	0xba	(186)
.aivt	0x84	0x108	0xf8	0xba	(186)
.text	0x100	0x200	0xae	0x831	(2097)
.const	0x676	0xcec	0x10	0xc	(12)
.dinit	0x67e	0xcfc	0x104	0xc3	(195)
.text	0x700	0xe00	0x14	0xf	(15)
.isr	0x70a	0xe14	0x4	0x3	(3)
Total program memory used (bytes):				0xa8c	(2700)

15.4 xc16-nm Utility

15.4.1 Introduction

The `xc16-nm` utility produces a list of symbols from object files. Each item in the list consists of the symbol value, symbol type and symbol name.

15.4.2 Highlights

The following topics are covered in this section:

- [Input/Output Files](#)
- [Syntax](#)
- [Options](#)
- [Output Formats](#)

15.4.3 Input/Output Files

- Input: Object archive files
- Output: Object archive files. If no object files are listed as arguments, `xc16-nm` assumes the file `a.out`.

15.4.4 Syntax

Command line syntax is:

```
xc16-nm [ -A | -o | --print-file-name ]
        [ -a | --debug-syms ] [ -B ]
        [ --defined-only ] [ -u | --undefined-only ]
        [ -f format | --format=format ] [ -g | --extern-only ]
        [ --help ] [-l | --line-numbers ]
        [ -n | -v | --numeric-sort ] [-omf=format]
        [ -p | --no-sort ]
        [ -P | --portability ] [ -r | --reverse-sort ]
        [ -s --print-armap ] [ --size-sort ]
        [ -t radix | --radix=radix ] [ -V | --version ]
        [ OBJFILE... ]
```

15.4.5 Options

Long and short forms of options, shown in [Table 15-3](#) as alternatives, are equivalent.

Table 15-3: xc16-nm Options

Option	Function
-A -o --print-file-name	Precede each symbol by the name of the input file (or archive member) in which it was found, rather than identifying the input file once only, before all of its symbols.
-a --debug-syms	Display all symbols, even debugger-only symbols; normally these are not listed.
-B --defined-only	The same as --format=bsd. Display only defined symbols for each object file.
-u --undefined-only	Display only undefined symbols (those external to each object file).
-f <i>format</i> --format= <i>format</i>	Use the output format <i>format</i> , which can be <i>bsd</i> , <i>sysv</i> or <i>posix</i> . The default is <i>bsd</i> . Only the first character of <i>format</i> is significant; it can be either upper or lower case.
-g --extern-only	Display only external symbols.
--help	Show a summary of the options to <i>xc16-nm</i> and exit.
-l --line-numbers	For each symbol, use debugging information to try to find a filename and line number. For a defined symbol, look for the line number of the address of the symbol. For an undefined symbol, look for the line number of a relocation entry that refers to the symbol. If line number information can be found, print it after the other symbol information.
-n -v --numeric-sort	Sort symbols numerically by their addresses, rather than alphabetically by their names.
-omf= <i>format</i>	Specify object file format. The following formats are supported: ELF, COFF. Format names are case-insensitive. ELF is the default.
-p --no-sort	Do not bother to sort the symbols in any order; print them in the order encountered.
-P --portability	Use the POSIX.2 standard output format instead of the default format. Equivalent to -f <i>posix</i> .
-r --reverse-sort	Reverse the order of the sort (whether numeric or alphabetic); let the last come first.
-s --print-arnmap	When listing symbols from archive members, include the index: a mapping (stored in the archive by <i>xc16-ar</i> or <i>xc16-ranlib</i>) of which modules contain definitions for which names.
--size-sort	Sort symbols by size. The size is computed as the difference between the value of the symbol and the value of the symbol with the next higher value. The size of the symbol is printed, rather than the value.
-t <i>radix</i> --radix= <i>radix</i>	Use <i>radix</i> as the radix for printing the symbol values. It must be <i>d</i> for decimal, <i>o</i> for octal or <i>x</i> for hexadecimal.
-V --version	Show the version number of <i>xc16-nm</i> and exit.

16-Bit Assembler, Linker and Utilities User's Guide

15.4.6 Output Formats

The symbol value is in the radix selected by the options, or hexadecimal by default.

If the symbol type is lowercase, the symbol is local; if uppercase, the symbol is global (external). Table 15-4 shows the symbol types.

Table 15-4: Symbol Types

Symbol	Description
A	The symbol's value is absolute, and will not be changed by further linking.
B	The symbol is in the uninitialized data section (known as BSS).
C	The symbol is common. Common symbols are uninitialized data. When linking, multiple common symbols may appear with the same name. If the symbol is defined anywhere, the common symbols are treated as undefined references.
D	The symbol is in the initialized data section.
N	The symbol is a debugging symbol.
R	The symbol is in a read only data section.
T	The symbol is in the text (code) section.
U	The symbol is undefined.
V	The symbol is a weak object. When a weak defined symbol is linked with a normal defined symbol, the normal defined symbol is used with no error. When a weak undefined symbol is linked and the symbol is not defined, the value of the weak symbol becomes zero with no error.
W	The symbol is a weak symbol that has not been specifically tagged as a weak object symbol. When a weak defined symbol is linked with a normal defined symbol, the normal defined symbol is used with no error. When a weak undefined symbol is linked and the symbol is not defined, the value of the weak symbol becomes zero with no error.
?	The symbol type is unknown, or object file format specific.

Example 15-3: xc16-nm output

```
00000474 T _fclose
0000023e T _fputc
000001b2 T _fputs
0000051e T _free
00000700 T _main
000003bc T _malloc
00000334 T _memcpy
00000198 T _puts
0000061a W _remove
0000062c W _sbrk
00000326 T _strlen
00000310 T _strchr
000005a0 W _write
```

15.5 xc16-objdump Utility

15.5.1 Introduction

The `xc16-objdump` utility displays information about one or more object files. The options control what particular information to display.

15.5.2 Highlights

The following topics are covered in this section:

- [Input/Output Files](#)
- [Syntax](#)
- [Options](#)

15.5.3 Input/Output Files

- Input: Object archive files
- Output: Object archive files. If no object files are listed as arguments, `xc16-nm` assumes the file `a.out`.

15.5.4 Syntax

Command line syntax is:

```
xc16-objdump [ -a | --archive-headers ]
               [ -d | --disassemble ]
               [ -D | --disassemble-all ]
               [ -EB | -EL | --endian={big | little} ]
               [ -f | --file-headers ]
               [ --file-start-context ]
               [ -g | --debugging ]
               [ -h | --section-headers | --headers ]
               [ -H | --help ]
               [ -j name | --section=name ]
               [ -l | --line-numbers ]
               [ -M options | --disassembler-options=options ]
               [-omf=format]
               [ --prefix-addresses ]
               [ --psrd-psrd-check ]
               [ -r | --reloc ]
               [ -s | --full-contents ]
               [ -S | --source ]
               [ --[no-]show-raw-insn ]
               [ --start-address=address ]
               [ --stop-address=address ]
               [ -t | --syms ]
               [ -V | --version ]
               [ -w | --wide ]
               [ -x | --all-headers ]
               [ -z | --disassemble-zeroes ]
OBJFILE...
```

OBJFILE... are the object files to be examined. When you specify archives, `xc16-objdump` shows information on each of the member object files.

16-Bit Assembler, Linker and Utilities User's Guide

15.5.5 Options

The long and short forms of options, shown in [Table 15-5](#), as alternatives, are equivalent. At least one of the following options `-a`, `-d`, `-D`, `-f`, `-g`, `-G`, `-h`, `-H`, `-p`, `-r`, `-R`, `-S`, `-t`, `-T`, `-V` or `-x` must be given.

Table 15-5: xc16-objdump Options

Option	Function
<code>-a</code> <code>--archive-header</code>	If any of the OBJFILE files are archives, display the archive header information (in a format similar to <code>ls -l</code>). Besides the information you could list with <code>xc16-ar tv</code> , <code>xc16-objdump -a</code> shows the object file format of each archive member.
<code>-d</code> <code>--disassemble</code>	Display the assembler mnemonics for the machine instructions from OBJFILE. This option only disassembles those sections that are expected to contain instructions.
<code>-D</code> <code>--disassemble-all</code>	Like <code>-d</code> , but disassemble the contents of all sections, not just those expected to contain instructions.
<code>-EB</code> <code>-EL</code> <code>--endian={big little}</code>	Specify the endianness of the object files. This only affects disassembly. This can be useful when disassembling a file format that does not describe endianness information, such as S-records.
<code>-f</code> <code>--file-header</code>	Display summary information from the overall header of each of the OBJFILE files.
<code>--file-start-context</code>	Specify that when displaying inter-listed source code/disassembly (assumes <code>'-S'</code>) from a file that has not yet been displayed, extend the context to the start of the file.
<code>-g</code> <code>--debugging</code>	Display debugging information. This attempts to parse debugging information stored in the file and print it out using a C like syntax. Only certain types of debugging information have been implemented.
<code>-h</code> <code>--section-header</code> <code>--header</code>	Display summary information from the section headers of the object file.
<code>-H</code> <code>--help</code>	Print a summary of the options to <code>xc16-objdump</code> and exit.
<code>-j name</code> <code>--section=name</code>	Display information only for section <i>name</i> .
<code>-l</code> <code>--line-numbers</code>	Label the display (using debugging information) with the filename and source line numbers corresponding to the object code or relocs shown. Only useful with <code>-d</code> , <code>-D</code> or <code>-r</code> .
<code>-M options</code> <code>--disassembler-options=options</code>	Pass target specific information to the disassembler. The dsPIC30F device supports the following target specific options: <code>symbolic</code> - Will perform symbolic disassembly.

Table 15-5: xc16-objdump Options (Continued)

Option	Function
<code>-omf=format</code>	Specify object file format. The following formats are supported: ELF, COFF. Format names are case-insensitive. ELF is the default.
<code>--prefix-addresses</code>	When disassembling, print the complete address on each line. This is the older disassembly format.
<code>--psrd-psrd-check</code> <code>[=library]</code>	Check for back-to-back data flash reads. Specifying the optional <code>=library</code> should be used on unlinked object files (such as a library). This option can be combined with <code>-d</code> to get a disassembly listing with additional information.
<code>-r</code> <code>--reloc</code>	Print the relocation entries of the file. If used with <code>-d</code> or <code>-D</code> , the relocations are printed interspersed with the disassembly.
<code>-s</code> <code>--full-contents</code>	Display the full contents of any sections requested.
<code>-S</code> <code>--source</code>	Display source code intermixed with disassembly, if possible. Implies <code>-d</code> .
<code>--show-raw-insn</code>	When disassembling instructions, print the instruction in hex, as well as in symbolic form. This is the default except when <code>--prefix-addresses</code> is used.
<code>--no-show-raw-insn</code>	When disassembling instructions, do not print the instruction bytes. This is the default when <code>--prefix-addresses</code> is used.
<code>--start-address=address</code> <code>ss</code>	Start displaying data at the specified address. This affects the output of the <code>-d</code> , <code>-r</code> and <code>-s</code> options.
<code>--stop-address=address</code> <code>s</code>	Stop displaying data at the specified address. This affects the output of the <code>-d</code> , <code>-r</code> and <code>-s</code> options.
<code>-t</code> <code>--syms</code>	Print the symbol table entries of the file. This is similar to the information provided by the <code>xc16-nm</code> program.
<code>-V</code> <code>--version</code>	Print the version number of <code>xc16-objdump</code> and exit.
<code>-w</code> <code>--wide</code>	Format some lines for output devices that have more than 80 columns.
<code>-x</code> <code>--all-header</code>	Display all available header information, including the symbol table and relocation entries. Using <code>-x</code> is equivalent to specifying all of <code>-a -f -h -r -t</code> .
<code>-z</code> <code>--disassemble-zeroes</code>	Normally, the disassembly output will skip blocks of zeroes. This option directs the disassembler to disassemble those blocks, just like any other data.

16-Bit Assembler, Linker and Utilities User's Guide

Example 15-4: -h OUTPUT

hello.out: file format coff-pic30

Sections:

Idx	Name	Size	VMA	LMA	File off	Algn
0	.reset	00000004	00000000	00000000	00000288	2**1
	CONTENTS, ALLOC, LOAD, CODE					
1	.text	00000576	00000100	00000100	00000290	2**1
	CONTENTS, ALLOC, LOAD, CODE					
2	.comment	0000005e	00000000	00000000	00000d7c	2**1
	CONTENTS, NEVER_LOAD					
3	.ivt	0000007c	00000004	00000004	00000e38	2**1
	CONTENTS, ALLOC, LOAD, CODE					
4	.aivt	0000007c	00000084	00000084	00000f30	2**1
	CONTENTS, ALLOC, LOAD, CODE					
5	__c30_signature	0000007e	0000005e	0000005e	00001028	2**1
	CONTENTS, DEBUGGING					
6	.data	0000008e	00000800	00000800	00001124	2**1
	CONTENTS, ALLOC, DATA, NEVER_LOAD					
7	.bss	00000002	0000088e	0000088e	00000000	2**1
	ALLOC					
8	.data	00000002	00000890	00000890	00001240	2**1
	CONTENTS, ALLOC, DATA, NEVER_LOAD					
9	.bss	00000002	00000892	00000892	00000000	2**1
	ALLOC					
10	.heap	00000080	00000894	00000894	00000000	2**1
	ALLOC					
11	.const	00000008	00008676	00000676	00001244	2**1
	CONTENTS, ALLOC, LOAD, READONLY, PSV					
12	.dinit	00000082	0000067e	0000067e	00001254	2**1
	CONTENTS, ALLOC, LOAD, CODE					
13	.text	0000000a	00000700	00000700	00001358	2**1
	CONTENTS, ALLOC, LOAD, CODE					
14	.isr	00000002	0000070a	0000070a	0000136c	2**1
	CONTENTS, ALLOC, LOAD, CODE					

15.6 xc16-ranlib Utility

15.6.1 Introduction

The `xc16-ranlib` utility generates an index to the contents of an archive and stores it in the archive. The index lists each symbol defined by a member of an archive that is a relocatable object file. You may use `xc16-nm -s` or `xc16-nm --print-armap` to list this index. An archive with such an index speeds up linking to the library and allows routines in the library to call each other without regard to their placement in the archive.

Running `xc16-ranlib` is completely equivalent to executing `xc16-ar -s` (i.e., the 16-bit archiver/librarian with the `-s` option).

15.6.2 Highlights

The following topics are covered in this section:

- [Input/Output Files](#)
- [Syntax](#)
- [Options](#)

15.6.3 Input/Output Files

- Input: Archive files
- Output: Archive files

15.6.4 Syntax

Command line syntax is:

```
xc16-ranlib [-omf=format] [-v | -V | --version] ARCHIVE
```

15.6.5 Options

The long and short forms of options, shown in [Table 15-6](#) as alternatives, are equivalent.

Table 15-6: xc16-ranlib Options

Option	Function
<code>-omf=<i>format</i></code>	Specify object file format. The following formats are supported: ELF, COFF. Format names are case-insensitive. ELF in the default.
<code>-v</code> <code>-V</code> <code>--version</code>	Show the version number of <code>xc16-ranlib</code>

15.7 xc16-strings Utility

15.7.1 Introduction

For each file given, the `xc16-strings` utility prints the printable character sequences that are at least 4 characters long (or the number given in the options) and are followed by an unprintable character. By default, it only prints the strings from the initialized and loaded sections of object files; for other types of files, it prints the strings from the whole file.

`xc16-strings` is mainly useful for determining the contents of non-text files.

15.7.2 Highlights

The following topics are covered in this section:

- [Input/Output Files](#)
- [Syntax](#)
- [Options](#)

15.7.3 Input/Output Files

- Input: Any files
- Output: Standard output

15.7.4 Syntax

Command line syntax is:

```
xc16-strings [-a | --all | -] [-f | --print-file-name]
              [--help] [-min-len | -n min-len | --bytes=min-len]
              [-omf=format] [-t radix | --radix=radix]
              [-v | --version] FILE...
```

15.7.5 Options

The long and short forms of options, shown in [Table 15-7](#) as alternatives, are equivalent.

Table 15-7: xc16-strings Options

Option	Function
-a --all -	Do not scan only the initialized and loaded sections of object files; scan the whole files.
-f --print-file-name	Print the name of the file before each string.
--help	Print a summary of the program usage on the standard output and exit.
-min-len -n min-len --bytes=min-len	Print sequences of characters that are at least <i>-min-len</i> characters long, instead of the default 4.
-omf=format	Specify object file format. The following formats are supported: ELF, COFF. Format names are case-insensitive. ELF is the default.
-t radix --radix=radix	Print the offset within the file before each string. The single character argument specifies the radix of the offset: <i>o</i> for octal, <i>x</i> for hexadecimal, or <i>d</i> for decimal.
-v --version	Print the program version number on the standard output and exit.

15.8 xc16-strip Utility

15.8.1 Introduction

The `xc16-strip` utility discards all symbols from the object and archive files specified. At least one file must be given. `xc16-strip` modifies the files named in its argument, rather than writing modified copies under different names.

15.8.2 Highlights

The following topics are covered in this section:

- [Input/Output Files](#)
- [Syntax](#)
- [Options](#)

15.8.3 Input/Output Files

- Input: Object or archive files
- Output: Object or archive files. If no object or archive files are listed as arguments, `xc16-strip` assumes the file `a.out`.

15.8.4 Syntax

Command line syntax is:

```
xc16-strip [ -g | -S | --strip-debug ] [ --help ]
          [ -K symbolname | --keep-symbol=symbolname ]
          [ -N symbolname | --strip-symbol=symbolname ]
          [ -o file ] [-omf=format]
          [ -p | --preserve-dates ]
          [ -R sectionname | --remove-section=sectionname ]
          [ -s | --strip-all ] [--strip-unneeded]
          [ -v | --verbose ] [ -V | --version ]
          [ -x | --discard-all ] [ -X | --discard-locals ]
          OBJFILE...
```

15.8.5 Options

The long and short forms of options, shown in [Table 15-8](#) as alternatives, are equivalent.

Table 15-8: xc16-strip Options

Option	Function
-g -S --strip-debug	Remove debugging symbols only.
--help	Show a summary of the options to <code>xc16-strip</code> and exit.
-K <i>symbolname</i> --keep-symbol= <i>symbolname</i>	Keep only symbol <i>symbolname</i> from the source file. This option may be given more than once.
-N <i>symbolname</i> --strip-symbol= <i>symbolname</i>	Remove symbol <i>symbolname</i> from the source file. This option may be given more than once, and may be combined with strip options other than -K.
-o <i>file</i>	Put the stripped output in <i>file</i> , rather than replacing the existing file. When this argument is used, only one <code>OBJFILE</code> argument may be specified.
-omf= <i>format</i>	Specify object file format. The following formats are supported: ELF, COFF. Format names are case-insensitive. ELF is the default.
-p --preserve-dates	Preserve the access and modification dates of the file.
-R <i>sectionname</i> --remove-section= <i>sectionname</i>	Remove any section named <i>sectionname</i> from the output file. This option may be given more than once. Note that using this option inappropriately may make the output file unusable.
-s --strip-all	Remove all symbols.
--strip-unneeded	Remove all symbols that are not needed for relocation processing.
-v --verbose	Verbose output: list all object files modified. In the case of archives, <code>xc16-strip -v</code> lists all members of the archive.
-V --version	Show the version number for <code>xc16-strip</code> .
-x --discard-all	Remove non-global symbols.
-X --discard-locals	Remove compiler-generated local symbols (these usually start with <code>_L</code> or <code>."</code>).

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:



Part 4 – Appendices

Appendix A. Deprecated Features	257
Appendix B. Useful Tables	259
Appendix C. GNU Free Documentation License	261
Appendix D. Document Revision History	269

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Appendix A. Deprecated Features

A.1 Introduction

The features described below are considered to be obsolete and have been replaced with more advanced functionality. Projects which depend on deprecated features will work properly with versions of the language tools cited. The use of a deprecated feature will result in a warning; programmers are encouraged to revise their projects in order to eliminate any dependency on deprecated features. Support for these features may be removed entirely in future versions of the language tools.

A.2 Highlights

Topics covered in this appendix are:

- [Assembler Directives that Define Sections](#)
- [Reserved Section Names with Implied Attributes](#)
- [Environmental Variables](#)

A.3 Assembler Directives that Define Sections

The following `.section` directive format was deprecated in v1.30. The new directive format may be found in [Section 4.2 "Directives that Define Sections"](#).

A.3.1 `.section name [, "flags"]`

Definition

Assembles the following code into a section named `name`. If the optional argument is quoted, it is taken as flags to use for the section. Each flag is a single character. The following flags are recognized:

- `b` bss section (uninitialized data)
- `n` Section is not loaded
- `d` Data section (initialized data)
- `r` Read-only data section (PSV window)
- `x` Executable section

If the `n` flag is used by itself, the section defaults to uninitialized data.

If no flags are specified, the default flags depend on the section name. If the section name is not recognized, the default will be for the section to be loadable data.

16-Bit Assembler, Linker and Utilities User's Guide

The following section names are recognized:

Table A-1: Section Names

Section Name	Default Flag
.text	x
.data	d
.bss	b

Note: Ensure that double quotes are used around flags. If the optional argument to the `.section` directive is not quoted, it is taken as a sub-section number. Remember, a single character in single quotes (i.e., 'b') is converted by the preprocessor to a number.

Example

```
.section .const, "r"
; The following symbols (C1 and C2) will be placed
; in the named section ".const".
C1: .word 0x1234
C2: .word 0x5678
```

A.4 Reserved Section Names with Implied Attributes

Implied attributes for the section names in the table below were deprecated in v1.30.

Reserved Name	Implied Attribute(s)
.xbss	bss, xmemory
.xdata	data, xmemory
.nbss	bss, near
.ndata	data, near
.ndconst	data, near
.pbss	bss, persist
.dconst	data
.ybss	bss, ymemory
.ydata	data, ymemory
.const	psv
.eedata	eedata

See [Section 4.2 "Directives that Define Sections"](#) for more information.

A.5 Environmental Variables

The environment variable `PIC30_OMF` was used to specify object file format for the 16-bit language tools. Now use `XC16_OMF`.

Appendix B. Useful Tables

B.1 INTRODUCTION

Some useful tables are included for reference here.

B.2 HIGHLIGHTS

The tables are:

- [ASCII Character Set](#)
- [Hexadecimal to Decimal Conversion](#)

B.3 ASCII CHARACTER SET

This table shows the ASCII character set in nibbles.

		Most Significant Nibbles							
Least Significant Nibbles	Hex	0	1	2	3	4	5	6	7
	0	NUL	DLE	Space	0	@	P	`	p
	1	SOH	DC1	!	1	A	Q	a	q
	2	STX	DC2	"	2	B	R	b	r
	3	ETX	DC3	#	3	C	S	c	s
	4	EOT	DC4	\$	4	D	T	d	t
	5	ENQ	NAK	%	5	E	U	e	u
	6	ACK	SYN	&	6	F	V	f	v
	7	Bell	ETB	'	7	G	W	g	w
	8	BS	CAN	(8	H	X	h	x
	9	HT	EM)	9	I	Y	i	y
	A	LF	SUB	*	:	J	Z	j	z
	B	VT	ESC	+	;	K	[k	{
	C	FF	FS	,	<	L	\	l	
	D	CR	GS	–	=	M]	m	}
	E	SO	RS	.	>	N	^	n	~
	F	SI	US	/	?	O	_	o	DEL

16-Bit Assembler, Linker and Utilities User's Guide

B.4 HEXADECIMAL TO DECIMAL CONVERSION

This appendix describes how to convert hexadecimal to decimal. For each hex digit, find the associated decimal value. Add the numbers together.

High Byte				Low Byte			
Hex 1000	Dec	Hex 100	Dec	Hex 10	Dec	Hex 1	Dec
0	0	0	0	0	0	0	0
1	4096	1	256	1	16	1	1
2	8192	2	512	2	32	2	2
3	12288	3	768	3	48	3	3
4	16384	4	1024	4	64	4	4
5	20480	5	1280	5	80	5	5
6	24576	6	1536	6	96	6	6
7	28672	7	1792	7	112	7	7
8	32768	8	2048	8	128	8	8
9	36864	9	2304	9	144	9	9
A	40960	A	2560	A	160	A	10
B	45056	B	2816	B	176	B	11
C	49152	C	3072	C	192	C	12
D	53248	D	3328	D	208	D	13
E	57344	E	3584	E	224	E	14
F	61440	F	3840	F	240	F	15

For example, hex A38F converts to 41871 as follows:

Hex 1000's Digit	Hex 100's Digit	Hex 10's Digit	Hex 1's Digit	Result
40960	768	128	15	41871 Decimal

Appendix C. GNU Free Documentation License

Version 1.3, 3 November 2008

Copyright (C) 2000, 2001, 2002, 2007, 2008 Free Software Foundation, Inc.

<<http://fsf.org/>>

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

C.1 Preamble

The purpose of this License is to make a manual, textbook, or other functional and useful document “free” in the sense of freedom: to assure everyone the effective freedom to copy and redistribute it, with or without modifying it, either commercially or non-commercially. Secondly, this License preserves for the author and publisher a way to get credit for their work, while not being considered responsible for modifications made by others.

This License is a kind of “copyleft,” which means that derivative works of the document must themselves be free in the same sense. It complements the GNU General Public License, which is a copyleft license designed for free software.

We have designed this License in order to use it for manuals for free software, because free software needs free documentation: a free program should come with manuals providing the same freedoms that the software does. But this License is not limited to software manuals; it can be used for any textual work, regardless of subject matter or whether it is published as a printed book. We recommend this License principally for works whose purpose is instruction or reference.

C.2 Applicability and Definitions

This License applies to any manual or other work, in any medium, that contains a notice placed by the copyright holder saying it can be distributed under the terms of this License. Such a notice grants a world-wide, royalty-free license, unlimited in duration, to use that work under the conditions stated herein. The “Document” below refers to any such manual or work. Any member of the public is a licensee and is addressed as “you.” You accept the license if you copy, modify or distribute the work in a way requiring permission under copyright law.

A “Modified Version” of the Document means any work containing the Document or a portion of it, either copied verbatim, or with modifications and/or translated into another language.

A “Secondary Section” is a named appendix or a front-matter section of the Document that deals exclusively with the relationship of the publishers or authors of the Document to the Document's overall subject (or to related matters) and contains nothing that could fall directly within that overall subject. (Thus, if the Document is in part a textbook of mathematics, a Secondary Section may not explain any mathematics.) The relationship could be a matter of historical connection with the subject or with related matters, or of legal, commercial, philosophical, ethical or political position regarding them.

The “Invariant Sections” are certain Secondary Sections whose titles are designated, as being those of Invariant Sections, in the notice that says that the Document is released under this License. If a section does not fit the above definition of Secondary then it is not allowed to be designated as Invariant. The Document may contain zero Invariant Sections. If the Document does not identify any Invariant Sections then there are none.

The “Cover Texts” are certain short passages of text that are listed, as Front-Cover Texts or Back-Cover Texts, in the notice that says that the Document is released under this License. A Front-Cover Text may be at most 5 words and a Back-Cover Text may be at most 25 words.

A “Transparent” copy of the Document means a machine-readable copy, represented in a format whose specification is available to the general public, that is suitable for revising the document straightforwardly with generic text editors or (for images composed of pixels) generic paint programs or (for drawings) some widely available drawing editor, and that is suitable for input to text formatters or for automatic translation to a variety of formats suitable for input to text formatters. A copy made in an otherwise Transparent file format whose markup, or absence of markup, has been arranged to thwart or discourage subsequent modification by readers is not Transparent. An image format is not Transparent if used for any substantial amount of text. A copy that is not “Transparent” is called “Opaque.”

Examples of suitable formats for Transparent copies include plain ASCII without markup, Texinfo input format, LaTeX input format, SGML or XML using a publicly available DTD, and standard-conforming simple HTML, PostScript or PDF designed for human modification. Examples of transparent image formats include PNG, XCF and JPG. Opaque formats include proprietary formats that can be read and edited only by proprietary word processors, SGML or XML for which the DTD and/or processing tools are not generally available, and the machine-generated HTML, PostScript or PDF produced by some word processors for output purposes only.

The “Title Page” means, for a printed book, the title page itself, plus such following pages as are needed to hold, legibly, the material this License requires to appear in the title page. For works in formats which do not have any title page as such, “Title Page” means the text near the most prominent appearance of the work's title, preceding the beginning of the body of the text.

The “publisher” means any person or entity that distributes copies of the Document to the public.

A section “Entitled XYZ” means a named subunit of the Document whose title either is precisely XYZ or contains XYZ in parentheses following text that translates XYZ in another language. Here XYZ stands for a specific section name mentioned below, such as “Acknowledgments,” “Dedications,” “Endorsements,” or “History.” To “Preserve the Title” of such a section when you modify the Document means that it remains a section “Entitled XYZ” according to this definition.

The Document may include Warranty Disclaimers next to the notice which states that this License applies to the Document. These Warranty Disclaimers are considered to be included by reference in this License, but only as regards disclaiming warranties: any other implication that these Warranty Disclaimers may have is void and has no effect on the meaning of this License.

C.3 Verbatim Copying

You may copy and distribute the Document in any medium, either commercially or non-commercially, provided that this License, the copyright notices, and the license notice saying this License applies to the Document are reproduced in all copies and that you add no other conditions whatsoever to those of this License. You may not use

technical measures to obstruct or control the reading or further copying of the copies you make or distribute. However, you may accept compensation in exchange for copies. If you distribute a large enough number of copies you must also follow the conditions in section 3.

You may also lend copies, under the same conditions stated above and you may publicly display copies.

C.4 Copying in Quantity

If you publish printed copies (or copies in media that commonly have printed covers) of the Document, numbering more than 100, and the Document's license notice requires Cover Texts, you must enclose the copies in covers that carry, clearly and legibly, all these Cover Texts: Front-Cover Texts on the front cover, and Back-Cover Texts on the back cover. Both covers must also clearly and legibly identify you as the publisher of these copies. The front cover must present the full title with all words of the title equally prominent and visible. You may add other material on the covers in addition. Copying with changes limited to the covers, as long as they preserve the title of the Document and satisfy these conditions, can be treated as verbatim copying in other respects.

If the required texts for either cover are too voluminous to fit legibly, you should put the first ones listed (as many as fit reasonably) on the actual cover and continue the rest onto adjacent pages.

If you publish or distribute Opaque copies of the Document numbering more than 100, you must either include a machine-readable Transparent copy along with each Opaque copy, or state in or with each Opaque copy a computer-network location from which the general network-using public has access to download using public-standard network protocols a complete Transparent copy of the Document, free of added material. If you use the latter option, you must take reasonably prudent steps, when you begin distribution of Opaque copies in quantity, to ensure that this Transparent copy will remain thus accessible at the stated location until at least one year after the last time you distribute an Opaque copy (directly or through your agents or retailers) of that edition to the public.

It is requested, but not required, that you contact the authors of the Document well before redistributing any large number of copies, to give them a chance to provide you with an updated version of the Document.

C.5 Modifications

You may copy and distribute a Modified Version of the Document under the conditions of sections 2 and 3 above, provided that you release the Modified Version under precisely this License, with the Modified Version filling the role of the Document, thus licensing distribution and modification of the Modified Version to whoever possesses a copy of it. In addition, you must do these things in the Modified Version:

- a) Use in the Title Page (and on the covers, if any) a title distinct from that of the Document, and from those of previous versions (which should, if there were any, be listed in the History section of the Document). You may use the same title as a previous version if the original publisher of that version gives permission.
- b) List on the Title Page, as authors, one or more persons or entities responsible for authorship of the modifications in the Modified Version, together with at least five of the principal authors of the Document (all of its principal authors, if it has fewer than five), unless they release you from this requirement.
- c) State on the Title page, the name of the publisher of the Modified Version, as the publisher.
- d) Preserve all the copyright notices of the Document.
- e) Add an appropriate copyright notice for your modifications adjacent to the other copyright notices.
- f) Include, immediately after the copyright notices, a license notice giving the public permission to use the Modified Version under the terms of this License, in the form shown in the Addendum below.
- g) Preserve in that license notice the full lists of Invariant Sections and required Cover Texts given in the Document's license notice.
- h) Include an unaltered copy of this License.
- i) Preserve the section entitled "History," as well as its Title, and add an item stating at least the title, year, new authors, and publisher of the Modified Version as given on the Title Page. If there is no section entitled "History" in the Document, create one stating the title, year, authors, and publisher of the Document as given on its Title Page, then add an item describing the Modified Version as stated in the previous sentence.
- j) Preserve the network location, if any, given in the Document for public access to a Transparent copy of the Document, and likewise the network locations given in the Document for previous versions it was based on. These may be placed in the "History" section. You may omit a network location for a work that was published at least four years before the Document itself, or if the original publisher of the version it refers to gives permission.
- k) For any section entitled "Acknowledgments" or "Dedications," preserve the Title of the section, and preserve in the section all the substance and tone of each of the contributor acknowledgments and/or dedications given therein.
- l) Preserve all the Invariant Sections of the Document, unaltered in their text and in their titles. Section numbers or the equivalent are not considered part of the section titles.
- m) Delete any section entitled "Endorsements." Such a section may not be included in the Modified Version.
- n) Do not retitle any existing section to be entitled "Endorsements" or to conflict in title with any Invariant Section. Preserve any Warranty Disclaimers.

If the Modified Version includes new front-matter sections or appendices that qualify as Secondary Sections and contain no material copied from the Document, you may at your option designate some or all of these sections as invariant. To do this, add their titles to the list of Invariant Sections in the Modified Version's license notice. These titles must be distinct from any other section titles.

You may add a section entitled “Endorsements,” provided it contains nothing but endorsements of your Modified Version by various parties—for example, statements of peer review or that the text has been approved by an organization as the authoritative definition of a standard.

You may add a passage of up to five words as a Front-Cover Text, and a passage of up to 25 words as a Back-Cover Text, to the end of the list of Cover Texts in the Modified Version. Only one passage of Front-Cover Text and one of Back-Cover Text may be added by (or through arrangements made by) any one entity. If the Document already includes a cover text for the same cover, previously added by you or by arrangement made by the same entity you are acting on behalf of, you may not add another; but you may replace the old one, on explicit permission from the previous publisher that added the old one.

The author(s) and publisher(s) of the Document do not by this License give permission to use their names for publicity for or to assert or imply endorsement of any Modified Version.

C.6 Combining Documents

You may combine the Document with other documents released under this License, under the terms defined in section 4 above for modified versions, provided that you include in the combination all of the Invariant Sections of all of the original documents, unmodified, and list them all as Invariant Sections of your combined work in its license notice, and that you preserve all their Warranty Disclaimers.

The combined work need only contain one copy of this License, and multiple identical Invariant Sections may be replaced with a single copy. If there are multiple Invariant Sections with the same name but different contents, make the title of each such section unique by adding at the end of it, in parentheses, the name of the original author or publisher of that section if known, or else a unique number. Make the same adjustment to the section titles in the list of Invariant Sections in the license notice of the combined work.

In the combination, you must combine any sections entitled “History” in the various original documents, forming one section entitled “History.” Likewise combine any sections entitled “Acknowledgments” and any sections entitled “Dedications.” You must delete all sections entitled “Endorsements.”

C.7 Collections of Documents

You may make a collection consisting of the Document and other documents released under this License, and replace the individual copies of this License in the various documents with a single copy that is included in the collection, provided that you follow the rules of this License for verbatim copying of each of the documents in all other respects.

You may extract a single document from such a collection, and distribute it individually under this License, provided you insert a copy of this License into the extracted document, and follow this License in all other respects regarding verbatim copying of that document.

C.8 Aggregation with Independent Works

A compilation of the Document or its derivatives with other separate and independent documents or works, in or on a volume of a storage or distribution medium, is called an “aggregate” if the copyright resulting from the compilation is not used to limit the legal rights of the compilation's users beyond what the individual works permit. When the Document is included in an aggregate, this License does not apply to the other works in the aggregate which are not themselves derivative works of the Document.

If the Cover Text requirement of section 3 is applicable to these copies of the Document, then if the Document is less than one half of the entire aggregate, the Document's Cover Texts may be placed on covers that bracket the Document within the aggregate, or the electronic equivalent of covers if the Document is in electronic form. Otherwise they must appear on printed covers that bracket the whole aggregate.

C.9 Translation

Translation is considered a kind of modification, so you may distribute translations of the Document under the terms of section 4. Replacing Invariant Sections with translations requires special permission from their copyright holders, but you may include translations of some or all Invariant Sections in addition to the original versions of these Invariant Sections. You may include a translation of this License, and all the license notices in the Document, and any Warranty Disclaimers, provided that you also include the original English version of this License and the original versions of those notices and disclaimers. In case of a disagreement between the translation and the original version of this License or a notice or disclaimer, the original version will prevail.

If a section in the Document is entitled “Acknowledgments,” “Dedications,” or “History,” the requirement (section 4) to Preserve its Title (section 1) will typically require changing the actual title.

C.10 Termination

You may not copy, modify, sublicense, or distribute the Document except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense, or distribute it is void, and will automatically terminate your rights under this License.

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, receipt of a copy of some or all of the same material does not give you any rights to use it.

C.11 Future Revisions of this License

The Free Software Foundation may publish new, revised versions of the GNU Free Documentation License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns. See <http://www.gnu.org/copyleft/>.

Each version of the License is given a distinguishing version number. If the Document specifies that a particular numbered version of this License “or any later version” applies to it, you have the option of following the terms and conditions either of that specified version or of any later version that has been published (not as a draft) by the Free Software Foundation. If the Document does not specify a version number of this License, you may choose any version ever published (not as a draft) by the Free Software Foundation. If the Document specifies that a proxy can decide which future versions of this License can be used, that proxy's public statement of acceptance of a version permanently authorizes you to choose that version for the Document.

C.12 Relicensing

“Massive Multi-author Collaboration Site” (or “MMC Site”) means any World Wide Web server that publishes copyrightable works and also provides prominent facilities for anybody to edit those works. A public wiki that anybody can edit is an example of such a server. A “Massive Multi-author Collaboration” (or “MMC”) contained in the site means any set of copyrightable works thus published on the MMC site.

“CC-BY-SA” means the Creative Commons Attribution-Share Alike 3.0 license published by Creative Commons Corporation, a not-for-profit corporation with a principal place of business in San Francisco, California, as well as future copyleft versions of that license published by that same organization.

“Incorporate” means to publish or republish a Document, in whole or in part, as part of another Document.

An MMC is “eligible for relicensing” if it is licensed under this License and if all works that were first published under this License somewhere other than this MMC, and subsequently incorporated in whole or in part into the MMC, (1) had no cover texts or invariant sections, and (2) were thus incorporated prior to November 1, 2008.

The operator of an MMC Site may republish an MMC contained in the site under CC-BY-SA on the same site at any time before August 1, 2009, provided the MMC is eligible for relicensing.

16-Bit Assembler, Linker and Utilities User's Guide

NOTES:

Appendix D. Document Revision History

Revision A (September 2013)

Initial release of this document.

Revision B (December 2014)

- [Section 4.2 “Directives that Define Sections”](#) - Added notes to `.bss` and `.data` to warn against using all data memory for symbols so there is no room for stack.
- [Chapter 8. “Linker Command Line Options”](#) - Changed title from “Linker Command Line Interface”. Also placed option lists into tables to better highlight option/no option pairs.
- [Section 8.4.24 “-mreserve”](#) - Added this section.
- [Section 8.4.48 “--wrap symbol”](#) - Added more to this section.
- [Section 8.5.6 “--local-stack”](#) and [Section 8.5.7 “--no-local-stack”](#) - Added these sections.
- [Section 10.10 “Stack Allocation”](#) - Added text to warn against using all data memory for symbols so there is no room for stack.
- [Section 13.3 “Errors”](#) - Updated definition for “undefined reference to ‘symbol’”.

Revision C (August 2016)

- [Section 4.2 “Directives that Define Sections”](#) - Added to “Attributes that Represent Section Types”: `packedflash`; added to “Attributes that Modify Section Types”: `shared`, `preserved`, `update`, and `priority`.
- [Section 8.4 “Options that Control Output File Creation”](#) - Added options `--application-id`, `--coresident`, `--no-ivt`, `--pad-flash`, `--preserve`, `--preserve-all`, `--reserve-const`; Updated option `--no-isr`.
- [Section 8.6 “Options that Control Informational Output”](#) - Added option `--memory-usage`.
- [Section 10.15 “Co-resident Application Linking”](#) - Co-resident applications information and options.

16-Bit Assembler, Linker and Utilities User's Guide

Revision D (February 2018)

- [Section 3.9 “Special Operators”](#) - [Table 3-8](#) updated with footnote that two operators cannot be used in an expression.
- [Chapter 4. “Assembler Directives”](#) - Numbered all directive sections for better reference.
- [Section 4.2.7.2 “Attributes that Modify Section Types”](#) - [Table 4-2](#) updated for ymemory and dma attributes also applicable to dsPIC33EP devices.
- [Section 8.4 “Options that Control Output File Creation”](#) - added `--ivt` and `--no-ivt` linker options.
- [Section 8.6 “Options that Control Informational Output”](#) - added `--no-psrd-psrd-check` linker option.
- [Section 9.7.5.7 “Output Section Data”](#) - removed reference to QUAD; not used.
- [Section 10.14.5 “Example of Simple Bootloader Application”](#) - added example.
- [Chapter 12. “Linker Map File”](#) - updated linker map file.
- [Section 15.5 “xc16-objdump Utility”](#) - added `--psrd-psrd-check[=library]` option.

Revision E (December 2019)

- [Section 8.4.1 “--add-flags-code=, --add-flags-data=, --add-flags-const=”](#) : Added options.
- [Section 8.4.48 “--wrap symbol”](#): Updated `__wrap_symbol` to `_wrap_symbol` and `__real_symbol` to `_real_symbol`. Includes updated example.

Revision F (February 2021)

- [Section 8.4.23 “--msecondary-id, --msecondary-id-location”](#) - Added options related to main/secondary core devices.
- [Section 4.2.7.3 “Reserved Section Names with Implied Attributes”](#) - Changed footnote “33” from “dsPIC33F” to “dsPIC33x”.
- [Section 10.10 “Stack Allocation”](#) - Added description of Stage Usage Guidance analysis tool.

Index

Symbols

__CODE_LENGTH	204
__DATA_LENGTH	204
__RAM_USAGE	127
__reset	188
__ROM_USAGE	127
__main	188
__PROGRAM_END	204
.....	41
.abort	82
.align	72, 179
.apline	82
.ascii	64
.asciz	65
.bss	54, 68
.bss section	122, 158, 186
.byte	65
.comm	68
.comm symbol, length	68
.const section	179, 188, 189
.data	54
.data section	122, 186
.def	84
.dim	84
.dinit section	187, 188
.double	65
.eject	76
.else	77
.elseif	77
.end	82
.endif	84
.endif	77
.endm	80
.endr	79, 80, 81
.equ	40, 69
.equiv	40, 69
.err	77
.error	77
.exitm	79
.extern	68
.fail	82
.file	84
.fill	73
.fillupper	62
.fillvalue	62
.fixed	66
.float	66
.global	68
.globl	68
.handle	125
.handle section	141, 175, 185

.hword	67
.icd section	145
.ident	82
.if	77
.ifdecl	77
.ifdef	77
.ifndecl	77
.ifndef	78
.ifnotdecl	77
.ifnotdef	78
.incbin	82
.include	31, 33, 83
.init section	141
.int	67
.irp	79
.irpc	80
.isr section	141
.lcomm	68
.lib* section	141
.libc section	141
.libdsp section	141
.libm section	141
.libperi section	141
.line	84
.list	76
.ln	82
.loc	83
.long	67
.macro	80
.memory	55
.nolist	76
.org	74
.palign	73
.pascii	64, 65
.pasciz	65
.pbss section	125
.pbyte	65, 176, 189
.pfill	73
.pfillvalue	62
.pincbn	83
.popsection	55
.porg	74
.print	83
.psize	76
.pskip	75
.pspace	75
.pstring	67
.purgem	81
.pushsection	55
.pword	67, 189
.rept	81
.reset section	140

16-Bit Assembler, Linker and Utilities User's Guide

.sbtll	76	Directives	35, 53
.scl	84	Overview	15
.section name	55, 257	Source	88
.set	40, 41, 69	ASSERT	165, 170
.short	67	Assigning Output Sections to Regions	179
.single	66	Assigning Values	152
.size	84	auto_psv	210, 213
.sizeof	50	auxflash	56
.skip	74	auxpsv	56
.sleb128	85	B	
.space	74	Base Memory Addresses	139
.startof	50	bin2hex utility	240
.string	67	Binary File	111
.struct	75	BLOCK	170
.tag	85	--boot	131
.text	61	boot	57
.text section	122, 141, 186	Bootloader Example	200
.title	76	bss	56
.type	85	Building the Output File	175
.uleb128	85	Built-in Functions	169
.user_init section	141	ADDR	170
.val	85	ALIGN	170
.version	83	ASSERT	170
.weak	68, 184	BLOCK	170
.word	67	DEFINED	171
-(-)	116	LOADADDR	171
\$	41	MAX	171
A		MIN	171
-A	116	NEXT	171
-a	20	SIZEOF	171
a.out	17, 30, 120	C	
-a=file	27	Character Constants	39
-ac	20	Characters	39
Accessing Data	47	--check-sections	127
Accumulator Select	36	code	56
-ad	21	CODE_LENGTH	204
--add-flags-code	116	Command-Line Information	
ADDR	170	Linker Scripts	136
address	57	Command-Line Interface	
-ah	22	Archiver/Librarian	235
-ai	23	Assembler	19
Aiwt Region	138	Linker	113
-al	24	Comments	36, 149
ALIGN	170	Computing Absolute Addresses	175
align	57	Condition Codes	35
alignment gaps	194	Conditional Assembly Directives	
Allocating Memory	174	else	77
Allocating Unmapped Sections	180	elseif	77
-am	24	endif	77
-an	25	err	77
--application-id	116	error	77
ar utility	233	if	77
--architecture	116	ifdecl	77
Archiver	233	ifdef	77
Command-Line Interface	235	ifndekl	77
Scripts	237	ifndef	78
Arguments	36	ifnotdecl	77
-as	27	ifnotdef	78
ASCII Character Set	259	Configuration Region	138
Assembler		Constant Data	47
Command-Line Interface	19		

Constants	166	Define Symbols	69
Fixed-Point Numbers	39	Fill	62
Floating-Point Numbers	38	Initialization	64
Integer	38	Miscellaneous	82
Locating in Program Memory	210	Output Listing	76
Numeric	38	Section	54
COPY	161	Substitution/Expansion	79
Co-resident Applications	202	--discard-all	117
Creating Special Sections	175	--discard-locals	117
--cref	130	dma	57
CRT0	188	Documentation	
CRT1	188	Conventions	9
Current Address	41	Layout	7
Custom Linker Script	149	DOT Symbol	41
Customer Support	12	Dot Variable	167
D		-dp	117
-d	117	DSECT	161
data	56	E	
Data Initialization Template	187	eds	57
Data Memory	47, 176	eedata	56
Data Region	137	EEDATA Memory Region	139
DATA_LENGTH	204	EEPROM, Locating and Accessing	211
--data-init	125	Empty Expressions	44
-dc	117	--end-group	116
Debug Information Directives		ENTRY	165
.def	84	Escape Characters	39
.dim	84	Evaluation	168
.endif	84	Examples, Linker	205
.file	84	EXCLUDE_FILE	156
.line	84	Expressions	44
.scl	84	Expressions, Empty	44
.size	84	Expressions, Integer	44
.sleb128	85	EXTERN	165
.tag	85	F	
.type	85	--fatal-warnings	28
.uleb128	85	File Commands, Linker Scripts	
.val	85	GROUP	151
Debugger Memory	145	INCLUDE	150
Declare Symbols Directives		INPUT	150
.bss	68	OUTPUT	151
.comm	68	SEARCH_DIR	151
.extern	68	STARTUP	151
.global	68	File Extensions	
.globl	68	Assembler	16, 110
.lcomm	68	Linker	110
.weak	68	File Registers	35
Define Symbols Directives		Files	
.equ	69	Library	110
.equiv	69	Linker Output	111
.set	69	Linker Script	110
DEFINED	171	Listing	17
--defsym	31, 117	Map	111
Destination Select	36	Object	17, 110
Directive	35	Source	16
Directives		Fill Directives	
Alignment	70	.fillupper	62
Assembler	53	.fillvalue	62
Conditional	77	.pfillvalue	62
Debug Information	84	--fill-upper	118
Declare Symbols	68		

16-Bit Assembler, Linker and Utilities User's Guide

Fixed-Point Numbers.....	39	--warn-common	129
Floating-Point Numbers.....	38	--warn-once	129
FORCE_COMMON_ALLOCATION	165	--warn-section-align.....	130
--force-exe-suffix	118	-y	128
--force-link	118	Initialization Directives	
Functions, Locating	207	.ascii	64
G		.asciz	65
gaps, alignment	194	.byte	65
--gc-sections	118	.double	65
--general	131	.fixed	66
Getting a Grip	185	.float	66
Global Symbols	184	.hword	67
GROUP	151	.int	67
H		.long	67
handle()	49, 175, 185	.pascii	64, 65
Handles	185	.pasciz	65
--handles	125	.pbyte	65
Header	88	.pstring	67
--heap	125	.pword	67
heap	56	.short	67
Heap Allocation	192	.single	66
--help	29, 127	.string	67
Hexadecimal to Decimal Conversion	260	.word	67
High-level Source	22, 88	Initialized Data	186
I		INPUT	150
-l	31	Input Section	
-i	121	Common Symbols	158
ICD Memory	145	Example	158
INCLUDE	150	Wildcard Patterns	157
Infix Operators	45	Input/Output Section Map	140
INFO	161	Integer Expressions	44
info	57	Integers	38
Informational Output Options, Assembler		Internal Preprocessor	33
--fatal-warnings	28	Internet Address, Microchip	11
--help	29	Interrupt	
-J	29	Handlers	192
--no-warn	29	Vector Tables	147, 192
--target-help	29	--isr	118
-v	29	--ivt	118
--verbose	29	Ivt Region	138
--version	29	J	
-W	29	-J	29
--warn	29	K	
Informational Output Options, Linker		K Suffix	166
--check-sections	127	--keep-locals	29
--help	127	L	
--memory-usage	127	-L	29, 119
--no-check-sections	127	-l	119
--no-psrd-psrd-check	127	Label	34, 40, 41, 176, 185
--no-warn-mismatch	128	-legacy-libc	119
--report-mem	128	LENGTH	154
-t	128	libpic30-coff.a	188
--trace	128	libpic30-elf.a	188
--trace-symbol	128	Librarian	233
-V	128	Command-Line Interface	235
-v	128	Scripts	237
--verbose	128	--library	119
--version	128	Library Files	110

--library-path	119	.space	74
Link Map Options, Linkerstruct	75
--cref	130	M	
-M	130	-M	130
-Map	130	M Suffix	166
--print-map	130	-Map	130
Linker		Map File	111
Allocation	178	Mapping Sections	178
Command-Line Interface	113	MAX	171
Examples	205	-MD	30
File Extensions	110	memory	56
Output File	111	Memory Addressing	176
Overview	109	MEMORY Command	154
Processing	173	Memory Region Information	137
Linker Scripts	135	--memory-usage	127
Command Language	149	merge	57
Command-Line Information	136	MIN	171
Concepts	149	Miscellaneous Directives	
Contents	136	.abort	82
Custom	149	.apline	82
Expressions	166	.end	82
File	110	.fail	82
File Commands	150	.incbin	82
Other Commands	165	.include	83
Listing Files	17	.indent	82
Listing Output Options, Assembler	20	.In	82
-a=file	27	.loc	83
-ac	20	.pincbin	83
-ad	21	.print	83
-ah	22	.version	83
-ai	23	Mnemonic	34
-al	24	Modification Options, Archiver/Librarian	
-am	24	a	236
-an	25, 27	b	236
--listing-cont-lines	28	c	236
--listing-lhs-width	27	f	236
--listing-lhs-width2	27	i	236
--listing-rhs-width	28	l	236
--listing-cont-lines	28	N	236
--listing-lhs-width	27	o	236
--listing-lhs-width2	27	P	236
--listing-rhs-width	28	S	236
Literals	35	s	236
LMA	150, 162, 171	u	236
Load Memory Address	150, 162, 171	V	236
LOADADDR	171	v	236
Loading Input Files	174	-mreserve	120
Local Symbols	40	myMicrochip Personalized Notification Service	11
--local-stack	125	N	
Location Counter	142, 167	near	57
Location Counter Directives		NEXT	171
.align	72	nm utility	242
.fill	73	--no-check-sections	127
.org	74	NOCROSSREFS	165
.palign	73	--no-data-init	125
.pfill	73	--no-force-link	118
.porg	74	--no-handles	125
.pskip	75	--noinhibit-exec	120
.pspace	75	--no-isr	118
.skip	74		

16-Bit Assembler, Linker and Utilities User's Guide

--no-ivt	118	--no-sort	243
--no-keep-memory	120	--numeric-sort	243
NOLOAD	161	-o	243
noload	57	-omf	243
--no-local-stack	126	-P	243
--no-pack-data	126	-p	243
--no-psrd-psrd-check	127	--portability	243
--no-relax	30	--print-armap	243
--no-select-objects	122	--print-file-name	243
--no-undefined	123	-r	243
--no-warn	29	--radix=	243
--no-warn-mismatch	128	--reverse-sort	243
Numeric Constants	38	-s	243
O		--size-sort	243
-o	30, 120	-t	243
objdump utility	245	-u	243
Object Files	17, 110	--undefined-only	243
-omf	17, 30, 120, 241, 243, 247, 249, 251, 253	-V	243
Operands	35	-v	243
Operators	44, 168	--version	243
Infix	45	Options, pic30-objdump	
Prefix	44	-a	246
Optimize	120	--all-header	247
Options, Archiver/Librarian		--archive-header	246
d	235	-D	246
m	235	-d	246
p	235	--debugging	246
q	235	--disassemble	246
r	235	--disassemble-all	246
t	235	--disassembler-options=	246
x	235	--disassemble-zeroes	247
Options, Assembler		-EB	246
Informational Output	28	-EL	246
Listing Output	19	--endian=	246
Other	31	-f	246
Output File Creation	29	--file-header	246
Options, Linker		--file-start-context	246
Informational Output	127	--full-contents	247
Link Map Output	130	-g	246
Output File Creation	115	-H	246
Run-time Initialization	125	-h	246
Options, pic30-bin2hex		--header	246
-a	241	--help	246
-omf	241	-j	246
-u	241	-l	246
-v	241	--line-numbers	246
Options, pic30-nm		-M	246
-A	243	--no-show-raw-insn	247
-a	243	-omf	247
-B	243	--prefix-addresses	247
--debug-syms	243	--psrd-psrd-check	247
--defined-only	243	-r	247
--extern-only	243	--reloc	247
-f	243	-S	247
--format=	243	-s	247
-g	243	--section=	246
--help	243	--section-header	246
-l	243	--show-raw-insn	247
--line-numbers	243	--source	247
-n	243	--start-address=	247

--stop-address=	247	FORCE_COMMON_ALLOCATION	165
--syms	247	NOCROSSREFS	165
-t	247	OUTPUT_ARCH	165
-V	247	OUTPUT_FORMAT	165
--version	247	TARGET	165
-w	247	Other Options, Assembler	
--wide	247	--defsym	31
-x	247	-l	31
-Z	247	-p	31
Options, pic30-ranlib		--processor	31
-omf	249	OUTPUT	151
-V	249	--output	120
-v	249	Output File Creation Options, Assembler	
--version	249	--keep-locals	29
Options, pic30-strings		-L	29
-	251	-MD	30
-a	251	--no-relax	30
--all	251	-o	30
--bytes=	251	-omf	30
-f	251	-R	30
--help	251	--relax	30
-n	251	-Z	30
-omf	251	Output File Format	136
--print-file-name	251	Output File Options, Linker	
--radix=	251	-(-)	116
-t	251	-A	116
-v	251	--add-flags-code	
--version	251	--add-flags-data	116
Options, pic30-strip		--add-flags-data	116
--discard-all	253	--application-id	116
--discard-locals	253	--architecture	116
-g	253	-d	117
--help	253	-dc	117
-K	253	--defsym	117
--keep-symbol=	253	--discard-all	117
-N	253	--discard-locals	117
-o	253	-dp	117
-omf	253	--end-group	116
-p	253	--fill-upper	118
--preserve-dates	253	--force-exe-suffix	118
-R	253	--force-link	118
--remove-section=	253	--gc-sections	118
-S	253	-i	121
-s	253	--isr	118
--strip-all	253	--ivt	118
--strip-debug	253	-L	119
--strip-symbol=	253	-l	119
--strip-unneeded	253	-legacy-libc	119
-V	253	--library	119
-v	253	--library-path	119
--verbose	253	-mreserve	120
--version	253	--no-force-link	118
-X	253	--noinhibit-exec	120
-x	253	--no-isr	118
ORG	154	--no-ivt	118
ORIGIN	154	--no-keep-memory	120
Other Linker Script Commands		--no-select-objects	122
ASSERT	165	--no-undefined	123
ENTRY	165	-o	120
EXTERN	165	-omf	120

16-Bit Assembler, Linker and Utilities User's Guide

--output	120	INFO	161
--p	121	NOLOAD	161
--pad-flash	121	OVERLAY	161
--preserve	121	Output Sections in	
-preserve-all	121	Configuration Memory	144
--processor	121	OUTPUT_ARCH	165
-r	121	OUTPUT_FORMAT	165
--relocateable	121	OVERLAY	161
--reserve-const	121	Overlay Description	163
--retain-symbols-file	121	Overview	
-S	122	Assembler	15
-s	122	Linker	109
--script	121	P	
--select-objects	121	--p	121
--start-group	116	-p	31
--strip-all	122	--pack-data	126
--strip-debug	122	packedflash	56
-T	121	paddr()	49
-Tbss	122	--pad-flash	121
-Tdata	122	page	57
-Ttext	122	Page Size	48
-u	123	persist	56
--undefined	123	pic30-ar utility	233
-Ur	121	pic30-bin2hex utility	240
--wrap	123	pic30-nm utility	242
-X	117	pic30-objdump utility	245
-x	117	pic30-ranlib utility	249
Output Formats, pic30-nm		pic30-strings utility	250
?	244	pic30-strip utility	252
A	244	Pointer	47
B	244	Precedence	45, 168
C	244	Prefix Operators	44
D	244	Preprocessor, Internal	33
N	244	--preserve	121
R	244	-preserve-all	121
T	244	preserved	57
U	244	--print-map	130
V	244	priority	57
W	244	Processing, Linker	173
Output Listing Directives		--processor	31, 121
.eject	76	Program Address	49
.list	76	Program Memory	47, 176
.nolist	76	Program Memory, Locating and Reserving	214
.psize	76	Program Region	138
.sbttl	76	Program Space Visibility Window	48, 163, 177, 179, 188
.title	76	PROGRAM_END	204
Output Section		PROVIDE	153
.reset	140	psv	56
.text	141	PSV Window	48, 163, 177, 179, 188, 257
Address	159	psvoffset()	48
Attributes	161	psvpage()	48
Data	160	R	
Description	159	-R	30
Discarding	160	-r	121
Fill	162	ranlib utility	249
LMA	162	Reading, Recommended	10
Region	162	Read-Only Data	189
Type	161	Registers	35
COPY	161		
DSECT	161		

Relative Branches	30	Section of an Expression	169
Relative Calls	30	SECTIONS Command	155
--relax	30	--secure	131
relocatable	17	secure	57
Relocatable Code	206	Security Feature Options, Linker	
--relocateable	121	--boot	131
--report-mem	128	--general	131
--reserve-const	121	--secure	131
Reserved Names	40	--select-objects	121
RESET	192	SFR Addresses	148
Reset Region	138	SFRs	137, 148, 174
Resolving Symbols	174	shared	57
--retain-symbols-file	121	Simple Assignments	152
reverse	57	SIZEOF	171
Run-time Initialization Options, Linker		Source Code	34
--data-init	125	Source Files	16
--handles	125	Special Function Registers	137, 148, 174
--heap	125	Special Operators	46
--local-stack	125	.sizeof	46
--no-data-init	125	.startof	46
--no-handles	125	boot	46
--no-local-stack	126	dmaoffset	46
--no-pack-data	126	dmapage	46
--pack-data	126	edsoffset	46
--stack	126	edspage	46
--stackguard	126	handle	46
Run-time Library Support	188	paddr	46
S		psvoffset	46
-S	122	psvpage	46
-s	122	secure	46
--script	121	tbloffset	46
Scripts		tblpage	46
Librarian	237	SPLIM	188, 191
Scripts, Archiver/Librarian		--stack	126
ADDLIB	237	stack	56
ADDMOD	237	Stack Allocation	191
CLEAR	237	Stack Pointer	188, 191
CREATE	237, 238	Stack Pointer Limit Register	188, 191
DELETE	238	Stack, Locating	213
DIRECTORY	238	--stackguard	126
END	238	Standard Data Section Names	186
EXTRACT	238	--start-group	116
LIST	238	Starting Address	50
OPEN	237, 238	STARTUP	151
REPLACE	238	Start-up Code	187
SAVE	237, 238	Start-up Module	188
VERBOSE	238	Startup Modules	136
SEARCH_DIR	151	Statement Format	34
Section Attributes		strings utility	250
Modify Section Types	57	strip utility	252
Represent Section Types	56	--strip-all	122
Reserved Section Names	59	--strip-debug	122
Section Directives		Substitution/Expansion Directives	
.bss	54	.endm	80
.data	54	.endr	79, 80, 81
.memory	55	.exitm	79
.popsection	55	.irpc	80
.pushsection	55	.macro	80
.section name	55, 257	.purgem	81
.text	61	.rept	81

16-Bit Assembler, Linker and Utilities User's Guide

irp	79
Subtitle	76, 88
Symbol Names	166
Symbol Table	27, 84, 85, 88, 150, 171
Syntax	
Archiver/Librarian	234
Assembler	19
Linker	113
pic30-bin2hex	240
pic30-nm	242
pic30-objdump	245
pic30-ranlib	249
pic30-strings	250
pic30-strip	252

T

-T	121
-t	128
Table Access Instructions	176
TARGET	165
--target-help	29
tboffset()	47, 176
tblpage()	47, 176
-Tbss	122
-Tdata	122
Title	76
Title Line	88
--trace	128
--trace-symbol	128
-Ttext	122

U

-u	123
--undefined	123
update	57
-Ur	121
User-Defined Section in Data Memory	146
User-Defined Section in Program Memory	142
Utilities	231

V

-V	128
-v	29, 128
Variables, Locating	207
--verbose	29, 128
--version	29, 128
Virtual Memory Address	150, 162
VMA	150, 162

W

-W	29
W15	188, 191
--warn	29
--warn-common	129
--warn-once	129
--warn-section-align	130
Watchdog Timer, Disabling	144
Weak Symbols	184
Web Site, Microchip	11
White Space	34
--wrap	123

X

-X	117
-x	117
X Memory, Creating a Modulo Buffer	212
xc16	246, 253
xmemory	57

Y

-y	128
Y Memory, Creating a Modulo Buffer	213
ymemory	57

Z

-Z	30
----------	----

Worldwide Sales and Service

AMERICAS

Corporate Office
2355 West Chandler Blvd.
Chandler, AZ 85224-6199
Tel: 480-792-7200
Fax: 480-792-7277
Technical Support:
<http://www.microchip.com/support>
Web Address:
www.microchip.com

Atlanta
Duluth, GA
Tel: 678-957-9614
Fax: 678-957-1455

Austin, TX
Tel: 512-257-3370

Boston
Westborough, MA
Tel: 774-760-0087
Fax: 774-760-0088

Chicago
Itasca, IL
Tel: 630-285-0071
Fax: 630-285-0075

Dallas
Addison, TX
Tel: 972-818-7423
Fax: 972-818-2924

Detroit
Novi, MI
Tel: 248-848-4000

Houston, TX
Tel: 281-894-5983

Indianapolis
Noblesville, IN
Tel: 317-773-8323
Fax: 317-773-5453
Tel: 317-536-2380

Los Angeles
Mission Viejo, CA
Tel: 949-462-9523
Fax: 949-462-9608
Tel: 951-273-7800

Raleigh, NC
Tel: 919-844-7510

New York, NY
Tel: 631-435-6000

San Jose, CA
Tel: 408-735-9110
Tel: 408-436-4270

Canada - Toronto
Tel: 905-695-1980
Fax: 905-695-2078

ASIA/PACIFIC

Australia - Sydney
Tel: 61-2-9868-6733

China - Beijing
Tel: 86-10-8569-7000

China - Chengdu
Tel: 86-28-8665-5511

China - Chongqing
Tel: 86-23-8980-9588

China - Dongguan
Tel: 86-769-8702-9880

China - Guangzhou
Tel: 86-20-8755-8029

China - Hangzhou
Tel: 86-571-8792-8115

China - Hong Kong SAR
Tel: 852-2943-5100

China - Nanjing
Tel: 86-25-8473-2460

China - Qingdao
Tel: 86-532-8502-7355

China - Shanghai
Tel: 86-21-3326-8000

China - Shenyang
Tel: 86-24-2334-2829

China - Shenzhen
Tel: 86-755-8864-2200

China - Suzhou
Tel: 86-186-6233-1526

China - Wuhan
Tel: 86-27-5980-5300

China - Xian
Tel: 86-29-8833-7252

China - Xiamen
Tel: 86-592-2388138

China - Zhuhai
Tel: 86-756-3210040

ASIA/PACIFIC

India - Bangalore
Tel: 91-80-3090-4444

India - New Delhi
Tel: 91-11-4160-8631

India - Pune
Tel: 91-20-4121-0141

Japan - Osaka
Tel: 81-6-6152-7160

Japan - Tokyo
Tel: 81-3-6880-3770

Korea - Daegu
Tel: 82-53-744-4301

Korea - Seoul
Tel: 82-2-554-7200

Malaysia - Kuala Lumpur
Tel: 60-3-7651-7906

Malaysia - Penang
Tel: 60-4-227-8870

Philippines - Manila
Tel: 63-2-634-9065

Singapore
Tel: 65-6334-8870

Taiwan - Hsin Chu
Tel: 886-3-577-8366

Taiwan - Kaohsiung
Tel: 886-7-213-7830

Taiwan - Taipei
Tel: 886-2-2508-8600

Thailand - Bangkok
Tel: 66-2-694-1351

Vietnam - Ho Chi Minh
Tel: 84-28-5448-2100

EUROPE

Austria - Wels
Tel: 43-7242-2244-39
Fax: 43-7242-2244-393

Denmark - Copenhagen
Tel: 45-4485-5910
Fax: 45-4485-2829

Finland - Espoo
Tel: 358-9-4520-820

France - Paris
Tel: 33-1-69-53-63-20
Fax: 33-1-69-30-90-79

Germany - Garching
Tel: 49-8931-9700

Germany - Haan
Tel: 49-2129-3766400

Germany - Heilbronn
Tel: 49-7131-72400

Germany - Karlsruhe
Tel: 49-721-625370

Germany - Munich
Tel: 49-89-627-144-0
Fax: 49-89-627-144-44

Germany - Rosenheim
Tel: 49-8031-354-560

Israel - Ra'anana
Tel: 972-9-744-7705

Italy - Milan
Tel: 39-0331-742611
Fax: 39-0331-466781

Italy - Padova
Tel: 39-049-7625286

Netherlands - Drunen
Tel: 31-416-690399
Fax: 31-416-690340

Norway - Trondheim
Tel: 47-7288-4388

Poland - Warsaw
Tel: 48-22-3325737

Romania - Bucharest
Tel: 40-21-407-87-50

Spain - Madrid
Tel: 34-91-708-08-90
Fax: 34-91-708-08-91

Sweden - Gothenberg
Tel: 46-31-704-60-40

Sweden - Stockholm
Tel: 46-8-5090-4654

UK - Wokingham
Tel: 44-118-921-5800
Fax: 44-118-921-5820